

The OpenCL[™] Extension Specification

Khronos® OpenCL Working Group

Version v3.0.8, Wed, 30 Jun 2021 20:00:00 +0000: from git branch: master commit: 09130de814688ec7b463cb089986b807c628ead3

Table of Contents

1. Extensions Overview	2
2. Installable Client Drivers	6
3. Byte Addressable Stores	12
4. Writing to 3D Image Objects	
5. Half Precision Floating-Point	
6. Double Precision Floating-Point	52
7. 32-bit Atomics	74
8. 64-bit Atomics	78
9. Selecting the Rounding Mode (DEPRECATED)	81
10. Creating an OpenCL Context from an OpenGL Context or Share Group	83
11. Creating OpenCL Memory Objects from OpenGL Objects	91
12. Creating OpenCL Event Objects from OpenGL Sync Objects	103
13. Creating OpenCL Memory Objects from Direct3D 10 Buffers and Textures	108
14. Creating OpenCL Memory Objects from Direct3D 11 Buffers and Textures	124
15. Creating OpenCL Memory Objects from DirectX 9 Media Surfaces	139
16. Depth Images	152
17. Sharing OpenGL and OpenGL ES Depth and Depth-Stencil Images	158
18. Creating OpenCL Memory Objects from OpenGL MSAA Textures	160
19. Creating OpenCL Event Objects from EGL Sync Objects	167
20. Creating OpenCL Memory Objects from EGL Images	171
21. Creating a 2D Image From A Buffer	178
22. Local and Private Memory Initialization	180
23. Terminating OpenCL contexts	182
24. Standard Portable Intermediate Representation Binaries	185
25. Intermediate Language Programs	187
26. Creating Command Queues with Properties	192
27. Device Enqueue Local Argument Types	195
28. Subgroups	196
29. Mipmaps	206
30. sRGB Image Writes	218
31. Priority Hints	219
32. Throttle Hints	220
33. Named Barriers for Subgroups	221
34. Extended Async Copies (Provisional)	222
35. Async Work Group Copy Fence (Provisional)	226
36. Unique Device Identifiers	228
37. Extended versioning	230
38. Extended Subgroup Functions	235

39. PCI Bus Information Query	257
40. Extended Bit Operations	259
41. Suggested Local Work Size Query	262
42. Integer dot product	265
43. Extensions to the OpenCL SPIR-V Environment.	269
Index	270
Appendix A: Extensions Promoted to Core Features	271
Appendix B: Deprecated Extensions	272
Appendix C: Quick Reference	273

Copyright 2008-2021 The Khronos Group.

This specification is protected by copyright laws and contains material proprietary to the Khronos Group, Inc. Except as described by these terms, it or any components may not be reproduced, republished, distributed, transmitted, displayed, broadcast or otherwise exploited in any manner without the express prior written permission of Khronos Group.

Khronos Group grants a conditional copyright license to use and reproduce the unmodified specification for any purpose, without fee or royalty, EXCEPT no licenses to any patent, trademark or other intellectual property rights are granted under these terms. Parties desiring to implement the specification and make use of Khronos trademarks in relation to that implementation, and receive reciprocal patent license protection under the Khronos IP Policy must become Adopters and confirm the implementation as conformant under the process defined by Khronos for this specification; see https://www.khronos.org/adopters.

Khronos Group makes no, and expressly disclaims any, representations or warranties, express or implied, regarding this specification, including, without limitation: merchantability, fitness for a particular purpose, non-infringement of any intellectual property, correctness, accuracy, completeness, timeliness, and reliability. Under no circumstances will the Khronos Group, or any of its Promoters, Contributors or Members, or their respective partners, officers, directors, employees, agents or representatives be liable for any damages, whether direct, indirect, special or consequential damages for lost revenues, lost profits, or otherwise, arising from or in connection with these materials.

Vulkan and Khronos are registered trademarks, and OpenXR, SPIR, SPIR-V, SYCL, WebGL, WebCL, OpenVX, OpenVG, EGL, COLLADA, glTF, NNEF, OpenKODE, OpenKCAM, StreamInput, OpenWF, OpenSL ES, OpenMAX, OpenMAX AL, OpenMAX IL, OpenMAX DL, OpenML and DevU are trademarks of the Khronos Group Inc. ASTC is a trademark of ARM Holdings PLC, OpenCL is a trademark of Apple Inc. and OpenGL and OpenML are registered trademarks and the OpenGL ES and OpenGL SC logos are trademarks of Hewlett Packard Enterprise used under license by Khronos. All other product names, trademarks, and/or company names are used solely for identification and belong to their respective owners.

Chapter 1. Extensions Overview

This document describes the list of optional features supported by OpenCL. Optional extensions are not required to be supported by a conformant OpenCL implementation, but are expected to be widely available, and in some cases may define functionality that is likely to be required in a future revision of the OpenCL specification.

This document describes all extensions that have been approved by the OpenCL working group. It is a *unified* specification, meaning that the extensions described in this document are not specific to a specific core OpenCL specification version.

OpenCL extensions approved by the OpenCL working group may be *promoted* to core features in later revisions of OpenCL. When this occurs, the feature described by the extension specification is merged into the core OpenCL specification. The extension will continue to be documented in this specification, both for backwards compatibility and for devices that wish to support the feature but that are unable to support the newer core OpenCL version.

1.1. Naming Convention for Optional Extensions

OpenCL extensions approved by the OpenCL working group use the following naming convention:

- A unique *name string* of the form "cl_khr_<name>" is associated with each extension. If the extension is supported by an implementation, this string will be present in the implementation's CL_PLATFORM_EXTENSIONS string or CL_DEVICE_EXTENSIONS string.
- All API functions defined by the extension will have names of the form cl<function_name
 >KHR.
- All enumerants defined by the extension will have names of the form CL_<enum_name>_KHR.

Functions and enumerants defined by extensions that are promoted to core features will have their **KHR** affix removed. OpenCL implementations of such later revisions must also export the name strings of promoted extensions in the CL_PLATFORM_EXTENSIONS or CL_DEVICE_EXTENSIONS string, and support the **KHR**-affixed versions of functions and enumerants as a transition aid.

Vendor extensions are strongly encouraged to follow a similar naming convention:

- A unique *name string* of the form "cl_<vendor_name>_<name>" is associated with each extension. If the extension is supported by an implementation, this string will be present in the implementation's CL_PLATFORM_EXTENSIONS string or CL_DEVICE_EXTENSIONS string.
- All API functions defined by the vendor extension will have names of the form cl<function_name><vendor_name>.
- All enumerants defined by the vendor extension will have names of the form **CL_**< *enum_name*>_<*vendor_name*>.

1.2. Compiler Directives for Optional Extensions

The #pragma OPENCL EXTENSION directive controls the behavior of the OpenCL compiler with

respect to extensions. The **#pragma OPENCL EXTENSION** directive is defined as:

```
#pragma OPENCL EXTENSION <extension_name> : <behavior>
#pragma OPENCL EXTENSION all : <behavior>
```

where <code>extension_name</code> is the name of the extension. The <code>extension_name</code> will have names of the form <code>cl_khr_<name></code> for an extension approved by the OpenCL working group and will have names of the form <code>cl_<vendor_name>_<name></code> for vendor extensions. The token <code>all</code> means that the behavior applies to all extensions supported by the compiler. The <code>behavior</code> can be set to one of the following values given by the table below.

behavior	Description
enable	Behave as specified by the extension <i>extension_name</i> .
	Report an error on the #pragma OPENCL EXTENSION if the <i>extension_name</i> is not supported, or if all is specified.
disable	Behave (including issuing errors and warnings) as if the extension <i>extension_name</i> is not part of the language definition.
	If all is specified, then behavior must revert back to that of the non-extended core version of the language being compiled to.
	Warn on the #pragma OPENCL EXTENSION if the extension <i>extension_name</i> is not supported.

The **#pragma OPENCL EXTENSION** directive is a simple, low-level mechanism to set the behavior for each extension. It does not define policies such as which combinations are appropriate; those must be defined elsewhere. The order of directives matter in setting the behavior for each extension. Directives that occur later override those seen earlier. The **all** variant sets the behavior for all extensions, overriding all previously issued extension directives, but only if the *behavior* is set to **disable**.

The initial state of the compiler is as if the directive

```
#pragma OPENCL EXTENSION all : disable
```

was issued, telling the compiler that all error and warning reporting must be done according to this specification, ignoring any extensions.

Every extension which affects the OpenCL language semantics, syntax or adds built-in functions to the language must create a preprocessor #define that matches the extension name string. This #define would be available in the language if and only if the extension is supported on a given implementation.

Example:

An extension which adds the extension string "cl_khr_3d_image_writes" should also add a preprocessor #define called cl_khr_3d_image_writes. A kernel can now use this preprocessor #define to do something like:

```
#ifdef cl_khr_3d_image_writes
  // do something using the extension
#else
  // do something else or #error!
#endif
```

1.3. Getting OpenCL API Extension Function Pointers

The function

returns the address of the extension function named by *funcname* for a given *platform* The pointer returned should be cast to a function pointer type matching the extension function's definition defined in the appropriate extension specification and header file. A return value of NULL indicates that the specified function does not exist for the implementation or *platform* is not a valid platform. A non-NULL return value for **clGetExtensionFunctionAddressForPlatform** does not guarantee that an extension function is actually supported by the platform. The application must also make a corresponding query using **clGetPlatformInfo**(platform, CL_PLATFORM_EXTENSIONS, ...) or **clGetDeviceInfo**(device, CL_DEVICE_EXTENSIONS, ...) to determine if an extension is supported by the OpenCL implementation.

Since there is no way to qualify the query with a device, the function pointer returned must work for all implementations of that extension on different devices for a platform. The behavior of calling a device extension function on a device not supporting that extension is undefined.

clGetExtensionFunctionAddressForPlatform may not be be used to query for core (non-extension) functions in OpenCL. For extension functions that may be queried using **clGetExtensionFunctionAddressForPlatform**, implementations may also choose to export those functions statically from the object libraries implementing those functions, however, portable applications cannot rely on this behavior.

Function pointer typedefs must be declared for all extensions that add API entrypoints. These typedefs are a required part of the extension interface, to be provided in an appropriate header (such as cl_ext.h if the extension is an OpenCL extension, or cl_gl_ext.h if the extension is an OpenCL / OpenGL sharing extension).

The following convention must be followed for all extensions affecting the host API:

where TAG can be KHR, EXT or vendor-specific.

Consider, for example, the **cl_khr_gl_sharing** extension. This extension would add the following to cl_gl_ext.h:

```
#ifndef cl_khr_gl_sharing
#define cl khr gl sharing 1
// all data typedefs, token #defines, prototypes, and
// function pointer typedefs for this extension
#define CL INVALID GL SHAREGROUP REFERENCE KHR -1000
#define CL_CURRENT_DEVICE_FOR_GL_CONTEXT_KHR
                                                 0x2006
#define CL DEVICES FOR GL CONTEXT KHR
                                                 0x2007
#define CL_GL_CONTEXT_KHR
                                                 0x2008
#define CL_EGL_DISPLAY_KHR
                                                 0x2009
#define CL GLX DISPLAY KHR
                                                 0x200A
#define CL WGL HDC KHR
                                                 0x200B
#define CL_CGL_SHAREGROUP_KHR
                                                 0x200C
// function pointer typedefs must use the
// following naming convention
typedef cl_int
        (CL_API_CALL *clGetGLContextInfoKHR_fn)(
            const cl_context_properties * /* properties */,
            cl_gl_context_info /* param_name */,
            size_t /* param_value_size */,
            void * /* param_value */,
            size t * /*param value size ret*/);
#endif // cl_khr_gl_sharing
```

Chapter 2. Installable Client Drivers

2.1. Overview

This section describes a platform extension which defines a simple mechanism through which the Khronos OpenCL installable client driver loader (ICD Loader) may expose multiple separate vendor installable client drivers (Vendor ICDs) for OpenCL. An application written against the ICD Loader will be able to access all cl_platform_ids exposed by all vendor implementations with the ICD Loader acting as a demultiplexor.

This is a platform extension, so if this extension is supported by an implementation, the string **cl_khr_icd** will be present in the **CL_PLATFORM_EXTENSIONS** string.

2.2. General information

2.2.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

2.3. Inferring Vendors from Function Call Arguments

At every OpenCL function call, the ICD Loader infers the vendor ICD function to call from the arguments to the function. An object is said to be ICD compatible if it is of the following structure:

```
struct _cl_<object>
{
    struct _cl_icd_dispatch *dispatch;
    // ... remainder of internal data
};
```

<object> is one of platform_id, device_id, context, command_queue, mem, program, kernel, event, or sampler.

The structure <u>cl_icd_dispatch</u> is a function pointer dispatch table which is used to direct calls to a particular vendor implementation. All objects created from ICD compatible objects must be ICD compatible.

The definition for _cl_icd_dispatch is provided along with the OpenCL headers. Existing members can never be removed from that structure but new members can be appended.

Functions which do not have an argument from which the vendor implementation may be inferred have been deprecated and may be ignored.

2.4. ICD Data

A Vendor ICD is defined by two pieces of data:

- The Vendor ICD library specifies a library which contains the OpenCL entry points for the vendor's OpenCL implementation. The vendor ICD's library file name should include the vendor name, or a vendor-specific implementation identifier.
- The Vendor ICD extension suffix is a short string which specifies the default suffix for extensions implemented only by that vendor. The vendor suffix string is optional.

2.5. ICD Loader Vendor Enumeration on Windows

To enumerate Vendor ICDs on Windows, the ICD Loader will first scan for REG_SZ string values in the "Display Adapter" and "Software Components" HKR registry keys. The exact registry keys to scan should be obtained via PnP Configuration Manager APIs, but will look like:

For 64-bit ICDs:

HKLM\SYSTEM\CurrentControlSet\Control\Class\
{Display Adapter GUID}\{Instance ID}\OpenCLDriverName, or

HKLM\SYSTEM\CurrentControlSet\Control\Class\
{Software Component GUID}\{Instance ID}\OpenCLDriverName

For 32-bit ICDs:

HKLM\SYSTEM\CurrentControlSet\Control\Class\
{Display Adapter GUID}\{Instance ID}\OpenCLDriverNameWoW, or

HKLM\SYSTEM\CurrentControlSet\Control\Class\
{Software Component GUID}\{Instance ID}\OpenCLDriverNameWoW

These registry values contain the path to the Vendor ICD library. For example, if the registry contains the value:

[HKLM\SYSTEM\CurrentControlSet\Control\Class\{GUID}\{Instance}]
"OpenCLDriverName"="c:\\vendor a\\vndra_ocl.dll"

Then the ICD Loader will open the Vendor ICD library:

c:\vendor a\vndra_ocl.dll

The ICD Loader will also scan for REG_DWORD values in the registry key:

HKLM\SOFTWARE\Khronos\OpenCL\Vendors

For each registry value in this key which has data set to 0, the ICD Loader will open the Vendor ICD library specified by the name of the registry value.

For example, if the registry contains the value:

```
[HKLM\SOFTWARE\Khronos\OpenCL\Vendors]
"c:\vendor a\vndra_ocl.dll"=dword:00000000
```

Then the ICD will open the Vendor ICD library:

c:\vendor a\vndra_ocl.dll

2.6. ICD Loader Vendor Enumeration on Linux

To enumerate vendor ICDs on Linux, the ICD Loader scans the files in the path /etc/OpenCL/vendors. For each file in this path, the ICD Loader opens the file as a text file. The expected format for the file is a single line of text which specifies the Vendor ICD's library. The ICD Loader will attempt to open that file as a shared object using dlopen(). Note that the library specified may be an absolute path or just a file name.

For example, if the following file exists

```
/etc/OpenCL/vendors/VendorA.icd
```

and contains the text

libVendorAOpenCL.so

then the ICD Loader will load the library libVendorAOpenCL.so.

2.7. ICD Loader Vendor Enumeration on Android

To enumerate vendor ICDs on Android, the ICD Loader scans the files in the path /system/vendor/Khronos/OpenCL/vendors. For each file in this path, the ICD Loader opens the file as a text file. The expected format for the file is a single line of text which specifies the Vendor ICD's library. The ICD Loader will attempt to open that file as a shared object using dlopen(). Note that the library specified may be an absolute path or just a file name.

For example, if the following file exists

/system/vendor/Khronos/OpenCL/vendors/VendorA.icd

and contains the text

```
libVendorAOpenCL.so
```

then the ICD Loader will load the library libVendorAOpenCL.so.

2.8. Adding a Vendor Library

Upon successfully loading a Vendor ICD's library, the ICD Loader queries the following functions from the library: **clIcdGetPlatformIDsKHR**, **clGetPlatformInfo**, and **clGetExtensionFunctionAddress** (note: **clGetExtensionFunctionAddress** has been deprecated, but is still required for the ICD loader). If any of these functions are not present then the ICD Loader will close and ignore the library.

Next the ICD Loader queries available ICD-enabled platforms in the library using **clicdGetPlatformIDsKHR**. For each of these platforms, the ICD Loader queries the platform's extension string to verify that **cl_khr_icd** is supported, then queries the platform's Vendor ICD extension suffix using **clGetPlatformInfo** with the value CL_PLATFORM_ICD_SUFFIX_KHR.

If any of these steps fail, the ICD Loader will ignore the Vendor ICD and continue on to the next.

2.9. New Procedures and Functions

2.10. New Tokens

Accepted as *param_name* to the function **clGetPlatformInfo**:

```
CL_PLATFORM_ICD_SUFFIX_KHR
```

Returned by **clGetPlatformIDs** when no platforms are found:

```
CL_PLATFORM_NOT_FOUND_KHR
```

2.11. Additions to Chapter 4 of the OpenCL 2.2 Specification

In section 4.1, replace the description of the return values of **clGetPlatformIDs** with:

"clGetPlatformIDs* returns CL_SUCCESS if the function is executed successfully and there are a non zero number of platforms available. It returns CL_PLATFORM_NOT_FOUND_KHR if zero platforms are available. It returns CL_INVALID_VALUE if *num_entries* is equal to zero and *platforms* is not NULL or if both *num_platforms* and *platforms* are NULL."

In section 4.1, add the following after the description of **clGetPlatformIDs**:

"The list of platforms accessible through the Khronos ICD Loader can be obtained using the following function:

num_entries is the number of cl_platform_id entries that can be added to *platforms*. If *platforms* is not NULL, then *num_entries* must be greater than zero.

platforms returns a list of OpenCL platforms available for access through the Khronos ICD Loader. The cl_platform_id values returned in platforms are ICD compatible and can be used to identify a specific OpenCL platform. If the platforms argument is NULL, then this argument is ignored. The number of OpenCL platforms returned is the minimum of the value specified by num_entries or the number of OpenCL platforms available.

num_platforms returns the number of OpenCL platforms available. If *num_platforms* is NULL, then this argument is ignored.

clicdGetPlatformIDsKHR returns CL_SUCCESS if the function is executed successfully and there are a non zero number of platforms available. It returns CL_PLATFORM_NOT_FOUND_KHR if zero platforms are available. It returns CL_INVALID_VALUE if *num_entries* is equal to zero and *platforms* is not NULL or if both *num_platforms* and *platforms* are NULL."

Add the following to table 4.1:

cl_platform_info enum	Return Type	Description
CL_PLATFORM_ICD_SUFFIX_KHR	char[]	The function name suffix used to
		identify extension functions to be
		directed to this platform by the ICD
		Loader.

2.12. Source Code

The official source for the ICD loader is available on github, at:

https://github.com/KhronosGroup/OpenCL-ICD-Loader

The complete <u>cl_icd_dispatch</u> structure is defined in the header <u>cl_icd.h</u>, which is available as a part of the OpenCL headers.

2.13. Issues

1. Some OpenCL functions do not take an object argument from which their vendor library may be identified (e.g, clUnloadCompiler), how will they be handled?

RESOLVED: Such functions will be a noop for all calls through the ICD.

2. How are OpenCL extension to be handled?

RESOLVED: OpenCL extension functions may be added to the ICD as soon as they are implemented by any vendor. The suffix mechanism provides access for vendor extensions which are not yet added to the ICD.

3. How will the ICD handle a NULL cl_platform_id?

RESOLVED: The ICD will by default choose the first enumerated platform as the NULL platform. The user can override this default by setting an environment variable OPENCL_ICD_DEFAULT_PLATFORM to the desired platform index. The API calls that deal with platforms will return CL_INVALID_PLATFORM if the index is not between zero and (number of platforms - 1), both inclusive.

4. There exists no mechanism to unload the ICD, should there be one?

RESOLVED: As there is no standard mechanism for unloading a vendor implementation, do not add one for the ICD.

5. How will the ICD loader handle NULL objects passed to the OpenCL functions?

RESOLVED: The ICD loader will check for NULL objects passed to the OpenCL functions without trying to dereference the NULL objects for obtaining the ICD dispatch table. On detecting a NULL object it will return one of the CL_INVALID_* error values corresponding to the object in question.

Chapter 3. Byte Addressable Stores

This section describes the **cl_khr_byte_addressable_store** extension. This extension relaxes restrictions on pointers to **char**, **uchar**2, **uchar**2, **uchar**2, **short**, **ushort** and **half** that were present in *Section 6.8m: Restrictions* of the OpenCL 1.0 specification. With this extension, applications are able to read from and write to pointers to these types.

This extension became a core feature in OpenCL 1.1.

3.1. General information

3.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

Chapter 4. Writing to 3D Image Objects

This section describes the cl_khr_3d_image_writes extension.

This extension adds built-in functions that allow a kernel to write to 3D image objects in addition to 2D image objects.

This extension became a core feature in OpenCL 2.0.

4.1. General information

4.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

The new built-in functions are described in the table below:

Table 1. 3D Image Built-in Image Write Functions

Function void write_imagef (image3d t image, int4 coord, float4 *color*) void write_imagei (image3d_t image, int4 coord, int4 color) void write_imageui (image3d_t image, int4 coord, uint4 color) stored.

Description

Write *color* value to the location specified by coordinate (x, y, z) in the 3D image specified by *image*. The appropriate data format conversion to the specified image format is done before writing the color value. coord.x, coord.y, and coord.z are considered to be unnormalized coordinates and must be in the range 0 ... image width - 1, 0 ... image height - 1, and 0 ... image depth - 1.

write_imagef can only be used with image objects created with *image_channel_data_type* set to one of the pre-defined packed formats or set to CL_SNORM_INT8, CL_UNORM_INT8, CL_SNORM_INT16, CL_UNORM_INT16, CL HALF FLOAT, or CL FLOAT. Appropriate data format conversion will be done to convert the channel data from a floating-point value to the actual data format in which the channels are

write imagei can only be used with image objects created with *image_channel_data_type* set to one of the following values: CL SIGNED INT8, CL_SIGNED_INT16, or CL_SIGNED_INT32.

write_imageui can only be used with image objects created with image_channel_data_type set to one of the following values: CL_UNSIGNED_INT8, CL_UNSIGNED_INT16, or CL_UNSIGNED_INT32.

The behavior of write_imagef, write_imagei, and write_imageui for image objects created with image channel data type values not specified in the description above, or with (x, y,z) coordinate values that are not in the range (0 ... image width - 1, 0 ... image height - 1, 0 ... image depth - 1) respectively, is undefined.

Chapter 5. Half Precision Floating-Point

This section describes the **cl_khr_fp16** extension. This extension adds support for half scalar and vector types as built-in types that can be used for arithmetic operations, conversions etc.

5.1. General information

5.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

5.2. Additions to Chapter 6 of the OpenCL 2.0 C Specification

The list of built-in scalar, and vector data types defined in *tables 6.1*, and *6.2* are extended to include the following:

Туре	Description
half2	A 2-component half-precision floating-point vector.
half3	A 3-component half-precision floating-point vector.
half4	A 4-component half-precision floating-point vector.
half8	A 8-component half-precision floating-point vector.
half16	A 16-component half-precision floating-point vector.

The built-in vector data types for halfn are also declared as appropriate types in the OpenCL API (and header files) that can be used by an application. The following table describes the built-in vector data types for halfn as defined in the OpenCL C programming language and the corresponding data type available to the application:

Type in OpenCL Language	API type for application
half2	cl_half2
half3	cl_half3
half4	cl_half4
half8	cl_half8
half16	cl_half16

The relational, equality, logical and logical unary operators described in *section 6.3* can be used with half scalar and halfn vector types and shall produce a scalar int and vector shortn result respectively.

The OpenCL compiler accepts an h and H suffix on floating point literals, indicating the literal is

typed as a half.

5.2.1. Conversions

The implicit conversion rules specified in *section 6.2.1* now include the half scalar and halfn vector data types.

The explicit casts described in *section 6.2.2* are extended to take a half scalar data type and a halfn vector data type.

The explicit conversion functions described in *section 6.2.3* are extended to take a half scalar data type and a halfn vector data type.

The as_typen() function for re-interpreting types as described in *section 6.2.4.2* is extended to allow conversion-free casts between shortn, ushortn, and halfn scalar and vector data types.

5.2.2. Math Functions

The built-in math functions defined in *table 6.8* (also listed below) are extended to include appropriate versions of functions that take half and half{2|3|4|8|16} as arguments and return values. gentype now also includes half, half2, half3, half4, half8, and half16.

For any specific use of a function, the actual type has to be the same for all arguments and the return type.

Table 2. Half Precision Built-in Math Functions

Function	Description
gentype acos (gentype <i>x</i>)	Arc cosine function.
gentype acosh (gentype x)	Inverse hyperbolic cosine.
gentype acospi (gentype x)	Compute acos (x) / π .
gentype asin (gentype x)	Arc sine function.
gentype asinh (gentype <i>x</i>)	Inverse hyperbolic sine.
gentype asinpi (gentype x)	Compute asin (x) / π .
gentype atan (gentype <i>y_over_x</i>)	Arc tangent function.
gentype atan2 (gentype <i>y</i> , gentype <i>x</i>)	Arc tangent of y / x .
gentype atanh (gentype <i>x</i>)	Hyperbolic arc tangent.
gentype atanpi (gentype <i>x</i>)	Compute atan (x) / π .
gentype atan2pi (gentype <i>y</i> , gentype <i>x</i>)	Compute atan2 $(y, x) / \pi$.
gentype cbrt (gentype x)	Compute cube-root.
gentype ceil (gentype x)	Round to integral value using the round to positive infinity rounding mode.
gentype copysign (gentype <i>x</i> , gentype <i>y</i>)	Returns <i>x</i> with its sign changed to match the sign of <i>y</i> .

Function	Description
gentype cos (gentype <i>x</i>)	Compute cosine.
gentype cosh (gentype <i>x</i>)	Compute hyperbolic cosine.
gentype cospi (gentype <i>x</i>)	Compute $\cos (\pi x)$.
gentype erfc (gentype x)	Complementary error function.
gentype erf (gentype <i>x</i>)	Error function encountered in integrating the normal distribution.
gentype exp (gentype <i>x</i>)	Compute the base- e exponential of x .
gentype exp2 (gentype <i>x</i>)	Exponential base 2 function.
gentype exp10 (gentype x)	Exponential base 10 function.
gentype expm1 (gentype x)	Compute e^x - 1.0.
gentype fabs (gentype x)	Compute absolute value of a floating-point number.
gentype fdim (gentype x, gentype y)	x - y if $x > y$, +0 if x is less than or equal to y.
gentype floor (gentype <i>x</i>)	Round to integral value using the round to negative infinity rounding mode.
gentype fma (gentype <i>a</i> , gentype <i>b</i> , gentype <i>c</i>)	Returns the correctly rounded floating-point representation of the sum of c with the infinitely precise product of a and b . Rounding of intermediate products shall not occur. Edge case behavior is per the IEEE 754-2008 standard.
gentype fmax (gentype <i>x</i> , gentype <i>y</i>) gentype fmax (gentype <i>x</i> , half <i>y</i>)	Returns y if $x < y$, otherwise it returns x . If one argument is a NaN, fmax() returns the other argument. If both arguments are NaNs, fmax() returns a NaN.
gentype fmin (gentype <i>x</i> , gentype <i>y</i>) gentype fmin (gentype <i>x</i> , half <i>y</i>)	Returns <i>y</i> if <i>y</i> < <i>x</i> , otherwise it returns <i>x</i> . If one argument is a NaN, fmin() returns the other argument. If both arguments are NaNs, fmin() returns a NaN.
gentype fmod (gentype <i>x</i> , gentype <i>y</i>)	Modulus. Returns $x - y * \mathbf{trunc} (x/y)$.

Function	Description
gentype fract (gentype x ,global gentype * $iptr$) gentype fract (gentype x ,local gentype * $iptr$) gentype fract (gentype x ,private gentype * $iptr$)	Returns fmin (x - floor (x), 0x1.ffcp-1f). floor (x) is returned in $iptr$.
For OpenCL C 2.0 or with theopencl_c_generic_address_space feature macro:	
gentype fract (gentype <i>x</i> , gentype * <i>iptr</i>)	
halfn frexp (halfn x,global intn *exp) half frexp (half x,global int *exp)	Extract mantissa and exponent from x . For each component the mantissa returned is a float with magnitude in the interval [1/2, 1) or 0. Each component of x equals mantissa returned * 2^{exp} .
half frexp (half x,local int *exp) half frexp (half x,local int *exp)	
half <i>n</i> frexp (half <i>n x</i> ,private int <i>n</i> *exp) half frexp (half <i>x</i> ,private int *exp)	
For OpenCL C 2.0 or with theopencl_c_generic_address_space feature macro:	
half n frexp (half n x , int n *exp) half frexp (half x , int *exp)	
gentype hypot (gentype <i>x</i> , gentype <i>y</i>)	Compute the value of the square root of $x^2 + y^2$ without undue overflow or underflow.
int <i>n</i> ilogb (half <i>n x</i>) int ilogb (half <i>x</i>)	Return the exponent as an integer value.
half n ldexp (half n x , int n k) half n ldexp (half n x , int n k) half ldexp (half n x , int n k)	Multiply <i>x</i> by 2 to the power <i>k</i> .

Function	Description
gentype lgamma (gentype x)	Log gamma function. Returns the natural logarithm of the absolute value of the gamma function. The sign of the gamma function is
half <i>n</i> lgamma_r (half <i>n x</i> ,global int <i>n</i> * <i>signp</i>) half lgamma_r (half <i>x</i> ,global int * <i>signp</i>)	returned in the <i>signp</i> argument of lgamma_r .
halfn lgamma_r (halfn x,local intn *signp) half lgamma_r (half x,local int *signp)	
half <i>n</i> lgamma_r (half <i>n x</i> ,private int <i>n *signp</i>) half lgamma_r (half <i>x</i> ,private int *signp)	
For OpenCL C 2.0 or with theopencl_c_generic_address_space feature macro:	
halfn lgamma_r (halfn x, intn *signp) half lgamma_r (half x, int *signp)	
gentype log (gentype x)	Compute natural logarithm.
gentype log2 (gentype <i>x</i>)	Compute a base 2 logarithm.
gentype log10 (gentype <i>x</i>)	Compute a base 10 logarithm.
gentype log1p (gentype x)	Compute $\log_{e}(1.0 + x)$.
gentype logb (gentype x)	Compute the exponent of x , which is the integral part of $\log_r x $.
gentype ${\bf mad}$ (gentype a , gentype b , gentype c)	mad computes $a*b+c$. The function may compute $a*b+c$ with reduced accuracy in the embedded profile. See the OpenCL SPIR-V Environment Specification for details. On some hardware the mad instruction may provide better performance than expanded computation of $a*b+c$.
	Note: For some usages, e.g. mad (a, b, -a*b), the half precision definition of mad () is loose enough that almost any result is allowed from mad () for some values of a and b.
gentype maxmag (gentype x, gentype y)	Returns x if $ x > y $, y if $ y > x $, otherwise fmax (x , y).

Function	Description
gentype minmag (gentype <i>x</i> , gentype <i>y</i>)	Returns x if $ x < y $, y if $ y < x $, otherwise fmin (x , y).
gentype modf (gentype x ,global gentype *iptr) gentype modf (gentype x ,local gentype *iptr) gentype modf (gentype x ,private gentype * iptr)	Decompose a floating-point number. The modf function breaks the argument <i>x</i> into integral and fractional parts, each of which has the same sign as the argument. It stores the integral part in the object pointed to by <i>iptr</i> .
For OpenCL C 2.0 or with theopencl_c_generic_address_space feature macro:	
gentype modf (gentype <i>x</i> , gentype * <i>iptr</i>)	
halfn nan (ushortn nancode) half nan (ushort nancode)	Returns a quiet NaN. The <i>nancode</i> may be placed in the significand of the resulting NaN.
gentype nextafter (gentype <i>x</i> , gentype <i>y</i>)	Computes the next representable half-precision floating-point value following x in the direction of y . Thus, if y is less than x , nextafter () returns the largest representable floating-point number less than x .
gentype pow (gentype <i>x</i> , gentype <i>y</i>)	Compute <i>x</i> to the power <i>y</i> .
half <i>n</i> pown (half <i>n x</i> , int <i>n y</i>) half pown (half <i>x</i> , int <i>y</i>)	Compute x to the power y , where y is an integer.
gentype powr (gentype <i>x</i> , gentype <i>y</i>)	Compute x to the power y , where x is ≥ 0 .
gentype remainder (gentype <i>x</i> , gentype <i>y</i>)	Compute the value r such that $r = x - n*y$, where n is the integer nearest the exact value of x/y . If there are two integers closest to x/y , n shall be the even one. If r is zero, it is given the same sign as x .

Function	Description
halfn remquo (halfn x, halfn y,global intn *quo) half remquo (half x, half y,global int *quo) halfn remquo (halfn x, halfn y,local intn *quo) halfn remquo (halfn x, halfn y,local int *quo) halfn remquo (halfn x, halfn y,private intn *quo) half remquo (half x, half y,private int *quo) For OpenCL C 2.0 or with theopencl_c_generic_address_space feature macro: halfn remquo (halfn x, halfn y, intn *quo) half remquo (half x, half y, int *quo)	The remquo function computes the value r such that $r = x - k^*y$, where k is the integer nearest the exact value of x/y . If there are two integers closest to x/y , k shall be the even one. If r is zero, it is given the same sign as x . This is the same value that is returned by the remainder function. remquo also calculates the lower seven bits of the integral quotient x/y , and gives that value the same sign as x/y . It stores this signed value in the object pointed to by quo .
gentype rint (gentype <i>x</i>)	Round to integral value (using round to nearest even rounding mode) in floating-point format. Refer to section 7.1 for description of rounding modes.
half <i>n</i> rootn (half <i>n x</i> , int <i>n y</i>) half rootn (half <i>x</i> , int <i>y</i>)	Compute <i>x</i> to the power 1/ <i>y</i> .
gentype round (gentype x)	Return the integral value nearest to <i>x</i> rounding halfway cases away from zero, regardless of the current rounding direction.
gentype rsqrt (gentype <i>x</i>)	Compute inverse square root.
gentype sin (gentype <i>x</i>)	Compute sine.

Function	Description
gentype sincos (gentype x,global gentype *cosval) gentype sincos (gentype x,local gentype *cosval) gentype sincos (gentype x,private gentype *cosval) For OpenCL C 2.0 or with theopencl_c_generic_address_space feature macro:	Compute sine and cosine of x. The computed sine is the return value and computed cosine is returned in <i>cosval</i> .
gentype sincos (gentype <i>x</i> , gentype * <i>cosval</i>)	
gentype sinh (gentype x)	Compute hyperbolic sine.
gentype sinpi (gentype <i>x</i>)	Compute $\sin (\pi x)$.
gentype sqrt (gentype <i>x</i>)	Compute square root.
gentype tan (gentype <i>x</i>)	Compute tangent.
gentype tanh (gentype <i>x</i>)	Compute hyperbolic tangent.
gentype tanpi (gentype <i>x</i>)	Compute $tan (\pi x)$.
gentype tgamma (gentype <i>x</i>)	Compute the gamma function.
gentype trunc (gentype x)	Round to integral value using the round to zero rounding mode.

The **FP_FAST_FMA_HALF** macro indicates whether the **fma()** family of functions are fast compared with direct code for half precision floating-point. If defined, the **FP_FAST_FMA_HALF** macro shall indicate that the **fma()** function generally executes about as fast as, or faster than, a multiply and an add of **half** operands.

The macro names given in the following list must use the values specified. These constant expressions are suitable for use in #if preprocessing directives.

```
#define HALF_DIG
                             3
#define HALF_MANT_DIG
                            11
#define HALF_MAX_10_EXP
                            +4
#define HALF_MAX_EXP
                            +16
#define HALF_MIN_10_EXP
                             -4
#define HALF_MIN_EXP
                             -13
#define HALF_RADIX
#define HALF_MAX
                            0x1.ffcp15h
#define HALF_MIN
                            0x1.0p-14h
#define HALF_EPSILON
                            0x1.0p-10h
```

The following table describes the built-in macro names given above in the OpenCL C programming language and the corresponding macro names available to the application.

Macro in OpenCL Language	Macro for application
HALF_DIG	CL_HALF_DIG
HALF_MANT_DIG	CL_HALF_MANT_DIG
HALF_MAX_10_EXP	CL_HALF_MAX_10_EXP
HALF_MAX_EXP	CL_HALF_MAX_EXP
HALF_MIN_10_EXP	CL_HALF_MIN_10_EXP
HALF_MIN_EXP	CL_HALF_MIN_EXP
HALF_RADIX	CL_HALF_RADIX
HALF_MAX	CL_HALF_MAX
HALF_MIN	CL_HALF_MIN
HALF_EPSILSON	CL_HALF_EPSILON

The following constants are also available. They are of type half and are accurate within the precision of the half type.

Constant	Description
M_E_H	Value of e
M_LOG2E_H	Value of log ₂ e
M_LOG10E_H	Value of log ₁₀ e
M_LN2_H	Value of log _e 2
M_LN10_H	Value of log _e 10
M_PI_H	Value of π
M_PI_2_H	Value of π / 2
M_PI_4_H	Value of π / 4
M_1_PI_H	Value of 1 / π
M_2_PI_H	Value of 2 / π
M_2_SQRTPI_H	Value of 2 / √π
M_SQRT2_H	Value of √2
M_SQRT1_2_H	Value of 1 / √2

5.2.3. Common Functions

The built-in common functions defined in *table 6.12* (also listed below) are extended to include appropriate versions of functions that take half and half{2|3|4|8|16} as arguments and return values. gentype now also includes half, half2, half3, half4, half8 and half16. These are described below.

Table 3. Half Precision Built-in Common Functions

Function	Description
gentype clamp (gentype x, gentype minval, gentype maxval) gentype clamp (gentype x, half minval, half maxval)	Returns $fmin(fmax(x, minval), maxval)$. Results are undefined if $minval > maxval$.
gentype degrees (gentype <i>radians</i>)	Converts $radians$ to degrees, i.e. $(180 / \pi) * radians$.
gentype max (gentype <i>x</i> , gentype <i>y</i>) gentype max (gentype <i>x</i> , half <i>y</i>)	Returns y if $x < y$, otherwise it returns x . If x and y are infinite or NaN, the return values are undefined.
gentype min (gentype <i>x</i> , gentype <i>y</i>) gentype min (gentype <i>x</i> , half <i>y</i>)	Returns y if $y < x$, otherwise it returns x . If x and y are infinite or NaN, the return values are undefined.
gentype mix (gentype <i>x</i> , gentype <i>y</i> , gentype <i>a</i>) gentype mix (gentype <i>x</i> , gentype <i>y</i> , half <i>a</i>)	Returns the linear blend of x and y implemented as: $x + (y - x) * a$ a must be a value in the range 0.0 1.0. If a is not in the range 0.0 1.0, the return values are undefined. Note: The half precision mix function can be implemented using contractions such as mad or fma .
gentype radians (gentype <i>degrees</i>)	Converts degrees to radians, i.e. $(\pi / 180)$ * degrees.
gentype step (gentype <i>edge</i> , gentype <i>x</i>) gentype step (half <i>edge</i> , gentype <i>x</i>)	Returns 0.0 if $x < edge$, otherwise it returns 1.0.

Function	Description
gentype smoothstep (gentype <i>edge0</i> , gentype <i>edge1</i> , gentype <i>x</i>) gentype smoothstep (half <i>edge0</i> , half <i>edge1</i> , gentype <i>x</i>)	Returns 0.0 if $x \le edge0$ and 1.0 if $x \ge edge1$ and performs smooth Hermite interpolation between 0 and 1 when $edge0 < x < edge1$. This is useful in cases where you would want a threshold function with a smooth transition.
	This is equivalent to:
	gentype t; t = clamp ((x - edge0) / (edge1 - edge0), 0, 1); return t * t * (3 - 2 * t);
	Results are undefined if <i>edge0</i> >= <i>edge1</i> . Note: The half precision smoothstep function can be implemented using contractions such as mad or fma .
gentype sign (gentype <i>x</i>)	Returns 1.0 if $x > 0$, -0.0 if $x = -0.0$, +0.0 if $x = +0.0$, or -1.0 if $x < 0$. Returns 0.0 if x is a NaN.

5.2.4. Geometric Functions

The built-in geometric functions defined in *table 6.13* (also listed below) are extended to include appropriate versions of functions that take half and half{2|3|4} as arguments and return values. gentype now also includes half, half2, half3 and half4. These are described below.

Note: The half precision geometric functions can be implemented using contractions such as **mad** or **fma**.

Table 4. Half Precision Built-in Geometric Functions

Function	Description
half4 cross (half4 <i>p0</i> , half4 <i>p1</i>) half3 cross (half3 <i>p0</i> , half3 <i>p1</i>)	Returns the cross product of $p0.xyz$ and $p1.xyz$. The w component of the result will be 0.0.
half dot (gentype <i>p0</i> , gentype <i>p1</i>)	Compute the dot product of $p0$ and $p1$.
half distance (gentype $p0$, gentype $p1$)	Returns the distance between $p0$ and $p1$. This is calculated as length ($p0 - p1$).
half length (gentype <i>p</i>)	Return the length of vector x, i.e., sqrt($p.x^2 + p.y^2 +$)
gentype normalize (gentype <i>p</i>)	Returns a vector in the same direction as p but with a length of 1.

5.2.5. Relational Functions

The scalar and vector relational functions described in *table 6.14* are extended to include versions that take half, half2, half3, half4, half8 and half16 as arguments.

The relational and equality operators (<, <=, >, >=, !=, ==) can be used with halfn vector types and shall produce a vector shortn result as described in *section 6.3*.

The functions **isequal**, **isnotequal**, **isgreater**, **isgreaterequal**, **isless**, **islessequal**, **islessgreater**, **isfinite**, **isinf**, **isnan**, **isnormal**, **isordered**, **isunordered** and **signbit** shall return a 0 if the specified relation is *false* and a 1 if the specified relation is true for scalar argument types. These functions shall return a 0 if the specified relation is *false* and a -1 (i.e. all bits set) if the specified relation is *true* for vector argument types.

The relational functions **isequal**, **isgreater**, **isgreaterequal**, **isless**, **islessequal**, and **islessgreater** always return 0 if either argument is not a number (NaN). **isnotequal** returns 1 if one or both arguments are not a number (NaN) and the argument type is a scalar and returns -1 if one or both arguments are not a number (NaN) and the argument type is a vector.

The functions described in *table 6.14* are extended to include the halfn vector types.

Table 5. Half Precision Relational Functions

Function	Description
int isequal (half x , half y) short n isequal (half n x , half n y)	Returns the component-wise compare of $x == y$.
int isnotequal (half <i>x</i> , half <i>y</i>) short <i>n</i> isnotequal (half <i>n x</i> , half <i>n y</i>)	Returns the component-wise compare of $x = y$.
int isgreater (half <i>x</i> , half <i>y</i>) short <i>n</i> isgreater (half <i>n x</i> , half <i>n y</i>)	Returns the component-wise compare of $x > y$.
int isgreaterequal (half x, half y) shortn isgreaterequal (halfn x, halfn y)	Returns the component-wise compare of $x \ge y$.
int isless (half <i>x</i> , half <i>y</i>) short <i>n</i> isless (half <i>n x</i> , half <i>n y</i>)	Returns the component-wise compare of $x < y$.
int islessequal (half x, half y) shortn islessequal (halfn x, halfn y)	Returns the component-wise compare of $x \le y$.
int islessgreater (half <i>x</i> , half <i>y</i>) short <i>n</i> islessgreater (half <i>n x</i> , half <i>n y</i>)	Returns the component-wise compare of $(x < y)$ $(x > y)$.
int isfinite (half) short <i>n</i> isfinite (half <i>n</i>)	Test for finite value.
int isinf (half) short <i>n</i> isinf (half <i>n</i>)	Test for infinity value (positive or negative) .
int isnan (half) short <i>n</i> isnan (half <i>n</i>)	Test for a NaN.

Function	Description
int isnormal (half) short <i>n</i> isnormal (half <i>n</i>)	Test for a normal value.
int isordered (half <i>x</i> , half <i>y</i>) short <i>n</i> isordered (half <i>n x</i> , half <i>n y</i>)	Test if arguments are ordered. isordered () takes arguments <i>x</i> and <i>y</i> , and returns the result isequal (<i>x</i> , <i>x</i>) && isequal (<i>y</i> , <i>y</i>).
int isunordered (half x, half y) shortn isunordered (halfn x, halfn y)	Test if arguments are unordered. isunordered () takes arguments <i>x</i> and <i>y</i> , returning non-zero if <i>x</i> or <i>y</i> is a NaN, and zero otherwise.
int signbit (half) short <i>n</i> signbit (half <i>n</i>)	Test for sign bit. The scalar version of the function returns a 1 if the sign bit in the half is set else returns 0. The vector version of the function returns the following for each component in halfn: -1 (i.e all bits set) if the sign bit in the half is set else returns 0.
half <i>n</i> bitselect (half <i>n a</i> , half <i>n b</i> , half <i>n c</i>)	Each bit of the result is the corresponding bit of a if the corresponding bit of c is 0. Otherwise it is the corresponding bit of b .
half n select (half n a , half n b , short n c) half n select (half n a , half n b , ushort n c)	For each component, result[i] = if MSB of c[i] is set ? b[i] : a[i].

5.2.6. Vector Data Load and Store Functions

The vector data load (**vloadn**) and store (**vstoren**) functions described in *table 6.13* (also listed below) are extended to include versions that read or write half vector values. The generic type gentype is extended to include half. The generic type gentypen is extended to include half2, half3, half4, half8, and half16.

Note: **vload3** reads x, y, z components from address (p + (*offset* * 3)) into a 3-component vector and **vstore3** writes x, y, z components from a 3-component vector to address (p + (*offset* * 3)).

Table 6. Half Precision Vector Data Load and Store Functions

Function	Description
gentypen vloadn (size_t <i>offset</i> , constglobal gentype *p) gentypen vloadn (size_t <i>offset</i> , constlocal gentype *p) gentypen vloadn (size_t <i>offset</i> , constconstant gentype *p) gentypen vloadn (size_t <i>offset</i> , constprivate gentype *p)	Return sizeof (gentype n) bytes of data read from address (p + (offset * n)). If gentype is half, the read address computed as (p + (offset * n)) must be 16-bit aligned.
For OpenCL C 2.0 or with theopencl_c_generic_address_space feature macro: gentypen vloadn(size_t offset, const gentype *p)	
void vstoren (gentypen data, size_t offset,global gentype *p) void vstoren (gentypen data, size_t offset,local gentype *p) void vstoren (gentypen data, size_t offset,private gentype *p)	Write sizeof (gentypen) bytes given by $data$ to address (p + ($offset * n$)). If gentype is half, the write address computed as (p + ($offset * n$)) must be 16-bit aligned.
For OpenCL C 2.0 or with theopencl_c_generic_address_space feature macro:	
void vstore <i>n</i> (gentype <i>n data</i> , size_t <i>offset</i> , gentype * <i>p</i>)	

5.2.7. Async Copies from Global to Local Memory, Local to Global Memory, and Prefetch

The OpenCL C programming language implements the following functions that provide asynchronous copies between global and local memory and a prefetch from global memory.

The generic type gentype is extended to include half, half2, half3, half4, half8, and half16.

Table 7. Half Precision Built-in Async Copy and Prefetch Functions

Function	Description
event_t async_work_group_copy (Perform an async copy of <i>num_gentypes</i> gentype
_local gentype *dst,	elements from <i>src</i> to <i>dst</i> . The async copy is
constglobal gentype *src,	performed by all work-items in a work-group
size_t num_gentypes, event_t event)	and this built-in function must therefore be
	encountered by all work-items in a work-group
event_t async_work_group_copy (executing the kernel with the same argument
global gentype *dst,	values; otherwise the results are undefined.
const _local gentype *src,	
size_t num_gentypes, event_t event)	Returns an event object that can be used by
	wait_group_events to wait for the async copy to
	finish. The <i>event</i> argument can also be used to
	associate the async_work_group_copy with a
	previous async copy allowing an event to be
	shared by multiple async copies; otherwise <i>event</i>
	should be zero.
	Te abia at
	If event argument is not zero, the event object
	supplied in <i>event</i> argument will be returned.
	This function does not perform any implicit
	synchronization of source data such as using a
	barrier before performing the copy.
	1 0 17

Function	Description
event_t async_work_group_strided_copy (local gentype *dst, constglobal gentype *src, size_t num_gentypes, size_t src_stride, event_t event) event_t async_work_group_strided_copy (global gentype *dst, constlocal gentype *src, size_t num_gentypes, size_t dst_stride, event_t event)	Perform an async gather of <i>num_gentypes</i> gentype elements from <i>src</i> to <i>dst</i> . The <i>src_stride</i> is the stride in elements for each gentype element read from <i>src</i> . The async gather is performed by all work-items in a work-group and this built-in function must therefore be encountered by all work-items in a work-group executing the kernel with the same argument values; otherwise the results are undefined. Returns an event object that can be used by wait_group_events to wait for the async copy to finish. The <i>event</i> argument can also be used to associate the async_work_group_strided_copy with a previous async copy allowing an event to be shared by multiple async copies; otherwise <i>event</i> should be zero. If <i>event</i> argument is not zero, the event object supplied in <i>event</i> argument will be returned. This function does not perform any implicit synchronization of source data such as using a barrier before performing the copy. The behavior of async_work_group_strided_copy is undefined if <i>src_stride</i> or <i>dst_stride</i> is 0, or if the <i>src_stride</i> or <i>dst_stride</i> values cause the <i>src</i> or <i>dst</i> pointers to exceed the upper bounds of the address space during the copy.
void wait_group_events (int num_events, event_t *event_list)	Wait for events that identify the async_work_group_copy operations to complete. The event objects specified in event_list will be released after the wait is performed. This function must be encountered by all work-items in a work-group executing the kernel with the same num_events and event objects specified in event_list; otherwise the results are undefined.

Function	Description
void prefetch (Prefetch num_gentypes * sizeof(gentype) bytes
constglobal gentype *p, size_t num_gentypes)	into the global cache. The prefetch instruction is
	applied to a work-item in a work-group and does
	not affect the functional behavior of the kernel.

5.2.8. Image Read and Write Functions

The image read and write functions defined in *tables 6.23*, *6.24* and *6.25* are extended to support image color values that are a half type.

5.2.9. Built-in Image Read Functions

Table 8. Half Precision Built-in Image Read Functions

Function	Description
half4 read_imageh (Use the coordinate (coord.x, coord.y) to do an
read_only image2d_t image,	element lookup in the 2D image object specified
sampler_t sampler,	by image.
int2 coord)	
	read_imageh returns half precision floating-
half4 read_imageh (point values in the range [0.0 1.0] for image
read_only image2d_t image,	objects created with image_channel_data_type
sampler_t sampler,	set to one of the pre-defined packed formats,
float2 coord)	CL_UNORM_INT8, or CL_UNORM_INT16.
	read_imageh returns half precision floating-
	point values in the range [-1.0 1.0] for image
	objects created with image_channel_data_type
	set to CL_SNORM_INT8, or CL_SNORM_INT16.
	read_imageh returns half precision floating-
	point values for image objects created with
	image_channel_data_type set to CL_HALF_FLOAT.
	The read_imageh calls that take integer
	coordinates must use a sampler with filter mode
	set to CLK_FILTER_NEAREST, normalized
	coordinates set to
	CLK_NORMALIZED_COORDS_FALSE and
	addressing mode set to
	CLK_ADDRESS_CLAMP_TO_EDGE,
	CLK_ADDRESS_CLAMP or CLK_ADDRESS_NONE;
	otherwise the values returned are undefined.
	Values returned by read_imageh for image
	objects with image_channel_data_type values not
	specified in the description above are undefined.

Function	Description
half4 read_imageh (Use the coordinate (coord.x, coord.y, coord.z) to
read_only image3d_t image,	do an elementlookup in the 3D image object
sampler_t sampler,	specified by <i>image</i> . <i>coord.w</i> is ignored.
int4 coord)	
	read_imageh returns half precision floating-
half4 read_imageh (point values in the range [0.0 1.0] for image
read_only image3d_t image,	objects created with image_channel_data_type
sampler_t sampler,	set to one of the pre-defined packed formats or
float4 coord)	CL_UNORM_INT8, or CL_UNORM_INT16.
	read_imageh returns half precision floating-
	point values in the range [-1.0 1.0] for image
	objects created with image_channel_data_type
	set to CL_SNORM_INT8, or CL_SNORM_INT16.
	read_imagehreturns half precision floating-
	point values for image objects created with
	image_channel_data_type set to CL_HALF_FLOAT.
	The read_imageh calls that take integer
	coordinates must use a sampler with filter mode
	set to CLK_FILTER_NEAREST, normalized
	coordinates set to
	CLK_NORMALIZED_COORDS_FALSE and
	addressing mode set to
	CLK_ADDRESS_CLAMP_TO_EDGE,
	CLK_ADDRESS_CLAMP or CLK_ADDRESS_NONE;
	otherwise the values returned are undefined.
	Values returned by read_imageh for image
	objects with image_channel_data_type values not
	specified in the description are undefined.

Description
Use <i>coord.xy</i> to do an element lookup in the 2D
image identified by coord.z in the 2D image
array specified by <i>image</i> .
read_imageh returns half precision floating-
point values in the range [0.0 1.0] for image
objects created with image_channel_data_type
set to one of the pre-defined packed formats or
CL_UNORM_INT8, or CL_UNORM_INT16.
read_imageh returns half precision floating-
point values in the range [-1.0 1.0] for image
objects created with image_channel_data_type
set to CL_SNORM_INT8, or CL_SNORM_INT16.
read_imageh returns half precision floating-
point values for image objects created with
image_channel_data_type set to
CL_HALF_FLOAT.
The read_imageh calls that take integer
coordinates must use a sampler with filter mode
set to CLK_FILTER_NEAREST, normalized
coordinates set to
CLK_NORMALIZED_COORDS_FALSE and
addressing mode set to
CLK_ADDRESS_CLAMP_TO_EDGE,
CLK_ADDRESS_CLAMP or CLK_ADDRESS_NONE;
otherwise the values returned are undefined.
Values returned by read_imageh for image
objects with image_channel_data_type values
not specified in the description above are
undefined.

Function	Description
half4 read_imageh (Use <i>coord</i> to do an element lookup in the 1D
read_only image1d_t <i>image</i> ,	image object specified by image.
sampler_t <i>sampler</i> ,	
int <i>coord</i>)	read_imageh returns half precision floating-
	point values in the range [0.0 1.0] for image
half4 read_imageh (objects created with image_channel_data_type
read_only image1d_t <i>image</i> ,	set to one of the pre-defined packed formats or
sampler_t <i>sampler</i> ,	CL_UNORM_INT8, or CL_UNORM_INT16.
float <i>coord</i>)	
	read_imageh returns half precision floating-
	point values in the range [-1.0 1.0] for image
	objects created with image_channel_data_type
	set to CL_SNORM_INT8, or CL_SNORM_INT16.
	read_imageh returns half precision floating-
	point values for image objects created with
	image_channel_data_type set to CL_HALF_FLOAT
	The read_imageh calls that take integer
	coordinates must use a sampler with filter mode
	set to CLK_FILTER_NEAREST, normalized
	coordinates set to
	CLK_NORMALIZED_COORDS_FALSE and
	addressing mode set to
	CLK_ADDRESS_CLAMP_TO_EDGE,
	CLK_ADDRESS_CLAMP or CLK_ADDRESS_NONE
	otherwise the values returned are undefined.
	Values returned by read_imageh for image
	objects with image_channel_data_type values no
	specified in the description above are undefined

Function	Description
half4 read_imageh (Use <i>coord.x</i> to do an element lookup in the 1D
read_only image1d_array_t image,	image identified by <i>coord.y</i> in the 1D image
sampler_t sampler,	array specified by <i>image</i> .
int2 coord)	
	read_imageh returns half precision floating-
half4 read_imageh (point values in the range [0.0 1.0] for image
read_only image1d_array_t image,	objects created with image_channel_data_type
sampler_t sampler,	set to one of the pre-defined packed formats or
float2 coord)	CL_UNORM_INT8, or CL_UNORM_INT16.
	read_imageh returns half precision floating-
	point values in the range [-1.0 1.0] for image
	objects created with image_channel_data_type
	set to CL_SNORM_INT8, or CL_SNORM_INT16.
	read_imageh returns half precision floating-
	point values for image objects created with
	image_channel_data_type set to
	CL_HALF_FLOAT.
	The read_imageh calls that take integer
	coordinates must use a sampler with filter mode
	set to CLK_FILTER_NEAREST, normalized
	coordinates set to
	CLK_NORMALIZED_COORDS_FALSE and
	addressing mode set to
	CLK_ADDRESS_CLAMP_TO_EDGE,
	CLK_ADDRESS_CLAMP or CLK_ADDRESS_NONE;
	otherwise the values returned are undefined.
	Values returned by read_imageh for image
	objects with image_channel_data_type values
	not specified in the description above are
	undefined.

5.2.10. Built-in Image Sampler-less Read Functions

aQual in Table 6.24 refers to one of the access qualifiers. For sampler-less read functions this may be $read_only$ or $read_write$.

Table 9. Half Precision Built-in Image Sampler-less Read Functions

Function	Description
half4 read_imageh (aQual image2d_t image, int2 coord)	Use the coordinate (coord.x, coord.y) to do an element lookup in the 2D image object specified by image.
	read_imageh returns half precision floating-point values in the range [0.0 1.0] for image objects created with <i>image_channel_data_type</i> set to one of the pre-defined packed formats or CL_UNORM_INT8, or CL_UNORM_INT16.
	read_imageh returns half precision floating-point values in the range [-1.0 1.0] for image objects created with <i>image_channel_data_type</i> set to CL_SNORM_INT8, or CL_SNORM_INT16.
	read_imageh returns half precision floating- point values for image objects created with image_channel_data_type set to CL_HALF_FLOAT.
	Values returned by read_imageh for image objects with <i>image_channel_data_type</i> values not specified in the description above are undefined.
half4 read_imageh (aQual image3d_t image, int4 coord)	Use the coordinate (coord.x, coord.y, coord.z) to do an element lookup in the 3D image object specified by image. coord.w is ignored.
	read_imageh returns half precision floating-point values in the range [0.0 1.0] for image objects created with <i>image_channel_data_type</i> set to one of the pre-defined packed formats or CL_UNORM_INT8, or CL_UNORM_INT16.
	read_imageh returns half precision floating-point values in the range [-1.0 1.0] for image objects created with <i>image_channel_data_type</i> set to CL_SNORM_INT8, or CL_SNORM_INT16.
	read_imageh returns half precision floating- point values for image objects created with image_channel_data_type set to CL_HALF_FLOAT.
	Values returned by read_imageh for image objects with <i>image_channel_data_type</i> values not specified in the description are undefined.

Function	Description
half4 read_imageh (aQual image2d_array_t image, int4 coord)	Use coord.xy to do an element lookup in the 2D image identified by coord.z in the 2D image array specified by image. read_imageh returns half precision floating-point values in the range [0.0 1.0] for image objects created with image_channel_data_type set to one of the pre-defined packed formats or CL_UNORM_INT8, or CL_UNORM_INT16. read_imageh returns half precision floating-point values in the range [-1.0 1.0] for image objects created with image_channel_data_type set to CL_SNORM_INT8, or CL_SNORM_INT16. read_imageh returns half precision floating-point values for image objects created with image_channel_data_type set to CL_HALF_FLOAT. Values returned by read_imageh for image objects with image_channel_data_type values not specified in the description above are undefined.
half4 read_imageh (aQual image1d_t image, int coord) half4 read_imageh (aQual image1d_buffer_t image, int coord)	Use coord to do an element lookup in the 1D image or 1D image buffer object specified by image. read_imageh returns half precision floating-point values in the range [0.0 1.0] for image objects created with image_channel_data_type set to one of the pre-defined packed formats or CL_UNORM_INT8, or CL_UNORM_INT16. read_imageh returns half precision floating-point values in the range [-1.0 1.0] for image objects created with image_channel_data_type set to CL_SNORM_INT8, or CL_SNORM_INT16. read_imageh returns half precision floating-point values for image objects created with image_channel_data_type set to CL_HALF_FLOAT. Values returned by read_imageh for image objects with image_channel_data_type values not specified in the description above are undefined.

Function	Description
half4 read_imageh (aQual image1d_array_t image, int2 coord)	Use <i>coord.x</i> to do an element lookup in the 2D image identified by <i>coord.y</i> in the 2D image array specified by <i>image</i> .
	read_imageh returns half precision floating-point values in the range [0.0 1.0] for image objects created with <i>image_channel_data_type</i> set to one of the pre-defined packed formats or CL_UNORM_INT8, or CL_UNORM_INT16.
	read_imageh returns half precision floating-point values in the range [-1.0 1.0] for image objects created with <i>image_channel_data_type</i> set to CL_SNORM_INT8, or CL_SNORM_INT16.
	read_imageh returns half precision floating- point values for image objects created with image_channel_data_type set to CL_HALF_FLOAT.
	Values returned by read_imageh for image objects with <i>image_channel_data_type</i> values not specified in the description above are undefined.

5.2.11. Built-in Image Write Functions

aQual in Table 6.25 refers to one of the access qualifiers. For write functions this may be $write_only$ or $read_write$.

Table 10. Half Precision Built-in Image Write Functions

Function	Description
void write_imageh (Write <i>color</i> value to location specified by
aQual image2d_t image,	coord.xy in the 2D image specified by image.
int2 coord,	
half4 color)	Appropriate data format conversion to the
	specified image format is done before writing
	the color value. <i>x</i> & <i>y</i> are considered to be
	unnormalized coordinates and must be in the
	range 0 width - 1, and 0 height - 1.
	write_imageh can only be used with image
	objects created with image_channel_data_type
	set to one of the pre-defined packed formats or
	set to CL_SNORM_INT8, CL_UNORM_INT8,
	CL_SNORM_INT16, CL_UNORM_INT16 or
	CL_HALF_FLOAT.
	The behavior of write_imageh for image objects
	created with image_channel_data_type values
	not specified in the description above or with (<i>x</i> ,
	y) coordinate values that are not in the range (0
	width - 1, 0 height - 1) respectively, is
	undefined.

Function	Description
void write_imageh (aQual image2d_array_t image, int4 coord, half4 color)	Write <i>color</i> value to location specified by <i>coord.xy</i> in the 2D image identified by <i>coord.z</i> in the 2D image array specified by <i>image</i> . Appropriate data format conversion to the specified image format is done before writing the color value. <i>coord.x</i> , <i>coord.y</i> and <i>coord.z</i> are considered to be unnormalized coordinates and must be in the range 0 image width - 1, 0 image height - 1 and 0 image number of layers - 1.
	write_imageh can only be used with image objects created with image_channel_data_type set to one of the pre-defined packed formats or set to CL_SNORM_INT8, CL_UNORM_INT8, CL_SNORM_INT16, CL_UNORM_INT16 or CL_HALF_FLOAT. The behavior of write_imageh for image objects created with image_channel_data_type values not specified in the description above or with (x,
	y, z) coordinate values that are not in the range (0 image width - 1, 0 image height - 1, 0 image number of layers - 1), respectively, is undefined.

Function	Description
void write_imageh (Write <i>color</i> value to location specified by <i>coord</i>
aQual image1d_t image,	in the 1D image or 1D image buffer object
int coord,	specified by image. Appropriate data format
half4 color)	conversion to the specified image format is done
	before writing the color value. coord is
void write_imageh (considered to be unnormalized coordinates and
aQual image1d_buffer_t image,	must be in the range 0 image width - 1.
int coord,	
half4 color)	write_imageh can only be used with image
	objects created with image_channel_data_type
	set to one of the pre-defined packed formats or
	set to CL_SNORM_INT8, CL_UNORM_INT8,
	CL_SNORM_INT16, CL_UNORM_INT16 or
	CL_HALF_FLOAT. Appropriate data format
	conversion will be done to convert channel data
	from a floating-point value to actual data format
	in which the channels are stored.
	The behavior of write_imageh for image objects created with <i>image_channel_data_type</i> values not specified in the description above or with coordinate values that is not in the range (0 image width - 1), is undefined.

Function	Description
void write_imageh (aQual image1d_array_t image, int2 coord, half4 color)	Write <i>color</i> value to location specified by <i>coord.x</i> in the 1D image identified by <i>coord.y</i> in the 1D image array specified by <i>image</i> . Appropriate data format conversion to the specified image format is done before writing the color value. <i>coord.x</i> and <i>coord.y</i> are considered to be unnormalized coordinates and must be in the range 0 image width - 1 and 0 image number of layers - 1.
	write_imageh can only be used with image objects created with image_channel_data_type set to one of the pre-defined packed formats or set to CL_SNORM_INT8, CL_UNORM_INT8, CL_SNORM_INT16, CL_UNORM_INT16 or CL_HALF_FLOAT. Appropriate data format conversion will be done to convert channel data from a floating-point value to actual data format in which the channels are stored.
	created with <i>image_channel_data_type</i> values not specified in the description above or with (<i>x</i> , <i>y</i>) coordinate values that are not in the range (0 image width - 1, 0 image number of layers - 1), respectively, is undefined.

Function	Description
void write_imageh (aQual image3d_t image,	Write color value to location specified by coord.xyz in the 3D image object specified by
int4 coord,	image.
half4 color)	
	Appropriate data format conversion to the specified image format is done before writing the color value. coord.x, coord.y and coord.z are considered to be unnormalized coordinates and must be in the range 0 image width - 1, 0 image height - 1 and 0 image depth - 1.
	mage height - I and 0 mage deput - 1.
	write_imageh can only be used with image objects created with image_channel_data_type set to one of the pre-defined packed formats or set to CL_SNORM_INT8, CL_UNORM_INT8, CL_SNORM_INT16, CL_UNORM_INT16 or CL_HALF_FLOAT.
	The behavior of write_imageh for image objects created with image_channel_data_type values not specified in the description above or with (x, y, z) coordinate values that are not in the range (0 image width - 1, 0 image height - 1, 0 image depth - 1), respectively, is undefined.
	Note: This built-in function is only available if the cl_khr_3d_image_writes extension is also supported by the device.

5.2.12. IEEE754 Compliance

The following table entry describes the additions to *table 4.3*, which allows applications to query the configuration information using **clGetDeviceInfo** for an OpenCL device that supports half precision floating-point.

Op-code	Return Type	Description
CL_DEVICE_HALF_FP_ CONFIG	cl_device_fp_config	Describes half precision floating-point capability of the OpenCL device. This is a bit-field that describes one or more of the following values:
		CL_FP_DENORM — denorms are supported
		CL_FP_INF_NAN — INF and NaNs are supported
		CL_FP_ROUND_TO_NEAREST — round to nearest even rounding mode supported
		CL_FP_ROUND_TO_ZERO — round to zero rounding mode supported
		CL_FP_ROUND_TO_INF — round to positive and negative infinity rounding modes supported
		CL_FP_FMA — IEEE754-2008 fused multiply-add is supported
		CL_FP_SOFT_FLOAT — Basic floating-point operations (such as addition, subtraction, multiplication) are implemented in software.
		The required minimum half precision floating- point capability as implemented by this extension is:
		CL_FP_ROUND_TO_ZERO, or CL_FP_ROUND_TO_NEAREST CL_FP_INF_NAN.

5.2.13. Rounding Modes

If CL_FP_ROUND_TO_NEAREST is supported, the default rounding mode for half-precision floating-point operations will be round to nearest even; otherwise the default rounding mode will be round to zero.

Conversions to half floating point format must be correctly rounded using the indicated convert operator rounding mode or the default rounding mode for half-precision floating-point operations if no rounding mode is specified by the operator, or a C-style cast is used.

Conversions from half to integer format shall correctly round using the indicated convert operator rounding mode, or towards zero if no rounding mode is specified by the operator or a C-style cast is used. All conversions from half to floating point formats are exact.

5.2.14. Relative Error as ULPs

In this section we discuss the maximum relative error defined as *ulp* (units in the last place).

Addition, subtraction, multiplication, fused multiply-add operations on half types are required to be correctly rounded using the default rounding mode for half-precision floating-point operations.

The following table describes the minimum accuracy of half precision floating-point arithmetic operations given as ULP values. 0 ULP is used for math functions that do not require rounding. The reference value used to compute the ULP value of an arithmetic operation is the infinitely precise result.

Table 11. ULP Values for Half Precision Floating-Point Arithmetic Operations

Function	Min Accuracy - Full Profile	Min Accuracy - Embedded Profile
x + y	Correctly rounded	Correctly rounded
x - y	Correctly rounded	Correctly rounded
<i>x</i> * <i>y</i>	Correctly rounded	Correctly rounded
1.0 / x	Correctly rounded	<= 1 ulp
<i>x</i> / <i>y</i>	Correctly rounded	<= 1 ulp
acos	<= 2 ulp	<= 3 ulp
acosh	<= 2 ulp	<= 3 ulp
acospi	<= 2 ulp	<= 3 ulp
asin	<= 2 ulp	<= 3 ulp
asinh	<= 2 ulp	<= 3 ulp
asinpi	<= 2 ulp	<= 3 ulp
atan	<= 2 ulp	<= 3 ulp
atanh	<= 2 ulp	<= 3 ulp
atanpi	<= 2 ulp	<= 3 ulp
atan2	<= 2 ulp	<= 3 ulp
atan2pi	<= 2 ulp	<= 3 ulp
cbrt	<= 2 ulp	<= 2 ulp
ceil	Correctly rounded	Correctly rounded
clamp	0 ulp	0 ulp
copysign	0 ulp	0 ulp
cos	<= 2 ulp	<= 2 ulp
cosh	<= 2 ulp	<= 3 ulp
cospi	<= 2 ulp	<= 2 ulp

Function	Min Accuracy - Full Profile	Min Accuracy - Embedded Profile
cross	absolute error tolerance of 'max * max * (3 * HLF_EPSILON)' per vector component, where max is the maximum input operand magnitude	
degrees	<= 2 ulp	<= 2 ulp
distance	<= 2n ulp, for gentype with vector width <i>n</i>	Implementation-defined
dot	absolute error tolerance of 'max * max * (2n - 1) * HLF_EPSILON', for vector width n and maximum input operand magnitude max across all vector components	
erfc	<= 4 ulp	<= 4 ulp
erf	<= 4 ulp	<= 4 ulp
exp	<= 2 ulp	<= 3 ulp
exp2	<= 2 ulp	<= 3 ulp
exp10	<= 2 ulp	<= 3 ulp
expm1	<= 2 ulp	<= 3 ulp
fabs	0 ulp	0 ulp
fdim	Correctly rounded	Correctly rounded
floor	Correctly rounded	Correctly rounded
fma	Correctly rounded	Correctly rounded
fmax	0 ulp	0 ulp
fmin	0 ulp	0 ulp
fmod	0 ulp	0 ulp
fract	Correctly rounded	Correctly rounded
frexp	0 ulp	0 ulp
hypot	<= 2 ulp	<= 3 ulp
ilogb	0 ulp	0 ulp
ldexp	Correctly rounded	Correctly rounded
length	<= 0.25 + 0.5n ulp, for gentype with vector width n	Implementation-defined
log	<= 2 ulp	<= 3 ulp
log2	<= 2 ulp	<= 3 ulp

Function	Min Accuracy - Full Profile	Min Accuracy - Embedded Profile
log10	<= 2 ulp	<= 3 ulp
log1p	<= 2 ulp	<= 3 ulp
logb	0 ulp	0 ulp
mad	Implementation-defined	Implementation-defined
max	0 ulp	0 ulp
maxmag	0 ulp	0 ulp
min	0 ulp	0 ulp
minmag	0 ulp	0 ulp
mix	Implementation-defined	Implementation-defined
modf	0 ulp	0 ulp
nan	0 ulp	0 ulp
nextafter	0 ulp	0 ulp
normalize	<= 1 + n ulp, for gentype with vector width n	Implementation-defined
pow(x, y)	<= 4 ulp	<= 5 ulp
pown(x, y)	<= 4 ulp	<= 5 ulp
powr(x, y)	<= 4 ulp	<= 5 ulp
radians	<= 2 ulp	<= 2 ulp
remainder	0 ulp	0 ulp
remquo	0 ulp for the remainder, at least the lower 7 bits of the integral quotient	0 ulp for the remainder, at least the lower 7 bits of the integral quotient
rint	Correctly rounded	Correctly rounded
rootn	<= 4 ulp	<= 5 ulp
round	Correctly rounded	Correctly rounded
rsqrt	<=1 ulp	<=1 ulp
sign	0 ulp	0 ulp
sin	<= 2 ulp	<= 2 ulp
sincos	<= 2 ulp for sine and cosine values	<= 2 ulp for sine and cosine values
sinh	<= 2 ulp	<= 3 ulp
sinpi	<= 2 ulp	<= 2 ulp
smoothstep	Implementation-defined	Implementation-defined
sqrt	Correctly rounded	<= 1 ulp

Function	Min Accuracy - Full Profile	Min Accuracy - Embedded Profile
step	0 ulp	0 ulp
tan	<= 2 ulp	<= 3 ulp
tanh	<= 2 ulp	<= 3 ulp
tanpi	<= 2 ulp	<= 3 ulp
tgamma	<= 4 ulp	<= 4 ulp
trunc	Correctly rounded	Correctly rounded

Note: Implementations may perform floating-point operations on half scalar or vector data types by converting the half values to single precision floating-point values and performing the operation in single precision floating-point. In this case, the implementation will use the half scalar or vector data type as a storage only format.

5.3. Additions to Chapter 8 of the OpenCL 2.0 C Specification

Add new sub-sections to section 8.3.1. Conversion rules for normalized integer channel data types:

5.3.1. Converting normalized integer channel data types to half precision floating-point values

For images created with image channel data type of CL_UNORM_INT8 and CL_UNORM_INT16, read_imagef will convert the channel values from an 8-bit or 16-bit unsigned integer to normalized half precision floating-point values in the range [0.0h, 1.0h].

For images created with image channel data type of CL_SNORM_INT8 and CL_SNORM_INT16, read_imagef will convert the channel values from an 8-bit or 16-bit signed integer to normalized half precision floating-point values in the range [-1.0h, 1.0h].

These conversions are performed as follows:

```
CL_UNORM_INT8 (8-bit unsigned integer) → half
    normalized half value = round_to_half(c / 255)

CL_UNORM_INT_101010 (10-bit unsigned integer) → half
    normalized half value = round_to_half(c / 1023)

CL_UNORM_INT16 (16-bit unsigned integer) → half
    normalized half value = round_to_half(c / 65535)

CL_SNORM_INT8 (8-bit signed integer) → half
    normalized half value = max(-1.0h, round_to_half(c / 127))
```

```
CL_SNORM_INT16 (16-bit signed integer) → half
   normalized half value = max(-1.0h, round to half(c / 32767))
The accuracy of the above conversions must be <= 1.5 ulp except for the following cases.
For CL UNORM INT8
   0 must convert to 0.0h and
   255 must convert to 1.0h
For CL_UNORM_INT_101010
   0 must convert to 0.0h and
   1023 must convert to 1.0h
For CL_UNORM_INT16
   0 must convert to 0.0h and
   65535 must convert to 1.0h
For CL_SNORM_INT8
   -128 and -127 must convert to -1.0h,
   0 must convert to 0.0h and
   127 must convert to 1.0h
For CL SNORM INT16
   -32768 and -32767 must convert to -1.0h,
   0 must convert to 0.0h and
   32767 must convert to 1.0h
```

5.3.2. Converting half precision floating-point values to normalized integer channel data types

For images created with image channel data type of CL_UNORM_INT8 and CL_UNORM_INT16, write_imagef will convert the floating-point color value to an 8-bit or 16-bit unsigned integer.

For images created with image channel data type of CL_SNORM_INT8 and CL_SNORM_INT16, write_imagef will convert the floating-point color value to an 8-bit or 16-bit signed integer.

The preferred conversion uses the round to nearest even (_rte) rounding mode, but OpenCL implementations may choose to approximate the rounding mode used in the conversions described below. When approximate rounding is used instead of the preferred rounding, the result of the conversion must satisfy the bound given below.

```
half → CL_UNORM_INT8 (8-bit unsigned integer)
```

```
Let f_{\text{exact}} = \max(0, \min(f * 255, 255))
    Let f_{preferred} = convert\_uchar\_sat\_rte(f * 255.0f)
    Let f_{approx} = convert\_uchar\_sat\_<impl-rounding-mode>(f * 255.0f)
    fabs(f_{exact} - f_{approx}) must be <= 0.6
half → CL UNORM INT 101010 (10-bit unsigned integer)
    Let f_{exact} = max(0, min(f * 1023, 1023))
    Let f_{preferred} = min(convert\_ushort\_sat\_rte(f * 1023.0f), 1023)
    Let f_{approx} = convert\_ushort\_sat\_<impl-rounding-mode>(f * 1023.0f)
    fabs(f_{exact} - f_{approx}) must be <= 0.6
half → CL_UNORM_INT16 (16-bit unsigned integer)
    Let f_{\text{exact}} = \max(0, \min(f * 65535, 65535))
    Let f_{preferred} = convert\_ushort\_sat\_rte(f * 65535.0f)
    Let f_{approx} = convert\_ushort\_sat\_<impl-rounding-mode>(f * 65535.0f)
    fabs(f_{exact} - f_{approx}) must be <= 0.6
half → CL_SNORM_INT8 (8-bit signed integer)
    Let f_{\text{exact}} = \text{max}(-128, \text{min}(f * 127, 127))
    Let f_{preferred} = convert\_char\_sat\_rte(f * 127.0f)
    Let f_{approx} = convert\_char\_sat\_<impl\_rounding\_mode>(f * 127.0f)
    fabs(f_{exact} - f_{approx}) must be <= 0.6
half → CL_SNORM_INT16 (16-bit signed integer)
    Let f_{exact} = max(-32768, min(f * 32767, 32767))
    Let f_{preferred} = convert\_short\_sat\_rte(f * 32767.0f)
    Let f_{approx} = convert\_short\_sat\_<impl-rounding-mode>(f * 32767.0f)
    fabs(f_{exact} - f_{approx}) must be <= 0.6
```

Chapter 6. Double Precision Floating-Point

This section describes the **cl_khr_fp64** extension. This extension became an optional core feature in OpenCL 1.2.

6.1. General information

6.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

6.2. Additions to Chapter 6

The list of built-in scalar, and vector data types defined in *tables 6.1* and *6.2* are extended to include the following:

Туре	Description
double	A double precision float.
double2	A 2-component double-precision floating-point vector.
double3	A 3-component double-precision floating-point vector.
double4	A 4-component double-precision floating-point vector.
double8	A 8-component double-precision floating-point vector.
double16	A 16-component double-precision floating-point vector.

The built-in scalar and vector data types for doublen are also declared as appropriate types in the OpenCL API (and header files) that can be used by an application. The following table describes the built-in scalar and vector data types for doublen as defined in the OpenCL C programming language and the corresponding data type available to the application:

Type in OpenCL Language	API type for application
double	cl_double
double2	cl_double2
double3	cl_double3
double4	cl_double4
double8	cl_double8
double16	cl_double16

The double data type must conform to the IEEE-754 double precision storage format.

The following text is added to *Section 6.1.1.1 The half data type*:

Conversions from double to half are correctly rounded. Conversions from half to double are lossless.

6.2.1. Conversions

The implicit conversion rules specified in *section 6.2.1* now include the double scalar and doublen vector data types.

The explicit casts described in *section 6.2.2* are extended to take a double scalar data type and a doublen vector data type.

The explicit conversion functions described in *section 6.2.3* are extended to take a double scalar data type and a doublen vector data type.

The as_typen() function for re-interpreting types as described in *section 6.2.4.2* is extended to allow conversion-free casts between longn, ulongn and doublen scalar and vector data types.

6.2.2. Math Functions

The built-in math functions defined in *table 6.8* (also listed below) are extended to include appropriate versions of functions that take double and double{2|3|4|8|16} as arguments and return values. gentype now also includes double, double2, double3, double4, double8 and double16.

For any specific use of a function, the actual type has to be the same for all arguments and the return type.

Table 12. Double Precision Built-in Math Functions

Function	Description
gentype acos (gentype <i>x</i>)	Arc cosine function.
gentype acosh (gentype x)	Inverse hyperbolic cosine.
gentype acospi (gentype x)	Compute acos (x) / π .
gentype asin (gentype x)	Arc sine function.
gentype asinh (gentype x)	Inverse hyperbolic sine.
gentype asinpi (gentype x)	Compute asin (x) / π .
gentype atan (gentype <i>y_over_x</i>)	Arc tangent function.
gentype atan2 (gentype <i>y</i> , gentype <i>x</i>)	Arc tangent of y / x .
gentype atanh (gentype <i>x</i>)	Hyperbolic arc tangent.
gentype atanpi (gentype <i>x</i>)	Compute atan (x) / π .
gentype atan2pi (gentype <i>y</i> , gentype <i>x</i>)	Compute atan2 $(y, x) / \pi$.
gentype cbrt (gentype <i>x</i>)	Compute cube-root.
gentype ceil (gentype x)	Round to integral value using the round to positive infinity rounding mode.

Function	Description
gentype copysign (gentype <i>x</i> , gentype <i>y</i>)	Returns <i>x</i> with its sign changed to match the sign of <i>y</i> .
gentype cos (gentype <i>x</i>)	Compute cosine.
gentype cosh (gentype <i>x</i>)	Compute hyperbolic cosine.
gentype cospi (gentype <i>x</i>)	Compute $\cos (\pi x)$.
gentype erfc (gentype <i>x</i>)	Complementary error function.
gentype erf (gentype x)	Error function encountered in integrating the normal distribution.
gentype exp (gentype x)	Compute the base- e exponential of x .
gentype exp2 (gentype <i>x</i>)	Exponential base 2 function.
gentype exp10 (gentype x)	Exponential base 10 function.
gentype expm1 (gentype x)	Compute e^x - 1.0.
gentype fabs (gentype x)	Compute absolute value of a floating-point number.
gentype fdim (gentype <i>x</i> , gentype <i>y</i>)	x - y if $x > y$, +0 if x is less than or equal to y.
gentype floor (gentype <i>x</i>)	Round to integral value using the round to negative infinity rounding mode.
gentype fma (gentype a , gentype b , gentype c)	Returns the correctly rounded floating-point representation of the sum of c with the infinitely precise product of a and b . Rounding of intermediate products shall not occur. Edge case behavior is per the IEEE 754-2008 standard.
gentype fmax (gentype <i>x</i> , gentype <i>y</i>) gentype fmax (gentype <i>x</i> , double <i>y</i>)	Returns y if $x < y$, otherwise it returns x . If one argument is a NaN, fmax() returns the other argument. If both arguments are NaNs, fmax() returns a NaN.
gentype fmin (gentype <i>x</i> , gentype <i>y</i>) gentype fmin (gentype <i>x</i> , double <i>y</i>)	Returns <i>y</i> if <i>y</i> < <i>x</i> , otherwise it returns <i>x</i> . If one argument is a NaN, fmin() returns the other argument. If both arguments are NaNs, fmin() returns a NaN.
gentype fmod (gentype <i>x</i> , gentype <i>y</i>)	Modulus. Returns $x - y * trunc (x/y)$.
gentype fract (gentype x ,global gentype * $iptr$) gentype fract (gentype x ,local gentype * $iptr$) gentype fract (gentype x ,private gentype * $iptr$)	Returns fmin (x - floor (x), 0x1. ffffffffffffp-1). floor (x) is returned in $iptr$.

Function	Description
doublen frexp (doublen x,global intn *exp) doublen frexp (doublen x,local intn *exp) doublen frexp (doublen x,private intn *exp) double frexp (double x,global int *exp) double frexp (double x,local int *exp) double frexp (double x,private int *exp)	Extract mantissa and exponent from x . For each component the mantissa returned is a float with magnitude in the interval [1/2, 1) or 0. Each component of x equals mantissa returned * 2^{exp} .
gentype hypot (gentype <i>x</i> , gentype <i>y</i>)	Compute the value of the square root of $x^2 + y^2$ without undue overflow or underflow.
int <i>n</i> ilogb (double <i>n x</i>) int ilogb (double <i>x</i>)	Return the exponent as an integer value.
double n ldexp (double n x , int n k) double n ldexp (double n x , int n k) double ldexp (double n x , int n k)	Multiply <i>x</i> by 2 to the power <i>k</i> .
gentype lgamma (gentype x) doublen lgamma_r (doublen x,global intn *signp) doublen lgamma_r (doublen x,local intn *signp) doublen lgamma_r (doublen x,private intn *signp) double lgamma_r (double x,global int *signp) double lgamma_r (double x,local int *signp) double lgamma_r (double x,private int * signp)	Log gamma function. Returns the natural logarithm of the absolute value of the gamma function. The sign of the gamma function is returned in the <i>signp</i> argument of lgamma_r .
gentype log (gentype x)	Compute natural logarithm.
gentype log2 (gentype <i>x</i>)	Compute a base 2 logarithm.
gentype log10 (gentype <i>x</i>)	Compute a base 10 logarithm.
gentype log1p (gentype <i>x</i>)	Compute $\log_{e}(1.0 + x)$.
gentype logb (gentype x)	Compute the exponent of x , which is the integral part of $\log_r x $.
gentype ${\bf mad}$ (gentype a , gentype b , gentype c)	mad computes $a * b + c$. The function may compute $a * b + c$ with reduced accuracy in the embedded profile. See the OpenCL SPIR-V Environment Specification for details. On some hardware the mad instruction may provide better performance than expanded computation of $a * b + c$.
gentype maxmag (gentype <i>x</i> , gentype <i>y</i>)	Returns x if $ x > y $, y if $ y > x $, otherwise fmax (x , y).
gentype minmag (gentype <i>x</i> , gentype <i>y</i>)	Returns x if $ x < y $, y if $ y < x $, otherwise fmin (x , y).

Function	Description
gentype modf (gentype x ,global gentype *iptr) gentype modf (gentype x ,local gentype *iptr) gentype modf (gentype x ,private gentype * iptr)	Decompose a floating-point number. The modf function breaks the argument <i>x</i> into integral and fractional parts, each of which has the same sign as the argument. It stores the integral part in the object pointed to by <i>iptr</i> .
double <i>n</i> nan (ulong <i>n nancode</i>) double nan (ulong <i>nancode</i>)	Returns a quiet NaN. The <i>nancode</i> may be placed in the significand of the resulting NaN.
gentype nextafter (gentype <i>x</i> , gentype <i>y</i>)	Computes the next representable double-precision floating-point value following x in the direction of y . Thus, if y is less than x , nextafter () returns the largest representable floating-point number less than x .
gentype pow (gentype <i>x</i> , gentype <i>y</i>)	Compute <i>x</i> to the power <i>y</i> .
double n pown (double n x , int n y) double pown (double x , int y)	Compute x to the power y , where y is an integer.
gentype powr (gentype <i>x</i> , gentype <i>y</i>)	Compute x to the power y , where x is >= 0.
gentype remainder (gentype <i>x</i> , gentype <i>y</i>)	Compute the value r such that $r = x - n^*y$, where n is the integer nearest the exact value of x/y . If there are two integers closest to x/y , n shall be the even one. If r is zero, it is given the same sign as x .
doublen remquo (doublen x, doublen y,global intn *quo) doublen remquo (doublen x, doublen y,local intn *quo) doublen remquo (doublen x, doublen y,private intn *quo) double remquo (double x, double y,global int *quo) double remquo (double x, double y,local int *quo) double remquo (double x, double y,private int *quo)	The remquo function computes the value r such that $r = x - k^*y$, where k is the integer nearest the exact value of x/y . If there are two integers closest to x/y , k shall be the even one. If r is zero, it is given the same sign as x . This is the same value that is returned by the remainder function. remquo also calculates the lower seven bits of the integral quotient x/y , and gives that value the same sign as x/y . It stores this signed value in the object pointed to by y
gentype rint (gentype <i>x</i>)	Round to integral value (using round to nearest even rounding mode) in floating-point format. Refer to section 7.1 for description of rounding modes.
double <i>n</i> rootn (double <i>n x</i> , int <i>n y</i>) double rootn (double <i>x</i> , int <i>y</i>)	Compute <i>x</i> to the power 1/ <i>y</i> .
gentype round (gentype <i>x</i>)	Return the integral value nearest to <i>x</i> rounding halfway cases away from zero, regardless of the current rounding direction.

Function	Description
gentype rsqrt (gentype <i>x</i>)	Compute inverse square root.
gentype sin (gentype <i>x</i>)	Compute sine.
gentype sincos (gentype <i>x</i> ,global gentype *cosval) gentype sincos (gentype <i>x</i> ,local gentype *cosval) gentype sincos (gentype <i>x</i> ,private gentype *cosval)	Compute sine and cosine of x. The computed sine is the return value and computed cosine is returned in <i>cosval</i> .
gentype sinh (gentype <i>x</i>)	Compute hyperbolic sine.
gentype sinpi (gentype <i>x</i>)	Compute $\sin (\pi x)$.
gentype sqrt (gentype <i>x</i>)	Compute square root.
gentype tan (gentype <i>x</i>)	Compute tangent.
gentype tanh (gentype <i>x</i>)	Compute hyperbolic tangent.
gentype tanpi (gentype <i>x</i>)	Compute $tan (\pi x)$.
gentype tgamma (gentype <i>x</i>)	Compute the gamma function.
gentype trunc (gentype <i>x</i>)	Round to integral value using the round to zero rounding mode.

In addition, the following symbolic constant will also be available:

HUGE_VAL - A positive double expression that evaluates to infinity. Used as an error value returned by the built-in math functions.

The **FP_FAST_FMA** macro indicates whether the **fma()** family of functions are fast compared with direct code for double precision floating-point. If defined, the **FP_FAST_FMA** macro shall indicate that the **fma()** function generally executes about as fast as, or faster than, a multiply and an add of **double** operands.

The macro names given in the following list must use the values specified. These constant expressions are suitable for use in #if preprocessing directives.

```
#define DBL DIG
                            15
#define DBL_MANT_DIG
                            53
#define DBL MAX 10 EXP
                            +308
#define DBL_MAX_EXP
                            +1024
#define DBL_MIN_10_EXP
                            -307
#define DBL MIN EXP
                            -1021
#define DBL RADIX
#define DBL_MAX
                            0x1.ffffffffffffp1023
#define DBL MIN
                            0x1.0p-1022
#define DBL EPSILON
                            0x1.0p-52
```

The following table describes the built-in macro names given above in the OpenCL C programming

language and the corresponding macro names available to the application.

Macro in OpenCL Language	Macro for application
DBL_DIG	CL_DBL_DIG
DBL_MANT_DIG	CL_DBL_MANT_DIG
DBL_MAX_10_EXP	CL_DBL_MAX_10_EXP
DBL_MAX_EXP	CL_DBL_MAX_EXP
DBL_MIN_10_EXP	CL_DBL_MIN_10_EXP
DBL_MIN_EXP	CL_DBL_MIN_EXP
DBL_RADIX	CL_DBL_RADIX
DBL_MAX	CL_DBL_MAX
DBL_MIN	CL_DBL_MIN
DBL_EPSILSON	CL_DBL_EPSILON

The following constants are also available. They are of type double and are accurate within the precision of the double type.

Constant	Description
M_E	Value of e
M_LOG2E	Value of log₂e
M_LOG10E	Value of log ₁₀ e
M_LN2	Value of log _e 2
M_LN10	Value of log _e 10
M_PI	Value of π
M_PI_2	Value of π / 2
M_PI_4	Value of π / 4
M_1_PI	Value of 1 / π
M_2_PI	Value of 2 / π
M_2_SQRTPI	Value of 2 / $\sqrt{\pi}$
M_SQRT2	Value of √2
M_SQRT1_2	Value of 1 / √2

6.2.3. Common Functions

The built-in common functions defined in *table 6.12* (also listed below) are extended to include appropriate versions of functions that take double and double{2|3|4|8|16} as arguments and return values. gentype now also includes double, double2, double3, double4, double8 and double16. These are described below.

Table 13. Double Precision Built-in Common Functions

Function	Description
gentype clamp (gentype <i>x</i> , gentype <i>minval</i> , gentype <i>maxval</i>) gentype clamp (gentype <i>x</i> , double <i>minval</i> , double <i>maxval</i>)	Returns $fmin(fmax(x, minval), maxval)$. Results are undefined if $minval > maxval$.
gentype degrees (gentype <i>radians</i>)	Converts $radians$ to degrees, i.e. $(180 / \pi) * radians$.
gentype max (gentype <i>x</i> , gentype <i>y</i>) gentype max (gentype <i>x</i> , double <i>y</i>)	Returns y if $x < y$, otherwise it returns x . If x and y are infinite or NaN, the return values are undefined.
gentype min (gentype <i>x</i> , gentype <i>y</i>) gentype min (gentype <i>x</i> , double <i>y</i>)	Returns y if $y < x$, otherwise it returns x . If x and y are infinite or NaN, the return values are undefined.
gentype mix (gentype <i>x</i> , gentype <i>y</i> , gentype <i>a</i>) gentype mix (gentype <i>x</i> , gentype <i>y</i> , double <i>a</i>)	Returns the linear blend of x and y implemented as: $x + (y - x) * a$ a must be a value in the range 0.0 1.0. If a is not in the range 0.0 1.0, the return values are undefined. Note: The double precision mix function can be implemented using contractions such as mad or fma .
gentype radians (gentype <i>degrees</i>)	Converts degrees to radians, i.e. $(\pi / 180)$ * degrees.
gentype step (gentype <i>edge</i> , gentype <i>x</i>) gentype step (double <i>edge</i> , gentype <i>x</i>)	Returns 0.0 if $x < edge$, otherwise it returns 1.0.

Function	Description
gentype smoothstep (gentype <i>edge1</i> , gentype <i>x</i>)	Returns 0.0 if $x \le edge0$ and 1.0 if $x \ge edge1$ and performs smooth Hermite interpolation between 0 and 1 when $edge0 < x < edge1$. This is useful in cases where you would want a
gentype smoothstep (threshold function with a smooth transition.
double <i>edge0</i> , double <i>edge1</i> , gentype <i>x</i>)	This is equivalent to: gentype t ; $t = \text{clamp } ((x - edge0) / (edge1 - edge0), 0, 1);$ return $t * t * (3 - 2 * t);$
	Results are undefined if <i>edge0</i> >= <i>edge1</i> . Note: The double precision smoothstep function can be implemented using contractions such as mad or fma .
gentype sign (gentype <i>x</i>)	Returns 1.0 if $x > 0$, -0.0 if $x = -0.0$, +0.0 if $x = +0.0$, or -1.0 if $x < 0$. Returns 0.0 if x is a NaN.

6.2.4. Geometric Functions

The built-in geometric functions defined in *table 6.13* (also listed below) are extended to include appropriate versions of functions that take double and double{2|3|4} as arguments and return values. gentype now also includes double, double2, double3 and double4. These are described below.

Note: The double precision geometric functions can be implemented using contractions such as **mad** or **fma**.

Table 14. Double Precision Built-in Geometric Functions

Function	Description
double4 cross (double4 <i>p0</i> , double4 <i>p1</i>) double3 cross (double3 <i>p0</i> , double3 <i>p1</i>)	Returns the cross product of $p0.xyz$ and $p1.xyz$. The w component of the result will be 0.0.
double dot (gentype <i>p0</i> , gentype <i>p1</i>)	Compute the dot product of $p0$ and $p1$.
double distance (gentype <i>p0</i> , gentype <i>p1</i>)	Returns the distance between $p0$ and $p1$. This is calculated as length ($p0 - p1$).
double length (gentype <i>p</i>)	Return the length of vector x, i.e., sqrt($p.x^2 + p.y^2 +$)
gentype normalize (gentype <i>p</i>)	Returns a vector in the same direction as p but with a length of 1.

6.2.5. Relational Functions

The scalar and vector relational functions described in *table 6.14* are extended to include versions that take double, double3, double4, double8 and double16 as arguments.

The relational and equality operators (<, <=, >, >=, !=, ==) can be used with doublen vector types and shall produce a vector longn result as described in *section 6.3*.

The functions **isequal**, **isnotequal**, **isgreater**, **isgreaterequal**, **isless**, **islessequal**, **islessgreater**, **isfinite**, **isinf**, **isnan**, **isnormal**, **isordered**, **isunordered** and **signbit** shall return a 0 if the specified relation is *false* and a 1 if the specified relation is true for scalar argument types. These functions shall return a 0 if the specified relation is *false* and a -1 (i.e. all bits set) if the specified relation is *true* for vector argument types.

The relational functions **isequal**, **isgreater**, **isgreaterequal**, **isless**, **islessequal**, and **islessgreater** always return 0 if either argument is not a number (NaN). **isnotequal** returns 1 if one or both arguments are not a number (NaN) and the argument type is a scalar and returns -1 if one or both arguments are not a number (NaN) and the argument type is a vector.

The functions described in *table 6.14* are extended to include the doublen' vector types.

Table 15. Double Precision Relational Functions

Function	Description
int isequal (double <i>x</i> , double <i>y</i>) long <i>n</i> isequal (double <i>n x</i> , double <i>n y</i>)	Returns the component-wise compare of $x == y$.
int isnotequal (double <i>x</i> , double <i>y</i>) long <i>n</i> isnotequal (double <i>n x</i> , double <i>n y</i>)	Returns the component-wise compare of $x = y$.
int isgreater (double <i>x</i> , double <i>y</i>) long <i>n</i> isgreater (double <i>n x</i> , double <i>n y</i>)	Returns the component-wise compare of $x > y$.
int isgreaterequal (double <i>x</i> , double <i>y</i>) long <i>n</i> isgreaterequal (double <i>n x</i> , double <i>n y</i>)	Returns the component-wise compare of $x \ge y$.
int isless (double <i>x</i> , double <i>y</i>) long <i>n</i> isless (double <i>n x</i> , double <i>n y</i>)	Returns the component-wise compare of $x < y$.
int islessequal (double <i>x</i> , double <i>y</i>) long <i>n</i> islessequal (double <i>n x</i> , double <i>n y</i>)	Returns the component-wise compare of $x \le y$.
int islessgreater (double <i>x</i> , double <i>y</i>) long <i>n</i> islessgreater (double <i>n x</i> , double <i>n y</i>)	Returns the component-wise compare of $(x < y)$ $(x > y)$.
int isfinite (double) long <i>n</i> isfinite (double <i>n</i>)	Test for finite value.
int isinf (double) long <i>n</i> isinf (double <i>n</i>)	Test for infinity value (positive or negative) .
int isnan (double) long <i>n</i> isnan (double <i>n</i>)	Test for a NaN.

Function	Description
int isnormal (double) long <i>n</i> isnormal (double <i>n</i>)	Test for a normal value.
int isordered (double <i>x</i> , double <i>y</i>) long <i>n</i> isordered (double <i>n x</i> , double <i>n y</i>)	Test if arguments are ordered. isordered () takes arguments <i>x</i> and <i>y</i> , and returns the result isequal (<i>x</i> , <i>x</i>) && isequal (<i>y</i> , <i>y</i>).
int isunordered (double <i>x</i> , double <i>y</i>) long <i>n</i> isunordered (double <i>n x</i> , double <i>n y</i>)	Test if arguments are unordered. isunordered () takes arguments <i>x</i> and <i>y</i> , returning non-zero if <i>x</i> or <i>y</i> is a NaN, and zero otherwise.
int signbit (double) long <i>n</i> signbit (double <i>n</i>)	Test for sign bit. The scalar version of the function returns a 1 if the sign bit in the double is set else returns 0. The vector version of the function returns the following for each component in double n: -1 (i.e all bits set) if the sign bit in the double is set else returns 0.
double n bitselect (double n a , double n b , double n c)	Each bit of the result is the corresponding bit of a if the corresponding bit of c is 0. Otherwise it is the corresponding bit of b .
double <i>n</i> select (double <i>n a</i> , double <i>n b</i> , long <i>n c</i>) double <i>n</i> select (double <i>n a</i> , double <i>n b</i> , ulong <i>n c</i>)	For each component, result[i] = if MSB of c[i] is set ? b[i] : a[i].

6.2.6. Vector Data Load and Store Functions

The vector data load (**vloadn**) and store (**vstoren**) functions described in *table 6.13* (also listed below) are extended to include versions that read from or write to double scalar or vector values. The generic type gentype is extended to include double. The generic type gentypen is extended to include double2, double3, double4, double8 and double16. The **vstore_half**, **vstore_half** and **vstorea_half** functions are extended to allow a double precision scalar or vector value to be written to memory as half values.

Note: **vload3** reads (x,y,z) components from address (p + (offset * 3)) into a 3-component vector. **vstore3**, and **vstore_half3** write (x,y,z) components from a 3-component vector to address (p + (offset * 3)). In addition, **vloada_half3** reads (x,y,z) components from address (p + (offset * 4)) into a 3-component vector and **vstorea_half3** writes (x,y,z) components from a 3-component vector to address (p + (offset * 4)). Whether **vloada_half3** and **vstorea_half3** read/write padding data between the third vector element and the next alignment boundary is implementation defined. **vloada_** and **vstoreaa_** variants are provided to access data that is aligned to the size of the vector, and are intended to enable performance on hardware that can take advantage of the increased alignment.

Table 16. Double Precision Vector Data Load and Store Functions

Function	Description
gentypen vloadn (size_t <i>offset</i> , constglobal gentype *p)	Return sizeof (gentype n) bytes of data read from address (p + (offset * n)). If gentype is double, the read address computed as (p + (offset * n)) must
gentype <i>n</i> vload <i>n</i> (size_t <i>offset</i> , constlocal gentype * <i>p</i>)	be 64-bit aligned.
gentype <i>n</i> vload <i>n</i> (size_t <i>offset</i> , constconstant gentype * <i>p</i>)	
gentype <i>n</i> vload <i>n</i> (size_t <i>offset</i> , constprivate gentype * <i>p</i>)	
void vstoren (gentypen data, size_t offset,global gentype *p)	Write sizeof (gentypen) bytes given by <i>data</i> to address ($p + (offset * n)$). If gentype is double, the write address computed as ($p + (offset * n)$) must
void vstoren (gentypen data, size_t offset,local gentype *p)	be 64-bit aligned.
void vstore <i>n</i> (gentype <i>n data</i> , size_t <i>offset</i> ,private gentype *p)	

Function Description void **vstore_half**(double *data*, size_t *offset*, The double value given by *data* is first converted to a half value using the appropriate rounding global half *p) void vstore_half_rte(double data, size_t offset, mode. The half value is then written to the global half *p) address computed as (p + offset). The address void **vstore_half_rtz**(double *data*, size_t *offset*, computed as (p + offset) must be 16-bit aligned. $_{\rm global\ half\ *p)}$ void **vstore_half_rtp**(double *data*, size_t *offset*, vstore_half uses the current rounding mode. __global half *p) The default current rounding mode is round to void **vstore_half_rtn**(double data, size_t offset, nearest even. $_{global}$ half *p) void **vstore_half**(double data, size_t offset, $_$ local half *p) void vstore_half_rte(double data, size_t offset, local half *p) void **vstore_half_rtz**(double *data*, size_t *offset*, local half *p) void **vstore_half_rtp**(double *data*, size_t *offset*, local half *p) void **vstore_half_rtn**(double *data*, size_t *offset*, _local half *p) void **vstore_half**(double data, size_t offset, __private half *p) void **vstore_half_rte**(double *data*, size_t *offset*, __private half *p) void **vstore_half_rtz**(double *data*, size_t *offset*, __private half *p) void **vstore_half_rtp**(double *data*, size_t *offset*, __private half *p) void **vstore_half_rtn**(double *data*, size_t *offset*,

_private half *p)

Function

void vstore_halfn(doublen data, size_t offset,
 __global half *p)

void vstore_halfn_rte(doublen data, size_t
offset, __global half *p)

void **vstore_halfn_rtz**(double*n data*, size_t *offset*, __global half **p*)

void vstore_halfn_rtp(doublen data, size_t
offset, __global half *p)

void vstore_halfn_rtn(doublen data, size_t
offset, __global half *p)

void vstore_halfn(doublen data, size_t offset,
_local half *p)

void vstore_halfn_rte(doublen data, size_t
offset, __local half *p)

void vstore_halfn_rtz(doublen data, size_t
offset, __local half *p)

void vstore_halfn_rtp(doublen data, size_t
offset, local half *p)

void vstore_halfn_rtn(doublen data, size_t
offset, _local half *p)

void vstore_halfn(doublen data, size_t offset,
__private half *p)

void vstore_halfn_rte(doublen data, size_t
offset, __private half *p)

void vstore_halfn_rtz(doublen data, size_t
offset, __private half *p)

void **vstore_halfn_rtp**(double*n data*, size_t *offset*, __private half **p*)

void vstore_halfn_rtn(doublen data, size_t
offset, __private half *p)

Description

The double*n* value given by *data* is converted to a half*n* value using the appropriate rounding mode. The half*n* value is then written to the address computed as (p + (offset * n)). The address computed as (p + (offset * n)) must be 16-bit aligned.

vstore_half*n* uses the current rounding mode. The default current rounding mode is round to nearest even.

Function

void vstorea_halfn(doublen data, size_t offset,
 _global half *p)

void **vstorea_halfn_rte**(double*n data*, size_t *offset*, __global half **p*)

void vstorea_halfn_rtz(doublen data, size_t
offset, __global half *p)

void **vstorea_halfn_rtp**(double*n data*, size_t *offset*, __global half **p*)

void **vstorea_half***n_***rtn**(double*n data*, size_t *offset*, __global half **p*)

void vstorea_halfn(doublen data, size_t offset,
 __local half *p)

void vstorea_halfn_rte(doublen data, size_t
offset, __local half *p)

void vstorea_halfn_rtz(doublen data, size_t
offset, __local half *p)

void vstorea_halfn_rtp(doublen data, size_t
offset, local half *p)

void vstorea_halfn_rtn(doublen data, size_t
offset, local half *p)

void vstorea_halfn(doublen data, size_t offset,
 _private half *p)

void **vstorea_halfn_rte**(double*n data*, size_t *offset*, __private half **p*)

void vstorea_halfn_rtz(doublen data, size_t
offset, __private half *p)

void vstorea_halfn_rtp(doublen data, size_t
offset, __private half *p)

void vstorea_halfn_rtn(doublen data, size_t
offset, __private half *p)

Description

The double*n* value is converted to a half*n* value using the appropriate rounding mode.

For n = 1, 2, 4, 8 or 16, the half n value is written to the address computed as (p + (offset * n)). The address computed as (p + (offset * n)) must be aligned to size of (half n) bytes.

For n = 3, the half3 value is written to the address computed as (p + (offset * 4)). The address computed as (p + (offset * 4)) must be aligned to size of (half) * 4 bytes.

vstorea_half*n* uses the current rounding mode. The default current rounding mode is round to nearest even.

6.2.7. Async Copies from Global to Local Memory, Local to Global Memory, and Prefetch

The OpenCL C programming language implements the following functions that provide asynchronous copies between global and local memory and a prefetch from global memory.

The generic type gentype is extended to include double, double2, double3, double4, double8 and double16.

Table 17. Double Precision Built-in Async Copy and Prefetch Functions

Function	Description
event_t async_work_group_copy (Perform an async copy of <i>num_gentypes</i> gentype
_local gentype *dst,	elements from <i>src</i> to <i>dst</i> . The async copy is
constglobal gentype *src,	performed by all work-items in a work-group
size_t num_gentypes, event_t event)	and this built-in function must therefore be
	encountered by all work-items in a work-group
event_t async_work_group_copy (executing the kernel with the same argument
global gentype *dst,	values; otherwise the results are undefined.
const _local gentype *src,	
size_t num_gentypes, event_t event)	Returns an event object that can be used by
	wait_group_events to wait for the async copy to
	finish. The event argument can also be used to
	associate the async_work_group_copy with a
	previous async copy allowing an event to be
	shared by multiple async copies; otherwise <i>event</i>
	should be zero.
	Te abia at
	If event argument is not zero, the event object
	supplied in <i>event</i> argument will be returned.
	This function does not perform any implicit
	synchronization of source data such as using a
	barrier before performing the copy.
	1 0 17

Function	Description
event_t async_work_group_strided_copy (local gentype *dst, constglobal gentype *src, size_t num_gentypes, size_t src_stride, event_t event) event_t async_work_group_strided_copy (global gentype *dst, constlocal gentype *src, size_t num_gentypes, size_t dst_stride, event_t event)	Perform an async gather of num_gentypes gentype elements from src to dst. The src_stride is the stride in elements for each gentype element read from src. The async gather is performed by all work-items in a work-group and this built-in function must therefore be encountered by all work-items in a work-group executing the kernel with the same argument values; otherwise the results are undefined. Returns an event object that can be used by wait_group_events to wait for the async copy to finish. The event argument can also be used to associate the async_work_group_strided_copy with a previous async copy allowing an event to be shared by multiple async copies; otherwise event should be zero. If event argument is not zero, the event object
	supplied in <i>event</i> argument will be returned. This function does not perform any implicit synchronization of source data such as using a barrier before performing the copy. The behavior of async_work_group_strided_copy is undefined if src_stride or dst_stride is 0, or if the src_stride or dst_stride values cause the src or dst pointers to exceed the upper bounds of the address space during the copy.
void wait_group_events (int num_events, event_t *event_list)	Wait for events that identify the async_work_group_copy operations to complete. The event objects specified in event_list will be released after the wait is performed. This function must be encountered by all work-items in a work-group executing the kernel with the same num_events and event objects specified in event_list; otherwise the results are undefined.

Function	Description
<pre>void prefetch (constglobal gentype *p, size_t num_gentypes)</pre>	Prefetch <i>num_gentypes</i> * sizeof(gentype) bytes into the global cache. The prefetch instruction is
constgrobat gentlype p, size_t num_gentlypes)	applied to a work-item in a work-group and does
	not affect the functional behavior of the kernel.

6.2.8. IEEE754 Compliance

The following table entry describes the additions to *table 4.3*, which allows applications to query the configuration information using **clGetDeviceInfo** for an OpenCL device that supports double precision floating-point.

Op-code	Return Type	Description
CL_DEVICE_DOUBLE_ FP_CONFIG	cl_device_fp_config	Describes double precision floating-point capability of the OpenCL device. This is a bit-field that describes one or more of the following values:
		CL_FP_DENORM — denorms are supported
		CL_FP_INF_NAN — INF and NaNs are supported
		CL_FP_ROUND_TO_NEAREST — round to nearest even rounding mode supported
		CL_FP_ROUND_TO_ZERO — round to zero rounding mode supported
		CL_FP_ROUND_TO_INF — round to positive and negative infinity rounding modes supported
		CL_FP_FMA — IEEE754-2008 fused multiply-add is supported
		CL_FP_SOFT_FLOAT — Basic floating-point operations (such as addition, subtraction, multiplication) are implemented in software.
		The required minimum double precision floating-point capability as implemented by this extension is:
		CL_FP_FMA CL_FP_ROUND_TO_NEAREST CL_FP_ROUND_TO_ZERO CL_FP_ROUND_TO_INF CL_FP_INF_NAN CL_FP_DENORM.

IEEE754 fused multiply-add, denorms, INF and NaNs are required to be supported for double precision floating-point numbers and operations on double precision floating-point numbers.

6.2.9. Relative Error as ULPs

In this section we discuss the maximum relative error defined as *ulp* (units in the last place).

Addition, subtraction, multiplication, fused multiply-add and conversion between integer and a floating-point format are IEEE 754 compliant and are therefore correctly rounded using round-to-nearest even rounding mode.

The following table describes the minimum accuracy of double precision floating-point arithmetic operations given as ULP values. 0 ULP is used for math functions that do not require rounding. The reference value used to compute the ULP value of an arithmetic operation is the infinitely precise result.

Table 18. ULP Values for Double Precision Floating-Point Arithmetic Operations

Function	Min Accuracy
<i>x</i> + <i>y</i>	Correctly rounded
<i>x</i> - <i>y</i>	Correctly rounded
x * y	Correctly rounded
1.0 / x	Correctly rounded
x / y	Correctly rounded
acos	<= 4 ulp
acosh	<= 4 ulp
acospi	<= 5 ulp
asin	<= 4 ulp
asinh	<= 4 ulp
asinpi	<= 5 ulp
atan	<= 5 ulp
atanh	<= 5 ulp
atanpi	<= 5 ulp
atan2	<= 6 ulp
atan2pi	<= 6 ulp
cbrt	<= 2 ulp
ceil	Correctly rounded
clamp	0 ulp
copysign	0 ulp
cos	<= 4 ulp
cosh	<= 4 ulp
cospi	<= 4 ulp
cross	absolute error tolerance of 'max * max * (3 * FLT_EPSILON)' per vector component, where <i>max</i> is the maximum input operand magnitude
degrees	<= 2 ulp
distance	<= 5.5 + 2n ulp, for gentype with vector width n

Function	Min Accuracy	
dot	absolute error tolerance of 'max * max * (2n - 1) * FLT_EPSILON', for vector width <i>n</i> and maximum input operand magnitude <i>max</i> across all vector components	
erfc	<= 16 ulp	
erf	<= 16 ulp	
exp	<= 3 ulp	
exp2	<= 3 ulp	
exp10	<= 3 ulp	
expm1	<= 3 ulp	
fabs	0 ulp	
fdim	Correctly rounded	
floor	Correctly rounded	
fma	Correctly rounded	
fmax	0 ulp	
fmin	0 ulp	
fmod	0 ulp	
fract	Correctly rounded	
frexp	0 ulp	
hypot	<= 4 ulp	
ilogb	0 ulp	
ldexp	Correctly rounded	
length	\leq 5.5 + n ulp, for gentype with vector width n	
log	<= 3 ulp	
log2	<= 3 ulp	
log10	<= 3 ulp	
log1p	<= 2 ulp	
logb	0 ulp	
mad	Implementation-defined	
max	0 ulp	
maxmag	0 ulp	
min	0 ulp	
minmag	0 ulp	
mix	Implementation-defined	

Function	Min Accuracy
modf	0 ulp
nan	0 ulp
nextafter	0 ulp
normalize	<= 4.5 + n ulp, for gentype with vector width n
pow(x, y)	<= 16 ulp
pown(x, y)	<= 16 ulp
powr(x, y)	<= 16 ulp
radians	<= 2 ulp
remainder	0 ulp
remquo	0 ulp for the remainder, at least the lower 7 bits of the integral quotient
rint	Correctly rounded
rootn	<= 16 ulp
round	Correctly rounded
rsqrt	<= 2 ulp
sign	0 ulp
sin	<= 4 ulp
sincos	<= 4 ulp for sine and cosine values
sinh	<= 4 ulp
sinpi	<= 4 ulp
smoothstep	Implementation-defined
sqrt	Correctly rounded
step	0 ulp
tan	<= 5 ulp
tanh	<= 5 ulp
tanpi	<= 6 ulp
tgamma	<= 16 ulp
trunc	Correctly rounded

Chapter 7. 32-bit Atomics

This section describes the extensions cl_khr_global_int32_base_atomics, cl_khr_global_int32_extended_atomics, cl_khr_local_int32_base_atomics, and cl_khr_local_int32_extended_atomics. These extensions allow atomic operations to be performed on 32-bit signed and unsigned integers in global and local memory.

These extensions became core features in OpenCL 1.1, except the built-in atomic function names are changed to use the **atomic_** prefix instead of **atom_** and the volatile qualifier was added to the pointer parameter *p*.

7.1. General information

7.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

7.2. Global Atomics for 32-bit Integers

7.2.1. Base Atomics

Table 19. Built-in Atomic Functions for cl_khr_global_int32_base_atomics

Function	Description
int atom_add (volatileglobal int *p, int val) uint atom_add (volatileglobal uint *p, uint val)	Read the 32-bit value (referred to as old) stored at location pointed by p . Compute ($old + val$) and store result at location pointed by p . The function returns old .
int atom_sub (volatileglobal int *p, int val) uint atom_sub (volatileglobal uint *p, uint val)	Read the 32-bit value (referred to as <i>old</i>) stored at location pointed by <i>p</i> . Compute (<i>old - val</i>) and store result at location pointed by <i>p</i> . The function returns <i>old</i> .
int atom_xchg (volatileglobal int *p, int val) uint atom_xchg (volatileglobal uint *p, uint val)	Swaps the <i>old</i> value stored at location <i>p</i> with new value given by <i>val</i> . Returns <i>old</i> value.
int atom_inc (volatileglobal int *p) uint atom_inc (volatileglobal uint *p)	Read the 32-bit value (referred to as old) stored at location pointed by p . Compute ($old + 1$) and store result at location pointed by p . The function returns old .

Function	Description
int atom_dec (volatileglobal int *p) uint atom_dec (volatileglobal uint *p)	Read the 32-bit value (referred to as <i>old</i>) stored at location pointed by <i>p</i> . Compute (<i>old - 1</i>) and store result at location pointed by <i>p</i> . The function returns <i>old</i> .
int atom_cmpxchg (volatileglobal int *p, int cmp, int val) uint atom_cmpxchg (volatileglobal uint *p, uint cmp, uint val)	Read the 32-bit value (referred to as old) stored at location pointed by p . Compute ($old == cmp$) ? val : old and store result at location pointed by p . The function returns old .

7.2.2. Extended Atomics

 $\textit{Table 20. Built-in Atomic Functions for \textbf{cl_khr_global_int32_extended_atomics}}$

Function	Description
int atom_min (volatileglobal int *p, int val) uint atom_min (volatileglobal uint *p, uint val)	Read the 32-bit value (referred to as <i>old</i>) stored at location pointed by <i>p</i> . Compute min (<i>old</i> , <i>val</i>) and store minimum value at location pointed by <i>p</i> . The function returns <i>old</i> .
int atom_max (volatileglobal int *p, int val) uint atom_max (volatileglobal uint *p, uint val)	Read the 32-bit value (referred to as <i>old</i>) stored at location pointed by <i>p</i> . Compute max (<i>old</i> , <i>val</i>) and store maximum value at location pointed by <i>p</i> . The function returns <i>old</i> .
int atom_and (volatileglobal int *p, int val) uint atom_and (volatileglobal uint *p, uint val)	Read the 32-bit value (referred to as <i>old</i>) stored at location pointed by <i>p</i> . Compute (<i>old</i> & val) and store result at location pointed by <i>p</i> . The function returns <i>old</i> .
int atom_or (volatileglobal int *p, int val) uint atom_or (volatileglobal uint *p, uint val)	Read the 32-bit value (referred to as old) stored at location pointed by p . Compute ($old \mid val$) and store result at location pointed by p . The function returns old .
int atom_xor (volatileglobal int *p, int val) uint atom_xor (volatileglobal uint *p, uint val)	Read the 32-bit value (referred to as old) stored at location pointed by p . Compute ($old \land val$) and store result at location pointed by p . The function returns old .

7.3. Local Atomics for 32-bit Integers

7.3.1. Base Atomics

Table 21. Built-in Atomic Functions for cl_khr_local_int32_base_atomics

Function	Description
int atom_add (volatilelocal int *p, int val) uint atom_add (volatilelocal uint *p, uint val)	Read the 32-bit value (referred to as old) stored at location pointed by p . Compute ($old + val$) and store result at location pointed by p . The function returns old .
int atom_sub (volatilelocal int *p, int val) uint atom_sub (volatilelocal uint *p, uint val)	Read the 32-bit value (referred to as <i>old</i>) stored at location pointed by <i>p</i> . Compute (<i>old - val</i>) and store result at location pointed by <i>p</i> . The function returns <i>old</i> .
int atom_xchg (volatilelocal int *p, int val) uint atom_xchg (volatilelocal uint *p, uint val)	Swaps the <i>old</i> value stored at location <i>p</i> with new value given by <i>val</i> . Returns <i>old</i> value.
int atom_inc (volatilelocal int *p) uint atom_inc (volatilelocal uint *p)	Read the 32-bit value (referred to as old) stored at location pointed by p . Compute ($old + 1$) and store result at location pointed by p . The function returns old .
int atom_dec (volatilelocal int *p) uint atom_dec (volatilelocal uint *p)	Read the 32-bit value (referred to as <i>old</i>) stored at location pointed by <i>p</i> . Compute (<i>old</i> - 1) and store result at location pointed by <i>p</i> . The function returns <i>old</i> .
int atom_cmpxchg (volatilelocal int *p, int cmp, int val) uint atom_cmpxchg (volatilelocal uint *p, uint cmp, uint val)	Read the 32-bit value (referred to as old) stored at location pointed by p . Compute ($old == cmp$) ? val : old and store result at location pointed by p . The function returns old .

7.3.2. Extended Atomics

 $\textit{Table 22. Built-in Atomic Functions for \textbf{cl_khr_local_int32_extended_atomics}}$

Function	Description
int atom_min (volatilelocal int *p, int val) uint atom_min (volatilelocal uint *p, uint val)	Read the 32-bit value (referred to as <i>old</i>) stored at location pointed by <i>p</i> . Compute min (<i>old</i> , <i>val</i>) and store minimum value at location pointed by <i>p</i> . The function returns <i>old</i> .
int atom_max (volatilelocal int *p, int val) uint atom_max (volatilelocal uint *p, uint val)	Read the 32-bit value (referred to as <i>old</i>) stored at location pointed by <i>p</i> . Compute max (<i>old</i> , <i>val</i>) and store maximum value at location pointed by <i>p</i> . The function returns <i>old</i> .
int atom_and (volatilelocal int *p, int val) uint atom_and (volatilelocal uint *p, uint val)	Read the 32-bit value (referred to as <i>old</i>) stored at location pointed by <i>p</i> . Compute (<i>old</i> & val) and store result at location pointed by <i>p</i> . The function returns <i>old</i> .
int atom_or (volatilelocal int *p, int val) uint atom_or (volatilelocal uint *p, uint val)	Read the 32-bit value (referred to as old) stored at location pointed by p . Compute ($old \mid val$) and store result at location pointed by p . The function returns old .
int atom_xor (volatilelocal int *p, int val) uint atom_xor (volatilelocal uint *p, uint val)	Read the 32-bit value (referred to as old) stored at location pointed by p . Compute ($old \land val$) and store result at location pointed by p . The function returns old .

Chapter 8. 64-bit Atomics

This section describes the **cl_khr_int64_base_atomics** and **cl_khr_int64_extended_atomics** extensions. These extensions allow atomic operations to be performed on 64-bit signed and unsigned integers in global and local memory.

8.1. General information

8.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

Table 23. Built-in Atomic Functions for cl_khr_int64_base_atomics

Function	Description		
long atom_add (volatileglobal long *p, long val) long atom_add (volatilelocal long *p, long val) ulong atom_add (volatileglobal ulong *p, ulong val) ulong atom_add (volatilelocal ulong *p, ulong val)	Read the 64-bit value (referred to as old) stored at location pointed by p . Compute ($old + val$) and store result at location pointed by p . The function returns old .		
long atom_sub (volatileglobal long *p, long val) long atom_sub (volatilelocal long *p, long val) ulong atom_sub (volatileglobal ulong *p, ulong val) ulong atom_sub (volatilelocal ulong *p, ulong val)	Read the 64-bit value (referred to as <i>old</i>) stored at location pointed by <i>p</i> . Compute (<i>old</i> - <i>val</i>) and store result at location pointed by <i>p</i> . The function returns <i>old</i> .		
long atom_xchg (volatileglobal long *p, long val) long atom_xchg (volatilelocal long *p, long val) ulong atom_xchg (volatileglobal ulong *p, ulong val) ulong atom_xchg (volatilelocal ulong *p, ulong val)	Swaps the <i>old</i> value stored at location <i>p</i> with new value given by <i>val</i> . Returns <i>old</i> value.		
long atom_inc (volatileglobal long *p) long atom_inc (volatilelocal long *p) ulong atom_inc (volatileglobal ulong *p) ulong atom_inc (volatilelocal ulong *p)	Read the 64-bit value (referred to as old) stored at location pointed by p . Compute ($old + 1$) and store result at location pointed by p . The function returns old .		
long atom_dec (volatileglobal long *p) long atom_dec (volatilelocal long *p) ulong atom_dec (volatileglobal ulong *p) ulong atom_dec (volatilelocal ulong *p)	Read the 64-bit value (referred to as <i>old</i>) stored at location pointed by <i>p</i> . Compute (<i>old</i> - 1) and store result at location pointed by <i>p</i> . The function returns <i>old</i> .		

Function	Description
long atom_cmpxchg (volatileglobal long *p, long cmp, long val)	Read the 64-bit value (referred to as <i>old</i>) stored at location pointed
long atom_cmpxchg (volatilelocal long *p, long cmp, long val)	by <i>p</i> . Compute (old == cmp) ? val : old and store result at location pointed by <i>p</i> . The function returns
ulong atom_cmpxchg (volatileglobal ulong *p, ulong cmp, ulong val) ulong atom_cmpxchg (volatilelocal ulong *p, ulong cmp, ulong val)	old.

 $\textit{Table 24. Built-in Atomic Functions for \textbf{cl_khr_int64_extended_atomics}}$

Function	Description
long atom_min (volatileglobal long *p, long val) long atom_min (volatilelocal long *p, long val) ulong atom_min (volatileglobal ulong *p, ulong val) ulong atom_min (volatilelocal ulong *p, ulong val)	Read the 64-bit value (referred to as old) stored at location pointed by p . Compute $\min(old, val)$ and store minimum value at location pointed by p . The function returns old .
long atom_max (volatileglobal long *p, long val) long atom_max (volatilelocal long *p, long val) ulong atom_max (volatileglobal ulong *p, ulong val) ulong atom_max (volatilelocal ulong *p, ulong val) ulong atom_max (volatilelocal ulong *p, ulong val)	Read the 64-bit value (referred to as old) stored at location pointed by p . Compute $max(old, val)$ and store maximum value at location pointed by p . The function returns old .
long atom_and (volatileglobal long *p, long val) long atom_and (volatilelocal long *p, long val) ulong atom_and (volatileglobal ulong *p, ulong val) ulong atom_and (volatilelocal ulong *p, ulong val)	Read the 64-bit value (referred to as old) stored at location pointed by p . Compute (old & val) and store result at location pointed by p . The function returns old .
long atom_or (volatileglobal long *p, long val) long atom_or (volatilelocal long *p, long val) ulong atom_or (volatileglobal ulong *p, ulong val) ulong atom_or (volatilelocal ulong *p, ulong val) ulong atom_or (volatilelocal ulong *p, ulong val)	Read the 64-bit value (referred to as old) stored at location pointed by p . Compute ($old \mid val$) and store result at location pointed by p . The function returns old .

Function	Description
long atom_xor (volatileglobal long *p, long val) long atom_xor (volatilelocal long *p, long val)	Read the 64-bit value (referred to as old) stored at location pointed by p . Compute ($old \land val$) and store result at location pointed by p . The function returns old .
ulong atom_xor (volatileglobal ulong *p, ulong val) ulong atom_xor (volatilelocal ulong *p, ulong val)	

Note: Atomic operations on 64-bit integers and 32-bit integers (and float) are also atomic w.r.t. each other.

Chapter 9. Selecting the Rounding Mode (DEPRECATED)

This section describes the **cl_khr_select_fprounding_mode** extension. It allows an application to specify the rounding mode for an instruction or group of instructions in the program source.

This extension was deprecated in OpenCL 1.1 and its use is not recommended.

9.1. General information

9.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

9.2. Changes to OpenCL C specification

With this extension, the rounding mode may be specified using the following **#pragma** in the OpenCL program source:

```
#pragma OPENCL SELECT_ROUNDING_MODE <rounding-mode>
```

The <rounding-mode> may be one of the following values:

- rte round to nearest even
- · rtz round to zero
- rtp round to positive infinity
- rtn round to negative infinity

If this extensions is supported then the OpenCL implementation must support all four rounding modes for single precision floating-point.

The **#pragma** sets the rounding mode for all instructions that operate on floating-point types (scalar or vector types) or produce floating-point values that follow this pragma in the program source until the next **#pragma**. Note that the rounding mode specified for a block of code is known at compile time. When inside a compound statement, the pragma takes effect from its occurrence until another **#pragma** is encountered (including within a nested compound statement), or until the end of the compound statement; at the end of a compound statement the state for the pragma is restored to its condition just before the compound statement. Except where otherwise documented, the callee functions do not inherit the rounding mode of the caller function.

If this extension is enabled, the <u>__ROUNDING_MODE__</u> preprocessor symbol shall be defined to be one of the following according to the current rounding mode:

```
#define __ROUNDING_MODE__ rte
#define __ROUNDING_MODE__ rtz
#define __ROUNDING_MODE__ rtp
#define __ROUNDING_MODE__ rtz
```

This is intended to enable remapping foo() to foo_rte() by the preprocessor by using:

```
#define foo foo ## __ROUNDING_MODE__
```

The default rounding mode is round to nearest even. The built-in math functions described in *section 6.11.2*, the common functions described in *section 6.11.4* and the geometric functions described in *section 6.11.5* are implemented with the round to nearest even rounding mode. Various built-in conversions and the **vstore_half** and **vstorea_half** built-in functions that do not specify a rounding mode inherit the current rounding mode. Conversions from floating-point to integer type always use rtz mode, except where the user specifically asks for another rounding mode.

Chapter 10. Creating an OpenCL Context from an OpenGL Context or Share Group

10.1. Overview

This section describes functionality in the **cl_khr_gl_sharing** extension to associate an OpenCL context with an OpenGL context or share group object. Once an OpenCL context is associated with an OpenGL context or share group object, the functionality described in the section Creating OpenCL Memory Objects from OpenGL Objects may be used to share OpenGL buffer, texture, and renderbuffer objects with the OpenCL context.

An OpenGL implementation supporting buffer objects and sharing of texture and buffer object images with OpenCL is required by this extension.

10.2. General information

10.2.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

10.3. New Procedures and Functions

10.4. New Tokens

Returned by **clCreateContext**, **clCreateContextFromType**, and **clGetGLContextInfoKHR** when an invalid OpenGL context or share group object handle is specified in *properties*:

```
CL_INVALID_GL_SHAREGROUP_REFERENCE_KHR
```

Accepted as the *param_name* argument of **clGetGLContextInfoKHR**:

```
CL_CURRENT_DEVICE_FOR_GL_CONTEXT_KHR
CL_DEVICES_FOR_GL_CONTEXT_KHR
```

Accepted as an attribute name in the *properties* argument of **clCreateContext** and **clCreateContextFromType**:

CL_GL_CONTEXT_KHR
CL_EGL_DISPLAY_KHR
CL_GLX_DISPLAY_KHR
CL_WGL_HDC_KHR
CL_CGL_SHAREGROUP_KHR

10.5. Additions to Chapter 4 of the OpenCL 2.2 Specification

In section 4.4, replace the description of properties under clCreateContext with:

"`properties points to an attribute list, which is a array of ordered <attribute name, value> pairs terminated with zero. If an attribute is not specified in *properties*, then its default value (listed in *table 4.5*) is used (it is said to be specified implicitly). If *properties* is NULL or empty (points to a list whose first value is zero), all attributes take on their default values.

Attributes control sharing of OpenCL memory objects with OpenGL buffer, texture, and renderbuffer objects. Depending on the platform-specific API used to bind OpenGL contexts to the window system, the following attributes may be set to identify an OpenGL context:

- When the CGL binding API is supported, the attribute CL_CGL_SHAREGROUP_KHR should be set to a CGLShareGroup handle to a CGL share group object.
- When the EGL binding API is supported, the attribute CL_GL_CONTEXT_KHR should be set to an EGLContext handle to an OpenGL ES or OpenGL context, and the attribute CL_EGL_DISPLAY_KHR should be set to the EGLDisplay handle of the display used to create the OpenGL ES or OpenGL context.
- When the GLX binding API is supported, the attribute CL_GL_CONTEXT_KHR should be set to a GLXContext handle to an OpenGL context, and the attribute CL_GLX_DISPLAY_KHR should be set to the Display handle of the X Window System display used to create the OpenGL context.
- When the WGL binding API is supported, the attribute CL_GL_CONTEXT_KHR should be set to an HGLRC handle to an OpenGL context, and the attribute CL_WGL_HDC_KHR should be set to the HDC handle of the display used to create the OpenGL context.

Memory objects created in the context so specified may be shared with the specified OpenGL or OpenGL ES context (as well as with any other OpenGL contexts on the share list of that context, according to the description of sharing in the GLX 1.4 and EGL 1.4 specifications, and the WGL documentation for OpenGL implementations on Microsoft Windows), or with the explicitly identified OpenGL share group for CGL. If no OpenGL or OpenGL ES context or share group is specified in the attribute list, then memory objects may not be shared, and calling any of the commands described in Creating OpenCL Memory Objects from OpenGL Objects will result in a CL_INVALID_GL_SHAREGROUP_REFERENCE_KHR error.`"

OpenCL / OpenGL sharing does not support the CL_CONTEXT_INTEROP_USER_SYNC property

defined in *table 4.5*. Specifying this property when creating a context with OpenCL / OpenGL sharing will return an appropriate error.

Add to *table 4.5*:

Table 25. OpenGL Sharing Context Creation Attributes

Attribute Name	Allowed Values (Default value is in bold)	Description
CL_GL_CONTEXT_KHR	0 , OpenGL context handle	OpenGL context to associated the OpenCL context with
CL_CGL_SHAREGROUP_KHR	0 , CGL share group handle	CGL share group to associate the OpenCL context with
CL_EGL_DISPLAY_KHR	EGL_NO_DISPLAY , EGLDisplay handle	EGLDisplay an OpenGL context was created with respect to
CL_GLX_DISPLAY_KHR	None, X handle	X Display an OpenGL context was created with respect to
CL_WGL_HDC_KHR	0 , HDC handle	HDC an OpenGL context was created with respect to

Replace the first error in the list for **clCreateContext** with:

- A context was specified for an EGL-based OpenGL ES or OpenGL implementation by setting the attributes CL_GL_CONTEXT_KHR and CL_EGL_DISPLAY_KHR.
- A context was specified for a GLX-based OpenGL implementation by setting the attributes CL_GL_CONTEXT_KHR and CL_GLX_DISPLAY_KHR.
- A context was specified for a WGL-based OpenGL implementation by setting the attributes CL_GL_CONTEXT_KHR and CL_WGL_HDC_KHR

and any of the following conditions hold:

- The specified display and context attributes do not identify a valid OpenGL or OpenGL ES context.
- The specified context does not support buffer and renderbuffer objects.
- The specified context is not compatible with the OpenCL context being created (for example, it exists in a physically distinct address space, such as another hardware device; or it does not support sharing data with OpenCL due to implementation restrictions).

errcode_ret returns CL_INVALID_GL_SHAREGROUP_REFERENCE_KHR if a share group was specified for a CGL-based OpenGL implementation by setting the attribute CL_CGL_SHAREGROUP_KHR, and the specified share group does not identify a valid CGL share group object.

errcode_ret returns CL_INVALID_OPERATION if a context was specified as described above and any
of the following conditions hold:

[&]quot;`errcode_ret returns CL_INVALID_GL_SHAREGROUP_REFERENCE_KHR if a context was specified by any of the following means:

- A context or share group object was specified for one of CGL, EGL, GLX, or WGL and the OpenGL implementation does not support that window-system binding API.
- More than one of the attributes CL_CGL_SHAREGROUP_KHR, CL_EGL_DISPLAY_KHR, CL_GLX_DISPLAY_KHR, and CL_WGL_HDC_KHR is set to a non-default value.
- Both of the attributes CL_CGL_SHAREGROUP_KHR and CL_GL_CONTEXT_KHR are set to nondefault values.
- Any of the devices specified in the *devices* argument cannot support OpenCL objects which share the data store of an OpenGL object.

errcode_ret returns CL_INVALID_PROPERTY if an attribute name other than those specified in table 4.5 or if CL_CONTEXT_INTEROP_USER_SYNC is specified in properties.`"

Replace the description of *properties* under **clCreateContextFromType** with:

"_properties_ points to an attribute list whose format and valid contents are identical to the **properties** argument of **clCreateContext**."

Replace the first error in the list for **clCreateContextFromType** with the same two new errors described above for **clCreateContext**.

10.6. Additions to Chapter 5 of the OpenCL 2.2 Specification

Add a new section to describe the new API for querying OpenCL devices that support sharing with OpenGL:

"`OpenCL device(s) corresponding to an OpenGL context may be queried. Such a device may not always exist (for example, if an OpenGL context is specified on a GPU not supporting OpenCL command queues, but which does support shared CL/GL objects), and if it does exist, may change over time. When such a device does exist, acquiring and releasing shared CL/GL objects may be faster on a command queue corresponding to this device than on command queues corresponding to other devices available to an OpenCL context.

To query the currently corresponding device, use the function

properties points to an attribute list whose format and valid contents are identical to the properties argument of clCreateContext. properties must identify a single valid GL context or GL share group object.

param_name is a constant that specifies the device types to query, and must be one of the values shown in the table below.

param_value is a pointer to memory where the result of the query is returned as described in the table below. If *param_value* is NULL, it is ignored.

param_value_size specifies the size in bytes of memory pointed to by param_value. This size must be greater than or equal to the size of the return type described in the table below.

param_value_size_ret returns the actual size in bytes of data being queried by param_value. If param_value_size_ret is NULL, it is ignored.

Table 26. Supported Device Types for clGetGLContextInfoKHR

param_name	Return Type	Information returned in param_value
CL_CURRENT_DEVICE_FOR_GL_CONTE XT_KHR	cl_device_id	Return the OpenCL device currently associated with the specified OpenGL context.
CL_DEVICES_FOR_GL_CONTEXT_KHR	cl_device_id[]	Return all OpenCL devices which may be associated with the specified OpenGL context.

clGetGLContextInfoKHR returns CL_SUCCESS if the function is executed successfully. If no device(s) exist corresponding to *param_name*, the call will not fail, but the value of *param value size ret* will be zero.

clGetGLContextInfoKHR returns CL_INVALID_GL_SHAREGROUP_REFERENCE_KHR if a context was specified by any of the following means:

- A context was specified for an EGL-based OpenGL ES or OpenGL implementation by setting the attributes CL_GL_CONTEXT_KHR and CL_EGL_DISPLAY_KHR.
- A context was specified for a GLX-based OpenGL implementation by setting the attributes CL_GL_CONTEXT_KHR and CL_GLX_DISPLAY_KHR.
- A context was specified for a WGL-based OpenGL implementation by setting the attributes CL_GL_CONTEXT_KHR and CL_WGL_HDC_KHR.

and any of the following conditions hold:

- The specified display and context attributes do not identify a valid OpenGL or OpenGL ES context.
- The specified context does not support buffer and renderbuffer objects.
- The specified context is not compatible with the OpenCL context being created (for example, it exists in a physically distinct address space, such as another hardware device; or it does not support sharing data with OpenCL due to implementation restrictions).

clGetGLContextInfoKHR returns CL_INVALID_GL_SHAREGROUP_REFERENCE_KHR if a share group was specified for a CGL-based OpenGL implementation by setting the attribute CL_CGL_SHAREGROUP_KHR, and the specified share group does not identify a valid CGL share group object.

clGetGLContextInfoKHR returns CL_INVALID_OPERATION if a context was specified as described above and any of the following conditions hold:

- A context or share group object was specified for one of CGL, EGL, GLX, or WGL and the OpenGL implementation does not support that window-system binding API.
- More than one of the attributes CL_CGL_SHAREGROUP_KHR, CL_EGL_DISPLAY_KHR, CL_GLX_DISPLAY_KHR, and CL_WGL_HDC_KHR is set to a non-default value.
- Both of the attributes CL_CGL_SHAREGROUP_KHR and CL_GL_CONTEXT_KHR are set to non-default values.
- Any of the devices specified in the <devices> argument cannot support OpenCL objects which share the data store of an OpenGL object.

clGetGLContextInfoKHR returns CL_INVALID_VALUE if an attribute name other than those specified in *table 4.5* is specified in *properties*.

Additionally, **clGetGLContextInfoKHR** returns CL_INVALID_VALUE if *param_name* is not one of the values listed in the table *GL context information that can be queried with clGetGLContextInfoKHR, or if the size in bytes specified by <i>param_value_size* is less than the size of the return type shown in the table and *param_value* is not a NULL value; CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device; or CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.`"

10.7. Issues

1. How should the OpenGL context be identified when creating an associated OpenCL context?

RESOLVED: by using a (display,context handle) attribute pair to identify an arbitrary OpenGL or OpenGL ES context with respect to one of the window-system binding layers EGL, GLX, or WGL, or a share group handle to identify a CGL share group. If a context is specified, it need not be current to the thread calling clCreateContext*.

A previously suggested approach would use a single boolean attribute CL_USE_GL_CONTEXT_KHR to allow creating a context associated with the currently bound OpenGL context. This may still be implemented as a separate extension, and might allow more efficient acquire/release behavior in the special case where they are being executed in the same thread as the bound GL context used to create the CL context.

2. What should the format of an attribute list be?

After considerable discussion, we think we can live with a list of <attribute name,value> pairs terminated by zero. The list is passed as 'cl_context_properties *properties', where cl_context_properties is typedefed to be 'intptr_t' in cl.h.

This effectively allows encoding all scalar integer, pointer, and handle values in the host API into the argument list and is analogous to the structure and type of EGL attribute lists. NULL attribute lists are also allowed. Again as for EGL, any attributes not explicitly passed in the list will take on a defined default value that does something reasonable.

Experience with EGL, GLX, and WGL has shown attribute lists to be a sufficiently flexible and general mechanism to serve the needs of management calls such as context creation. It is not completely general (encoding floating-point and non-scalar attribute values is not straightforward), and other approaches were suggested such as opaque attribute lists with getter/setter methods, or arrays of varadic structures.

3. What's the behavior of an associated OpenGL or OpenCL context when using resources defined by the other associated context, and that context is destroyed?

RESOLVED: OpenCL objects place a reference on the data store underlying the corresponding GL object when they're created. The GL name corresponding to that data store may be deleted, but the data store itself remains so long as any CL object has a reference to it. However, destroying all GL contexts in the share group corresponding to a CL context results in implementation-dependent behavior when using a corresponding CL object, up to and including program termination.

4. How about sharing with D3D?

Sharing between D3D and OpenCL should use the same attribute list mechanism, though obviously with different parameters, and be exposed as a similar parallel OpenCL extension. There may be an interaction between that extension and this one since it's not yet clear if it will be possible to create a CL context simultaneously sharing GL and D3D objects.

5. Under what conditions will context creation fail due to sharing?

RESOLVED: Several cross-platform failure conditions are described (GL context or CGL share group doesn't exist, GL context doesn't support types of GL objects, GL context implementation doesn't allow sharing), but additional failures may result due to implementation-dependent reasons and should be added to this extension as such failures are discovered. Sharing between OpenCL and OpenGL requires integration at the driver internals level.

6. What command queues can **clEnqueueAcquire/ReleaseGLObjects** be placed on?

RESOLVED: All command queues. This restriction is enforced at context creation time. If any device passed to context creation cannot support shared CL/GL objects, context creation will fail with a CL_INVALID_OPERATION error.

7. How can applications determine which command queue to place an Acquire/Release on?

RESOLVED: The **clGetGLContextInfoKHR** returns either the CL device currently corresponding to a specified GL context (typically the display it's running on), or a list of all the CL devices the specified context might run on (potentially useful in multiheaded / "virtual screen" environments). This command is not simply placed in Creating OpenCL Memory Objects from OpenGL Objects because it relies on the same property-list method of specifying a GL context introduced by this extension.

If no devices are returned, it means that the GL context exists on an older GPU not capable of running OpenCL, but still capable of sharing objects between GL running on that GPU and CL running elsewhere.

8. What is the meaning of the CL_DEVICES_FOR_GL_CONTEXT_KHR query?

RESOLVED: The list of all CL devices that may ever be associated with a specific GL context. On platforms such as MacOS X, the "virtual screen" concept allows multiple GPUs to back a single virtual display. Similar functionality might be implemented on other windowing systems, such as a transparent heterogenous multiheaded X server. Therefore the exact meaning of this query is interpreted relative to the binding layer API in use.

Chapter 11. Creating OpenCL Memory Objects from OpenGL Objects

This section describes functionality in the **cl_khr_gl_sharing** extension to use OpenGL buffer, texture, and renderbuffer objects as OpenCL memory objects. OpenCL memory objects may be created from OpenGL objects if and only if the OpenCL context is associated with an OpenGL context or share group object. The section Creating an OpenCL Context from an OpenGL Context or Share Group describes how to create an OpenCL context associated with an OpenGL context or share group object.

An OpenCL image object may be created from an OpenGL texture or renderbuffer object. An OpenCL buffer object may be created from an OpenGL buffer object.

Any supported OpenGL object defined within the associated OpenGL context or share group object may be shared, with the exception of the default OpenGL objects (i.e. objects named zero), which may not be shared.

11.1. General information

11.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

11.2. Lifetime of Shared Objects

An OpenCL memory object created from an OpenGL object (hereinafter referred to as a "shared CL/GL object") remains valid as long as the corresponding GL object has not been deleted. If the GL object is deleted through the GL API (e.g. **glDeleteBuffers**, **glDeleteTextures**, or **glDeleteRenderbuffers**), subsequent use of the CL buffer or image object will result in undefined behavior, including but not limited to possible CL errors and data corruption, but may not result in program termination.

The CL context and corresponding command-queues are dependent on the existence of the GL share group object, or the share group associated with the GL context from which the CL context is created. If the GL share group object or all GL contexts in the share group are destroyed, any use of the CL context or command-queue(s) will result in undefined behavior, which may include program termination. Applications should destroy the CL command-queue(s) and CL context before destroying the corresponding GL share group or contexts

11.3. OpenCL Buffer Objects from OpenGL Buffer Objects

The function

creates an OpenCL buffer object from an OpenGL buffer object.

context is a valid OpenCL context created from an OpenGL context.

flags is a bit-field that is used to specify usage information. Refer to table 5.3 for a description of flags. Only CL_MEM_READ_ONLY, CL_MEM_WRITE_ONLY and CL_MEM_READ_WRITE values specified in table 5.3 can be used.

bufobj is the name of a GL buffer object. The data store of the GL buffer object must have have been previously created by calling **glBufferData**, although its contents need not be initialized. The size of the data store will be used to determine the size of the CL buffer object.

errcode_ret will return an appropriate error code as described below. If errcode_ret is NULL, no error code is returned.

clCreateFromGLBuffer returns a valid non-zero OpenCL buffer object and *errcode_ret* is set to CL_SUCCESS if the buffer object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode_ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context or was not created from a GL context.
- CL_INVALID_VALUE if values specified in *flags* are not valid.
- CL_INVALID_GL_OBJECT if *bufobj* is not a GL buffer object or is a GL buffer object but does not have an existing data store or the size of the buffer is 0.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The size of the GL buffer object data store at the time **clCreateFromGLBuffer** is called will be used as the size of buffer object returned by **clCreateFromGLBuffer**. If the state of a GL buffer object is modified through the GL API (e.g. **glBufferData**) while there exists a corresponding CL buffer object, subsequent use of the CL buffer object will result in undefined behavior.

The **clRetainMemObject** and **clReleaseMemObject** functions can be used to retain and release the buffer object.

The CL buffer object created using clCreateFromGLBuffer can also be used to create a CL 1D image buffer object.

11.4. OpenCL Image Objects from OpenGL Textures

The function

creates the following:

- an OpenCL 2D image object from an OpenGL 2D texture object or a single face of an OpenGL cubemap texture object,
- an OpenCL 2D image array object from an OpenGL 2D texture array object,
- an OpenCL 1D image object from an OpenGL 1D texture object,
- an OpenCL 1D image buffer object from an OpenGL texture buffer object,
- an OpenCL 1D image array object from an OpenGL 1D texture array object,
- an OpenCL 3D image object from an OpenGL 3D texture object.

context is a valid OpenCL context created from an OpenGL context.

flags is a bit-field that is used to specify usage information. Refer to *table 5.3* for a description of flags. Only CL_MEM_READ_ONLY, CL_MEM_WRITE_ONLY and CL_MEM_READ_WRITE values specified in *table 5.3* may be used.

texture_target must be one of GL_TEXTURE_1D, GL_TEXTURE_1D_ARRAY, GL_TEXTURE_BUFFER, GL_TEXTURE_2D, GL_TEXTURE_2D_ARRAY, GL TEXTURE 3D, GL_TEXTURE_CUBE_MAP_POSITIVE_X, GL TEXTURE CUBE MAP POSITIVE Y, GL_TEXTURE_CUBE_MAP_POSITIVE_Z, GL_TEXTURE_CUBE_MAP_NEGATIVE_X, GL_TEXTURE_CUBE_MAP_NEGATIVE_Y, GL_TEXTURE_CUBE_MAP_NEGATIVE_Z, GL_TEXTURE_RECTANGLE (Note: GL_TEXTURE_RECTANGLE requires OpenGL 3.1. Alternatively, GL_TEXTURE_RECTANGLE_ARB be specified if the OpenGL may extension **GL_ARB_texture_rectangle** is supported.). *texture_target* is used only to define the image type of texture. No reference to a bound GL texture object is made or implied by this parameter.

miplevel is the mipmap level to be used. If *texture_target* is GL_TEXTURE_BUFFER, *miplevel* must be 0. Note: Implementations may return CL_INVALID_OPERATION for miplevel values > 0.

texture is the name of a GL 1D, 2D, 3D, 1D array, 2D array, cubemap, rectangle or buffer texture object. The texture object must be a complete texture as per OpenGL rules on texture completeness. The *texture* format and dimensions defined by OpenGL for the specified *miplevel* of the texture will be used to create the OpenCL image memory object. Only GL texture objects with an internal format that maps to appropriate image channel order and data type specified in *tables 5.5* and *5.6* may be used to create the OpenCL image memory object.

errcode_ret will return an appropriate error code as described below. If *errcode_ret* is NULL, no error code is returned.

clCreateFromGLTexture returns a valid non-zero OpenCL image object and errcode_ret is set to

CL_SUCCESS if the image object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode_ret*:

- CL INVALID CONTEXT if context is not a valid context or was not created from a GL context.
- CL_INVALID_VALUE if values specified in *flags* are not valid or if value specified in *texture_target* is not one of the values specified in the description of *texture_target*.
- CL_INVALID_MIP_LEVEL if *miplevel* is less than the value of $level_{base}$ (for OpenGL implementations) or zero (for OpenGL ES implementations); or greater than the value of q (for both OpenGL and OpenGL ES). $level_{base}$ and q are defined for the texture in section 3.8.10 (Texture Completeness) of the OpenGL 2.1 specification and section 3.7.10 of the OpenGL ES 2.0.
- CL_INVALID_MIP_LEVEL if *miplevel* is greather than zero and the OpenGL implementation does not support creating from non-zero mipmap levels.
- CL_INVALID_GL_OBJECT if *texture* is not a GL texture object whose type matches *texture_target*, if the specified *miplevel* of *texture* is not defined, or if the width or height of the specified *miplevel* is zero or if the GL texture object is incomplete.
- CL_INVALID_IMAGE_FORMAT_DESCRIPTOR if the OpenGL texture internal format does not map to a supported OpenCL image format.
- CL_INVALID_OPERATION if *texture* is a GL texture object created with a border width value greater than zero.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

If the state of a GL texture object is modified through the GL API (e.g. glTexImage2D, glTexImage3D or the values of the texture parameters GL_TEXTURE_BASE_LEVEL or GL_TEXTURE_MAX_LEVEL are modified) while there exists a corresponding CL image object, subsequent use of the CL image object will result in undefined behavior.

The **clRetainMemObject** and **clReleaseMemObject** functions can be used to retain and release the image objects.

11.4.1. List of OpenGL and corresponding OpenCL Image Formats

The table below describes the list of OpenGL texture internal formats and the corresponding OpenCL image formats. If a OpenGL texture object with an internal format from the table below is successfully created by OpenGL, then there is guaranteed to be a mapping to one of the corresponding OpenCL image format(s) in that table. Texture objects created with other OpenGL internal formats may (but are not guaranteed to) have a mapping to an OpenCL image format; if such mappings exist, they are guaranteed to preserve all color components, data types, and at least the number of bits/component actually allocated by OpenGL for that format.

Table 27. OpenGL internal formats and corresponding OpenCL internal formats

GL internal format	CL image format (channel order, channel data type)
GL_RGBA8	CL_RGBA, CL_UNORM_INT8 or
	CL_BGRA, CL_UNORM_INT8
GL_SRGB8_ALPHA8	CL_sRGBA, CL_UNORM_INT8
GL_RGBA, GL_UNSIGNED_INT_8_8_8_8_REV	CL_RGBA, CL_UNORM_INT8
GL_BGRA, GL_UNSIGNED_INT_8_8_8_8_REV	CL_BGRA, CL_UNORM_INT8
GL_RGBA8I, GL_RGBA8I_EXT	CL_RGBA, CL_SIGNED_INT8
GL_RGBA16I, GL_RGBA16I_EXT	CL_RGBA, CL_SIGNED_INT16
GL_RGBA32I, GL_RGBA32I_EXT	CL_RGBA, CL_SIGNED_INT32
GL_RGBA8UI, GL_RGBA8UI_EXT	CL_RGBA, CL_UNSIGNED_INT8
GL_RGBA16UI, GL_RGBA16UI_EXT	CL_RGBA, CL_UNSIGNED_INT16
GL_RGBA32UI, GL_RGBA32UI_EXT	CL_RGBA, CL_UNSIGNED_INT32
GL_RGBA8_SNORM	CL_RGBA, CL_SNORM_INT8
GL_RGBA16	CL_RGBA, CL_UNORM_INT16
GL_RGBA16_SNORM	CL_RGBA, CL_SNORM_INT16
GL_RGBA16F, GL_RGBA16F_ARB	CL_RGBA, CL_HALF_FLOAT
GL_RGBA32F, GL_RGBA32F_ARB	CL_RGBA, CL_FLOAT
GL_R8	CL_R, CL_UNORM_INT8
GL_R8_SNORM	CL_R, CL_SNORM_INT8
GL_R16	CL_R, CL_UNORM_INT16
GL_R16_SNORM	CL_R, CL_SNORM_INT16
GL_R16F	CL_R, CL_HALF_FLOAT
GL_R32F	CL_R, CL_FLOAT
GL_R8I	CL_R, CL_SIGNED_INT8
GL_R16I	CL_R, CL_SIGNED_INT16
GL_R32I	CL_R, CL_SIGNED_INT32
GL_R8UI	CL_R, CL_UNSIGNED_INT8
GL_R16UI	CL_R, CL_UNSIGNED_INT16
GL_R32UI	CL_R, CL_UNSIGNED_INT32
GL_RG8	CL_RG, CL_UNORM_INT8
GL_RG8_SNORM	CL_RG, CL_SNORM_INT8

GL internal format	CL image format (channel order, channel data type)
GL_RG16	CL_RG, CL_UNORM_INT16
GL_RG16_SNORM	CL_RG, CL_SNORM_INT16
GL_RG16F	CL_RG, CL_HALF_FLOAT
GL_RG32F	CL_RG, CL_FLOAT
GL_RG8I	CL_RG, CL_SIGNED_INT8
GL_RG16I	CL_RG, CL_SIGNED_INT16
GL_RG32I	CL_RG, CL_SIGNED_INT32
GL_RG8UI	CL_RG, CL_UNSIGNED_INT8
GL_RG16UI	CL_RG, CL_UNSIGNED_INT16
GL_RG32UI	CL_RG, CL_UNSIGNED_INT32

11.5. OpenCL Image Objects from OpenGL Renderbuffers

The function

creates an OpenCL 2D image object from an OpenGL renderbuffer object.

context is a valid OpenCL context created from an OpenGL context.

flags is a bit-field that is used to specify usage information. Refer to table 5.3 for a description of flags. Only CL_MEM_READ_ONLY, CL_MEM_WRITE_ONLY and CL_MEM_READ_WRITE values specified in table 5.3 can be used.

renderbuffer is the name of a GL renderbuffer object. The renderbuffer storage must be specified before the image object can be created. The renderbuffer format and dimensions defined by OpenGL will be used to create the 2D image object. Only GL renderbuffers with internal formats that maps to appropriate image channel order and data type specified in tables 5.5 and 5.6 can be used to create the 2D image object.

errcode_ret will return an appropriate error code as described below. If *errcode_ret* is NULL, no error code is returned.

clCreateFromGLRenderbuffer returns a valid non-zero OpenCL image object and *errcode_ret* is set to CL_SUCCESS if the image object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode_ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context or was not created from a GL context.
- CL_INVALID_VALUE if values specified in *flags* are not valid.
- CL_INVALID_GL_OBJECT if *renderbuffer* is not a GL renderbuffer object or if the width or height of *renderbuffer* is zero.
- CL_INVALID_IMAGE_FORMAT_DESCRIPTOR if the OpenGL renderbuffer internal format does not map to a supported OpenCL image format.
- CL_INVALID_OPERATION if renderbuffer is a multi-sample GL renderbuffer object.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

If the state of a GL renderbuffer object is modified through the GL API (i.e. changes to the dimensions or format used to represent pixels of the GL renderbuffer using appropriate GL API calls such as **glRenderbufferStorage**) while there exists a corresponding CL image object, subsequent use of the CL image object will result in undefined behavior.

The **clRetainMemObject** and **clReleaseMemObject** functions can be used to retain and release the image objects.

The table *OpenGL internal formats and corresponding OpenCL internal formats* describes the list of OpenGL renderbuffer internal formats and the corresponding OpenCL image formats. If an OpenGL renderbuffer object with an internal format from the table is successfully created by OpenGL, then there is guaranteed to be a mapping to one of the corresponding OpenCL image format(s) in that table. Renderbuffer objects created with other OpenGL internal formats may (but are not guaranteed to) have a mapping to an OpenCL image format; if such mappings exist, they are guaranteed to preserve all color components, data types, and at least the number of bits/component actually allocated by OpenGL for that format.

11.6. Querying OpenGL object information from an OpenCL memory object

The OpenGL object used to create the OpenCL memory object and information about the object type i.e. whether it is a texture, renderbuffer or buffer object can be queried using the following function.

gl_object_type returns the type of GL object attached to memobj and can be CL_GL_OBJECT_BUFFER, CL_GL_OBJECT_TEXTURE2D, CL_GL_OBJECT_TEXTURE3D, CL_GL_OBJECT_TEXTURE2D_ARRAY, CL_GL_OBJECT_TEXTURE1D, CL_GL_OBJECT_TEXTURE1D_ARRAY, CL_GL_OBJECT_TEXTURE_BUFFER, or CL_GL_OBJECT_RENDERBUFFER. If gl_object_type is NULL, it is

ignored

gl_object_name returns the GL object name used to create memobj. If gl_object_name is NULL, it is ignored.

clGetGLObjectInfo returns CL_SUCCESS if the call was executed successfully. Otherwise, it returns one of the following errors:

- CL_INVALID_MEM_OBJECT if *memobj* is not a valid OpenCL memory object.
- CL_INVALID_GL_OBJECT if there is no GL object associated with *memobj*.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

returns additional information about the GL texture object associated with memobj.

param_name specifies what additional information about the GL texture object associated with memobj to query. The list of supported param_name types and the information returned in param_value by clGetGLTextureInfo is described in the table below.

param_value is a pointer to memory where the result being queried is returned. If *param_value* is NULL, it is ignored.

param_value_size is used to specify the size in bytes of memory pointed to by *param_value*. This size must be >= size of return type as described in the table below.

param_value_size_ret returns the actual size in bytes of data copied to param_value. If param_value_size_ret is NULL, it is ignored.

Table 28. OpenGL texture info that may be queried with clGetGLTextureInfo

cl_gl_texture_info	Return Type	Info. returned in param_value
CL_GL_TEXTURE_TARGET	GLenum	The texture_target argument specified in clCreateFromGLTexture.
CL_GL_MIPMAP_LEVEL	GLint	The <i>miplevel</i> argument specified in clCreateFromGLTexture .

clGetGLTextureInfo returns CL_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- CL_INVALID_MEM_OBJECT if *memobj* is not a valid OpenCL memory object.
- CL_INVALID_GL_OBJECT if there is no GL texture object associated with *memobj*.
- CL_INVALID_VALUE if *param_name* is not valid, or if size in bytes specified by *param_value_size* is less than the size of the return type as described in the table above and *param_value* is not NULL, or if *param_value* and *param_value_size_ret* are NULL.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

11.7. Sharing memory objects that map to GL objects between GL and CL contexts

The function

is used to acquire OpenCL memory objects that have been created from OpenGL objects. These objects need to be acquired before they can be used by any OpenCL commands queued to a command-queue. The OpenGL objects are acquired by the OpenCL context associated with command_queue and can therefore be used by all command-queues associated with the OpenCL context.

command_queue is a valid command-queue. All devices used to create the OpenCL context associated with command_queue must support acquiring shared CL/GL objects. This constraint is enforced at context creation time.

num_objects is the number of memory objects to be acquired in *mem_objects*.

mem_objects is a pointer to a list of CL memory objects that correspond to GL objects.

event_wait_list and num_events_in_wait_list specify events that need to complete before this particular command can be executed. If event_wait_list is NULL, then this particular command does not wait on any event to complete. If event_wait_list is NULL, num_events_in_wait_list must be 0. If event_wait_list is not NULL, the list of events pointed to by event_wait_list must be valid and num events in wait list must be greater than 0. The events specified in

event_wait_list act as synchronization points.

event returns an event object that identifies this command and can be used to query wait for this command to complete. If event is NULL or the enqueue is unsuccessful, no event will be created and therefore it will not be possible to query the status of this command or to wait for this command to complete. If event_wait_list and event are not NULL, event must not refer to an element of the event_wait_list array.

clEnqueueAcquireGLObjects returns CL_SUCCESS if the function is executed successfully. If *num_objects* is 0 and *mem_objects* is NULL the function does nothing and returns CL_SUCCESS. Otherwise, it returns one of the following errors:

- CL_INVALID_VALUE if num_objects is zero and mem_objects is not a NULL value or if num_objects
 o and mem_objects is NULL.
- CL_INVALID_MEM_OBJECT if memory objects in *mem_objects* are not valid OpenCL memory objects.
- CL_INVALID_COMMAND_QUEUE if command_queue is not a valid command-queue.
- CL_INVALID_CONTEXT if context associated with *command_queue* was not created from an OpenGL context
- CL_INVALID_GL_OBJECT if memory objects in *mem_objects* have not been created from a GL object(s).
- CL_INVALID_EVENT_WAIT_LIST if event_wait_list is NULL and num_events_in_wait_list > 0, or event_wait_list is not NULL and num_events_in_wait_list is 0, or if event objects in event_wait_list are not valid events.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

is used to release OpenCL memory objects that have been created from OpenGL objects. These objects need to be released before they can be used by OpenGL. The OpenGL objects are released by the OpenCL context associated with *command_queue*.

num_objects is the number of memory objects to be released in *mem_objects*.

mem objects is a pointer to a list of CL memory objects that correspond to GL objects.

event_wait_list and num_events_in_wait_list specify events that need to complete before this
command can be executed. If event_wait_list is NULL, then this particular command does not wait on

any event to complete. If <code>event_wait_list</code> is <code>NULL</code>, <code>num_events_in_wait_list</code> must be 0. If <code>event_wait_list</code> is not <code>NULL</code>, the list of events pointed to by <code>event_wait_list</code> must be valid and <code>num_events_in_wait_list</code> must be greater than 0. The events specified in <code>event_wait_list</code> act as synchronization points.

event returns an event object that identifies this command and can be used to query or wait for this command to complete. If event is NULL or the enqueue is unsuccessful, no event will be created and therefore it will not be possible to query the status of this command or to wait for this command to complete. If event_wait_list and event are not NULL, event must not refer to an element of the event_wait_list array.

clEnqueueReleaseGLObjects returns CL_SUCCESS if the function is executed successfully. If *num_objects* is 0 and *mem_objects* is NULL the function does nothing and returns CL_SUCCESS. Otherwise, it returns one of the following errors:

- CL_INVALID_VALUE if num_objects is zero and mem_objects is not a NULL value or if num_objects
 o and mem_objects is NULL.
- CL_INVALID_MEM_OBJECT if memory objects in mem_objects are not valid OpenCL memory objects.
- CL_INVALID_COMMAND_QUEUE if command_queue is not a valid command-queue.
- CL_INVALID_CONTEXT if context associated with *command_queue* was not created from an OpenGL context
- CL_INVALID_GL_OBJECT if memory objects in *mem_objects* have not been created from a GL object(s).
- CL_INVALID_EVENT_WAIT_LIST if event_wait_list is NULL and num_events_in_wait_list > 0, or event_wait_list is not NULL and num_events_in_wait_list is 0, or if event objects in event_wait_list are not valid events.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

11.7.1. Synchronizing OpenCL and OpenGL Access to Shared Objects

In order to ensure data integrity, the application is responsible for synchronizing access to shared CL/GL objects by their respective APIs. Failure to provide such synchronization may result in race conditions and other undefined behavior including non-portability between implementations.

Prior to calling **clEnqueueAcquireGLObjects**, the application must ensure that any pending GL operations which access the objects specified in *mem_objects* have completed. This may be accomplished portably by issuing and waiting for completion of a **glFinish** command on all GL contexts with pending references to these objects. Implementations may offer more efficient synchronization methods; for example on some platforms calling **glFlush** may be sufficient, or synchronization may be implicit within a thread, or there may be vendor-specific extensions that enable placing a fence in the GL command stream and waiting for completion of that fence in the CL command queue. Note that no synchronization methods other than **glFinish** are portable between OpenGL implementations at this time.

Similarly, after calling **clEnqueueReleaseGLObjects**, the application is responsible for ensuring that any pending OpenCL operations which access the objects specified in *mem_objects* have completed prior to executing subsequent GL commands which reference these objects. This may be accomplished portably by calling **clWaitForEvents** with the event object returned by **clEnqueueReleaseGLObjects**, or by calling **clFinish**. As above, some implementations may offer more efficient methods.

The application is responsible for maintaining the proper order of operations if the CL and GL contexts are in separate threads.

If a GL context is bound to a thread other than the one in which **clEnqueueReleaseGLObjects** is called, changes to any of the objects in *mem_objects* may not be visible to that context without additional steps being taken by the application. For an OpenGL 3.1 (or later) context, the requirements are described in Appendix D ("Shared Objects and Multiple Contexts") of the OpenGL 3.1 Specification. For prior versions of OpenGL, the requirements are implementation-dependent.

Attempting to access the data store of an OpenGL object after it has been acquired by OpenCL and before it has been released will result in undefined behavior. Similarly, attempting to access a shared CL/GL object from OpenCL before it has been acquired by the OpenCL command queue, or after it has been released, will result in undefined behavior.

11.7.2. Event Command Types for Sharing memory objects that map to GL objects

The following table describes the event command types for the OpenCL commands to acquire and release OpenCL memory objects that have been created from OpenGL objects:

Table 29. List of supported event command types

Events Created By	Event Command Type
clEnqueueAcquireGLObjects	CL_COMMAND_ACQUIRE_GL_OBJECTS
clEnqueueReleaseGLObjects	CL_COMMAND_RELEASE_GL_OBJECTS

Chapter 12. Creating OpenCL Event Objects from OpenGL Sync Objects

12.1. Overview

This section describes the **cl_khr_gl_event** extension. This extension allows creating OpenCL event objects linked to OpenGL fence sync objects, potentially improving efficiency of sharing images and buffers between the two APIs. The companion **GL_ARB_cl_event** extension provides the complementary functionality of creating an OpenGL sync object from an OpenCL event object.

In addition, this extension modifies the behavior of **clEnqueueAcquireGLObjects** and **clEnqueueReleaseGLObjects** to implicitly guarantee synchronization with an OpenGL context bound in the same thread as the OpenCL context.

12.2. General information

12.2.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

12.3. New Procedures and Functions

12.4. New Tokens

Returned by **clGetEventInfo** when *param_name* is CL_EVENT_COMMAND_TYPE:

```
CL_COMMAND_GL_FENCE_SYNC_OBJECT_KHR
```

12.5. Additions to Chapter 5 of the OpenCL 2.2 Specification

Add following to the fourth paragraph of *section 5.11* (prior to the description of **clWaitForEvents**):

"Event objects can also be used to reflect the status of an OpenGL sync object. The sync object in turn refers to a fence command executing in an OpenGL command stream. This provides another method of coordinating sharing of buffers and images between OpenGL and OpenCL." Add CL_COMMAND_GL_FENCE_SYNC_OBJECT_KHR to the valid *param_value* values returned by **clGetEventInfo** for *param_name* CL_EVENT_COMMAND_TYPE (in the third row and third column of *table 5.22*).

Add new subsection 5.11.1:

"`5.11.1 Linking Event Objects to OpenGL Synchronization Objects

An event object may be created by linking to an OpenGL **sync object**. Completion of such an event object is equivalent to waiting for completion of the fence command associated with the linked GL sync object.

The function

creates a linked event object.

context is a valid OpenCL context created from an OpenGL context or share group, using the cl_khr_gl_sharing extension.

sync is the name of a sync object in the GL share group associated with *context*.

clCreateEventFromGLsyncKHR returns a valid OpenCL event object and *errcode_ret* is set to CL_SUCCESS if the event object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode_ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context, or was not created from a GL context.
- CL_INVALID_GL_OBJECT if *sync* is not the name of a sync object in the GL share group associated with *context*.

The parameters of an event object linked to a GL sync object will return the following values when queried with **clGetEventInfo**:

- The CL_EVENT_COMMAND_QUEUE of a linked event is NULL, because the event is not associated with any OpenCL command queue.
- The CL_EVENT_COMMAND_TYPE of a linked event is CL_COMMAND_GL_FENCE_SYNC_OBJECT_KHR, indicating that the event is associated with a GL sync object, rather than an OpenCL command.
- The CL_EVENT_COMMAND_EXECUTION_STATUS of a linked event is either CL_SUBMITTED, indicating that the fence command associated with the sync object has not yet completed, or CL_COMPLETE, indicating that the fence command has completed.

clCreateEventFromGLsyncKHR performs an implicit **clRetainEvent** on the returned event object. Creating a linked event object also places a reference on the linked GL sync object. When the event object is deleted, the reference will be removed from the GL sync object.

Events returned from **clCreateEventFromGLsyncKHR** can be used in the *event_wait_list* argument to **clEnqueueAcquireGLObjects** and CL APIs that take a cl_event as an argument but do not enqueue commands. Passing such events to any other CL API that enqueues commands will generate a CL_INVALID_EVENT error.`"

12.6. Additions to the OpenCL Extension Specification

Add following the paragraph describing parameter *event* to **clEnqueueAcquireGLObjects**:

- "`If an OpenGL context is bound to the current thread, then any OpenGL commands which
- 1. affect or access the contents of a memory object listed in the mem_objects list, and
- 2. were issued on that OpenGL context prior to the call to clEnqueueAcquireGLObjects

will complete before execution of any OpenCL commands following the **clEnqueueAcquireGLObjects** which affect or access any of those memory objects. If a non-NULL *event* object is returned, it will report completion only after completion of such OpenGL commands.`"

Add following the paragraph describing parameter *event* to **clEnqueueReleaseGLObjects**:

- "`If an OpenGL context is bound to the current thread, then then any OpenGL commands which
- 1. affect or access the contents of the memory objects listed in the *mem_objects* list, and
- 2. are issued on that context after the call to clEnqueueReleaseGLObjects

will not execute until after execution of any OpenCL commands preceding the

clEnqueueReleaseGLObjects which affect or access any of those memory objects. If a non-NULL *event* object is returned, it will report completion before execution of such OpenGL commands.`"

Replace the second paragraph of Synchronizing OpenCL and OpenGL Access to Shared Objects with:

"`Prior to calling **clEnqueueAcquireGLObjects**, the application must ensure that any pending OpenGL operations which access the objects specified in *mem_objects* have completed.

If the **cl_khr_gl_event** extension is supported, then the OpenCL implementation will ensure that any such pending OpenGL operations are complete for an OpenGL context bound to the same thread as the OpenCL context. This is referred to as *implicit synchronization*.

If the **cl_khr_gl_event** extension is supported and the OpenGL context in question supports fence sync objects, completion of OpenGL commands may also be determined by placing a GL fence command after those commands using **glFenceSync**, creating an event from the resulting GL sync object using **clCreateEventFromGLsyncKHR**, and determining completion of that event object via **clEnqueueAcquireGLObjects**. This method may be considerably more efficient than calling **glFinish**, and is referred to as *explicit synchronization*. Explicit synchronization is most useful when an OpenGL context bound to another thread is accessing the memory objects.

If the cl_khr_gl_event extension is not supported, completion of OpenGL commands may be

determined by issuing and waiting for completion of a **glFinish** command on all OpenGL contexts with pending references to these objects. Some implementations may offer other efficient synchronization methods. If such methods exist they will be described in platform-specific documentation.

Note that no synchronization method other than **glFinish** is portable between all OpenGL implementations and all OpenCL implementations. While this is the only way to ensure completion that is portable to all platforms, **glFinish** is an expensive operation and its use should be avoided if the **cl_khr_gl_event** extension is supported on a platform.`"

12.7. Issues

1. How are references between CL events and GL syncs handled?

PROPOSED: The linked CL event places a single reference on the GL sync object. That reference is removed when the CL event is deleted. A more expensive alternative would be to reflect changes in the CL event reference count through to the GL sync.

2. How are linkages to synchronization primitives in other APIs handled?

UNRESOLVED. We will at least want to have a way to link events to EGL sync objects. There is probably no analogous DX concept. There would be an entry point for each type of synchronization primitive to be linked to, such as clCreateEventFromEGLSyncKHR.

An alternative is a generic clCreateEventFromExternalEvent taking an attribute list. The attribute list would include information defining the type of the external primitive and additional information (GL sync object handle, EGL display and sync object handle, etc.) specific to that type. This allows a single entry point to be reused.

These will probably be separate extensions following the API proposed here.

3. Should the CL_EVENT_COMMAND_TYPE correspond to the type of command (fence) or the type of the linked sync object?

PROPOSED: To the type of the linked sync object.

4. Should we support both explicit and implicit synchronization?

PROPOSED: Yes. Implicit synchronization is suitable when GL and CL are executing in the same application thread. Explicit synchronization is suitable when they are executing in different threads but the expense of glFinish is too high.

5. Should this be a platform or device extension?

PROPOSED: Platform extension. This may result in considerable under-the-hood work to implement the sync→event semantics using only the public GL API, however, when multiple drivers and devices with different GL support levels coexist in the same runtime.

6. Where can events generated from GL syncs be usable?

PROPOSED: Only with clEnqueueAcquireGLObjects, and attempting to use such an event

elsewhere will generate an error. There is no apparent use case for using such events elsewhere, and possibly some cost to supporting it, balanced by the cost of checking the source of events in all other commands accepting them as parameters.

Chapter 13. Creating OpenCL Memory Objects from Direct3D 10 Buffers and Textures

13.1. Overview

This section describes the **cl_khr_d3d10_sharing** extension. The goal of this extension is to provide interoperability between OpenCL and Direct3D 10.

13.2. General information

13.2.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

13.3. New Procedures and Functions

```
cl_int clGetDeviceIDsFromD3D10KHR(cl_platform_id platform,
                                  cl_d3d10_device_source_khr d3d_device_source,
                                  void *d3d object,
                                  cl_d3d10_device_set_khr d3d_device_set,
                                  cl_uint num_entries,
                                  cl device id *devices,
                                  cl_uint *num_devices)
cl mem clCreateFromD3D10BufferKHR(cl context context,
                                  cl_mem_flags flags,
                                  ID3D10Buffer *resource,
                                  cl_int *errcode_ret)
cl mem clCreateFromD3D10Texture2DKHR(cl context context,
                                     cl_mem_flags flags,
                                     ID3D10Texture2D *resource,
                                     UINT subresource,
                                     cl_int *errcode_ret)
cl mem clCreateFromD3D10Texture3DKHR(cl context context,
                                     cl_mem_flags flags,
                                     ID3D10Texture3D *resource,
                                     UINT subresource,
                                     cl_int *errcode_ret)
cl_int clEnqueueAcquireD3D10ObjectsKHR(cl_command_queue command_queue,
                                       cl_uint num_objects,
                                       const cl_mem *mem_objects,
                                       cl uint num events in wait list,
                                       const cl_event *event_wait_list,
                                       cl_event *event)
cl_int clEnqueueReleaseD3D100bjectsKHR(cl_command_queue command_queue,
                                       cl_uint num_objects,
                                       const cl_mem *mem_objects,
                                       cl_uint num_events_in_wait_list,
                                       const cl_event *event_wait_list,
                                       cl_event *event)
```

13.4. New Tokens

Accepted as a Direct3D 10 device source in the *d3d_device_source* parameter of **clGetDeviceIDsFromD3D10KHR**:

```
CL_D3D10_DEVICE_KHR
CL_D3D10_DXGI_ADAPTER_KHR
```

Accepted as a set of Direct3D 10 devices in the d3d_device_set parameter of

clGetDeviceIDsFromD3D10KHR:

```
CL_PREFERRED_DEVICES_FOR_D3D10_KHR
CL_ALL_DEVICES_FOR_D3D10_KHR
```

Accepted as a property name in the *properties* parameter of **clCreateContext** and **clCreateContextFromType**:

```
CL_CONTEXT_D3D10_DEVICE_KHR
```

Accepted as a property name in the *param_name* parameter of **clGetContextInfo**:

```
CL_CONTEXT_D3D10_PREFER_SHARED_RESOURCES_KHR
```

Accepted as the property being queried in the *param_name* parameter of **clGetMemObjectInfo**:

```
CL_MEM_D3D10_RESOURCE_KHR
```

Accepted as the property being queried in the *param_name* parameter of **clGetImageInfo**:

```
CL_IMAGE_D3D10_SUBRESOURCE_KHR
```

Returned in the *param_value* parameter of **clGetEventInfo** when *param_name* is CL_EVENT_COMMAND_TYPE:

```
CL_COMMAND_ACQUIRE_D3D10_OBJECTS_KHR
CL_COMMAND_RELEASE_D3D10_OBJECTS_KHR
```

Returned by **clCreateContext** and **clCreateContextFromType** if the Direct3D 10 device specified for interoperability is not compatible with the devices against which the context is to be created:

```
CL_INVALID_D3D10_DEVICE_KHR
```

Returned by **clCreateFromD3D10BufferKHR** when *resource* is not a Direct3D 10 buffer object, and by **clCreateFromD3D10Texture2DKHR** and **clCreateFromD3D10Texture3DKHR** when *resource* is not a Direct3D 10 texture object:

```
CL_INVALID_D3D10_RESOURCE_KHR
```

Returned by **clEnqueueAcquireD3D10ObjectsKHR** when any of *mem_objects* are currently acquired by OpenCL:

CL_D3D10_RESOURCE_ALREADY_ACQUIRED_KHR

Returned by **clEnqueueReleaseD3D10ObjectsKHR** when any of *mem_objects* are not currently acquired by OpenCL:

CL_D3D10_RESOURCE_NOT_ACQUIRED_KHR

13.5. Additions to Chapter 4 of the OpenCL 2.2 Specification

In section 4.4, replace the description of properties under clCreateContext with:

"_properties_ specifies a list of context property names and their corresponding values. Each property is followed immediately by the corresponding desired value. The list is terminated with zero. If a property is not specified in *properties*, then its default value (listed in *table 4.5*) is used (it is said to be specified implicitly). If *properties* is NULL or empty (points to a list whose first value is zero), all attributes take on their default values."

Add the following to *table 4.5*:

cl_context_properties enum	Property value	Description
CL_CONTEXT_D3D10_DEVICE_KHR	ID3D10Device *	Specifies the ID3D10Device * to use for Direct3D 10 interoperability. The default value is NULL.

Add to the list of errors for **clCreateContext**:

- CL_INVALID_D3D10_DEVICE_KHR if the value of the property CL_CONTEXT_D3D10_DEVICE_KHR is non-NULL and does not specify a valid Direct3D 10 device with which the *cl_device_ids* against which this context is to be created may interoperate.
- CL_INVALID_OPERATION if Direct3D 10 interoperability is specified by setting CL_INVALID_D3D10_DEVICE_KHR to a non-NULL value, and interoperability with another graphics API is also specified.

Add to the list of errors for **clCreateContextFromType** the same new errors described above for **clCreateContext**.

Add the following row to table 4.6:

cl_context_info	Return Type	Information returned in param_value
CL_CONTEXT_D3D10_PREFER_SHAR ED_RESOURCES_KHR	cl_bool	Returns CL_TRUE if Direct3D 10 resources created as shared by setting <i>MiscFlags</i> to include D3D10_RESOURCE_MISC_SHARED will perform faster when shared with OpenCL, compared with resources which have not set this flag. Otherwise returns CL_FALSE.

13.6. Additions to Chapter 5 of the OpenCL 2.2 Specification

Add to the list of errors for clGetMemObjectInfo:

• CL_INVALID_D3D10_RESOURCE_KHR if *param_name* is CL_MEM_D3D10_RESOURCE_KHR and *memobj* was not created by the function **clCreateFromD3D10BufferKHR**, **clCreateFromD3D10Texture2DKHR**, or **clCreateFromD3D10Texture3DKHR**.

Extend table 5.12 to include the following entry.

cl_mem_info	Return type	Info. returned in param_value
CL_MEM_D3D10_RESOURCE_KHR	ID3D10Resource *	If memobj was created using clCreateFromD3D10BufferKHR, clCreateFromD3D10Texture2DKHR, or clCreateFromD3D10Texture3DKHR, returns the resource argument specified when memobj was created.

Add to the list of errors for clGetImageInfo:

• CL_INVALID_D3D10_RESOURCE_KHR if *param_name* is CL_MEM_D3D10_SUBRESOURCE_KHR and *image* was not created by the function **clCreateFromD3D10Texture2DKHR**, or **clCreateFromD3D10Texture3DKHR**.

Extend table 5.9 to include the following entry.

cl_image_info	Return type	Info. returned in param_value
CL_MEM_D3D10_SUBRESOURCE_KH R	UINT	If image was created using clCreateFromD3D10Texture2DKHR, or clCreateFromD3D10Texture3DKHR, returns the subresource argument specified when image was created.

Add to *table 5.22* in the **Info returned in <param_value>** column for *cl_event_info* = CL_EVENT_COMMAND_TYPE:

```
CL_COMMAND_ACQUIRE_D3D10_OBJECTS_KHR
CL_COMMAND_RELEASE_D3D10_OBJECTS_KHR
```

13.7. Sharing Memory Objects with Direct3D 10 Resources

This section discusses OpenCL functions that allow applications to use Direct3D 10 resources as OpenCL memory objects. This allows efficient sharing of data between OpenCL and Direct3D 10. The OpenCL API may be used to execute kernels that read and/or write memory objects that are also Direct3D 10 resources. An OpenCL image object may be created from a Direct3D 10 texture resource. An OpenCL buffer object may be created from a Direct3D 10 buffer resource. OpenCL memory objects may be created from Direct3D 10 objects if and only if the OpenCL context has been created from a Direct3D 10 device.

13.7.1. Querying OpenCL Devices Corresponding to Direct3D 10 Devices

The OpenCL devices corresponding to a Direct3D 10 device may be queried. The OpenCL devices corresponding to a DXGI adapter may also be queried. The OpenCL devices corresponding to a Direct3D 10 device will be a subset of the OpenCL devices corresponding to the DXGI adapter against which the Direct3D 10 device was created.

The OpenCL devices corresponding to a Direct3D 10 device or a DXGI device may be queried using the function

platform refers to the platform ID returned by **clGetPlatformIDs**.

d3d_device_source specifies the type of d3d_object, and must be one of the values shown in the table below.

d3d_object specifies the object whose corresponding OpenCL devices are being queried. The type of d3d_object must be as specified in the table below.

d3d_device_set specifies the set of devices to return, and must be one of the values shown in the table below.

num_entries is the number of cl_device_id entries that can be added to devices. If devices is not NULL

then *num_entries* must be greater than zero.

devices returns a list of OpenCL devices found. The cl_device_id values returned in devices can be used to identify a specific OpenCL device. If devices is NULL, this argument is ignored. The number of OpenCL devices returned is the minimum of the value specified by num_entries and the number of OpenCL devices corresponding to d3d_object.

num_devices returns the number of OpenCL devices available that correspond to *d3d_object*. If *num_devices* is NULL, this argument is ignored.

clGetDeviceIDsFromD3D10KHR returns CL_SUCCESS if the function is executed successfully. Otherwise it may return

- CL_INVALID_PLATFORM if *platform* is not a valid platform.
- CL_INVALID_VALUE if *d3d_device_source* is not a valid value, *d3d_device_set* is not a valid value, *num_entries* is equal to zero and *devices* is not NULL, or if both *num_devices* and *devices* are NULL.
- CL_DEVICE_NOT_FOUND if no OpenCL devices that correspond to *d3d_object* were found.

Table 30. Direct3D 10 object types that may be used by clGetDeviceIDsFromD3D10KHR

cl_d3d_device_source_khr	Type of d3d_object	
CL_D3D10_DEVICE_KHR	ID3D10Device *	
CL_D3D10_DXGI_ADAPTER_KHR	IDXGIAdapter *	

Table 31. Sets of devices queriable using clGetDeviceIDsFromD3D10KHR

cl_d3d_device_set_khr	Devices returned in devices
CL_PREFERRED_DEVICES_FOR_D3D10_KHR	The preferred OpenCL devices associated with the specified Direct3D object.
CL_ALL_DEVICES_FOR_D3D10_KHR	All OpenCL devices which may interoperate with the specified Direct3D object. Performance of sharing data on these devices may be considerably less than on the preferred devices.

13.7.2. Lifetime of Shared Objects

An OpenCL memory object created from a Direct3D 10 resource remains valid as long as the corresponding Direct3D 10 resource has not been deleted. If the Direct3D 10 resource is deleted through the Direct3D 10 API, subsequent use of the OpenCL memory object will result in undefined behavior, including but not limited to possible OpenCL errors, data corruption, and program termination.

The successful creation of a cl_context against a Direct3D 10 device specified via the context create parameter CL_CONTEXT_D3D10_DEVICE_KHR will increment the internal Direct3D reference count on the specified Direct3D 10 device. The internal Direct3D reference count on that Direct3D 10 device will be decremented when the OpenCL reference count on the returned OpenCL context drops to zero.

The OpenCL context and corresponding command-queues are dependent on the existence of the Direct3D 10 device from which the OpenCL context was created. If the Direct3D 10 device is deleted through the Direct3D 10 API, subsequent use of the OpenCL context will result in undefined behavior, including but not limited to possible OpenCL errors, data corruption, and program termination.

13.7.3. Sharing Direct3D 10 Buffer Resources as OpenCL Buffer Objects

The function

creates an OpenCL buffer object from a Direct3D 10 buffer.

context is a valid OpenCL context created from a Direct3D 10 device.

flags is a bit-field that is used to specify usage information. Refer to table 5.3 for a description of flags. Only CL_MEM_READ_ONLY, CL_MEM_WRITE_ONLY and CL_MEM_READ_WRITE values specified in table 5.3 can be used.

resource is a pointer to the Direct3D 10 buffer to share.

errcode_ret will return an appropriate error code. If *errcode_ret* is NULL, no error code is returned.

clCreateFromD3D10BufferKHR returns a valid non-zero OpenCL buffer object and *errcode_ret* is set to CL_SUCCESS if the buffer object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode_ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context.
- CL_INVALID_VALUE if values specified in *flags* are not valid.
- CL_INVALID_D3D10_RESOURCE_KHR if *resource* is not a Direct3D 10 buffer resource, if *resource* was created with the D3D10_USAGE flag D3D10_USAGE_IMMUTABLE, if a cl_mem from *resource* has already been created using **clCreateFromD3D10BufferKHR**, or if *context* was not created against the same Direct3D 10 device from which *resource* was created.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The size of the returned OpenCL buffer object is the same as the size of *resource*. This call will increment the internal Direct3D reference count on *resource*. The internal Direct3D reference count on *resource* will be decremented when the OpenCL reference count on the returned OpenCL memory object drops to zero.

13.7.4. Sharing Direct3D 10 Texture and Resources as OpenCL Image Objects

The function

creates an OpenCL 2D image object from a subresource of a Direct3D 10 2D texture.

context is a valid OpenCL context created from a Direct3D 10 device.

flags is a bit-field that is used to specify usage information. Refer to *table 5.3* for a description of flags. Only CL_MEM_READ_ONLY, CL_MEM_WRITE_ONLY and CL_MEM_READ_WRITE values specified in *table 5.3* can be used.

resource is a pointer to the Direct3D 10 2D texture to share.

subresource is the subresource of *resource* to share.

errcode_ret will return an appropriate error code. If *errcode_ret* is NULL, no error code is returned.

clCreateFromD3D10Texture2DKHR returns a valid non-zero OpenCL image object and *errcode_ret* is set to CL_SUCCESS if the image object is created successfully. Otherwise, it returns a **NULL** value with one of the following error values returned in *errcode ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context.
- CL_INVALID_VALUE if values specified in *flags* are not valid or if *subresource* is not a valid subresource index for *resource*.
- CL_INVALID_D3D10_RESOURCE_KHR if *resource* is not a Direct3D 10 texture resource, if *resource* was created with the D3D10_USAGE flag D3D10_USAGE_IMMUTABLE, if *resource* is a multisampled texture, if a cl_mem from subresource *subresource* of *resource* has already been created using **clCreateFromD3D10Texture2DKHR**, or if *context* was not created against the same Direct3D 10 device from which *resource* was created.
- CL_INVALID_IMAGE_FORMAT_DESCRIPTOR if the Direct3D 10 texture format of *resource* is not listed in the table *Direct3D 10 formats and corresponding OpenCL image formats* or if the Direct3D 10 texture format of *resource* does not map to a supported OpenCL image format.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The width and height of the returned OpenCL 2D image object are determined by the width and height of subresource subresource of resource. The channel type and order of the returned OpenCL 2D image object is determined by the format of resource by the table Direct3D 10 formats and corresponding OpenCL image formats.

This call will increment the internal Direct3D reference count on *resource*. The internal Direct3D reference count on *resource* will be decremented when the OpenCL reference count on the returned OpenCL memory object drops to zero.

The function

creates an OpenCL 3D image object from a subresource of a Direct3D 10 3D texture.

context is a valid OpenCL context created from a Direct3D 10 device.

flags is a bit-field that is used to specify usage information. Refer to table 5.3 for a description of flags. Only CL_MEM_READ_ONLY, CL_MEM_WRITE_ONLY and CL_MEM_READ_WRITE values specified in table 5.3 can be used.

resource is a pointer to the Direct3D 10 3D texture to share.

subresource is the subresource of *resource* to share.

errcode_ret will return an appropriate error code. If *errcode_ret* is NULL, no error code is returned.

clCreateFromD3D10Texture3DKHR returns a valid non-zero OpenCL image object and *errcode_ret* is set to CL_SUCCESS if the image object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context.
- CL_INVALID_VALUE if values specified in *flags* are not valid or if *subresource* is not a valid subresource index for *resource*.
- CL_INVALID_D3D10_RESOURCE_KHR if *resource* is not a Direct3D 10 texture resource, if *resource* was created with the D3D10_USAGE flag D3D10_USAGE_IMMUTABLE, if *resource* is a multisampled texture, if a cl_mem from subresource *subresource* of *resource* has already been created using **clCreateFromD3D10Texture3DKHR**, or if *context* was not created against the same Direct3D 10 device from which *resource* was created.
- CL_INVALID_IMAGE_FORMAT_DESCRIPTOR if the Direct3D 10 texture format of *resource* is not listed in the table *Direct3D 10 formats and corresponding OpenCL image formats* or if the Direct3D 10 texture format of *resource* does not map to a supported OpenCL image format.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The width, height and depth of the returned OpenCL 3D image object are determined by the width, height and depth of subresource *subresource* of *resource*. The channel type and order of the returned OpenCL 3D image object is determined by the format of *resource* by the table *Direct3D 10 formats and corresponding OpenCL image formats*.

This call will increment the internal Direct3D reference count on *resource*. The internal Direct3D reference count on *resource* will be decremented when the OpenCL reference count on the returned OpenCL memory object drops to zero.

Table 32. Direct3D 10 formats and corresponding OpenCL image formats

DXGI format	CL image format (channel order, channel data type)	
DXGI_FORMAT_R32G32B32A32_FLOAT	CL_RGBA, CL_FLOAT	
DXGI_FORMAT_R32G32B32A32_UINT	CL_RGBA, CL_UNSIGNED_INT32	
DXGI_FORMAT_R32G32B32A32_SINT	CL_RGBA, CL_SIGNED_INT32	
DXGI_FORMAT_R16G16B16A16_FLOAT	CL_RGBA, CL_HALF_FLOAT	
DXGI_FORMAT_R16G16B16A16_UNORM	CL_RGBA, CL_UNORM_INT16	
DXGI_FORMAT_R16G16B16A16_UINT	CL_RGBA, CL_UNSIGNED_INT16	
DXGI_FORMAT_R16G16B16A16_SNORM	CL_RGBA, CL_SNORM_INT16	
DXGI_FORMAT_R16G16B16A16_SINT	CL_RGBA, CL_SIGNED_INT16	
DXGI_FORMAT_B8G8R8A8_UNORM	CL_BGRA, CL_UNORM_INT8	
DXGI_FORMAT_R8G8B8A8_UNORM	CL_RGBA, CL_UNORM_INT8	
DXGI_FORMAT_R8G8B8A8_UINT	CL_RGBA, CL_UNSIGNED_INT8	
DXGI_FORMAT_R8G8B8A8_SNORM	CL_RGBA, CL_SNORM_INT8	
DXGI_FORMAT_R8G8B8A8_SINT	CL_RGBA, CL_SIGNED_INT8	
DXGI_FORMAT_R32G32_FLOAT	CL_RG, CL_FLOAT	
DXGI_FORMAT_R32G32_UINT	CL_RG, CL_UNSIGNED_INT32	
DXGI_FORMAT_R32G32_SINT	CL_RG, CL_SIGNED_INT32	
DXGI_FORMAT_R16G16_FLOAT	CL_RG, CL_HALF_FLOAT	
DXGI_FORMAT_R16G16_UNORM	CL_RG, CL_UNORM_INT16	
DXGI_FORMAT_R16G16_UINT	CL_RG, CL_UNSIGNED_INT16	
DXGI_FORMAT_R16G16_SNORM	CL_RG, CL_SNORM_INT16	
DXGI_FORMAT_R16G16_SINT	CL_RG, CL_SIGNED_INT16	
DXGI_FORMAT_R8G8_UNORM	CL_RG, CL_UNORM_INT8	
DXGI_FORMAT_R8G8_UINT	CL_RG, CL_UNSIGNED_INT8	
DXGI_FORMAT_R8G8_SNORM	CL_RG, CL_SNORM_INT8	
DXGI_FORMAT_R8G8_SINT	CL_RG, CL_SIGNED_INT8	
DXGI_FORMAT_R32_FLOAT	CL_R, CL_FLOAT	
DXGI_FORMAT_R32_UINT	CL_R, CL_UNSIGNED_INT32	
DXGI_FORMAT_R32_SINT	CL_R, CL_SIGNED_INT32	
DXGI_FORMAT_R16_FLOAT	CL_R, CL_HALF_FLOAT	
DXGI_FORMAT_R16_UNORM	CL_R, CL_UNORM_INT16	

DXGI format	CL image format (channel order, channel data type)	
DXGI_FORMAT_R16_UINT	CL_R, CL_UNSIGNED_INT16	
DXGI_FORMAT_R16_SNORM	CL_R, CL_SNORM_INT16	
DXGI_FORMAT_R16_SINT	CL_R, CL_SIGNED_INT16	
DXGI_FORMAT_R8_UNORM	CL_R, CL_UNORM_INT8	
DXGI_FORMAT_R8_UINT	CL_R, CL_UNSIGNED_INT8	
DXGI_FORMAT_R8_SNORM	CL_R, CL_SNORM_INT8	
DXGI_FORMAT_R8_SINT	CL_R, CL_SIGNED_INT8	

13.7.5. Querying Direct3D properties of memory objects created from Direct3D 10 resources

Properties of Direct3D 10 objects may be queried using **clGetMemObjectInfo** and **clGetImageInfo** with *param_name* CL_MEM_D3D10_RESOURCE_KHR and

CL_IMAGE_D3D10_SUBRESOURCE_KHR respectively as described in sections 5.4.3 and 5.3.6.

13.7.6. Sharing memory objects created from Direct3D 10 resources between Direct3D 10 and OpenCL contexts

The function

is used to acquire OpenCL memory objects that have been created from Direct3D 10 resources. The Direct3D 10 objects are acquired by the OpenCL context associated with *command_queue* and can therefore be used by all command-queues associated with the OpenCL context.

OpenCL memory objects created from Direct3D 10 resources must be acquired before they can be used by any OpenCL commands queued to a command-queue. If an OpenCL memory object created from a Direct3D 10 resource is used while it is not currently acquired by OpenCL, the call attempting to use that OpenCL memory object will return CL_D3D10_RESOURCE_NOT_ACQUIRED_KHR.

If CL_CONTEXT_INTEROP_USER_SYNC is not specified as CL_TRUE during context creation, clEnqueueAcquireD3D10ObjectsKHR provides the synchronization guarantee that any Direct3D 10 calls involving the interop device(s) used in the OpenCL context made before clEnqueueAcquireD3D10ObjectsKHR is called will complete executing before *event* reports

completion and before the execution of any subsequent OpenCL work issued in *command_queue* begins. If the context was created with properties specifying CL_CONTEXT_INTEROP_USER_SYNC as CL_TRUE, the user is responsible for guaranteeing that any Direct3D 10 calls involving the interop device(s) used in the OpenCL context made before **clEnqueueAcquireD3D10ObjectsKHR** is called have completed before calling **clEnqueueAcquireD3D10ObjectsKHR**.

command_queue is a valid command-queue.

num_objects is the number of memory objects to be acquired in *mem_objects*.

mem_objects is a pointer to a list of OpenCL memory objects that were created from Direct3D 10 resources.

event_wait_list and num_events_in_wait_list specify events that need to complete before this particular command can be executed. If event_wait_list is NULL, then this particular command does not wait on any event to complete. If event_wait_list is NULL, num_events_in_wait_list must be 0. If event_wait_list is not NULL, the list of events pointed to by event_wait_list must be valid and num_events_in_wait_list must be greater than 0. The events specified in event_wait_list act as synchronization points.

event returns an event object that identifies this command and can be used to query or wait for this command to complete. If event is NULL or the enqueue is unsuccessful, no event will be created and therefore it will not be possible to query the status of this command or to wait for this command to complete. If event_wait_list and event are not NULL, event must not refer to an element of the event_wait_list array.

clEnqueueAcquireD3D10ObjectsKHR returns CL_SUCCESS if the function is executed successfully. If *num_objects* is 0 and *mem_objects* is NULL then the function does nothing and returns CL_SUCCESS. Otherwise it returns one of the following errors:

- CL_INVALID_VALUE if *num_objects* is zero and *mem_objects* is not a NULL value or if *num_objects* > 0 and *mem_objects* is NULL.
- CL_INVALID_MEM_OBJECT if memory objects in *mem_objects* are not valid OpenCL memory objects or if memory objects in *mem_objects* have not been created from Direct3D 10 resources.
- CL_INVALID_COMMAND_QUEUE if command_queue is not a valid command-queue.
- CL_INVALID_CONTEXT if context associated with *command_queue* was not created from an Direct3D 10 context.
- CL_D3D10_RESOURCE_ALREADY_ACQUIRED_KHR if memory objects in *mem_objects* have previously been acquired using **clEnqueueAcquireD3D10ObjectsKHR** but have not been released using **clEnqueueReleaseD3D10ObjectsKHR**.
- CL_INVALID_EVENT_WAIT_LIST if event_wait_list is NULL and num_events_in_wait_list > 0, or event_wait_list is not NULL and num_events_in_wait_list is 0, or if event objects in event_wait_list are not valid events.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

is used to release OpenCL memory objects that have been created from Direct3D 10 resources. The Direct3D 10 objects are released by the OpenCL context associated with *command_queue*.

OpenCL memory objects created from Direct3D 10 resources which have been acquired by OpenCL must be released by OpenCL before they may be accessed by Direct3D 10. Accessing a Direct3D 10 resource while its corresponding OpenCL memory object is acquired is in error and will result in undefined behavior, including but not limited to possible OpenCL errors, data corruption, and program termination.

If CL_CONTEXT_INTEROP_USER_SYNC is not specified as CL_TRUE during context creation, clEnqueueReleaseD3D10ObjectsKHR provides the synchronization guarantee that any calls to Direct3D 10 calls involving the interop device(s) used in the OpenCL context made after the call to clEnqueueReleaseD3D10ObjectsKHR will not start executing until after all events in event_wait_list are complete and all work already submitted to command_queue completes execution. If the created context was with properties specifying CL_CONTEXT_INTEROP_USER_SYNC as CL_TRUE, the user is responsible for guaranteeing that any Direct3D 10 calls involving the interop device(s) used in the OpenCL context made after clEnqueueReleaseD3D10ObjectsKHR will not start executing until after event returned by clEnqueueReleaseD3D10ObjectsKHR reports completion.

num_objects is the number of memory objects to be released in *mem_objects*.

mem_objects is a pointer to a list of OpenCL memory objects that were created from Direct3D 10 resources.

event_wait_list and num_events_in_wait_list specify events that need to complete before this particular command can be executed. If event_wait_list is NULL, then this particular command does not wait on any event to complete. If event_wait_list is NULL, num_events_in_wait_list must be 0. If event_wait_list is not NULL, the list of events pointed to by event_wait_list must be valid and num_events_in_wait_list must be greater than 0.

event returns an event object that identifies this command and can be used to query or wait for this command to complete. If event is NULL or the enqueue is unsuccessful, no event will be created and therefore it will not be possible to query the status of this command or to wait for this command to complete. If event_wait_list and event are not NULL, event must not refer to an element of the event_wait_list array.

clEnqueueReleaseD3D10ObjectsKHR returns CL_SUCCESS if the function is executed successfully. If *num_objects* is 0 and *mem_objects* is NULL the function does nothing and returns CL_SUCCESS. Otherwise it returns one of the following errors:

• CL_INVALID_VALUE if num_objects is zero and mem_objects is not a NULL value or if num_objects

- > 0 and mem_objects is NULL.
- CL_INVALID_MEM_OBJECT if memory objects in *mem_objects* are not valid OpenCL memory objects or if memory objects in *mem_objects* have not been created from Direct3D 10 resources.
- CL_INVALID_COMMAND_QUEUE if command_queue is not a valid command-queue.
- CL_INVALID_CONTEXT if context associated with command_queue was not created from a Direct3D 10 device.
- CL_D3D10_RESOURCE_NOT_ACQUIRED_KHR if memory objects in *mem_objects* have not previously been acquired using **clEnqueueAcquireD3D10ObjectsKHR**, or have been released using **clEnqueueReleaseD3D10ObjectsKHR** since the last time that they were acquired.
- CL_INVALID_EVENT_WAIT_LIST if event_wait_list is NULL and num_events_in_wait_list > 0, or event_wait_list is not NULL and num_events_in_wait_list> is 0, or if event objects in event_wait_list are not valid events.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

13.7.7. Event Command Types for Sharing memory objects that map to Direct3D 10 objects

The following table describes the event command types for the OpenCL commands to acquire and release OpenCL memory objects that have been created from Direct3D 10 objects:

Table 33. List of supported event command types

Events Created By	Event Command Type	
clEnqueueAcquireD3D10ObjectsKH R	CL_COMMAND_ACQUIRE_D3D10_OBJECTS_KHR	
clEnqueueReleaseD3D10ObjectsKHR	CL_COMMAND_RELEASE_D3D10_OBJECTS_KHR	

13.8. Issues

1. Should this extension be KHR or EXT?

PROPOSED: KHR. If this extension is to be approved by Khronos then it should be KHR, otherwise EXT. Not all platforms can support this extension, but that is also true of OpenGL interop.

RESOLVED: KHR.

2. Requiring SharedHandle on ID3D10Resource

Requiring this can largely simplify things at the DDI level and make some implementations faster. However, the DirectX spec only defines the shared handle for a subset of the resources we would like to support:

```
D3D10_RESOURCE_MISC_SHARED - Enables the sharing of resource data between two or more Direct3D devices.
The only resources that can be shared are 2D non-mipmapped textures.
```

PROPOSED A: Add wording to the spec about some implementations needing the resource setup as shared:

"Some implementations may require the resource to be shared on the D3D10 side of the API"

If we do that, do we need another enum to describe this failure case?

PROPOSED B: Require that all implementations support both shared and non-shared resources. The restrictions prohibiting multisample textures and the flag D3D10_USAGE_IMMUTABLE guarantee software access to all shareable resources.

RESOLVED: Require that implementations support both D3D10_RESOURCE_MISC_SHARED being set and not set. Add the query for CL_CONTEXT_D3D10_PREFER_SHARED_RESOURCES_KHR to determine on a per-context basis which method will be faster.

3. Texture1D support

There is not a matching CL type, so do we want to support this and map to buffer or Texture2D? If so the command might correspond to the 2D / 3D versions:

RESOLVED: We will not add support for ID3D10Texture1D objects unless a corresponding OpenCL 1D Image type is created.

4. CL/D3D10 queries

The GL interop has clGetGLObjectInfo and clGetGLTextureInfo. It is unclear if these are needed on the D3D10 interop side since the D3D10 spec makes these queries trivial on the D3D10 object itself. Also, not all of the semantics of the GL call map across.

PROPOSED: Add the **clGetMemObjectInfo** and **clGetImageInfo** parameter names CL_MEM_D3D10_RESOURCE_KHR and CL_IMAGE_D3D10_SUBRESOURCE_KHR to query the D3D10 resource from which a cl_mem was created. From this data, any D3D10 side information may be queried using the D3D10 API.

RESOLVED: We will use **clGetMemObjectInfo** and **clGetImageInfo** to access this information.

Chapter 14. Creating OpenCL Memory Objects from Direct3D 11 Buffers and Textures

14.1. Overview

This section describes the **cl_khr_d3d11_sharing** extension. The goal of this extension is to provide interoperability between OpenCL and Direct3D 11.

14.2. General information

14.2.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

14.3. New Procedures and Functions

```
cl_int clGetDeviceIDsFromD3D11KHR(cl_platform_id platform,
                                  cl_d3d11_device_source_khr d3d_device_source,
                                  void *d3d object,
                                  cl_d3d11_device_set_khr d3d_device_set,
                                  cl_uint num_entries,
                                  cl device id *devices,
                                  cl_uint *num_devices)
cl mem clCreateFromD3D11BufferKHR(cl context context,
                                  cl_mem_flags flags,
                                  ID3D11Buffer *resource,
                                  cl_int *errcode_ret)
cl mem clCreateFromD3D11Texture2DKHR(cl context context,
                                     cl_mem_flags flags,
                                     ID3D11Texture2D *resource,
                                     UINT subresource,
                                     cl_int *errcode_ret)
cl mem clCreateFromD3D11Texture3DKHR(cl context context,
                                     cl_mem_flags flags,
                                     ID3D11Texture3D *resource,
                                     UINT subresource,
                                     cl_int *errcode_ret)
cl_int clEnqueueAcquireD3D110bjectsKHR(cl_command_queue command_queue,
                                       cl_uint num_objects,
                                       const cl_mem *mem_objects,
                                       cl uint num events in wait list,
                                       const cl_event *event_wait_list,
                                       cl_event *event)
cl_int clEnqueueReleaseD3D110bjectsKHR(cl_command_queue command_queue,
                                       cl_uint num_objects,
                                       const cl_mem *mem_objects,
                                       cl_uint num_events_in_wait_list,
                                       const cl_event *event_wait_list,
                                       cl_event *event)
```

14.4. New Tokens

Accepted as a Direct3D 11 device source in the *d3d_device_source* parameter of **clGetDeviceIDsFromD3D11KHR**:

```
CL_D3D11_DEVICE_KHR
CL_D3D11_DXGI_ADAPTER_KHR
```

Accepted as a set of Direct3D 11 devices in the _d3d_device_set_parameter of

clGetDeviceIDsFromD3D11KHR:

```
CL_PREFERRED_DEVICES_FOR_D3D11_KHR
CL_ALL_DEVICES_FOR_D3D11_KHR
```

Accepted as a property name in the *properties* parameter of **clCreateContext** and **clCreateContextFromType**:

```
CL_CONTEXT_D3D11_DEVICE_KHR
```

Accepted as a property name in the *param_name* parameter of **clGetContextInfo**:

```
CL_CONTEXT_D3D11_PREFER_SHARED_RESOURCES_KHR
```

Accepted as the property being queried in the *param_name* parameter of **clGetMemObjectInfo**:

```
CL_MEM_D3D11_RESOURCE_KHR
```

Accepted as the property being queried in the *param_name* parameter of **clGetImageInfo**:

```
CL_IMAGE_D3D11_SUBRESOURCE_KHR
```

Returned in the *param_value* parameter of **clGetEventInfo** when *param_name* is CL_EVENT_COMMAND_TYPE:

```
CL_COMMAND_ACQUIRE_D3D11_OBJECTS_KHR
CL_COMMAND_RELEASE_D3D11_OBJECTS_KHR
```

Returned by **clCreateContext** and **clCreateContextFromType** if the Direct3D 11 device specified for interoperability is not compatible with the devices against which the context is to be created:

```
CL_INVALID_D3D11_DEVICE_KHR
```

Returned by **clCreateFromD3D11BufferKHR** when *resource* is not a Direct3D 11 buffer object, and by **clCreateFromD3D11Texture2DKHR** and **clCreateFromD3D11Texture3DKHR** when *resource* is not a Direct3D 11 texture object.

```
CL_INVALID_D3D11_RESOURCE_KHR
```

Returned by **clEnqueueAcquireD3D11ObjectsKHR** when any of *mem_objects* are currently acquired by OpenCL:

CL_D3D11_RESOURCE_ALREADY_ACQUIRED_KHR

Returned by **clEnqueueReleaseD3D11ObjectsKHR** when any of *mem_objects* are not currently acquired by OpenCL:

CL_D3D11_RESOURCE_NOT_ACQUIRED_KHR

14.5. Additions to Chapter 4 of the OpenCL 2.2 Specification

In section 4.4, replace the description of properties under clCreateContext with:

"_properties_ specifies a list of context property names and their corresponding values. Each property is followed immediately by the corresponding desired value. The list is terminated with zero. If a property is not specified in *properties*, then its default value (listed in *table 4.5*) is used (it is said to be specified implicitly). If *properties* is NULL or empty (points to a list whose first value is zero), all attributes take on their default values."

Add the following to *table 4.5*:

cl_context_properties enum	Property value	Description
CL_CONTEXT_D3D11_DEVICE_KHR	ID3D11Device *	Specifies the ID3D11Device * to use for Direct3D 11 interoperability.
		The default value is NULL.

Add to the list of errors for **clCreateContext**:

- CL_INVALID_D3D11_DEVICE_KHR if the value of the property CL_CONTEXT_D3D11_DEVICE_KHR is non-NULL and does not specify a valid Direct3D 11 device with which the *cl_device_ids* against which this context is to be created may interoperate.
- CL_INVALID_OPERATION if Direct3D 11 interoperability is specified by setting CL_INVALID_D3D11_DEVICE_KHR to a non-NULL value, and interoperability with another graphics API is also specified.

Add to the list of errors for **clCreateContextFromType** the same new errors described above for **clCreateContext**.

Add the following row to table 4.6:

cl_context_info	Return Type	Information returned in param_value
CL_CONTEXT_D3D11_PREFER_SHAR ED_RESOURCES_KHR	cl_bool	Returns CL_TRUE if Direct3D 11 resources created as shared by setting <i>MiscFlags</i> to include D3D11_RESOURCE_MISC_SHARED will perform faster when shared with OpenCL, compared with resources which have not set this flag. Otherwise returns CL_FALSE.

14.6. Additions to Chapter 5 of the OpenCL 2.2 Specification

Add to the list of errors for **clGetMemObjectInfo**:

• CL_INVALID_D3D11_RESOURCE_KHR if *param_name* is CL_MEM_D3D11_RESOURCE_KHR and *memobj* was not created by the function **clCreateFromD3D11BufferKHR**, **clCreateFromD3D11Texture2DKHR**, or **clCreateFromD3D11Texture3DKHR**.

Extend table 5.12 to include the following entry.

cl_mem_info	Return type	Info. returned in param_value
CL_MEM_D3D11_RESOURCE_KHR	ID3D11Resource *	If memobj was created using clCreateFromD3D11BufferKHR, clCreateFromD3D11Texture2DKHR, or clCreateFromD3D11Texture3DKHR, returns the resource argument specified when memobj was created.

Add to the list of errors for clGetImageInfo:

• CL_INVALID_D3D11_RESOURCE_KHR if *param_name* is CL_MEM_D3D11_SUBRESOURCE_KHR and *image* was not created by the function **clCreateFromD3D11Texture2DKHR**, or **clCreateFromD3D11Texture3DKHR**.

Extend table 5.9 to include the following entry.

cl_image_info	Return type	Info. returned in param_value
CL_MEM_D3D11_SUBRESOURCE_KH R	UINT	If image was created using clCreateFromD3D11Texture2DKHR, or
		clCreateFromD3D11Texture3DKHR, returns the <i>subresource</i> argument specified when <i>image</i> was created.

Add to *table 5.22* in the **Info returned in param_value** column for *cl_event_info* = CL_EVENT_COMMAND_TYPE:

```
CL_COMMAND_ACQUIRE_D3D11_OBJECTS_KHR
CL_COMMAND_RELEASE_D3D11_OBJECTS_KHR
```

14.7. Sharing Memory Objects with Direct3D 11 Resources

This section discusses OpenCL functions that allow applications to use Direct3D 11 resources as OpenCL memory objects. This allows efficient sharing of data between OpenCL and Direct3D 11. The OpenCL API may be used to execute kernels that read and/or write memory objects that are also Direct3D 11 resources. An OpenCL image object may be created from a Direct3D 11 texture resource. An OpenCL buffer object may be created from a Direct3D 11 buffer resource. OpenCL memory objects may be created from Direct3D 11 objects if and only if the OpenCL context has been created from a Direct3D 11 device.

14.7.1. Querying OpenCL Devices Corresponding to Direct3D 11 Devices

The OpenCL devices corresponding to a Direct3D 11 device may be queried. The OpenCL devices corresponding to a DXGI adapter may also be queried. The OpenCL devices corresponding to a Direct3D 11 device will be a subset of the OpenCL devices corresponding to the DXGI adapter against which the Direct3D 11 device was created.

The OpenCL devices corresponding to a Direct3D 11 device or a DXGI device may be queried using the function

platform refers to the platform ID returned by **clGetPlatformIDs**.

d3d_device_source specifies the type of d3d_object, and must be one of the values shown in the table below.

d3d_object specifies the object whose corresponding OpenCL devices are being queried. The type of d3d_object must be as specified in the table below.

d3d_device_set specifies the set of devices to return, and must be one of the values shown in the table below.

num_entries is the number of cl_device_id entries that can be added to devices. If devices is not NULL

then *num_entries* must be greater than zero.

devices returns a list of OpenCL devices found. The cl_device_id values returned in devices can be used to identify a specific OpenCL device. If devices is NULL, this argument is ignored. The number of OpenCL devices returned is the minimum of the value specified by num_entries and the number of OpenCL devices corresponding to d3d_object.

num_devices returns the number of OpenCL devices available that correspond to *d3d_object*. If *num_devices* is NULL, this argument is ignored.

clGetDeviceIDsFromD3D10KHR returns CL_SUCCESS if the function is executed successfully. Otherwise it may return

- CL_INVALID_PLATFORM if *platform* is not a valid platform.
- CL_INVALID_VALUE if *d3d_device_source* is not a valid value, *d3d_device_set* is not a valid value, *num_entries* is equal to zero and *devices* is not NULL, or if both *num_devices* and *devices* are NULL.
- CL_DEVICE_NOT_FOUND if no OpenCL devices that correspond to *d3d_object* were found.

Table 34. Direct3D 11 object types that may be used by clGetDeviceIDsFromD3D11KHR

cl_d3d_device_source_khr	Type of d3d_object
CL_D3D11_DEVICE_KHR	ID3D11Device *
CL_D3D11_DXGI_ADAPTER_KHR	IDXGIAdapter *

Table 35. Sets of devices queriable using clGetDeviceIDsFromD3D11KHR

cl_d3d_device_set_khr	Devices returned in devices
CL_PREFERRED_DEVICES_FOR_D3D11_KHR	The preferred OpenCL devices associated with the specified Direct3D object.
CL_ALL_DEVICES_FOR_D3D11_KHR	All OpenCL devices which may interoperate with the specified Direct3D object. Performance of sharing data on these devices may be considerably less than on the preferred devices.

14.7.2. Lifetime of Shared Objects

An OpenCL memory object created from a Direct3D 11 resource remains valid as long as the corresponding Direct3D 11 resource has not been deleted. If the Direct3D 11 resource is deleted through the Direct3D 11 API, subsequent use of the OpenCL memory object will result in undefined behavior, including but not limited to possible OpenCL errors, data corruption, and program termination.

The successful creation of a cl_context against a Direct3D 11 device specified via the context create parameter CL_CONTEXT_D3D11_DEVICE_KHR will increment the internal Direct3D reference count on the specified Direct3D 11 device. The internal Direct3D reference count on that Direct3D 11 device will be decremented when the OpenCL reference count on the returned OpenCL context drops to zero.

The OpenCL context and corresponding command-queues are dependent on the existence of the Direct3D 11 device from which the OpenCL context was created. If the Direct3D 11 device is deleted through the Direct3D 11 API, subsequent use of the OpenCL context will result in undefined behavior, including but not limited to possible OpenCL errors, data corruption, and program termination.

14.7.3. Sharing Direct3D 11 Buffer Resources as OpenCL Buffer Objects

The function

creates an OpenCL buffer object from a Direct3D 11 buffer.

context is a valid OpenCL context created from a Direct3D 11 device.

flags is a bit-field that is used to specify usage information. Refer to table 5.3 for a description of flags. Only CL_MEM_READ_ONLY, CL_MEM_WRITE_ONLY and CL_MEM_READ_WRITE values specified in table 5.3 can be used.

resource is a pointer to the Direct3D 11 buffer to share.

errcode_ret will return an appropriate error code. If *errcode_ret* is NULL, no error code is returned.

clCreateFromD3D11BufferKHR returns a valid non-zero OpenCL buffer object and *errcode_ret* is set to CL_SUCCESS if the buffer object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode_ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context.
- CL_INVALID_VALUE if values specified in *flags* are not valid.
- CL_INVALID_D3D11_RESOURCE_KHR if *resource* is not a Direct3D 11 buffer resource, if *resource* was created with the D3D11_USAGE flag D3D11_USAGE_IMMUTABLE, if a cl_mem from *resource* has already been created using **clCreateFromD3D11BufferKHR**, or if *context* was not created against the same Direct3D 11 device from which *resource* was created.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The size of the returned OpenCL buffer object is the same as the size of *resource*. This call will increment the internal Direct3D reference count on *resource*. The internal Direct3D reference count on *resource* will be decremented when the OpenCL reference count on the returned OpenCL memory object drops to zero.

14.7.4. Sharing Direct3D 11 Texture and Resources as OpenCL Image Objects

The function

creates an OpenCL 2D image object from a subresource of a Direct3D 11 2D texture.

context is a valid OpenCL context created from a Direct3D 11 device.

flags is a bit-field that is used to specify usage information. Refer to *table 5.3* for a description of flags. Only CL_MEM_READ_ONLY, CL_MEM_WRITE_ONLY and CL_MEM_READ_WRITE values specified in *table 5.3* can be used.

resource is a pointer to the Direct3D 11 2D texture to share.

subresource is the subresource of *resource* to share.

errcode_ret will return an appropriate error code. If *errcode_ret* is NULL, no error code is returned.

clCreateFromD3D11Texture2DKHR returns a valid non-zero OpenCL image object and *errcode_ret* is set to CL_SUCCESS if the image object is created successfully. Otherwise, it returns a **NULL** value with one of the following error values returned in *errcode ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context.
- CL_INVALID_VALUE if values specified in *flags* are not valid or if *subresource* is not a valid subresource index for *resource*.
- CL_INVALID_D3D11_RESOURCE_KHR if *resource* is not a Direct3D 11 texture resource, if *resource* was created with the D3D11_USAGE flag D3D11_USAGE_IMMUTABLE, if *resource* is a multisampled texture, if a cl_mem from subresource *subresource* of *resource* has already been created using **clCreateFromD3D11Texture2DKHR**, or if *context* was not created against the same Direct3D 10 device from which *resource* was created.
- CL_INVALID_IMAGE_FORMAT_DESCRIPTOR if the Direct3D 11 texture format of *resource* is not listed in the table *Direct3D 11 formats and corresponding OpenCL image formats* or if the Direct3D 11 texture format of *resource* does not map to a supported OpenCL image format.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The width and height of the returned OpenCL 2D image object are determined by the width and height of subresource subresource of resource. The channel type and order of the returned OpenCL 2D image object is determined by the format of resource by the table Direct3D 11 formats and corresponding OpenCL image formats.

This call will increment the internal Direct3D reference count on *resource*. The internal Direct3D reference count on *resource* will be decremented when the OpenCL reference count on the returned OpenCL memory object drops to zero.

The function

creates an OpenCL 3D image object from a subresource of a Direct3D 11 3D texture.

context is a valid OpenCL context created from a Direct3D 11 device.

flags is a bit-field that is used to specify usage information. Refer to table 5.3 for a description of flags. Only CL_MEM_READ_ONLY, CL_MEM_WRITE_ONLY and CL_MEM_READ_WRITE values specified in table 5.3 can be used.

resource is a pointer to the Direct3D 11 3D texture to share.

subresource is the subresource of *resource* to share.

errcode_ret will return an appropriate error code. If *errcode_ret* is NULL, no error code is returned.

clCreateFromD3D11Texture3DKHR returns a valid non-zero OpenCL image object and *errcode_ret* is set to CL_SUCCESS if the image object is created successfully. Otherwise, it returns a **NULL** value with one of the following error values returned in *errcode ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context.
- CL_INVALID_VALUE if values specified in *flags* are not valid or if *subresource* is not a valid subresource index for *resource*.
- CL_INVALID_D3D11_RESOURCE_KHR if *resource* is not a Direct3D 11 texture resource, if *resource* was created with the D3D11_USAGE flag D3D11_USAGE_IMMUTABLE, if *resource* is a multisampled texture, if a cl_mem from subresource *subresource* of *resource* has already been created using **clCreateFromD3D11Texture3DKHR**, or if *context* was not created against the same Direct3D 11 device from which *resource* was created.
- CL_INVALID_IMAGE_FORMAT_DESCRIPTOR if the Direct3D 11 texture format of *resource* is not listed in the table *Direct3D 11 formats and corresponding OpenCL image formats* or if the Direct3D 11 texture format of *resource* does not map to a supported OpenCL image format.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The width, height and depth of the returned OpenCL 3D image object are determined by the width, height and depth of subresource *subresource* of *resource*. The channel type and order of the returned OpenCL 3D image object is determined by the format of *resource* by the table *Direct3D 11 formats and corresponding OpenCL image formats*.

This call will increment the internal Direct3D reference count on *resource*. The internal Direct3D reference count on *resource* will be decremented when the OpenCL reference count on the returned OpenCL memory object drops to zero.

Table 36. Direct3D 11 formats and corresponding OpenCL image formats

DXGI format	CL image format (channel order, channel data type)
DXGI_FORMAT_R32G32B32A32_FLOAT	CL_RGBA, CL_FLOAT
DXGI_FORMAT_R32G32B32A32_UINT	CL_RGBA, CL_UNSIGNED_INT32
DXGI_FORMAT_R32G32B32A32_SINT	CL_RGBA, CL_SIGNED_INT32
DXGI_FORMAT_R16G16B16A16_FLOAT	CL_RGBA, CL_HALF_FLOAT
DXGI_FORMAT_R16G16B16A16_UNORM	CL_RGBA, CL_UNORM_INT16
DXGI_FORMAT_R16G16B16A16_UINT	CL_RGBA, CL_UNSIGNED_INT16
DXGI_FORMAT_R16G16B16A16_SNORM	CL_RGBA, CL_SNORM_INT16
DXGI_FORMAT_R16G16B16A16_SINT	CL_RGBA, CL_SIGNED_INT16
DXGI_FORMAT_B8G8R8A8_UNORM	CL_BGRA, CL_UNORM_INT8
DXGI_FORMAT_R8G8B8A8_UNORM	CL_RGBA, CL_UNORM_INT8
DXGI_FORMAT_R8G8B8A8_UINT	CL_RGBA, CL_UNSIGNED_INT8
DXGI_FORMAT_R8G8B8A8_SNORM	CL_RGBA, CL_SNORM_INT8
DXGI_FORMAT_R8G8B8A8_SINT	CL_RGBA, CL_SIGNED_INT8
DXGI_FORMAT_R32G32_FLOAT	CL_RG, CL_FLOAT
DXGI_FORMAT_R32G32_UINT	CL_RG, CL_UNSIGNED_INT32
DXGI_FORMAT_R32G32_SINT	CL_RG, CL_SIGNED_INT32
DXGI_FORMAT_R16G16_FLOAT	CL_RG, CL_HALF_FLOAT
DXGI_FORMAT_R16G16_UNORM	CL_RG, CL_UNORM_INT16
DXGI_FORMAT_R16G16_UINT	CL_RG, CL_UNSIGNED_INT16
DXGI_FORMAT_R16G16_SNORM	CL_RG, CL_SNORM_INT16
DXGI_FORMAT_R16G16_SINT	CL_RG, CL_SIGNED_INT16
DXGI_FORMAT_R8G8_UNORM	CL_RG, CL_UNORM_INT8
DXGI_FORMAT_R8G8_UINT	CL_RG, CL_UNSIGNED_INT8
DXGI_FORMAT_R8G8_SNORM	CL_RG, CL_SNORM_INT8
DXGI_FORMAT_R8G8_SINT	CL_RG, CL_SIGNED_INT8
DXGI_FORMAT_R32_FLOAT	CL_R, CL_FLOAT
DXGI_FORMAT_R32_UINT	CL_R, CL_UNSIGNED_INT32
DXGI_FORMAT_R32_SINT	CL_R, CL_SIGNED_INT32
DXGI_FORMAT_R16_FLOAT	CL_R, CL_HALF_FLOAT
DXGI_FORMAT_R16_UNORM	CL_R, CL_UNORM_INT16

DXGI format	CL image format (channel order, channel data type)
DXGI_FORMAT_R16_UINT	CL_R, CL_UNSIGNED_INT16
DXGI_FORMAT_R16_SNORM	CL_R, CL_SNORM_INT16
DXGI_FORMAT_R16_SINT	CL_R, CL_SIGNED_INT16
DXGI_FORMAT_R8_UNORM	CL_R, CL_UNORM_INT8
DXGI_FORMAT_R8_UINT	CL_R, CL_UNSIGNED_INT8
DXGI_FORMAT_R8_SNORM	CL_R, CL_SNORM_INT8
DXGI_FORMAT_R8_SINT	CL_R, CL_SIGNED_INT8

14.7.5. Querying Direct3D properties of memory objects created from Direct3D 11 resources

Properties of Direct3D 11 objects may be queried using **clGetMemObjectInfo** and **clGetImageInfo** with *param_name* CL_MEM_D3D11_RESOURCE_KHR and

CL_IMAGE_D3D11_SUBRESOURCE_KHR respectively as described in sections 5.4.3 and 5.3.6.

14.7.6. Sharing memory objects created from Direct3D 11 resources between Direct3D 11 and OpenCL contexts

The function

is used to acquire OpenCL memory objects that have been created from Direct3D 11 resources. The Direct3D 11 objects are acquired by the OpenCL context associated with *command_queue* and can therefore be used by all command-queues associated with the OpenCL context.

OpenCL memory objects created from Direct3D 11 resources must be acquired before they can be used by any OpenCL commands queued to a command-queue. If an OpenCL memory object created from a Direct3D 11 resource is used while it is not currently acquired by OpenCL, the call attempting to use that OpenCL memory object will return CL_D3D11_RESOURCE_NOT_ACQUIRED_KHR.

If CL_CONTEXT_INTEROP_USER_SYNC is not specified as CL_TRUE during context creation, clEnqueueAcquireD3D11ObjectsKHR provides the synchronization guarantee that any Direct3D 11 calls involving the interop device(s) used in the OpenCL context made before clEnqueueAcquireD3D11ObjectsKHR is called will complete executing before *event* reports

completion and before the execution of any subsequent OpenCL work issued in *command_queue* begins. If the context was created with properties specifying CL_CONTEXT_INTEROP_USER_SYNC as CL_TRUE, the user is responsible for guaranteeing that any Direct3D 11 calls involving the interop device(s) used in the OpenCL context made before **clEnqueueAcquireD3D11ObjectsKHR** is called have completed before calling **clEnqueueAcquireD3D11ObjectsKHR**.

command_queue is a valid command-queue.

num_objects is the number of memory objects to be acquired in *mem_objects*.

mem_objects is a pointer to a list of OpenCL memory objects that were created from Direct3D 11 resources.

event_wait_list and num_events_in_wait_list specify events that need to complete before this particular command can be executed. If event_wait_list is NULL, then this particular command does not wait on any event to complete. If event_wait_list is NULL, num_events_in_wait_list must be 0. If event_wait_list is not NULL, the list of events pointed to by event_wait_list must be valid and num_events_in_wait_list must be greater than 0. The events specified in event_wait_list act as synchronization points.

event returns an event object that identifies this command and can be used to query or wait for this command to complete. If event is NULL or the enqueue is unsuccessful, no event will be created and therefore it will not be possible to query the status of this command or to wait for this command to complete. If event_wait_list and event are not NULL, event must not refer to an element of the event_wait_list array.

clEnqueueAcquireD3D11ObjectsKHR returns CL_SUCCESS if the function is executed successfully. If *num_objects* is 0 and *mem_objects* is NULL then the function does nothing and returns CL_SUCCESS. Otherwise it returns one of the following errors:

- CL_INVALID_VALUE if *num_objects* is zero and *mem_objects* is not a NULL value or if *num_objects* > 0 and *mem_objects* is NULL.
- CL_INVALID_MEM_OBJECT if memory objects in *mem_objects* are not valid OpenCL memory objects or if memory objects in *mem_objects* have not been created from Direct3D 11 resources.
- CL_INVALID_COMMAND_QUEUE if command_queue is not a valid command-queue.
- CL_INVALID_CONTEXT if context associated with *command_queue* was not created from an Direct3D 11 context.
- CL_D3D11_RESOURCE_ALREADY_ACQUIRED_KHR if memory objects in *mem_objects* have previously been acquired using **clEnqueueAcquireD3D11ObjectsKHR** but have not been released using **clEnqueueReleaseD3D11ObjectsKHR**.
- CL_INVALID_EVENT_WAIT_LIST if event_wait_list is NULL and num_events_in_wait_list > 0, or event_wait_list is not NULL and num_events_in_wait_list is 0, or if event objects in event_wait_list are not valid events.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

is used to release OpenCL memory objects that have been created from Direct3D 11 resources. The Direct3D 11 objects are released by the OpenCL context associated with *command_queue*.

OpenCL memory objects created from Direct3D 11 resources which have been acquired by OpenCL must be released by OpenCL before they may be accessed by Direct3D 11. Accessing a Direct3D 11 resource while its corresponding OpenCL memory object is acquired is in error and will result in undefined behavior, including but not limited to possible OpenCL errors, data corruption, and program termination.

If CL_CONTEXT_INTEROP_USER_SYNC is not specified as CL_TRUE during context creation, clEnqueueReleaseD3D11ObjectsKHR provides the synchronization guarantee that any calls to Direct3D 11 calls involving the interop device(s) used in the OpenCL context made after the call to clEnqueueReleaseD3D11ObjectsKHR will not start executing until after all events in event_wait_list are complete and all work already submitted to command_queue completes execution. If the context created was with properties specifying CL_CONTEXT_INTEROP_USER_SYNC as CL_TRUE, the user is responsible for guaranteeing that any Direct3D 11 calls involving the interop device(s) used in the OpenCL context made after clEnqueueReleaseD3D11ObjectsKHR will not start executing until after event returned by clEnqueueReleaseD3D11ObjectsKHR reports completion.

num_objects is the number of memory objects to be released in *mem_objects*.

mem_objects is a pointer to a list of OpenCL memory objects that were created from Direct3D 11 resources.

event_wait_list and num_events_in_wait_list specify events that need to complete before this particular command can be executed. If event_wait_list is NULL, then this particular command does not wait on any event to complete. If event_wait_list is NULL, num_events_in_wait_list must be 0. If event_wait_list is not NULL, the list of events pointed to by event_wait_list must be valid and num_events_in_wait_list must be greater than 0.

event returns an event object that identifies this command and can be used to query or wait for this command to complete. If event is NULL or the enqueue is unsuccessful, no event will be created and therefore it will not be possible to query the status of this command or to wait for this command to complete. If event_wait_list and event are not NULL, event must not refer to an element of the event_wait_list array.

clEnqueueReleaseD3D11ObjectsKHR returns CL_SUCCESS if the function is executed successfully. If *num_objects* is 0 and *mem_objects* is NULL the function does nothing and returns CL_SUCCESS. Otherwise it returns one of the following errors:

• CL_INVALID_VALUE if num_objects is zero and mem_objects is not a NULL value or if num_objects

- > 0 and mem_objects is NULL.
- CL_INVALID_MEM_OBJECT if memory objects in *mem_objects* are not valid OpenCL memory objects or if memory objects in *mem_objects* have not been created from Direct3D 11 resources.
- CL_INVALID_COMMAND_QUEUE if command_queue is not a valid command-queue.
- CL_INVALID_CONTEXT if context associated with command_queue was not created from a Direct3D 11 device.
- CL_D3D11_RESOURCE_NOT_ACQUIRED_KHR if memory objects in *mem_objects* have not previously been acquired using **clEnqueueAcquireD3D11ObjectsKHR**, or have been released using **clEnqueueReleaseD3D11ObjectsKHR** since the last time that they were acquired.
- CL_INVALID_EVENT_WAIT_LIST if event_wait_list is NULL and num_events_in_wait_list > 0, or event_wait_list is not NULL and num_events_in_wait_list > is 0, or if event objects in event_wait_list are not valid events.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

14.7.7. Event Command Types for Sharing memory objects that map to Direct3D 11 objects

The following table describes the event command types for the OpenCL commands to acquire and release OpenCL memory objects that have been created from Direct3D 11 objects:

Table 37. List of supported event command types

Events Created By	Event Command Type
clEnqueueAcquireD3D11ObjectsKH R	CL_COMMAND_ACQUIRE_D3D11_OBJECTS_KHR
clEnqueueReleaseD3D11ObjectsKHR	CL_COMMAND_RELEASE_D3D11_OBJECTS_KHR

Chapter 15. Creating OpenCL Memory Objects from DirectX 9 Media Surfaces

15.1. Overview

This section describes the **cl_khr_dx9_media_sharing** extension. The goal of this extension is to allow applications to use media surfaces as OpenCL memory objects. This allows efficient sharing of data between OpenCL and selected adapter APIs (only DX9 for now). If this extension is supported, an OpenCL image object can be created from a media surface and the OpenCL API can be used to execute kernels that read and/or write memory objects that are media surfaces. Note that OpenCL memory objects may be created from the adapter media surface if and only if the OpenCL context has been created from that adapter.

15.2. General information

15.2.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

15.3. New Procedures and Functions

```
cl_int clGetDeviceIDsFromDX9MediaAdapterKHR(
                                     cl_platform_id platform,
                                     cl uint num media adapters,
                                     cl_dx9_media_adapter_type_khr *
media_adapters_type,
                                     void *media adapters,
                                     cl_dx9_media_adapter_set_khr media_adapter_set,
                                     cl_uint num_entries,
                                     cl device id *devices,
                                     cl_int *num_devices)
cl_mem clCreateFromDX9MediaSurfaceKHR(cl_context context,
                                       cl_mem_flags flags,
                                       cl dx9 media adapter type khr adapter type,
                                       void *surface_info,
                                       cl_uint plane,
                                       cl int *errcode ret)
cl_int clEnqueueAcquireDX9MediaSurfacesKHR(cl_command_queue command_queue,
                                            cl uint num objects,
                                            const cl_mem *mem_objects,
                                            cl_uint num_events_in_wait_list,
                                            const cl_event *event_wait_list,
                                            cl_event *event)
cl_int clEnqueueReleaseDX9MediaSurfacesKHR(cl_command_queue command_queue,
                                            cl_uint num_objects,
                                            const cl_mem *mem_objects,
                                            cl uint num events in wait list,
                                            const cl_event *event_wait_list,
                                            cl_event *event)
```

15.4. New Tokens

Accepted by the *media_adapter_type* parameter of **clGetDeviceIDsFromDX9MediaAdapterKHR**:

```
CL_ADAPTER_D3D9_KHR
CL_ADAPTER_D3D9EX_KHR
CL_ADAPTER_DXVA_KHR
```

Accepted by the *media_adapter_set* parameter of **clGetDeviceIDsFromDX9MediaAdapterKHR**:

```
CL_PREFERRED_DEVICES_FOR_DX9_MEDIA_ADAPTER_KHR
CL_ALL_DEVICES_FOR_DX9_MEDIA_ADAPTER_KHR
```

Accepted as a property name in the properties parameter of clCreateContext and

clCreateContextFromType:

```
CL_CONTEXT_ADAPTER_D3D9_KHR
CL_CONTEXT_ADAPTER_D3D9EX_KHR
CL_CONTEXT_ADAPTER_DXVA_KHR
```

Accepted as the property being queried in the *param_name* parameter of **clGetMemObjectInfo**:

```
CL_MEM_DX9_MEDIA_ADAPTER_TYPE_KHR
CL_MEM_DX9_MEDIA_SURFACE_INFO_KHR
```

Accepted as the property being queried in the *param_name* parameter of **clGetImageInfo**:

```
CL_IMAGE_DX9_MEDIA_PLANE_KHR
```

Returned in the *param_value* parameter of **clGetEventInfo** when *param_name* is CL_EVENT_COMMAND_TYPE:

```
CL_COMMAND_ACQUIRE_DX9_MEDIA_SURFACES_KHR
CL_COMMAND_RELEASE_DX9_MEDIA_SURFACES_KHR
```

Returned by **clCreateContext** and **clCreateContextFromType** if the media adapter specified for interoperability is not compatible with the devices against which the context is to be created:

```
CL_INVALID_DX9_MEDIA_ADAPTER_KHR
```

Returned by **clCreateFromDX9MediaSurfaceKHR** when *adapter_type* is set to a media adapter and the *surface_info* does not reference a media surface of the required type, or if *adapter_type* is set to a media adapter type and *surface_info* does not contain a valid reference to a media surface on that adapter, by **clGetMemObjectInfo** when *param_name* is a surface or handle when the image was not created from an appropriate media surface, and from **clGetImageInfo** when *param_name* is CL IMAGE_DX9_MEDIA_PLANE KHR and image was not created from an appropriate media surface.

```
CL_INVALID_DX9_MEDIA_SURFACE_KHR
```

Returned by **clEnqueueAcquireDX9MediaSurfacesKHR** when any of *mem_objects* are currently acquired by OpenCL:

```
CL_DX9_MEDIA_SURFACE_ALREADY_ACQUIRED_KHR
```

Returned by clEnqueueReleaseDX9MediaSurfacesKHR when any of mem_objects are not

CL_DX9_MEDIA_SURFACE_NOT_ACQUIRED_KHR

15.5. Additions to Chapter 4 of the OpenCL 2.2 Specification

In section 4.4, replace the description of properties under clCreateContext with:

"_properties_ specifies a list of context property names and their corresponding values. Each property is followed immediately by the corresponding desired value. The list is terminated with zero. If a property is not specified in *properties*, then its default value (listed in *table 4.5*) is used (it is said to be specified implicitly). If *properties* is NULL or empty (points to a list whose first value is zero), all attributes take on their default values."

Add the following to *table 4.5*:

cl_context_properties enum	Property value	Description
CL_CONTEXT_ADAPTER_D3D9 _KHR	IDirect3DDevice9 *	Specifies an IDirect3DDevice9 to use for D3D9 interop.
CL_CONTEXT_ADAPTER_D3D9 EX_KHR	IDirect3DDeviceEx*	Specifies an IDirect3DDevice9Ex to use for D3D9 interop.
CL_CONTEXT_ADAPTER_DXV A_KHR	IDXVAHD_Device *	Specifies an IDXVAHD_Device to use for DXVA interop.

Add to the list of errors for **clCreateContext**:

• CL_INVALID_ADAPTER_KHR if any of the values of the properties CL_CONTEXT_ADAPTER_D3D9_KHR, CL_CONTEXT_ADAPTER_D3D9EX_KHR or CL_CONTEXT_ADAPTER_DXVA_KHR is non-NULL and does not specify a valid media adapter with which the <code>cl_device_ids</code> against which this context is to be created may interoperate.

Add to the list of errors for **clCreateContextFromType** the same new errors described above for **clCreateContext**.

15.6. Additions to Chapter 5 of the OpenCL 2.2 Specification

Add to the list of errors for **clGetMemObjectInfo**:

• CL_INVALID_DX9_MEDIA_SURFACE_KHR if param_name is CL_MEM_DX9_MEDIA_SURFACE_INFO_KHR and memobj was not created by the function clCreateFromDX9MediaSurfaceKHR from a Direct3D9 surface.

Extend *table 5.12* to include the following entry:

cl_mem_info	Return type	Info. returned in param_value
CL_MEM_DX9_MEDIA_ADAPT ER_TYPE_KHR	cl_dx9_media_adapter_type_k hr	Returns the cl_dx9_media_adapter_type_khr argument value specified when memobj is created using clCreateFromDX9MediaSurfa ceKHR.
CL_MEM_DX9_MEDIA_SURFA CE_INFO_KHR	cl_dx9_surface_info_khr	Returns the cl_dx9_surface_info_khr argument value specified when memobj is created using clCreateFromDX9MediaSurfa ceKHR.

Add to the list of errors for **clGetImageInfo**:

• CL_INVALID_DX9_MEDIA_SURFACE_KHR if *param_name* is CL_IMAGE_DX9_MEDIA_PLANE_KHR and *image* was not created by the function **clCreateFromDX9MediaSurfaceKHR**.

Extend table 5.9 to include the following entry.

cl_image_info	Return type	Info. returned in param_value
CL_IMAGE_DX9_MEDIA_PLAN E_KHR	cl_uint	Returns the <i>plane</i> argument value specified when <i>memobj</i> is created using clCreateFromDX9MediaSurfa ceKHR .

Add to *table 5.22* in the **Info returned in param_value** column for *cl_event_info* = CL_EVENT_COMMAND_TYPE:

CL_COMMAND_ACQUIRE_DX9_MEDIA_SURFACES_KHR
CL_COMMAND_RELEASE_DX9_MEDIA_SURFACES_KHR

15.7. Sharing Media Surfaces with OpenCL

This section discusses OpenCL functions that allow applications to use media surfaces as OpenCL memory objects. This allows efficient sharing of data between OpenCL and media surface APIs. The OpenCL API may be used to execute kernels that read and/or write memory objects that are also media surfaces. An OpenCL image object may be created from a media surface. OpenCL memory objects may be created from media surfaces if and only if the OpenCL context has been created from a media adapter.

15.7.1. Querying OpenCL Devices corresponding to Media Adapters

Media adapters are an abstraction associated with devices that provide media capabilities.

The function

queries a media adapter for any associated OpenCL devices. Adapters with associated OpenCL devices can enable media surface sharing between the two.

platform refers to the platform ID returned by **clGetPlatformIDs**.

num_media_adapters specifies the number of media adapters.

media_adapters_type is an array of *num_media_adapters* entries. Each entry specifies the type of media adapter and must be one of the values described in the table below.

Table 38. cl_dx9_media_adapter_type_khr values

cl_dx9_media_adapter_type_khr	Type of media adapters
CL_ADAPTER_D3D9_KHR	IDirect3DDevice9 *
CL_ADAPTER_D3D9EX_KHR	IDirect3DDevice9Ex *
CL_ADAPTER_DXVA_KHR	IDXVAHD_Device *

Table 39. cl_dx9_media_adapter_set_khr values

cl_dx9_media_adapter_set_khr	Description
CL_PREFERRED_DEVICES_FOR_DX9_MEDIA_A DAPTER_KHR	The preferred OpenCL devices associated with the media adapter.
CL_ALL_DEVICES_FOR_MEDIA_DX9_ADAPTER _KHR	All OpenCL devices that may interoperate with the media adapter

media_adapters is an array of num_media_adapters entries. Each entry specifies the actual adapter whose type is specified by media_adapter_type. The media_adapters must be one of the types described in the table cl_dx9_media_adapter_type_khr values. media_adapter_set specifies the set of adapters to return and must be one of the values described in the table <<[[cl_khr_dx9_media_sharing-media-adapter-sets,cl_dx9_media_adapter_set_khr values>>.

num_entries is the number of cl_device_id entries that can be added to *devices*. If *devices* is not NULL, the *num_entries* must be greater than zero.

devices returns a list of OpenCL devices found that support the list of media adapters specified. The cl_device_id values returned in *devices* can be used to identify a specific OpenCL device. If *devices* argument is NULL, this argument is ignored. The number of OpenCL devices returned is the minimum of the value specified by *num_entries* or the number of OpenCL devices whose type matches *device_type*.

num_devices returns the number of OpenCL devices. If *num_devices* is NULL, this argument is ignored.

clGetDeviceIDsFromDX9MediaAdapterKHR returns CL_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- CL_INVALID_PLATFORM if *platform* is not a valid platform.
- CL_INVALID_VALUE if num_media_adapters is zero or if media_adapters_type is NULL or if media_adapters is NULL.
- CL_INVALID_VALUE if any of the entries in *media_adapters_type* or *media_adapters* is not a valid value.
- CL_INVALID_VALUE if *media_adapter_set* is not a valid value.
- CL_INVALID_VALUE if *num_entries* is equal to zero and *devices* is not NULL or if both *num_devices* and *devices* are NULL.
- CL_DEVICE_NOT_FOUND if no OpenCL devices that correspond to adapters specified in *media_adapters* and *media_adapters_type* were found.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

15.7.2. Creating Media Resources as OpenCL Image Objects

The function

creates an OpenCL image object from a media surface.

context is a valid OpenCL context created from a media adapter.

flags is a bit-field that is used to specify usage information. Refer to table 5.3 for a description of

flags. Only CL_MEM_READ_ONLY, CL_MEM_WRITE_ONLY and CL_MEM_READ_WRITE values specified in *table 5.3* can be used.

adapter_type is a value from enumeration of supported adapters described in the table <code>cl_dx9_media_adapter_type_khr</code> values. The type of <code>surface_info</code> is determined by the adapter type. The implementation does not need to support all adapter types. This approach provides flexibility to support additional adapter types in the future. Supported adapter types are <code>CL_ADAPTER_D3D9_KHR</code>, <code>CL_ADAPTER_D3D9EX_KHR</code> and <code>CL_ADAPTER_DXVA_KHR</code>.

If adapter_type is CL_ADAPTER_D3D9_KHR, CL_ADAPTER_D3D9EX_KHR and CL_ADAPTER_DXVA_KHR, the *surface_info* points to the following structure:

```
typedef struct _cl_dx9_surface_info_khr
{
    IDirect3DSurface9 *resource;
    HANDLE shared_handle;
} cl_dx9_surface_info_khr;
```

For DX9 surfaces, we need both the handle to the resource and the resource itself to have a sufficient amount of information to eliminate a copy of the surface for sharing in cases where this is possible. Elimination of the copy is driver dependent. *shared_handle* may be NULL and this may result in sub-optimal performance.

surface_info is a pointer to one of the structures defined in the *adapter_type* description above passed in as a void *.

plane is the plane of resource to share for planar surface formats. For planar formats, we use the plane parameter to obtain a handle to thie specific plane (Y, U or V for example). For non-planar formats used by media, *plane* must be 0.

errcode_ret will return an appropriate error code. If *errcode_ret* is NULL, no error code is returned.

clCreateFromDX9MediaSurfaceKHR returns a valid non-zero 2D image object and *errcode_ret* is set to CL_SUCCESS if the 2D image object is created successfully. Otherwise it returns a NULL value with one of the following error values returned in *errcode_ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context.
- CL_INVALID_VALUE if values specified in *flags* are not valid or if *plane* is not a valid plane of *resource* specified in *surface_info*.
- CL_INVALID_DX9_MEDIA_SURFACE_KHR if resource specified in surface_info is not a valid resource or is not associated with adapter_type (e.g., adapter_type is set to CL_ADAPTER_D3D9_KHR and resource is not a Direct3D 9 surface created in D3DPOOL DEFAULT).
- CL_INVALID_DX9_MEDIA_SURFACE_KHR if shared_handle specified in surface_info is not NULL or a valid handle value.
- CL_INVALID_IMAGE_FORMAT_DESCRIPTOR if the texture format of *resource* is not listed in *YUV* FourCC codes and corresponding OpenCL image format or Direct3D formats and corresponding

OpenCL image formats.

- CL_INVALID_OPERATION if there are no devices in *context* that support *adapter_type*.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The width and height of the returned OpenCL 2D image object are determined by the width and height of the plane of resource. The channel type and order of the returned image object is determined by the format and plane of resource and are described in the table YUV FourCC codes and corresponding OpenCL image format or Direct3D formats and corresponding OpenCL image formats.

This call will increment the internal media surface count on *resource*. The internal media surface reference count on *resource* will be decremented when the OpenCL reference count on the returned OpenCL memory object drops to zero.

15.7.3. Querying Media Surface Properties of Memory Objects created from Media Surfaces

Properties of media surface objects may be queried using **clGetMemObjectInfo** and **clGetImageInfo** with *param_name* CL_MEM_DX9_MEDIA_ADAPTER_TYPE_KHR, CL_MEM_DX9_MEDIA_SURFACE_INFO_KHR and CL_IMAGE_DX9_MEDIA_PLANE_KHR as described in *sections 5.4.3* and *5.3.6*.

15.7.4. Sharing Memory Objects created from Media Surfaces between a Media Adapter and OpenCL

The function

is used to acquire OpenCL memory objects that have been created from a media surface. The media surfaces are acquired by the OpenCL context associated with *command_queue* and can therefore be used by all command-queues associated with the OpenCL context.

OpenCL memory objects created from media surfaces must be acquired before they can be used by any OpenCL commands queued to a command-queue. If an OpenCL memory object created from a media surface is used while it is not currently acquired by OpenCL, the call attempting to use that OpenCL memory object will return CL_DX9_MEDIA_SURFACE_NOT_ACQUIRED_KHR.

If CL_CONTEXT_INTEROP_USER_SYNC is not specified as CL_TRUE during context creation,

clEnqueueAcquireDX9MediaSurfacesKHR provides the synchronization guarantee that any media adapter API calls involving the interop device(s) used in the OpenCL context made before clEnqueueAcquireDX9MediaSurfacesKHR is called will complete executing before event reports completion and before the execution of any subsequent OpenCL work issued in command_queue begins. If the context was created with properties specifying CL CONTEXT INTEROP USER SYNC as CL_TRUE, the user is responsible for guaranteeing that any media adapter API calls involving the device(s) the OpenCL before interop used in context made clEnqueueAcquireDX9MediaSurfacesKHR is called before calling have completed $cl Enqueue Acquire DX9 Media Surfaces KHR \ . \\$

command_queue is a valid command-queue.

num_objects is the number of memory objects to be acquired in mem_objects.

mem_objects is a pointer to a list of OpenCL memory objects that were created from media surfaces.

event_wait_list and num_events_in_wait_list specify events that need to complete before this particular command can be executed. If event_wait_list is NULL, then this particular command does not wait on any event to complete. If event_wait_list is NULL, num_events_in_wait_list must be 0. If event_wait_list is not NULL, the list of events pointed to by event_wait_list must be valid and num_events_in_wait_list must be greater than 0. The events specified in event_wait_list act as synchronization points.

event returns an event object that identifies this command and can be used to query or wait for this command to complete. If event is NULL or the enqueue is unsuccessful, no event will be created and therefore it will not be possible to query the status of this command or to wait for this command to complete. If event_wait_list and event are not NULL, event must not refer to an element of the event_wait_list array.

clEnqueueAcquireDX9MediaSurfacesKHR returns CL_SUCCESS if the function is executed successfully. If *num_objects* is 0 and *mem_objects* is NULL then the function does nothing and returns CL_SUCCESS. Otherwise it returns one of the following errors:

- CL_INVALID_VALUE if num_objects is zero and mem_objects is not a NULL value or if num_objects
 o and mem_objects is NULL.
- CL_INVALID_MEM_OBJECT if memory objects in *mem_objects* are not valid OpenCL memory objects or if memory objects in *mem_objects* have not been created from media surfaces.
- CL_INVALID_COMMAND_QUEUE if command_queue is not a valid command-queue.
- CL_INVALID_CONTEXT if context associated with *command_queue* was not created from a device that can share the media surface referenced by *mem_objects*.
- CL_DX9_MEDIA_SURFACE_ALREADY_ACQUIRED_KHR if memory objects in *mem_objects* have previously been acquired using **clEnqueueAcquireDX9MediaSurfacesKHR** but have not been released using **clEnqueueReleaseDX9MediaSurfacesKHR**.
- CL_INVALID_EVENT_WAIT_LIST if event_wait_list is NULL and num_events_in_wait_list > 0, or event_wait_list is not NULL and num_events_in_wait_list is 0, or if event objects in event_wait_list are not valid events.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL

implementation on the host.

The function

is used to release OpenCL memory objects that have been created from media surfaces. The media surfaces are released by the OpenCL context associated with *command_queue*.

OpenCL memory objects created from media surfaces which have been acquired by OpenCL must be released by OpenCL before they may be accessed by the media adapter API. Accessing a media surface while its corresponding OpenCL memory object is acquired is in error and will result in undefined behavior, including but not limited to possible OpenCL errors, data corruption, and program termination.

If CL CONTEXT INTEROP USER SYNC is not specified as CL TRUE during context creation, clEnqueueReleaseDX9MediaSurfacesKHR provides the synchronization guarantee that any calls to media adapter APIs involving the interop device(s) used in the OpenCL context made after the call to clEnqueueReleaseDX9MediaSurfacesKHR will not start executing until after all events in event_wait_list are complete and all work already submitted to command_queue completes execution. If the context was created with properties specifying CL_CONTEXT_INTEROP_USER_SYNC as CL_TRUE, the user is responsible for guaranteeing that any media adapter API calls involving the interop device(s) used in the OpenCL context made after clEnqueueReleaseDX9MediaSurfacesKHR will not start executing until after event returned by clEnqueueReleaseDX9MediaSurfacesKHR reports completion.

num_objects is the number of memory objects to be released in *mem_objects*.

mem_objects is a pointer to a list of OpenCL memory objects that were created from media surfaces.

event_wait_list and num_events_in_wait_list specify events that need to complete before this particular command can be executed. If event_wait_list is NULL, then this particular command does not wait on any event to complete. If event_wait_list is NULL, num_events_in_wait_list must be 0. If event_wait_list is not NULL, the list of events pointed to by event_wait_list must be valid and num_events_in_wait_list must be greater than 0.

event returns an event object that identifies this command and can be used to query or wait for this command to complete. If event is NULL or the enqueue is unsuccessful, no event will be created and therefore it will not be possible to query the status of this command or to wait for this command to complete. If event_wait_list and event are not NULL, event must not refer to an element of the event_wait_list array.

clEnqueueReleaseDX9MediaSurfaceKHR returns CL_SUCCESS if the function is executed successfully. If *num_objects* is 0 and *<mem_objects>* is **NULL** the function does nothing and returns

CL_SUCCESS. Otherwise it returns one of the following errors:

- CL_INVALID_VALUE if *num_objects* is zero and *mem_objects* is not a NULL value or if *num_objects* > 0 and *mem_objects* is NULL.
- CL_INVALID_MEM_OBJECT if memory objects in *mem_objects* are not valid OpenCL memory objects or if memory objects in *mem_objects* have not been created from valid media surfaces.
- CL_INVALID_COMMAND_QUEUE if command_queue is not a valid command-queue.
- CL_INVALID_CONTEXT if context associated with *command_queue* was not created from a media object.
- CL_DX9_MEDIA_SURFACE_NOT_ACQUIRED_KHR if memory objects in mem_objects have not previously been acquired using clEnqueueAcquireDX9MediaSurfacesKHR, or have been released using clEnqueueReleaseDX9MediaSurfacesKHR since the last time that they were acquired.
- CL_INVALID_EVENT_WAIT_LIST if event_wait_list is NULL and num_events_in_wait_list > 0, or event_wait_list is not NULL and num_events_in_wait_list > is 0, or if event objects in event_wait_list are not valid events.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

15.7.5. Event Command Types for Sharing Memory Objects created from Media Surfaces

The following table describes the event command types for the OpenCL commands to acquire and release OpenCL memory objects that have been created from media surfaces:

Table 40. List of supported event command types

Events Created By	Event Command Type
clEnqueueAcquireDX9MediaSurface sKHR	CL_COMMAND_ACQUIRE_DX9_MEDIA_SURFACES_KHR
clEnqueueReleaseDX9MediaSurface sKHR	CL_COMMAND_RELEASE_DX9_MEDIA_SURFACES_KHR

15.7.6. Surface formats for Media Surface Sharing

This section includes the D3D surface formats that are supported when the adapter type is one of the Direct 3D lineage . Using a D3D surface format not listed here is an error. To extend the use of this extension to support media adapters beyond DirectX9 tables similar to the ones in this section will need to be defined for the surface formats supported by the new media adapter. All implementations that support this extension are required to support the NV12 surface format, the other surface formats supported are the same surface formats that the adapter you are sharing with supports as long as they are listed in the table <code>YUV FourCC codes and corresponding OpenCL image formats</code>.

Table 41. YUV FourCC codes and corresponding OpenCL image format

FOUR CC code	CL image format (channel order, channel data type)
FOURCC('N','V','1','2'), Plane 0	CL_R, CL_UNORM_INT8
FOURCC('N','V','1','2'), Plane 1	CL_RG, CL_UNORM_INT8
FOURCC('Y','V','1','2'), Plane 0	CL_R, CL_UNORM_INT8
FOURCC('Y','V','1','2'), Plane 1	CL_R, CL_UNORM_INT8
FOURCC('Y','V','1','2'), Plane 2	CL_R, CL_UNORM_INT8

In the table *YUV FourCC codes and corresponding OpenCL image format* above, NV12 Plane 0 corresponds to the luminance (Y) channel and Plane 1 corresponds to the UV channels. The YV12 Plane 0 corresponds to the Y channel, Plane 1 corresponds to the V channel and Plane 2 corresponds to the U channel. Note that the YUV formats map to CL_R and CL_RG but do not perform any YUV to RGB conversion and vice-versa.

Table 42. Direct3D formats and corresponding OpenCL image formats

D3D format	CL image format (channel order, channel data type)
D3DFMT_R32F	CL_R, CL_FLOAT
D3DFMT_R16F	CL_R, CL_HALF_FLOAT
D3DFMT_L16	CL_R, CL_UNORM_INT16
D3DFMT_A8	CL_A, CL_UNORM_INT8
D3DFMT_L8	CL_R, CL_UNORM_INT8
D3DFMT_G32R32F	CL_RG, CL_FLOAT
D3DFMT_G16R16F	CL_RG, CL_HALF_FLOAT
D3DFMT_G16R16	CL_RG, CL_UNORM_INT16
D3DFMT_A8L8	CL_RG, CL_UNORM_INT8
D3DFMT_A32B32G32R32F	CL_RGBA, CL_FLOAT
D3DFMT_A16B16G16R16F	CL_RGBA, CL_HALF_FLOAT
D3DFMT_A16B16G16R16	CL_RGBA, CL_UNORM_INT16
D3DFMT_A8B8G8R8	CL_RGBA, CL_UNORM_INT8
D3DFMT_X8B8G8R8	CL_RGBA, CL_UNORM_INT8
D3DFMT_A8R8G8B8	CL_BGRA, CL_UNORM_INT8
D3DFMT_X8R8G8B8	CL_BGRA, CL_UNORM_INT8

Note: The D3D9 format names in the table above seem to imply that the order of the color channels are switched relative to OpenCL but this is not the case. For example, the layout of channels for each pixel for D3DFMT_A32FB32FG32FR32F is the same as CL_RGBA, CL_FLOAT.

Chapter 16. Depth Images

This section describes the cl_khr_depth_images extension.

This extension adds support for depth images.

This extension became a core feature in OpenCL 2.0.

16.1. General information

16.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

16.2. Additions to Chapter 5 of the OpenCL 1.2 Specification

This extension adds the following new image formats for depth images to *tables 5.6 and 5.7* of the OpenCL 1.2 specification.

Enum values that can be specified in channel_order CL_DEPTH. This format can only be used if channel data type = CL_UNORM_INT16 or CL_FLOAT.

Image Channel Data Type	Description
CL_UNORM_INT16	Each channel component is a normalized unsigned 16-bit integer value
CL_FLOAT	Each channel component is a single precision floating- point value

This extension adds the following new image format to the minimum list of supported image formats described in *table 5.8*:

Table 43. Required Image Formats for cl_khr_depth_images

num_channels	channel_order	channel_data_type
1	CL_DEPTH	CL_UNORM_INT CL_FLOAT

NOTE:

Depth image objects can be initialized, read and written using the appropriate CL APIs i.e. clEnqueueReadImage, clEnqueueWriteImage, clEnqueueCopyImage, clEnqueueCopyImage, clEnqueueCopyBufferToImage, clEnqueueMapImage and clEnqueueFillImage.

For clEnqueueFillImage, the fill color is a 4-component value where the R component refers to the depth value if the image format is CL_DEPTH. The fill color will be converted to the appropriate image channel format and order associated with image.

Update text that describes arg value argument to clSetKernelArg with the following:

If the kernel argument is declared to be of type image2d_depth_t or image2d_array_depth t, the arg_value entry will be a pointer to a depth image or depth image array object.

Add the following error condition for clSetKernelArg:

CL_INVALID_MEM_OBJECT for an argument declared to be a depth image or a depth image array and the argument value specified in arg_value does not follow the rules described above for a depth memory object or memory array object argument.

16.3. Additions to Chapter 6 of the OpenCL 1.2 Specification

Add the following new data types to *table 6.3* in *section 6.1.3* of the OpenCL 1.2 specification:

Туре	Description
image2d_depth_t	A 2D depth image. Refer to <i>section 6.12.14</i> for a detailed description of the built-in functions that use this type.
image2d_array_depth_t	A 2D depth image array. Refer to <i>section 6.12.14</i> for a detailed description of the built-in functions that use this type.

Add the following to the bulleted list in section 6.12.14.1.1 - Determining the border color:

• If the image channel order is CL_DEPTH, the border value is 0.0f.

Add the following built-in functions to section 6.12.14.2 - Built-in Image Read Functions:

Function	Description
float read_imagef(read_only image2d_depth_t image, sampler_t sampler, int2 coord) float read_imagef(read_only image2d_depth_t image, sampler_t sampler, float2 coord)	Use the coordinate (coord.x, coord.y) to do an element lookup in the 2D depth image object specified by image. read_imagef returns a floating-point value in the range [0.0, 1.0] for depth image objects created with image_channel_data_type set to CL_UNORM_INT16 or CL_UNORM_INT24. read_imagef returns a floating-point value for depth image objects created with image_channel_data_type set to CL_FLOAT. The read_imagef calls that take integer coordinates must use a sampler with filter mode set to CLK_FILTER_NEAREST, normalized coordinates set to CLK_NORMALIZED_COORDS_FALSE and addressing mode set to CLK_ADDRESS_CLAMP_TO_EDGE, CLK_ADDRESS_CLAMP or CLK_ADDRESS_NONE; otherwise the values returned are undefined. Values returned by read_imagef for depth image objects with image_channel_data_type values not specified in the description above are undefined.
float read_imagef(read_only image2d_array_depth_t image, sampler_t sampler, int4 coord) float read_imagef(read_only image2d_array_depth_t image, sampler_t sampler, float4 coord)	Use coord.xy to do an element lookup in the 2D image identified by coord.z in the 2D depth image array specified by image. read_imagef returns a floating-point value in the range [0.0, 1.0] for depth image objects created with image_channel_data_type set to CL_UNORM_INT16 or CL_UNORM_INT24. read_imagef returns a floating-point value for depth image objects created with image_channel_data_type set to CL_FLOAT. The read_imagef calls that take integer coordinates must use a sampler with filter mode set to CLK_FILTER_NEAREST, normalized coordinates set to CLK_NORMALIZED_COORDS_FALSE and addressing mode set to CLK_ADDRESS_CLAMP_TO_EDGE, CLK_ADDRESS_CLAMP or CLK_ADDRESS_NONE; otherwise the values returned are undefined. Values returned by read_imagef for image objects with image_channel_data_type values not specified in the description above are undefined.

Add the following built-in functions to section 6.12.14.3 - Built-in Image Sampler-less Read Functions:

Function	Description
float read_imagef (image2d_depth_t image, int2 coord)	Use the coordinate (<i>coord.x</i> , <i>coord.y</i>) to do an element lookup in the 2D depth image object specified by <i>image</i> .
	read_imagef returns a floating-point value in the range [0.0, 1.0] for depth image objects created with image_channel_data_type set to CL_UNORM_INT16 or CL_UNORM_INT24.
	read_imagef returns a floating-point value for depth image objects created with <i>image_channel_data_type</i> set to CL_FLOAT.
	Values returned by read_imagef for image objects with image_channel_data_type values not specified in the description above are undefined.
float read_imagef(image2d_array_depth_t image, int4 coord)	Use <i>coord.xy</i> to do an element lookup in the 2D image identified by <i>coord.z</i> in the 2D depth image array specified by <i>image</i> .
	read_imagef returns a floating-point value in the range [0.0, 1.0] for depth image objects created with image_channel_data_type set to CL_UNORM_INT16 or CL_UNORM_INT24.
	read_imagef returns a floating-point value for depth image objects created with <i>image_channel_data_type</i> set to CL_FLOAT.
	Values returned by read_imagef for image objects with image_channel_data_type values not specified in the description above are undefined.

Add the following built-in functions to section 6.12.14.4 – Built-in Image Write Functions:

Function	Description
void write_imagef (image2d_depth_t image, int2 coord, float depth)	Write <i>depth</i> value to location specified by <i>coord.xy</i> in the 2D depth image object specified by <i>image</i> . Appropriate data format conversion to the specified image format is done before writing the depth value. <i>coord.x</i> and <i>coord.y</i> are considered to be unnormalized coordinates, and must be in the range [0, image width-1], and [0, image height-1], respectively.
	write_imagef can only be used with image objects created with image_channel_data_type set to CL_UNORM_INT16, CL_UNORM_INT24 or CL_FLOAT. Appropriate data format conversion will be done to convert depth value from a floating-point value to actual data format associated with the image.
	The behavior of write_imagef , write_imagei and write_imageui for image objects created with <i>image_channel_data_type</i> values not specified in the description above or with (<i>x</i> , <i>y</i>) coordinate values that are not in the range [0, image width-1] and [0, image height-1], respectively, is undefined.
void write_imagef(image2d_array_depth_t image, int4 coord, float depth)	Write <i>depth</i> value to location specified by <i>coord.xy</i> in the 2D image identified by <i>coord.z</i> in the 2D depth image array specified by <i>image</i> . Appropriate data format conversion to the specified image format is done before writing the depth value. <i>coord.x</i> , <i>coord.y</i> and <i>coord.z</i> are considered to be unnormalized coordinates, and must be in the range [0, image width-1], [0, image height-1], and [0, image number of layers-1], respectively.
	write_imagef can only be used with image objects created with image_channel_data_type set to CL_UNORM_INT16, CL_UNORM_INT24 or CL_FLOAT. Appropriate data format conversion will be done to convert depth valye from a floating-point value to actual data format associated with the image.
	The behavior of write_imagef , write_imagei and write_imageui for image objects created with <i>image_channel_data_type</i> values not specified in the description above or with (<i>x</i> , <i>y</i> , <i>z</i>) coordinate values that are not in the range [0, image width-1], [0, image height-1], [0, image number of layers-1], respectively, is undefined.

Add the following built-in functions to section 6.12.14.5 – Built-in Image Query Functions:

Function	Description
<pre>int get_image_width(image2d_depth_t image) int get_image_width(image2d_array_dept h_t image)</pre>	
<pre>int get_image_height(image2d_depth_t image) int get_image_height(image2d_array_dep th_t image)</pre>	Return the image height in pixels.
int get_image_channel_data_type(image 2d_depth_t image) int get_image_channel_data_type(image 2d_array_depth_t image)	Return the channel data type. Valid values are: CLK_UNORM_INT16 CLK_FLOAT
<pre>int get_image_channel_order(image2d_d epth_t image) int get_image_channel_order(image2d_a rray_depth_t image)</pre>	Return the image channel order. Valid values are: CLK_DEPTH
<pre>int2 get_image_dim(image2d_depth_t image) int2 get_image_dim(image2d_array_depth _t image)</pre>	Return the 2D image width and height as an int2 type. The width is returned in the x component, and the height in the y component.
size_t get_image_array_size(image2d_array _depth_t image)	Return the number of images in the 2D image array.

Add the following text below the table in section 6.12.14.6 - Mapping image channels to color values returned by read_image and color values passed to write_image to image channels:

For CL_DEPTH images, a scalar value is returned by **read_imagef** or supplied to **write_imagef**.

Chapter 17. Sharing OpenGL and OpenGL ES Depth and Depth-Stencil Images

This section describes the **cl_khr_gl_depth_images** extension. The **cl_khr_gl_depth_images** extends OpenCL / OpenGL sharing (the cl_khr_gl_sharing_extension) defined in Creating OpenCL Memory Objects from OpenGL Objects to allow an OpenCL image to be created from an OpenGL depth or depth-stencil texture.

17.1. General information

17.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

17.2. Additions to Chapter 5 of the OpenCL 2.2 Specification

The **cl_khr_gl_depth_images** extension extends OpenCL / OpenGL sharing by allowing an OpenCL depth image to be created from an OpenGL depth or depth-stencil texture. Depth images with an image channel order of CL_DEPTH_STENCIL can only be created using the **clCreateFromGLTexture** API.

This extension adds the following new image format for depth-stencil images to *table 5.6 and 5.7* of the OpenCL 2.2 specification.

Enum values that can be specified in channel_order

CL_DEPTH_STENCIL. This format can only be used if channel data type = CL_UNORM_INT24 or CL_FLOAT.

Image Channel Data Type	Description
CL_UNORM_INT24	Each channel component is a normalized unsigned 24-bit integer value
CL_FLOAT	Each channel component is a single precision floating-point value

This extension adds the following new image format to the minimum list of supported image formats described in *tables 5.8.a* and *5.8.b*.

Table 44. Required Image Formats for cl_khr_gl_depth_images

1	CL_DEPTH_STENCIL	CL_UNORM_INT24	read only
		CL_FLOAT	

For the image format given by channel order of CL_DEPTH_STENCIL and channel data type of CL_UNORM_INT24, the depth is stored as an unsigned normalized 24-bit value.

For the image format given by channel order of CL DEPTH_STENCIL and channel data type of CL_FLOAT, each pixel is two 32-bit values. The depth is stored as a single precision floating-point value followed by the stencil which is stored as a 8-bit integer value.

The stencil value cannot be read or written using the **read_imagef** and **write_imagef** built-in functions in an OpenCL kernel.

Depth image objects with an image channel order equal to CL_DEPTH_STENCIL cannot be used as arguments to clEnqueueReadImage, clEnqueueWriteImage, clEnqueueCopyImage, clEnqueueCopyImage, clEnqueueMapImage and clEnqueueFillImage and will return a CL_INVALID_OPERATION error.

17.3. Additions to the OpenCL Extension Specification

The following new image formats are added to the table of OpenGL internal formats and corresponding OpenCL internal formats in the OpenCL extension specification. If an OpenGL texture object with an internal format in this table is successfully created by OpenGL, then there is guaranteed to be a mapping to one of the corresponding OpenCL image format(s) in that table.

GL internal format	CL image format (channel order, channel data type)
GL_DEPTH_COMPONENT32F	CL_DEPTH, CL_FLOAT
GL_DEPTH_COMPONENT16	CL_DEPTH, CL_UNORM_INT16
GL_DEPTH24_STENCIL8	CL_DEPTH_STENCIL, CL_UNORM_INT24
GL_DEPTH32F_STENCIL8	CL_DEPTH_STENCIL, CL_FLOAT

Chapter 18. Creating OpenCL Memory Objects from OpenGL MSAA Textures

This extension extends the OpenCL / OpenGL sharing (the cl_khr_gl_sharing_extension) defined in Creating OpenCL Memory Objects from OpenGL Objects to allow an OpenCL image to be created from an OpenGL multi-sampled (a.k.a. MSAA) texture (color or depth).

This extension name is **cl_khr_gl_msaa_sharing**. This extension requires **cl_khr_gl_depth_images**.

18.1. General information

18.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

18.2. Additions to the OpenCL Extension Specification

Allow *texture_target* argument to **clCreateFromGLTexture** to be GL_TEXTURE_2D_MULTISAMPLE or GL_TEXTURE_2D_MULTISAMPLE_ARRAY.

If texture_target is GL_TEXTURE_2D_MULTISAMPLE, clCreateFromGLTexture creates an OpenCL 2D multi-sample image object from an OpenGL 2D multi-sample texture.

If texture_target is GL_TEXTURE_2D_MULTISAMPLE_ARRAY, clCreateFromGLTexture creates an OpenCL 2D multi-sample array image object from an OpenGL 2D multi-sample texture.

Multi-sample OpenCL image objects can only be read from a kernel. Multi-sample OpenCL image objects cannot be used as arguments to clEnqueueReadImage , clEnqueueWriteImage, clEnqueueCopyImage, clEnqueueCopyImageToBuffer, clEnqueueCopyBufferToImage, clEnqueueMapImage and clEnqueueFillImage and will return a CL_INVALID_OPERATION error.

Add the following entry to the table describing OpenGL texture info that may be queried with clGetGLTextureInfo:

cl_gl_texture_info	Return Type	Info. returned in param_value
CL_GL_NUM_SAMPLES	GLsizei	The <i>samples</i> argument passed to glTexImage2DMultisample or glTexImage3DMultisample . If <i>image</i> is not a MSAA texture, 1 is returned.

18.3. Additions to Chapter 5 of the OpenCL 2.2 Specification

The formats described in tables 5.8.a and 5.8.b of the OpenCL 2.2 specification and the additional formats described in required image formats for cl_khr_gl_depth_images also support OpenCL images created from a OpenGL multi-sampled color or depth texture.

Update text that describes arg value argument to clSetKernelArg with the following:

"If the argument is a multi-sample 2D image, the *arg_value* entry must be a pointer to a multi-sample image object. If the argument is a multi-sample 2D depth image, the *arg_value* entry must be a pointer to a multi-sample depth image object. If the argument is a multi-sample 2D image array, the *arg_value* entry must be a pointer to a multi-sample image array object. If the argument is a multi-sample 2D depth image array, the *arg_value* entry must be a pointer to a multi-sample depth image array object."

Updated error code text for clSetKernelArg is:

Add the following text:

"CL_INVALID_MEM_OBJECT for an argument declared to be a multi-sample image, multi-sample image array, multi-sample depth image or a multi-sample depth image array and the argument value specified in *arg_value* does not follow the rules described above for a depth memory object or memory array object argument."

18.4. Additions to Chapter 6 of the OpenCL 2.2 Specification

Add the following new data types to table 6.3 in section 6.1.3 of the OpenCL 2.2 specification:

Туре	Description
image2d_msaa_t	A 2D multi-sample color image. Refer to <i>section</i> 6.13.14 for a detailed description of the built-in functions that use this type.
image2d_array_msaa_t	A 2D multi-sample color image array. Refer to section 6.13.14 for a detailed description of the built-in functions that use this type.
image2d_msaa_depth_t	A 2D multi-sample depth image. Refer to <i>section</i> 6.13.14 for a detailed description of the built-in functions that use this type.
image2d_array_msaa_depth_t	A 2D multi-sample depth image array. Refer to section 6.13.14 for a detailed description of the built-in functions that use this type.

Add the following built-in functions to section 6.13.14.3—Built-in Image Sampler-less Read Functions:

```
float4 read_imagef(
   image2d_msaa_t image,
   int2 coord,
   int sample)
```

Use the coordinate (coord.x, coord.y) and sample to do an element lookup in the 2D image object specified by image.

read_imagef returns floating-point values in the range [0.0 ... 1.0] for image objects created with *image_channel_data_type* set to one of the pre-defined packed formats or CL_UNORM_INT8, or CL_UNORM_INT16.

read_imagef returns floating-point values in the range [-1.0 ... 1.0] for image objects created with *image_channel_data_type* set to CL_SNORM_INT8, or CL_SNORM_INT16.

read_imagef returns floating-point values for image objects created with *image_channel_data_type* set to CL_HALF_FLOAT or CL_FLOAT.

Values returned by **read_imagef** for image objects with *image_channel_data_type* values not specified in the description above are undefined.

Use the coordinate (coord.x, coord.y) and sample to do an element lookup in the 2D image object specified by image.

read_imagei and **read_imageui** return unnormalized signed integer and unsigned integer values respectively. Each channel will be stored in a 32-bit integer.

read_imagei can only be used with image objects created with *image_channel_data_type* set to one of the following values:

- CL_SIGNED_INT8,
- CL_SIGNED_INT16, and
- CL_SIGNED_INT32.

If the *image_channel_data_type* is not one of the above values, the values returned by **read_imagei** are undefined.

read_imageui can only be used with image objects created with *image_channel_data_type* set to one of the following values:

- CL_UNSIGNED_INT8,
- CL_UNSIGNED_INT16, and
- CL UNSIGNED INT32.

If the *image_channel_data_type* is not one of the above values, the values returned by **read_imageui** are undefined.

Use *coord.xy* and *sample* to do an element lookup in the 2D image identified by *coord.z* in the 2D image array specified by *image*.

read_imagef returns floating-point values in the range [0.0 ... 1.0] for image objects created with *image_channel_data_type* set to one of the pre-defined packed formats or CL_UNORM_INT8, or CL_UNORM_INT16.

read_imagef returns floating-point values in the range [-1.0 ... 1.0] for image objects created with *image_channel_data_type* set to CL_SNORM_INT8, or CL_SNORM_INT16.

read_imagef returns floating-point values for image objects created with *image_channel_data_type* set to CL_HALF_FLOAT or CL_FLOAT.

Values returned by **read_imagef** for image objects with *image_channel_data_type* values not specified in the description above are undefined.

Use *coord.xy* and *sample* to do an element lookup in the 2D image identified by *coord.z* in the 2D image array specified by *image*.

read_imagei and **read_imageui** return unnormalized signed integer and unsigned integer values respectively. Each channel will be stored in a 32-bit integer.

read_imagei can only be used with image objects created with *image_channel_data_type* set to one of the following values:

- CL_SIGNED_INT8,
- CL_SIGNED_INT16, and
- CL_SIGNED_INT32.

If the *image_channel_data_type* is not one of the above values, the values returned by **read_imagei** are undefined.

read_imageui can only be used with image objects created with *image_channel_data_type* set to one of the following values:

- CL_UNSIGNED_INT8,
- CL_UNSIGNED_INT16, and
- CL_UNSIGNED_INT32.

If the *image_channel_data_type* is not one of the above values, the values returned by **read_imageui** are undefined.

Use the coordinate (coord.x, coord.y) and sample to do an element lookup in the 2D depth image object specified by image.

read_imagef returns a floating-point value in the range [0.0 ... 1.0] for depth image objects created with *image_channel_data_type* set to CL_UNORM_INT16 or CL_UNORM_INT24.

read_imagef returns a floating-point value for depth image objects created with *image_channel_data_type* set to CL_FLOAT.

Values returned by **read_imagef** for image objects with *image_channel_data_type* values not specified in the description above are undefined.

Use *coord.xy* and *sample* to do an element lookup in the 2D image identified by *coord.z* in the 2D depth image array specified by *image*.

read_imagef returns a floating-point value in the range [0.0 ... 1.0] for depth image objects created with *image_channel_data_type* set to CL_UNORM_INT16 or CL_UNORM_INT24.

read_imagef returns a floating-point value for depth image objects created with *image_channel_data_type* set to CL_FLOAT.

Values returned by **read_imagef** for image objects with *image_channel_data_type* values not specified in the description above are undefined.

Note: When a multisample image is accessed in a kernel, the access takes one vector of integers describing which pixel to fetch and an integer corresponding to the sample numbers describing which sample within the pixel to fetch. sample identifies the sample position in the multi-sample

image.

For best performance, we recommend that *sample* be a literal value so it is known at compile time and the OpenCL compiler can perform appropriate optimizations for multi-sample reads on the device.

No standard sampling instructions are allowed on the multisample image. Accessing a coordinate outside the image and/or a sample that is outside the number of samples associated with each pixel in the image is undefined

Add the following built-in functions to section 6.13.14.5 — Built-in Image Query Functions:

```
int get_image_width(image2d_msaa_t image)
int get_image_width(image2d_array_msaa_t image)
int get_image_width(image2d_msaa_depth_t image)
int get_image_width(image2d_array_msaa_depth_t image)
```

Return the image width in pixels.

```
int get_image_height(image2d_msaa_t image)
int get_image_height(image2d_array_msaa_t image)
int get_image_height(image2d_msaa_depth_t image)
int get_image_height(image2d_array_msaa_depth_t image)
```

Return the image height in pixels.

```
int get_image_channel_data_type(image2d_msaa_t image)
int get_image_channel_data_type(image2d_array_msaa_t image)
int get_image_channel_data_type(image2d_msaa_depth_t image)
int get_image_channel_data_type(image2d_array_msaa_depth_t image)
```

Return the channel data type.

```
int get_image_channel_order(image2d_msaa_t image)
int get_image_channel_order(image2d_array_msaa_t image)
int get_image_channel_order(image2d_msaa_depth_t image)
int get_image_channel_order(image2d_array_msaa_depth_t image)
```

Return the image channel order.

```
int2 get_image_dim(image2d_msaa_t image)
int2 get_image_dim(image2d_array_msaa_t image)
int2 get_image_dim(image2d_msaa_depth_t image)
int2 get_image_dim(image2d_array_msaa_depth_t image)
```

Return the 2D image width and height as an int2 type. The width is returned in the *x* component, and the height in the *y* component.

```
size_t get_image_array_size(image2d_array_msaa_depth_t image)
```

Return the number of images in the 2D image array.

```
int get_image_num_samples(image2d_msaa_t image)
int get_image_num_samples(image2d_array_msaa_t image)
int get_image_num_samples(image2d_msaa_depth_t image)
int get_image_num_samples(image2d_array_msaa_depth_t image)
```

Return the number of samples in the 2D MSAA image

Chapter 19. Creating OpenCL Event Objects from EGL Sync Objects

19.1. Overview

This section describes the **cl_khr_egl_event** extension. This extension allows creating OpenCL event objects linked to EGL fence sync objects, potentially improving efficiency of sharing images and buffers between the two APIs. The companion **EGL_KHR_cl_event** extension provides the complementary functionality of creating an EGL sync object from an OpenCL event object.

19.2. General information

19.2.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

19.3. New Procedures and Functions

19.4. New Tokens

Returned by clCreateEventFromEGLSyncKHR if *sync* is not a valid EGLSyncKHR handle created with respect to EGLDisplay *display*:

```
CL_INVALID_EGL_OBJECT_KHR
```

Returned by **clGetEventInfo** when *param_name* is CL_EVENT_COMMAND_TYPE:

```
CL_COMMAND_EGL_FENCE_SYNC_OBJECT_KHR
```

19.5. Additions to Chapter 5 of the OpenCL 2.2 Specification

Add following to the fourth paragraph of *section 5.11* (prior to the description of **clWaitForEvents**):

"Event objects can also be used to reflect the status of an EGL fence sync object. The sync object in turn refers to a fence command executing in an EGL client API command stream. This provides another method of coordinating sharing of EGL / EGL client API objects with OpenCL. Completion of EGL / EGL client API commands may be determined by placing an EGL fence command after commands using eglCreateSyncKHR, creating an event from the resulting EGL sync object using clCreateEventFromEGLSyncKHR and then specifying it in the <code>event_wait_list</code> of a clEnqueueAcquire*** command. This method may be considerably more efficient than calling operations like glFinish, and is referred to as <code>explicit synchronization</code>. The application is responsible for ensuring the command stream associated with the EGL fence is flushed to ensure the CL queue is submitted to the device. Explicit synchronization is most useful when an EGL client API context bound to another thread is accessing the memory objects."

Add CL_COMMAND_EGL_FENCE_SYNC_OBJECT_KHR to the valid *param_value* values returned by **clGetEventInfo** for *param_name* CL_EVENT_COMMAND_TYPE (in the third row and third column of *table 5.22*).

Add new *subsection 5.11.2*:

"`5.11.2 Linking Event Objects to EGL Synchronization Objects

An event object may be created by linking to an EGL **sync object**. Completion of such an event object is equivalent to waiting for completion of the fence command associated with the linked EGL sync object.

The function

creates a linked event object.

context is a valid OpenCL context created from an OpenGL context or share group, using the **cl_khr_gl_sharing** extension.

sync is the name of a sync object of type EGL_SYNC_FENCE_KHR created with respect to EGLDisplay *display*.

clCreateEventFromEGLSyncKHR returns a valid OpenCL event object and *errcode_ret* is set to CL_SUCCESS if the event object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode_ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context, or was not created from a GL context.
- CL_INVALID_EGL_OBJECT_KHR if *sync* is not a valid EGLSyncKHR object of type EGL_SYNC_FENCE_KHR created with respect to EGLDisplay *display*.

The parameters of an event object linked to an EGL sync object will return the following values when queried with **clGetEventInfo**:

- The CL_EVENT_COMMAND_QUEUE of a linked event is NULL, because the event is not associated with any OpenCL command queue.
- The CL_EVENT_COMMAND_TYPE of a linked event is CL_COMMAND_EGL_FENCE_SYNC_OBJECT_KHR, indicating that the event is associated with a EGL sync object, rather than an OpenCL command.
- The CL_EVENT_COMMAND_EXECUTION_STATUS of a linked event is either CL_SUBMITTED, indicating that the fence command associated with the sync object has not yet completed, or CL_COMPLETE, indicating that the fence command has completed.

clCreateEventFromEGLSyncKHR performs an implicit **clRetainEvent** on the returned event object. Creating a linked event object also places a reference on the linked EGL sync object. When the event object is deleted, the reference will be removed from the EGL sync object.

Events returned from **clCreateEventFromEGLSyncKHR** may only be consumed by **clEnqueueAcquire***** commands. Passing such events to any other CL API that enqueues commands will generate a CL_INVALID_EVENT error.`"

19.6. Additions to the OpenCL Extension Specification

Replace the second paragraph of Synchronizing OpenCL and OpenGL Access to Shared Objects with:

"`Prior to calling **clEnqueueAcquireGLObjects**, the application must ensure that any pending EGL or EGL client API operations which access the objects specified in *mem_objects* have completed.

If the **cl_khr_egl_event** extension is supported and the EGL context in question supports fence sync objects, *explicit synchronization* can be achieved as set out in *section 5.7.1*.

If the **cl_khr_egl_event** extension is not supported, completion of EGL client API commands may be determined by issuing and waiting for completion of commands such as glFinish or vgFinish on all client API contexts with pending references to these objects. Some implementations may offer other efficient synchronization methods. If such methods exist they will be described in platform-specific documentation.

Note that no synchronization methods other than glFinish and vgFinish are portable between all EGL client API implementations and all OpenCL implementations. While this is the only way to ensure completion that is portable to all platforms, these are expensive operation and their use should be avoided if the cl_khr_egl_event extension is supported on a platform.`"

19.7. Issues

Most issues are shared with **cl_khr_gl_event** and are resolved as described in that extension.

1. Should we support implicit synchronization?

RESOLVED: No, as this may be very difficult since the synchronization would not be with EGL, it would be with currently bound EGL client APIs. It would be necessary to know which client APIs might be bound, to validate that they're associated with the EGLDisplay associated with the

OpenCL context, and to reach into each such context.

2. Do we need to have typedefs to use EGL handles in OpenCL?

RESOLVED Using typedefs for EGL handles.

3. Should we restrict which CL APIs can be used with this cl_event?

RESOLVED Use is limited to clEnqueueAcquire*** calls only.

4. What is the desired behaviour for this extension when EGLSyncKHR is of a type other than EGL_SYNC_FENCE_KHR?

RESOLVED This extension only requires support for EGL_SYNC_FENCE_KHR. Support of other types is an implementation choice, and will result in CL_INVALID_EGL_OBJECT_KHR if unsupported.

Chapter 20. Creating OpenCL Memory Objects from EGL Images

20.1. Overview

This section describes the **cl_khr_egl_image** extension. This extension provides a mechanism to creating OpenCL memory objects from from EGLImages.

20.2. General information

20.2.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

20.3. New Procedures and Functions

```
cl_mem clCreateFromEGLImageKHR(cl_context context,
                               CLeglDisplayKHR display,
                               CLeglImageKHR image,
                               cl_mem_flags flags,
                               const cl_eql_image_properties_khr *properties,
                               cl_int *errcode_ret);
cl_int clEnqueueAcquireEGLObjectsKHR(cl_command_queue command_queue,
                                     cl_uint num_objects,
                                     const cl_mem *mem_objects,
                                     cl_uint num_events_in_wait_list,
                                     const cl_event *event_wait_list,
                                     cl_event *event)
cl_int clEnqueueReleaseEGLObjectsKHR(cl_command_queue command_queue,
                                     cl_uint num_objects,
                                     const cl_mem *mem_objects,
                                     cl_uint num_events_in_wait_list,
                                     const cl_event *event_wait_list,
                                     cl_event *event)
```

20.4. New Tokens

New error codes:

```
CL_EGL_RESOURCE_NOT_ACQUIRED_KHR
CL_INVALID_EGL_OBJECT_KHR
```

New command types:

```
CL_COMMAND_ACQUIRE_EGL_OBJECTS_KHR
CL_COMMAND_RELEASE_EGL_OBJECTS_KHR
```

20.5. Additions to Chapter 5 of the OpenCL 2.2 Specification

In section 5.2.4, add the following text after the paragraph defining clCreateImage:

"`The function

creates an EGLImage target of type cl_mem from the EGLImage source provided as image.

display should be of type EGLDisplay, cast into the type CLeglDisplayKHR.

image should be of type EGLImageKHR, cast into the type CLeglImageKHR. Assuming no errors are generated in this function, the resulting image object will be an EGLImage target of the specified EGLImage *image*. The resulting cl_mem is an image object which may be used normally by all OpenCL operations. This maps to an image2d t type in OpenCL kernel code.

flags is a bit-field that is used to specify usage information about the memory object being created.

The possible values for *flags* are: CL_MEM_READ_ONLY, CL_MEM_WRITE_ONLY and CL_MEM_READ_WRITE.

For OpenCL 1.2 flags also accepts: CL_MEM_HOST_WRITE_ONLY, CL_MEM_HOST_READ_ONLY or CL_MEM_HOST_NO_ACCESS.

This extension only requires support for CL_MEM _READ_ONLY, and for OpenCL 1.2 CL_MEM_HOST_NO_ACCESS. For OpenCL 1.1, a CL_INVALID_OPERATION will be returned for images which do not support host mapping.

If the value passed in *flags* is not supported by the OpenCL implementation it will return CL_INVALID_VALUE. The accepted *flags* may be dependent upon the texture format used.

properties specifies a list of property names and their corresponding values. Each property name is immediately followed by the corresponding desired value. The list is terminated with 0. No properties are currently supported with this version of the extension. *properties* can be NULL.

clCreateFromEGLImageKHR returns a valid non-zero OpenCL image object and *errcode_ret* is set to CL_SUCCESS if the image object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode_ret*:

- CL_INVALID_CONTEXT if *context* is not a valid OpenCL context.
- CL_INVALID_VALUE if *properties* contains invalid values, if *display* is not a valid display object or if *flags* are not in the set defined above.
- CL_INVALID_EGL_OBJECT_KHR if *image* is not a valid EGLImage object.
- CL_IMAGE_FORMAT_NOT_SUPPORTED if the OpenCL implementation is not able to create a cl_mem compatible with the provided CLeglImageKHR for an implementation-dependent reason (this could be caused by, but not limited to, reasons such as unsupported texture formats, etc).
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_INVALID_OPERATION if there are no devices in *context* that support images (i.e. CL_DEVICE_IMAGE_SUPPORT specified in table 4.3 is CL_FALSE) or if the flags passed are not supported for that image type. `"

20.5.1. Lifetime of Shared Objects

An OpenCL memory object created from an EGL image remains valid according to the lifetime behavior as described in EGL_KHR_image_base.

"Any EGLImage siblings exist in any client API context"

For OpenCL this means that while the application retains a reference on the cl_mem (the EGL sibling), the image remains valid.

20.5.2. Synchronizing OpenCL and EGL Access to Shared Objects

In order to ensure data integrity, the application is responsible for synchronizing access to shared CL/EGL objects by their respective APIs. Failure to provide such synchronization may result in race conditions and other undefined behavior including non-portability between implementations.

Prior to calling clEnqueueAcquireEGLObjectsKHR, the application must ensure that any pending operations which access the objects specified in mem_objects have completed. This may be accomplished in a portable way by ceasing all client operations on the resource, and issuing and waiting for completion of a glFinish command on all GL contexts with pending references to these objects. Implementations may offer more efficient synchronization methods, such as synchronization primitives or fence operations.

Similarly, after calling clEnqueueReleaseEGLImageObjects, the application is responsible for ensuring that any pending OpenCL operations which access the objects specified in mem_objects have completed prior to executing subsequent commands in other APIs which reference these objects. This may be accomplished in a portable way by calling clWaitForEvents with the event object returned by clEnqueueReleaseGLObjects, or by calling clFinish. As above, some implementations may offer more efficient methods.

Attempting to access the data store of an EGLImage object after it has been acquired by OpenCL and before it has been released will result in undefined behavior. Similarly, attempting to access a shared EGLImage object from OpenCL before it has been acquired by the OpenCL command queue or after it has been released, will result in undefined behavior.

20.5.3. Sharing memory objects created from EGL resources between EGLDisplays and OpenCL contexts

The function

is used to acquire OpenCL memory objects that have been created from EGL resources. The EGL objects are acquired by the OpenCL context associated with *command_queue* and can therefore be used by all command-queues associated with the OpenCL context.

OpenCL memory objects created from EGL resources must be acquired before they can be used by any OpenCL commands queued to a command-queue. If an OpenCL memory object created from a EGL resource is used while it is not currently acquired by OpenCL, the call attempting to use that OpenCL memory object will return CL_EGL_RESOURCE_NOT_ACQUIRED_KHR.

command_queue is a valid command-queue.

num_objects is the number of memory objects to be acquired in *mem_objects*.

mem_objects is a pointer to a list of OpenCL memory objects that were created from EGL resources, within the context associate with command_queue.

event_wait_list and num_events_in_wait_list specify events that need to complete before this particular command can be executed. If event_wait_list is NULL, then this particular command does not wait on any event to complete. If event_wait_list is NULL, num_events_in_wait_list must be 0. If event_wait_list is not NULL, the list of events pointed to by event_wait_list must be valid and num_events_in_wait_list must be greater than 0. The events specified in event_wait_list act as synchronization points.

event returns an event object that identifies this command and can be used to query or wait for this command to complete. If *event* is NULL or the enqueue is unsuccessful, no event will be created and

therefore it will not be possible to query the status of this command or to wait for this command to complete. If <code>event_wait_list</code> and <code>event</code> are not <code>NULL</code>, <code>event</code> must not refer to an element of the <code>event_wait_list</code> array.

clEnqueueAcquireEGLObjectsKHR returns CL_SUCCESS if the function is executed successfully. If *num_objects* is 0 and *mem_objects* is NULL then the function does nothing and returns CL_SUCCESS. Otherwise it returns one of the following errors:

- CL_INVALID_VALUE if num_objects is zero and mem_objects is not a NULL value or if num_objects
 o and mem_objects is NULL.
- CL_INVALID_MEM_OBJECT if memory objects in *mem_objects* are not valid OpenCL memory objects in the context associated with *command_queue*.
- CL_INVALID_EGL_OBJECT_KHR if memory objects in mem_objects have not been created from EGL resources.
- CL_INVALID_COMMAND_QUEUE if command_queue is not a valid command-queue.
- CL_INVALID_EVENT_WAIT_LIST if event_wait_list is NULL and num_events_in_wait_list > 0, or event_wait_list is not NULL and num_events_in_wait_list is 0, or if event objects in event_wait_list are not valid events.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

is used to release OpenCL memory objects that have been created from EGL resources. The EGL objects are released by the OpenCL context associated with <command_queue>.

OpenCL memory objects created from EGL resources which have been acquired by OpenCL must be released by OpenCL before they may be accessed by EGL or by EGL client APIs. Accessing a EGL resource while its corresponding OpenCL memory object is acquired is in error and will result in undefined behavior, including but not limited to possible OpenCL errors, data corruption, and program termination.

command_queue is a valid command-queue.

num_objects is the number of memory objects to be acquired in *mem_objects*.

mem_objects is a pointer to a list of OpenCL memory objects that were created from EGL resources,

within the context associate with command_queue.

event_wait_list and num_events_in_wait_list specify events that need to complete before this particular command can be executed. If event_wait_list is NULL, then this particular command does not wait on any event to complete. If event_wait_list is NULL, num_events_in_wait_list must be 0. If event_wait_list is not NULL, the list of events pointed to by event_wait_list must be valid and num_events_in_wait_list must be greater than 0. The events specified in event_wait_list act as synchronization points.

event returns an event object that identifies this command and can be used to query or wait for this command to complete. If event is NULL or the enqueue is unsuccessful, no event will be created and therefore it will not be possible to query the status of this command or to wait for this command to complete. If event_wait_list and event are not NULL, event must not refer to an element of the event_wait_list array.

clEnqueueReleaseEGLObjectsKHR returns CL_SUCCESS if the function is executed successfully. If *num_objects* is 0 and *mem_objects* is NULL then the function does nothing and returns CL_SUCCESS. Otherwise it returns one of the following errors:

- CL_INVALID_VALUE if *num_objects* is zero and *mem_objects* is not a NULL value or if num_objects > 0 and mem_objects is NULL.
- CL_INVALID_MEM_OBJECT if memory objects in *mem_objects* are not valid OpenCL memory objects in the context associated with *command_queue*.
- CL_INVALID_EGL_OBJECT_KHR if memory objects in *mem_objects* have not been created from EGL resources.
- CL_INVALID_COMMAND_QUEUE if command_queue is not a valid command-queue.
- CL_INVALID_EVENT_WAIT_LIST if event_wait_list is NULL and num_events_in_wait_list > 0, or event_wait_list is not NULL and num_events_in_wait_list is 0, or if event objects in event_wait_list are not valid events.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

20.5.4. Event Command Types for Sharing memory objects created from EGL resources

The following table describes the event command types for the OpenCL commands to acquire and release OpenCL memory objects that have been created from EGL resources:

Table 45. List of supported event command types

Events Created By	Event Command Type
clEnqueueAcquireEGLObjectsKHR	CL_COMMAND_ACQUIRE_EGL_OBJECTS_KHR
clEnqueueReleaseEGLObjectsKHR	CL_COMMAND_RELEASE_EGL_OBJECTS_KHR

20.6. Issues

- 1. This extension does not support reference counting of the images, so the onus is on the application to behave sensibly and not release the underlying cl_mem object while the EGLImage is still being used.
- 2. In order to ensure data integrity, the application is responsible for synchronizing access to shared CL/EGL image objects by their respective APIs. Failure to provide such synchronization may result in race conditions and other undefined behavior. This may be accomplished by calling clWaitForEvents with the event objects returned by any OpenCL commands which use the shared image object or by calling clFinish.
- 3. Currently CL_MEM_READ_ONLY is the only supported flag for *flags*.

RESOLVED: Implementation will now return an error if writing to a shared object that is not supported rather than disallowing it entirely.

- 4. Currently restricted to 2D image objects.
- 5. What should happen for YUV color-space conversion, multi plane images, and chroma-siting, and channel mapping?

RESOLVED: YUV is no longer explicitly described in this extension. Before this removal the behavior was dependent on the platform. This extension explicitly leaves the YUV layout to the platform and EGLImage source extension (i.e. is implementation specific). Colorspace conversion must be applied by the application using a color conversion matrix.

The expected extension path if YUV color-space conversion is to be supported is to introduce a YUV image type and provide overloaded versions of the read_image built-in functions.

Getting image information for a YUV image should return the original image size (non quantized size) when all of Y U and V are present in the image. If the planes have been separated then the actual dimensionality of the separated plane should be reported. For example with YUV 4:2:0 (NV12) with a YUV image of 256x256, the Y only image would return 256x256 whereas the UV only image would return 128x128.

6. Should an attribute list be used instead?

RESOLVED: function has been changed to use an attribute list.

7. What should happen for EGLImage extensions which introduce formats without a mapping to an OpenCL image channel data type or channel order?

RESOLVED: This extension does not define those formats. It is expected that as additional EGL extensions are added to create EGL images from other sources, an extension to CL will be introduced where needed to represent those image types.

8. What are the guarantees to synchronization behavior provided by the implementation?

The basic portable form of synchronization is to use a clFinish, as is the case for GL interop. In addition implementations which support the synchronization extensions cl_khr_egl_event and EGL_KHR_cl_event can interoperate more efficiently as described in those extensions.

Chapter 21. Creating a 2D Image From A Buffer

This section describes the **cl_khr_image2d_from_buffer** extension.

This extension allows a 2D image to be created from an existing OpenCL buffer memory object.

This extension became a core feature in OpenCL 2.0.

21.1. General information

21.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

21.2. Additions to Chapter 4 of the OpenCL 1.2 Specification

The following table entry describes the additions to *table 4.3*, which allows applications to query the configuration information using **clGetDeviceInfo** for an OpenCL device that supports creating a 2D image from a buffer.

cl_device_info	Return Type	Description
CL_DEVICE_IMAGE_ PITCH_ALIGNMENT_KHR	cl_uint	The row pitch alignment size in pixels for images created from a buffer. The value returned must be a power of 2. If the device does not support images, this value should be 0.
CL_DEVICE_IMAGE_BASE_ ADDRESS_ALIGNMENT_KHR	cl_uint	This query should be used when an image is created from a buffer which was created using CL_MEM_USE_HOST_PTR. The value returned must be a power of 2. This query specifies the minimum alignment in pixels of the host_ptr specified to clCreateBuffer. If the device does not support images, this value should be 0.

21.3. Additions to Chapter 5 of the OpenCL 1.2 Specification

Add to Section 5.3.1: Creating Image Objects:

A 2D image can be created from a buffer by specifying a *buffer* object in the *image_desc* passed to **clCreateImage** for an *image_type* equal to **CL_MEM_OBJECT_IMAGE2D**. When the 2D image from buffer is created, the client must specify the width, height and image format (i.e. channel order and channel data type). If these are not specified, **clCreateImage** returns a NULL value with *errcode_ret* set to **CL_INVALID_IMAGE_FORMAT_DESCRIPTOR**. The pitch can be optionally specified. If the pitch is not specified, the pitch is computed as width × bytes per pixel based on the image format.

The pitch specified (or computed if pitch specified is 0) must be a multiple of the maximum of the CL_DEVICE_IMAGE_PITCH_ALIGNMENT_KHR value for all devices in the context associated with the *buffer* that support images. Otherwise, **clCreateImage** returns a NULL value with *errcode_ret* set to CL_INVALID_IMAGE_FORMAT_DESCRIPTOR.

If the *buffer* was created with CL_MEM_USE_HOST_PTR, the *host_ptr* specified to **clCreateBuffer** must be aligned to the maximum of the CL_DEVICE_IMAGE_BASE_ADDRESS_ALIGNMENT_KHR value for all devices in the context associated with the *buffer* that support images. Otherwise, **clCreateImage** returns a NULL value with *errcode_ret* set to CL_INVALID_IMAGE_FORMAT_DESCRIPTOR.

The minimum list of supported image formats described in *table 5.8* of the OpenCL 1.2 specification must be supported for 2D images created from a buffer.

The OpenCL runtime APIs that operate on images (i.e. clEnqueueReadImage, clEnqueueWriteImage, clEnqueueFillImage, clEnqueueCopyImageToBuffer, clEnqueueCopyBufferToImage and clEnqueueMapImage) are supported for a 2D image created from a buffer.

When the contents of a buffer object data store are modified, those changes are reflected in the contents of the 2D image object and vice-versa at corresponding synchronization points. The <code>image_height × image_row_pitch</code> specified in <code>image_desc</code> must be less than or equal to the size of the buffer object data store.



Concurrent reading from, writing to, and copying between both a buffer object and the 2D image object associated with the buffer object is undefined. Only reading from both a buffer object and 2D image object associated with the buffer object is defined. A 2D image and a 2D image created from a buffer use the same image type in OpenCL C (image2d_t). The image built-ins functions described in section 6.12.14.2, 6.12.14.3, 6.12.14.4 and 6.12.14.5 for image2d_t behave the same way for a 2D image and a 2D image from a buffer.

Chapter 22. Local and Private Memory Initialization

Memory is allocated in various forms in OpenCL both explicitly (global memory) or implicitly (local, private memory). This allocation so far does not provide a straightforward mechanism to initialize the memory on allocation. In other words what is lacking is the equivalent of calloc for the currently supported malloc like capability. This functionality is useful for a variety of reasons including ease of debugging, application controlled limiting of visibility to previous contents of memory and in some cases, optimization.

This extension adds support for initializing local and private memory before a kernel begins execution. This extension name is **cl_khr_initialize_memory**.

22.1. General information

22.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

22.2. Additions to Chapter 4 of the OpenCL 2.2 Specification

Add a new context property to table 4.5 in section 4.4.

cl_context_properties enum	Property value	Description
CL_CONTEXT_MEMORY_INITI ALIZE_KHR	cl_context_memory_i nitialize_khr	Describes which memory types for the context must be initialized. This is a bit-field, where the following values are currently supported:
		CL_CONTEXT_MEMORY_INITIALIZE_LOCA L_KHR — Initialize local memory to zeros.
		CL_CONTEXT_MEMORY_INITIALIZE_PRIVA TE_KHR — Initialize private memory to zeros.

22.3. Additions to Chapter 6 of the OpenCL 2.2 Specification

Updates to section 6.9 — Restrictions

If the context is created with CL CONTEXT MEMORY INITIALIZE KHR, appropriate memory locations as specified by the bit-field is initialized with zeroes, prior to the start of execution of any kernel. The driver chooses when, prior to kernel execution, the initialization of local and/or private memory is performed. The only requirement is there should be no values set from outside the context, which can be read during a kernel execution.

Chapter 23. Terminating OpenCL contexts

Today, OpenCL provides an API to release a context. This operation is done only after all queues, memory object, programs and kernels are released, which in turn might wait for all ongoing operations to complete. However, there are cases in which a fast release is required, or release operation cannot be done, as commands are stuck in mid execution. An example of the first case can be program termination due to exception, or quick shutdown due to low power. Examples of the second case are when a kernel is running too long, or gets stuck, or it may result from user action which makes the results of the computation unnecessary.

In many cases, the driver or the device is capable of speeding up the closure of ongoing operations when the results are no longer required in a much more expedient manner than waiting for all previously enqueued operations to finish.

This extension implements a new query to check whether a device can terminate an OpenCL context and adds an API to terminate a context.

The extension name is **cl_khr_terminate_context**.

23.1. General information

23.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

23.2. Additions to Chapter 4 of the OpenCL 2.2 Specification

Add a new device property to table 4.3 in section 4.2.

cl_device_info	Return Type	Description
CL_DEVICE_TERMINATE_CAP ABILITY_KHR	cl_device_terminat e_capability_khr	Describes the termination capability of the OpenCL device. This is a bit-field, where the following values are currently supported:
		CL_DEVICE_TERMINATE_CAPABILITY_CON TEXT_KHR - Indicates that context termination is supported.

Add a new context property to *table 4.5* in *section 4.4*.

cl_context_properties enum	Property value	Description
CL_CONTEXT_TERMINATE_K HR	cl_bool	Specifies whether the context can be terminated. The default value is CL_FALSE.

CL_CONTEXT_TERMINATE_KHR can be specified in the context properties only if all devices associated with the context support the ability to support context termination (i.e. CL_DEVICE_TERMINATE_CAPABILITY_CONTEXT_KHR is set for CL_DEVICE_TERMINATE_CAPABILITY_KHR). Otherwise, context creation fails with error code of CL_INVALID_PROPERTY.

The new function

```
cl_int clTerminateContextKHR(cl context context)
```

terminates all pending work associated with the context and renders all data owned by the context invalid. It is the responsibility of the application to release all objects associated with the context being terminated.

When a context is terminated:

- The execution status of enqueued commands will be CL_TERMINATED_KHR. Event objects can be queried using **clGetEventInfo**. Event callbacks can be registered and registered event callbacks will be called with *event_command_status* set to CL_TERMINATED_KHR. **clWaitForEvents** will return as immediately for commands associated with event objects specified in event_list. The status of user events can be set. Event objects can be retained and released. **clGetEventProfilingInfo** returns CL_PROFILING_INFO_NOT_AVAILABLE.
- The context is considered to be terminated. A callback function registered when the context was created will be called. Only queries, retain and release operations can be performed on the context. All other APIs that use a context as an argument will return CL_CONTEXT_TERMINATED_KHR.
- The contents of the memory regions of the memory objects is undefined. Queries, registering a destructor callback, retain and release operations can be performed on the memory objects.
- Once a context has been terminated, all OpenCL API calls that create objects or enqueue commands will return CL_CONTEXT_TERMINATED_KHR. APIs that release OpenCL objects will continue to operate as though clTerminateContextKHR was not called.
- The behavior of callbacks will remain unchanged, and will report appropriate error, if executing after termination of context. This behavior is similar to enqueued commands, after the command queue has become invalid.

clTerminateContextKHR returns CL_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- CL_INVALID_CONTEXT if *context* is not a valid OpenCL context.
- CL_CONTEXT_TERMINATED_KHR if *context* has already been terminated.
- CL_INVALID_OPERATION if context was not created with CL_CONTEXT_TERMNATE_KHR set to

CL_TRUE.

- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

An implementation that supports this extension must be able to terminate commands currently executing on devices or queued across all command-queues associated with the context that is being terminated. The implementation cannot implement this extension by waiting for currently executing (or queued) commands to finish execution on devices associated with this context (i.e. doing a **clFinish**).

Chapter 24. Standard Portable Intermediate Representation Binaries

This extension adds the ability to create an OpenCL program object from a Standard Portable Intermediate Representation (SPIR) instance. A SPIR instance is a vendor-neutral non-source representation for OpenCL C programs.

The extension name is **cl_khr_spir**. This extension has been superseded by the SPIR-V intermediate representation, which is supported by the **cl_khr_il_program** extension, and is a core feature in OpenCL 2.1.

24.1. General information

24.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

24.2. Additions to Chapter 4 of the OpenCL 2.2 Specification

Add a new device property to table 4.3 in section 4.2:

cl_device_info	Return Type	Description
CL_DEVICE_SPIR_VERSIO NS	char[]	A space separated list of SPIR versions supported by the device.
		For example, returning "1.2" in this query implies that SPIR version 1.2 is supported by the implementation.

24.3. Additions to Chapter 5 of the OpenCL 2.2 Specification

Additions to *section 5.8.1* — Creating Program Objects:

"clCreateProgramWithBinary can be used to load a SPIR binary. Once a program object has been created from a SPIR binary, clBuildProgram can be called to build a program executable or clCompileProgram can be called to compile the SPIR binary."

Modify the CL_PROGRAM_BINARY_TYPE entry in *table 5.14* (**clGetProgramBuildInfo**) to add a potential value CL_PROGRAM_BINARY_TYPE_INTERMEDIATE:

cl_program_build_info	Return Type	Info. returned in param_value
CL_PROGRAM_BINARY_T YPE		CL_PROGRAM_BINARY_TYPE_INTERMEDIATE — An intermediate (non-source) representation for the program is loaded as a binary. The program must be further processed with clCompileProgram or clBuildProgram . If processed with clCompileProgram , the result will be a binary of type CL_PROGRAM_BINARY_TYPE_COMPILED_OBJECT or CL_PROGRAM_BINARY_TYPE_LIBRARY. If processed with clBuildProgram , the result will be a binary of type CL_PROGRAM_BINARY_TYPE_EXECUTABLE.

Additions to section 5.8.4 — Compiler Options:

"The compile option -x spir must be specified to indicate that the binary is in SPIR format, and the compile option -spir-std must be used to specify the version of the SPIR specification that describes the format and meaning of the binary. For example, if the binary is as described in SPIR version 1.2, then -spir-std=1.2 must be specified. Failing to specify these compile options may result in implementation defined behavior."

Additions to section 5.9.3 — Kernel Object Queries:

Modify following text in clGetKernelArgInfo from:

"Kernel argument information is only available if the program object associated with *kernel* is created with **clCreateProgramWithSource** and the program executable is built with the -cl-kernel -arg-info option specified in *options* argument to **clBuildProgram** or **clCompileProgram**."

to:

"Kernel argument information is only available if the program object associated with *kernel* is created with **clCreateProgramWithSource** and the program executable is built with the -cl-kernel -arg-info option specified in *options* argument to **clBuildProgram** or **clCompileProgram**, or if the program object associated with *kernel* is created with **clCreateProgramWithBinary** and the program executable is built with the -cl-kernel-arg-info and --x spir options specified in *options* argument to **clBuildProgram** or **clCompileProgram**."

Chapter 25. Intermediate Language Programs

This section describes the **cl_khr_il_program** extension.

This extension adds the ability to create programs with intermediate language (IL), usually SPIR-V. Further information about the format and contents of SPIR-V may be found in the SPIR-V Specification. Information about how SPIR-V modules behave in the OpenCL environment may be found in the OpenCL SPIR-V Environment Specification.

This functionality described by this extension is a core feature in OpenCL 2.1.

25.1. General information

25.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

25.2. New Procedures and Functions

25.3. New Tokens

Accepted as a new *param_name* argument to **clGetDeviceInfo**:

```
CL_DEVICE_IL_VERSION_KHR
```

Accepted as a new *param_name* argument to **clGetProgramInfo**:

```
CL_PROGRAM_IL_KHR
```

25.4. Additions to Chapter 3 of the OpenCL 2.0 Specification

In section 3.1, replace the fourth paragraph with:

"Programmers provide programs in the form of intermediate language binaries (usually SPIR-V), OpenCL C source strings, or implementation-defined binary objects. The OpenCL platform provides a compiler to translate programs represented as intermediate language binaries or OpenCL C source strings into device program executables. The compiler may be *online* or *offline*. An *online compiler* is available during host program execution using standard APIs. An *offline compiler* is invoked outside of host program control, using platform-specific methods. The OpenCL runtime allows developers to get a previously compiled device program executable and to load and execute a previously compiled device program executable."

25.5. Additions to Chapter 4 of the OpenCL 2.0 Specification

Add a new device property to **Table 4.3** OpenCL Device Queries:

cl_device_info	Return Type	Description
CL_DEVICE_IL_VERSION_KHR	char[]	The intermediate languages that are be supported by clCreateProgramWithILKHR for this device.
		Returns a space separated list of IL version strings of the form:
		<il_prefix>_<major_version>.<minor_version></minor_version></major_version></il_prefix>
		A device that supports the cl_khr_il_program extension must support the "SPIR-V" IL prefix.

25.6. Additions to Chapter 5 of the OpenCL 2.0 Specification

Add to Section 5.8.1: Creating Program Objects:

"The function

creates a new program object for *context* using the *length* bytes of intermediate language pointed to by *il*.

context must be a valid OpenCL context.

il is a pointer to a length-byte block of memory containing intermediate langage.

length is the length of the block of memory pointed to by il.

errcode_ret will return an appropriate error code. If errcode_ret is NULL, no error code is returned.

clCreateProgramWithILKHR returns a valid non-zero program object and *errcode_ret* is set to CL_SUCCESS if the program object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode_ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context
- CL_INVALID_VALUE if *il* is NULL or if *length* is zero.
- CL_INVALID_VALUE if the *length*-byte block of memory pointed to by *il* does not contain well-formed intermediate language.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host."

Add to Section 5.8.2: Building Program Executables:

Add the following to the description of the *options* parameter to **clBuildProgram**:

"Certain options are ignored when *program* is created with IL."

Additionally, replace the error:

• CL_INVALID_OPERATION if *program* was not created with **clCreateProgramWithSource** or **clCreateProgramWithBinary**.

with:

• CL_INVALID_OPERATION if *program* was not created with **clCreateProgramWithSource**, **clCreateProgramWithILKHR** or **clCreateProgramWithBinary**.

Add to Section 5.8.3: Separate Compilation and Linking of Programs:

Add the following to the description of the *options* parameter to **clCompileProgram**:

"Certain options are ignored when program is created with IL."

Additionally, replace the error:

• CL_INVALID_OPERATION if *program* has no source i.e. it has not been created with clCreateProgramWithSource.

with:

 CL_INVALID_OPERATION if program was not created with clCreateProgramWithSource or clCreateProgramWithILKHR. Add to Section 5.8.4.1: Preprocessor Options,

Add to Section 5.8.4.2: Math Intrinsic Options (for -cl-single-precision-constant-only),

Add to Section 5.8.4.3: Optimization Options,

Add to Section 5.8.4.4: Options to Request or Suppress Warnings, and

Add to Section 5.8.4.5: Options Controlling the OpenCL C Version:

"These options are ignored for programs created with IL."

Change one entry and add one new entry to **Table 5.17** *clGetProgramInfo parameter queries*:

cl_program_info	Return Type	Info returned in param_value
CL_PROGRAM_SOURCE	char[]	Return the program source code specified by clCreateProgramWithSource. The source string returned is a concatenation of all source strings specified to clCreateProgramWithSource with a null terminator. The concatenation strips any nulls in the original source strings. If program is created using clCreateProgramWithBinary, clCreateProgramWithBuiltIn Kernels,, or clCreateProgramWithILKHR a null string or the appropriate program source code is returned depending on whether or not the program source code is stored in the binary. The actual number of characters that represents the program source code including the null terminator is returned in param_value_size_ret.

cl_program_info	Return Type	Info returned in param_value
CL_PROGRAM_IL_KHR	unsigned char[]	Returns the program IL for programs created with clCreateProgramWithILKHR.
		If program is created with clCreateProgramWithSource, clCreateProgramWithBinary, or clCreateProgramWithBuiltIn Kernels, the memory pointed to
		by param_value will be unchanged and param_value_size_ret will be set to zero.

Chapter 26. Creating Command Queues with Properties

26.1. Overview

The section describes the **cl_khr_create_command_queue** extension.

This extension allows OpenCL 1.x devices to support an equivalent clCreateCommandQueueWithProperties API that was added in OpenCL 2.0. This allows OpenCL support other optional extensions or features clCreateCommandQueueWithProperties API to specify additional command queue properties that cannot be specified using the OpenCL 1.x clCreateCommandQueue API.

No new command queue properties are required by this extension. Applications may use the existing CL_DEVICE_QUEUE_PROPERTIES query to determine command queue properties that are supported by the device.

OpenCL 2.x devices may support this extension for compatibility. In this scenario, the function added by this extension will have the same capabilities as the core clCreateCommandQueueWithProperties API. Applications that only target OpenCL 2.x devices should use the core OpenCL 2.x clCreateCommandQueueWithProperties API instead of this extension API.

26.2. General information

26.2.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

26.3. New API Functions

```
cl_command_queue clCreateCommandQueueWithPropertiesKHR(
   cl_context context,
   cl_device_id device,
   const cl_queue_properties_khr *properties,
   cl_int *errcode_ret)
```

26.4. New API Types

```
typedef cl_properties cl_queue_properties_khr;
```

26.5. Modifications to the OpenCL 1.2 Specification

(Add to Table 5.2 for CL_QUEUE_PROPERTIES in Section 5.1)

Table 5.2 List of supported param_names by clGetCommandQueueInfo

cl_command_queue_info	Return Type	Information returned in param_value
CL_QUEUE_PROPERTIES	_	Returns the currently specified properties for the command-queue. These properties are specified by the <i>properties</i> argument in clCreateCommandQueue , or by the CL_QUEUE_PROPERTIES property value in clCreateCommandQueueWithPropertiesKH R .

(Add a new Section 5.1.1, Creating Command Queues With Properties)

The function

```
cl_command_queue clCreateCommandQueueWithPropertiesKHR(
    cl_context context,
    cl_device_id device,
    const cl_queue_properties_khr *properties,
    cl_int *errcode_ret)
```

allows creation of a command-queue from an array of properties for the specified device.

context must be a valid OpenCL context.

device must be a device or sub-device associated with *context*. It can either be in the list of devices and sub-devices specified when *context* is created using **clCreateContext** or be a root device with the same device type as specified when *context* is created using **clCreateContextFromType**.

properties specifies a list of properties for the command-queue and their corresponding values. Each property name is immediately followed by the corresponding desired value. The list is terminated with 0. The list of supported properties is described in the table below. If a supported property and its value is not specified in *properties*, its default value will be used. *properties* can be NULL in which case the default values for supported command-queue properties will be used.

Table X.Y List of supported cl_queue_properties_khr values and description

Queue Properties	Property Value	Description
CL_QUEUE_PROPERTIES	cl_bitfield	This is a bitfield and can be set to a combination of the following values:
		CL_QUEUE_OUT_OF_ORDER_ EXEC_MODE_ENABLE - Determines whether the commands queued in the command-queue are executed in-order or out-of-order. If set, the commands in the command-queue are executed out-of-order. Otherwise, commands are executed in-order.
		CL_QUEUE_PROFILING_ENABLE - Enable or disable profiling of commands in the command-queue. If set, the profiling of commands is enabled. Otherwise, profiling of commands is disabled.
		If CL_QUEUE_PROPERTIES is not specified an in-order command queue that does not support profiling of commands is created for the specified device.

errcode_ret will return an appropriate error code. If errcode_ret is NULL, no error code is returned.

clCreateCommandQueueWithPropertiesKHR returns a valid non-zero command-queue and *errcode_ret* is set to CL_SUCCESS if the command-queue is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode_ret*:

- CL_INVALID_CONTEXT if *context* is not a valid context.
- CL_INVALID_DEVICE if *device* is not a valid device or is not associated with *context*.
- CL_INVALID_VALUE if values specified in *properties* are not valid.
- CL_INVALID_QUEUE_PROPERTIES if values specified in *properties* are valid but are not supported by the device.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

Chapter 27. Device Enqueue Local Argument Types

This extension allows arguments to blocks that are passed to the **enqueue_kernel** built-in function to be pointers to any type (built-in or user-defined) in local memory, instead of requiring arguments to blocks to be pointers to void in local memory.

The name of this extension is **cl_khr_device_enqueue_local_arg_types**.

27.1. General information

27.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

27.2. Additions to Chapter 6 of the OpenCL 2.0 C Specification

Modify the second paragraph of Section 6.13.17: Enqueuing Kernels:

"The following table describes the list of built-in functions that can be used to enqueue a kernel. We use the generic type name gentype to indicate the built-in OpenCL C scalar or vector integer or floating-point data types, or any user defined type built from these scalar and vector data types, which can be used as the type of the pointee of the arguments of the kernel enqueue functions listed in table 6.31."

Then, replace all occurrences of local void * in table 6.31 with local gentype *. For example:

Additionally, replace all occurrences of local void* in table 6.33 with local gentype *. For example:

Chapter 28. Subgroups

This section describes the **cl_khr_subgroups** extension.

This extension adds support for implementation-controlled groups of work items, known as subgroups. Subgroups behave similarly to work groups and have their own sets of built-ins and synchronization primitives. Subgroups within a work group are independent, may make forward progress with respect to each other, and may map to optimized hardware structures where that makes sense.

Subgroups were promoted to a core feature in OpenCL 2.1, however note that:

- The subgroup OpenCL C built-in functions described by this extension must still be accessed as an OpenCL C extension in OpenCL 2.1.
- Subgroup independent forward progress is an optional device property in OpenCL 2.1, see CL_DEVICE_SUB_GROUP_INDEPENDENT_FORWARD_PROGRESS.

28.1. General information

28.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

28.2. Additions to Chapter 3 of the OpenCL 2.0 Specification

28.3. Additions to section 3.2 — Execution Model

Within a work-group work-items may be divided into sub-groups. The mapping of work-items to sub-groups is implementation-defined and may be queried at runtime. While sub-groups may be used in multi-dimensional work-groups, each subgroup is 1-dimensional and any given work-item may query which sub-group it is a member of.

Work items are mapped into subgroups through a combination of compile-time decisions and the parameters of the dispatch. The mapping to subgroups is invariant for the duration of a kernel's execution, across dispatches of a given kernel with the same launch parameters, and from one work-group to another within the dispatch (excluding the trailing edge work-groups in the presence of non-uniform work-group sizes). In addition, all sub-groups within a work-group will be the same size, apart from the sub-group with the maximum index which may be smaller if the size of the work-group is not evenly divisible by the size of the sub-group.

Sub-groups execute concurrently within a given work-group and make independent forward progress with respect to each other even in the absence of work-group barrier operations. Subgroups are able to internally synchronize using barrier operations without synchronizing with

each other.

In the degenerate case, with the extension enabled, a single sub-group must be supported for each work-group. In this situation all sub-group scope functions alias their work-group level equivalents.

28.4. Additions to Chapter 5 of the OpenCL 2.0 Specification

The function

returns information about the kernel object.

kernel specifies the kernel object being queried.

device identifies a specific device in the list of devices associated with *kernel*. The list of devices is the list of devices in the OpenCL context that is associated with *kernel*. If the list of devices associated with *kernel* is a single device, *device* can be a NULL value.

param_name specifies the information to query. The list of supported *param_name* types and the information returned in *param_value* by **clGetKernelSubGroupInfoKHR** is described in the Kernel Object Subgroup Queries table.

input_value_size is used to specify the size in bytes of memory pointed to by *input_value*. This size must be == size of input type as described in the table below.

input_value is a pointer to memory where the appropriate parameterization of the query is passed from. If *input_value* is NULL, it is ignored.

param_value is a pointer to memory where the appropriate result being queried is returned. If *param_value* is NULL, it is ignored.

 $param_value_size$ is used to specify the size in bytes of memory pointed to by $param_value$. This size must be \geq size of return type as described in the Kernel Object Subgroup Queries table.

param_value_size_ret returns the actual size in bytes of data being queried by param_name. If param_value_size_ret is NULL, it is ignored.

Table 46. clGetKernelSubGroupInfoKHR parameter queries

cl_kernel_sub_group_i nfo	Input Type	Return Type	Info. returned in param_value
CL_KERNEL_MAX_SUB_ GROUP_SIZE_FOR_ NDRANGE_KHR	size_t *	size_t	Returns the maximum sub-group size for this kernel. All sub-groups must be the same size, while the last subgroup in any work-group (i.e. the subgroup with the maximum index) could be the same or smaller size. The input_value must be an array of size_t values corresponding to the local work size parameter of the intended dispatch. The number of dimensions in the ND-range will be inferred from the value specified for input_value_size.

cl_kernel_sub_group_i nfo	Input Type	Return Type	Info. returned in param_value
CL_KERNEL_SUB_GROUP_ COUNT_FOR_NDRANGE_KHR	size_t *	size_t	Returns the number of sub-groups that will be present in each workgroup for a given local work size. All workgroups, apart from the last workgroup in each dimension in the presence of non-uniform work-group sizes, will have the same number of subgroups.
			The <i>input_value</i> must be an array of size_t values corresponding to the local work size parameter of the intended dispatch. The number of dimensions in the ND-range will be inferred from the value specified for <i>input_value_size</i> .

clGetKernelSubGroupInfoKHR returns CL_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- CL_INVALID_DEVICE if *device* is not in the list of devices associated with *kernel* or if *device* is NULL but there is more than one device associated with *kernel*.
- CL_INVALID_VALUE if *param_name* is not valid, or if size in bytes specified by *param_value_size* is < size of return type as described in the Kernel Object Subgroup Queries table and *param_value* is not NULL.
- CL_INVALID_VALUE if *param_name* is CL_KERNEL_MAX_SUB_GROUP_SIZE_FOR_NDRANGE_KHR and the size in bytes specified by *input_value_size* is not valid or if *input_value* is NULL.
- CL_INVALID_KERNEL if *kernel* is a not a valid kernel object.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

28.5. Additions to Chapter 6 of the OpenCL 2.0 C Specification

28.5.1. Additions to section 6.13.1 — Work Item Functions

Function	Description
uint get_sub_group_size ()	Returns the number of work items in the subgroup. This value is no more than the maximum subgroup size and is implementation-defined based on a combination of the compiled kernel and the dispatch dimensions. This will be a constant value for the lifetime of the subgroup.
uint get_max_sub_group_size ()	Returns the maximum size of a subgroup within the dispatch. This value will be invariant for a given set of dispatch dimensions and a kernel object compiled for a given device.
uint get_num_sub_groups ()	Returns the number of subgroups that the current work group is divided into. This number will be constant for the duration of a work group's execution. If the kernel is executed with a non-uniform work group size (i.e. the global_work_size values specified to clEnqueueNDRangeKernel are not evenly divisible by the local_work_size values for any dimension, calls to this built-in from some work groups may return different values than calls to this built-in from other work groups.
uint get_enqueued_num_sub_groups ()	Returns the same value as that returned by get_num_sub_groups if the kernel is executed with a uniform work group size. If the kernel is executed with a non-uniform work group size, returns the number of subgroups in each of the work groups that make up the uniform region of the global range.
uint get_sub_group_id ()	<pre>get_sub_group_id returns the subgroup ID which is a number from 0 get_num_sub_groups() - 1. For clEnqueueTask, this returns 0.</pre>

Function	Description
uint get_sub_group_local_id ()	Returns the unique work item ID within the current subgroup. The mapping from get_local_id (<i>dimindx</i>) to get_sub_group_local_id will be invariant for the lifetime of the work group.

${\bf 28.5.2.\ Additions\ to\ section\ 6.13.8-Synchronization\ Functions}$

Function	Description
void sub_group_barrier (cl_mem_fence_flags <i>flags</i>) void sub_group_barrier (cl_mem_fence_flags <i>flags</i> , memory_scope <i>scope</i>)	All work items in a subgroup executing the kernel on a processor must execute this function before any are allowed to continue execution beyond the subgroup barrier. This function must be encountered by all work items in a subgroup executing the kernel. These rules apply to ND-ranges implemented with uniform and non-uniform work groups.
memory_scope scope)	uniform work groups.
	If sub_group_barrier is inside a conditional statement, then all work items within the subgroup must enter the conditional if any work item in the subgroup enters the conditional statement and executes the sub_group_barrier.
	If sub_group_barrier is inside a loop, all work items within the subgroup must execute the sub_group_barrier for each iteration of the loop before any are allowed to continue execution beyond the sub_group_barrier.
	The sub_group_barrier function also queues a memory fence (reads and writes) to ensure correct ordering of memory operations to local or global memory.
	The flags argument specifies the memory address space and can be set to a combination of the following values:
	CLK_LOCAL_MEM_FENCE - The sub_group_barrier function will either flush any variables stored in local memory or queue a memory fence to ensure correct ordering of memory operations to local memory.
	CLK_GLOBAL_MEM_FENCE — The sub_group_barrier function will queue a memory fence to ensure correct ordering of memory operations to global memory. This can be useful when work items, for example, write to buffer objects and then want to read the updated data from these buffer objects.
	CLK_IMAGE_MEM_FENCE — The sub_group_barrier function will queue a memory fence to ensure correct ordering of memory operations to image objects. This can be useful when work items, for example, write to image objects and then want to read the updated data from these image objects.

28.5.3. Additions to section 6.13.11 — Atomic Functions

Add the following new value to the enumerated type memory_scope defined in section 6.13.11.4.

memory_scope_sub_group

The memory_scope_sub_group specifies that the memory ordering constraints given by memory_order apply to work items in a subgroup. This memory scope can be used when performing atomic operations to global or local memory.

28.5.4. Add a new section 6.13.X — Sub-Group Functions

The table below describes OpenCL C programming language built-in functions that operate on a subgroup level. These built-in functions must be encountered by all work items in the subgroup executing the kernel. For the functions below, the generic type name gentype may be the one of the supported built-in scalar data types int, uint, long, ulong, float, double (if double precision is supported), or half (if half precision is supported).

Function	Description
int sub_group_all (int <i>predicate</i>)	Evaluates <i>predicate</i> for all work items in the subgroup and returns a non-zero value if <i>predicate</i> evaluates to non-zero for all work items in the subgroup.
int sub_group_any (int predicate)	Evaluates <i>predicate</i> for all work items in the subgroup and returns a non-zero value if <i>predicate</i> evaluates to non-zero for any work items in the subgroup.
gentype sub_group_broadcast (gentype <i>x</i> , uint <i>sub_group_local_id</i>)	Broadcast the value of <i>x</i> for work item identified by $sub_group_local_id$ (value returned by $get_sub_group_local_id$) to all work items in the subgroup. $sub_group_local_id$ must be the same value for all work items in the subgroup.
gentype sub_group_reduce_<op></op> (gentype x)	Return result of reduction operation specified by < op > for all values of <i>x</i> specified by work items in a subgroup.
gentype sub_group_scan_exclusive_<op></op> (gentype <i>x</i>)	Do an exclusive scan operation specified by < op > of all values specified by work items in a subgroup. The scan results are returned for each work item.
	The scan order is defined by increasing subgroup local ID within the subgroup.

Function	Description
gentype sub_group_scan_inclusive_<op></op> (gentype <i>x</i>)	Do an inclusive scan operation specified by <op></op> of all values specified by work items in a subgroup. The scan results are returned for each work item. The scan order is defined by increasing subgroup local ID within the subgroup.

The <op> in sub_group_reduce_<op>, sub_group_scan_inclusive_<op> and sub_group_scan_exclusive_<op> defines the operator and can be add, min or max.

The exclusive scan operation takes a binary operator **op** with an identity I and n (where n is the size of the sub-group) elements $[a_0, a_1, ... a_{n-1}]$ and returns $[I, a_0, (a_0 \mathbf{op} a_1), ... (a_0 \mathbf{op} a_1 \mathbf{op} ... \mathbf{op} a_{n-2})]$.

The inclusive scan operation takes a binary operator **op** with n (where n is the size of the subgroup) elements $[a_0, a_1, ... a_{n-1}]$ and returns $[a_0, (a_0 \mathbf{op} a_1), ... (a_0 \mathbf{op} a_1 \mathbf{op} ... \mathbf{op} a_{n-1})]$.

If **op** = **add**, the identity I is 0. If **op** = **min**, the identity I is INT_MAX, UINT_MAX, LONG_MAX, ULONG_MAX, for int, uint, long, ulong types and is +INF for floating-point types. Similarly if **op** = max, the identity I is INT_MIN, 0, LONG_MIN, 0 and -INF.



The order of floating-point operations is not guaranteed for the sub_group_reduce_<op>, sub_group_scan_inclusive_<op> and
sub_group_scan_exclusive_<op> built-in functions that operate on half, float and double data types. The order of these floating-point operations is also non-deterministic for a given sub-group.

28.5.5. Additions to section 6.13.16 — Pipe Functions

The OpenCL C programming language implements the following built-in pipe functions that operate at a subgroup level. These built-in functions must be encountered by all work items in a subgroup executing the kernel with the same argument values; otherwise the behavior is undefined. We use the generic type name gentype to indicate the built-in OpenCL C scalar or vector integer or floating-point data types or any user defined type built from these scalar and vector data types can be used as the type for the arguments to the pipe functions listed in *table 6.29*.

Function	Description
reserve_id_t sub_group_reserve_read_pipe (read_only pipe gentype pipe, uint num_packets)	Reserve <i>num_packets</i> entries for reading from or writing to <i>pipe</i> . Returns a valid non-zero reservation ID if the reservation is successful and 0 otherwise.
reserve_id_t sub_group_reserve_write_pipe (write_only pipe gentype <i>pipe</i> , uint <i>num_packets</i>)	The reserved pipe entries are referred to by indices that go from 0 num_packets - 1.

Function	Description
<pre>void sub_group_commit_read_pipe (read_only pipe gentype pipe, reserve_id_t reserve_id)</pre>	Indicates that all reads and writes to num_packets associated with reservation reserve_id are completed.
void sub_group_commit_write_pipe (write_only pipe gentype <i>pipe</i> , reserve_id_t <i>reserve_id</i>)	

Note: Reservations made by a subgroup are ordered in the pipe as they are ordered in the program. Reservations made by different subgroups that belong to the same work group can be ordered using subgroup synchronization. The order of subgroup based reservations that belong to different work groups is implementation defined.

28.5.6. Additions to section 6.13.17.6 — Enqueuing Kernels (Kernel Query Functions)

Built-in Function	Description
uint get_kernel_sub_group_count_for_ndrange (const ndrange_t ndrange, void (^block)(void)); uint get_kernel_sub_group_count_for_ndrange (const ndrange_t ndrange, void (^block)(local void *,));	Returns the number of subgroups in each work group of the dispatch (except for the last in cases where the global size does not divide cleanly into work groups) given the combination of the passed ndrange and block. block specifies the block to be enqueued.
uint get_kernel_max_sub_group_size_for_ndrange (const ndrange_t ndrange, void (^block)(void)); uint get_kernel_max_sub_group_size_for_ndrange (const ndrange_t ndrange, void (^block)(local void *,));	Returns the maximum subgroup size for a block.

Chapter 29. Mipmaps

This section describes OpenCL support for mipmaps.

There are two optional mipmap extensions. The cl_khr_mipmap_image extension adds the ability to create a mip-mapped image, enqueue commands to read/write/copy/map/unmap a region of a mipmapped image, and built-in functions that can be used to read a mip-mapped image in an OpenCL C program. The cl_khr_mipmap_image_writes extension adds built-in functions that can write a mip-mapped image in an OpenCL C program. the cl_khr_mipmap_image_writes extension is supported by the OpenCL device, the cl_khr_mipmap_image extension must also be supported.

29.1. General information

29.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

29.2. Additions to Chapter 5 of the OpenCL 2.2 Specification

29.2.1. Additions to section 5.3 — Image Objects

A mip-mapped 1D image, 1D image array, 2D image, 2D image array or 3D image is created by specifying *num_mip_levels* to be a value greater than one in the *image_desc* passed to **clCreateImage**. The dimensions of a mip-mapped image can be a power of two or a non-power of two. Each successively smaller mipmap level is half the size of the previous level. If this half value is a fractional value, it is rounded down to the nearest integer.

Restrictions

The following restrictions apply when mip-mapped images are created with **clCreateImage**:

- CL_MEM_USE_HOST_PTR or CL_MEM_COPY_HOST_PTR cannot be specified if a mip-mapped image is created.
- The *host_ptr* argument to **clCreateImage** must be a **NULL** value.
- Mip-mapped images cannot be created for CL_MEM_OBJECT_IMAGE1D_BUFFER images, depth images or multi-sampled (i.e. msaa) images.

Calls to **clEnqueueReadImage**, **clEnqueueWriteImage** and **clEnqueueMapImage** can be used to read from or write to a specific mip-level of a mip-mapped image. If image argument is a 1D image, origin[1] specifies the mip-level to use. If image argument is a 1D image array, origin[2] specifies the mip-level to use. If image argument is a 2D image, origin[2] specifies the mip-level to use. If image array or a 3D image, origin[3] specifies the mip-level to use.

calls to clenqueueCopyImage, clenqueueCopyImageToBuffer and clenqueueCopyBufferToImage can also be used to copy from and to a specific mip-level of a mip-mapped image. If src_image argument is a 1D image, $src_origin[1]$ specifies the mip-level to use. If src_image argument is a 1D image array, $src_origin[2]$ specifies the mip-level to use. If src_image argument is a 2D image, $src_origin[2]$ specifies the mip-level to use. If src_image argument is a 2D image array or a 3D image, $src_origin[3]$ specifies the mip-level to use. If dst_image argument is a 1D image array, $dst_origin[1]$ specifies the mip-level to use. If dst_image argument is a 2D image, $dst_origin[2]$ specifies the mip-level to use. If dst_image argument is a 2D image, $dst_origin[3]$ specifies the mip-level to use. If dst_image argument is a 2D image, $dst_origin[3]$ specifies the mip-level to use.

If the mip level specified is not a valid value, these functions return the error CL_INVALID_MIP_LEVEL.

Calls to clEnqueueFillImage can be used to write to a specific mip-level of a mip-mapped image. If image argument is a 1D image, origin[1] specifies the mip-level to use. If image argument is a 1D image array, origin[2] specifies the mip-level to use. If image argument is a 2D image, origin[2] specifies the mip-level to use. If image argument is a 2D image array or a 3D image, origin[3] specifies the mip-level to use.

29.2.2. Additions to section 5.7 — Sampler Objects

Add the following sampler properties *to table 5.14* that can be specified when a sampler object is created using **clCreateSamplerWithProperties**.

cl_sampler_properties enum	Property Value	Default Value
CL_SAMPLER_MIP_FILTER_MODE_KHR	cl_filter_mode	CL_FILTER_NEAREST
CL_SAMPLER_LOD_MIN_KHR	cl_float	0.0f
CL_SAMPLER_LOD_MAX_KHR	cl_float	MAXFLOAT

Note: The sampler properties CL_SAMPLER_MIP_FILTER_MODE_KHR, CL_SAMPLER_LOD_MIN_KHR and CL_SAMPLER_LOD_MAX_KHR cannot be specified with any samplers initialized in the OpenCL program source. Only the default values for these properties will be used. To create a sampler with specific for these properties, a sampler object must be created values with clCreateSamplerWithProperties and passed as an argument to a kernel.

29.3. Additions to Chapter 6 of the OpenCL 2.0 Specification

29.3.1. Additions to section 6.13.14 - Image Read, Write and Query Functions

The image read and write functions described in *sections 6.13.14.2*, *6.13.14.3* and *6.13.14.4* read from and write to mip-level 0 if the image argument is a mip-mapped image.

The following new built-in functions are added to section 6.13.14.2.

```
float4 read_imagef(
    read_only image2d_t image,
    sampler_t sampler,
    float2 coord,
    float lod)
int4 read_imagei(
    read_only image2d_t image,
    sampler_t sampler,
    float2 coord,
    float lod)
uint4 read_imageui(
    read_only image2d_t image,
    sampler_t sampler,
    float2 coord,
    float lod)
float read_imagef(
    read_only image2d_depth_t image,
    sampler_t sampler,
    float2 coord,
    float lod)
```

Description

Use the coordinate *coord.xy* to do an element lookup in the mip-level specified by *lod* in the 2D image object specified by *image*.

```
float4 read_imagef(
    read only image2d t image,
    sampler_t sampler,
    float2 coord,
    float2 gradient_x,
    float2 gradient_y)
int4 read imagei(
    read_only image2d_t image,
    sampler_t sampler,
    float2 coord,
    float2 gradient_x,
    float2 gradient_y)
uint4 read_imageui(
    read_only image2d_t image,
    sampler_t sampler,
    float2 coord,
    float2 gradient_x,
    float2 gradient_y)
float read imagef(
    read_only image2d_depth_t image,
    sampler_t sampler,
    float2 coord,
    float2 gradient x,
    float2 gradient_y)
```

Description

Use the gradients to compute the lod and coordinate *coord.xy* to do an element lookup in the mip-level specified by the computed lod in the 2D image object specified by *image*.

```
float4 read_imagef(
    read_only image1d_t image,
    sampler_t sampler,
    float coord,
    float lod)

int4 read_imagei(
    read_only image1d_t image,
    sampler_t sampler,
    float coord,
    float lod)

uint4 read_imageui(
    read_only image1d_t image,
    sampler_t sampler,
    float coord,
    float coord,
    float lod)
```

Use the coordinate *coord* to do an element lookup in the mip-level specified by *lod* in the 1D image object specified by *image*.

```
float4 read_imagef(
    read only image1d t image,
    sampler_t sampler,
    float coord,
    float gradient_x,
    float gradient_y)
int4 read_imagei(
    read_only image1d_t image,
    sampler_t sampler,
    float coord,
    float gradient_x,
    float gradient_y)
uint4 read_imageui(
    read_only image1d_t image,
    sampler_t sampler,
    float coord,
    float gradient_x,
    float gradient_y)
```

Description

Use the gradients to compute the lod and coordinate *coord* to do an element lookup in the mip-level specified by the computed lod in the 1D image object specified by *image*.

```
float4 read_imagef(
    read_only image3d_t image,
    sampler_t sampler,
    float4 coord,
    float lod)

int4 read_imagei(
    read_only image3d_t image,
    sampler_t sampler,
    float4 coord,
    float lod)

uint4 read_imageui(
    read_only image3d_t image,
    sampler_t sampler,
    float4 coord,
    float4 coord,
    float lod)
```

Use the coordinate *coord.xyz* to do an element lookup in the mip-level specified by *lod* in the 3D image object specified by *image*.

```
float4 read_imagef(
    read only image3d t image,
    sampler_t sampler,
    float4 coord,
    float4 gradient_x,
    float4 gradient_y)
int4 read imagei(
    read_only image3d_t image,
    sampler_t sampler,
    float4 coord,
    float4 gradient_x,
    float4 gradient_y)
uint4 read_imageui(
    read_only image3d_t image,
    sampler_t sampler,
    float4 coord,
    float4 gradient_x,
    float4 gradient_y)
```

Description

Use the gradients to compute the lod and coordinate *coord.xyz* to do an element lookup in the mip-level specified by the computed lod in the 3D image object specified by *image*.

```
float4 read_imagef(
    read_only image1d_array_t image,
    sampler_t sampler,
    float2 coord,
    float lod)

int4 read_imagei(
    read_only image1d_array_t image,
    sampler_t sampler,
    float2 coord,
    float lod)

uint4 read_imageui(
    read_only image1d_array_t image,
    sampler_t sampler,
    float2 coord,
    float2 coord,
    float lod)
```

Use the coordinate *coord.x* to do an element lookup in the 1D image identified by *coord.x* and mip-level specified by *lod* in the 1D image array specified by *image*.

```
float4 read_imagef(
    read only image1d array t image,
    sampler_t sampler,
    float2 coord,
    float gradient x,
    float gradient_y)
int4 read imagei(
    read_only image1d_array_t image,
    sampler_t sampler,
    float2 coord,
    float gradient_x,
    float gradient_y)
uint4 read_imageui(
    read_only image1d_array_t image,
    sampler_t sampler,
    float2 coord,
    float gradient_x,
    float gradient_y)
```

Description

Use the gradients to compute the lod and coordinate *coord.x* to do an element lookup in the mip-level specified by the computed lod in the 1D image array specified by *image*.

```
float4 read imagef(
    read_only image2d_array_t image,
    sampler_t sampler,
    float4 coord,
    float lod)
int4 read_imagei(
    read_only image2d_array_t image,
    sampler_t sampler,
    float4 coord,
    float lod)
uint4 read imageui(
    read_only image2d_array_t image,
    sampler_t sampler,
    float4 coord,
    float lod)
float read_imagef(
    read_only image2d_array_depth_t image,
    sampler_t sampler,
    float4 coord,
    float lod)
```

Use the coordinate *coord.xy* to do an element lookup in the 2D image identified by *coord.z* and mip-level specified by *lod* in the 2D image array specified by *image*.

```
float4 read_imagef(
    read only image2d array t image,
    sampler_t sampler,
    float4 coord,
    float2 gradient_x,
    float2 gradient_y)
int4 read imagei(
    read_only image2d_array_t image,
    sampler_t sampler,
    float4 coord,
    float2 gradient_x,
    float2 gradient_y)
uint4 read_imageui(
    read_only image2d_array_t image,
    sampler_t sampler,
    float4 coord,
    float2 gradient_x,
    float2 gradient_y)
float read imagef(
    read_only image2d_array_depth_t image,
    sampler_t sampler,
    float4 coord,
    float2 gradient x,
    float2 gradient_y)
```

Description

Use the gradients to compute the lod coordinate and *coord.xy* to do an element lookup in the 2D image identified by *coord.z* and mip-level specified by the computed lod in the 2D image array specified by *image*.



CL_SAMPLER_NORMALIZED_COORDS must be CL_TRUE for built-in functions described in the table above that read from a mip-mapped image; otherwise the behavior is undefined. The value specified in the *lod* argument is clamped to the minimum of (actual number of mip-levels – 1) in the image or value specified for CL_SAMPLER_LOD_MAX.

The following new built-in functions are added to section 6.13.14.4.

```
void write_imagef(
    write only image2d t image,
    int2 coord,
    int lod,
    float4 color)
void write_imagei(
    write_only image2d_t image,
    int2 coord,
    int lod,
    int4 color)
void write_imageui(
    write_only image2d_t image,
    int2 coord,
    int lod,
    uint4 color)
void write imagef(
    write_only image2d_depth_t image,
    int2 coord,
    int lod,
    float depth)
```

Description

Write *color* value to location specified by *coord.xy* in the mip-level specified by *lod* in the 2D image object specified by *image*. Appropriate data format conversion to the specified image format is done before writing the color value. coord.x and coord.y are considered to be unnormalized coordinates and must be in the range 0 .. image width of mip-level specified by lod - 1, and 0 .. image height of mip-level specified by lod - 1.

The behavior of **write_imagef**, **write_imagei** and **write_imageui** if (x, y) coordinate values are not in the range (0 ... image width of the miplevel specified by <math>lod - 1, 0 ... image height of the mip-level specified by <math>lod - 1) or lod value exceeds the (number of mip-levels in the image – 1) is undefined.

```
void write_imagef(
    write_only image1d_t image,
    int coord,
    int lod,
    float4 color)

void write_imagei(
    write_only image1d_t image,
    int coord,
    int lod,
    int4 color)

void write_imageui(
    write_only image1d_t image,
    int coord,
    int lod,
    int coord,
    int lod,
    uint4 color)
```

Write *color* value to location specified by *coord* in the mip-level specified by *lod* in the 1D image object specified by *image*. Appropriate data format conversion to the specified image format is done before writing the color value. *coord* is considered to be unnormalized coordinates and must be in the range 0 .. image width of the mip-level specified by lod - 1.

The behavior of write_imagef, write_imagei and write_imageui if coordinate value is not in the range (0 .. image width of the mip-level specified by lod - 1) or lod value exceeds the (number of mip-levels in the image - 1), is undefined.

```
void write_imagef(
    write_only image1d_array_t image,
    int2 coord,
    int lod,
    float4 color)

void write_imagei(
    write_only image1d_array_t image,
    int2 coord,
    int lod,
    int4 color)

void write_imageui(
    write_only image1d_array_t image,
    int2 coord,
    int1 coord,
    int2 coord,
    int2 coord,
    int1 lod,
    uint4 color)
```

```
void write imagef(
    write_only image2d_array_t image,
    int4 coord,
    int lod,
    float4 color)
void write imagei(
    write_only image2d_array_t image,
    int4 coord,
    int lod,
    int4 color)
void write imageui(
    write_only image2d_array_t image,
    int4 coord,
    int lod,
    uint4 color)
void write imagef(
    write_only image2d_array_depth_t
image,
    int4 coord,
    int lod,
    float depth)
```

Description

Write *color* value to location specified by *coord.x* in the 1D image identified by *coord.y* and miplevel *lod* in the 1D image array specified by *image*. Appropriate data format conversion to the specified image format is done before writing the color value. *coord.x* and *coord.y* are considered to be unnormalized coordinates and must be in the range 0 .. image width of the miplevel specified by lod-1 and 0 .. image number of layers -1.

The behavior of **write_imagef**, **write_imagei** and **write_imageui** if (x, y) coordinate values are not in the range (0 ... image width of the miplevel specified by <math>lod - 1, 0 ... image number of layers <math>- 1), respectively or lod value exceeds the (number of mip-levels in the image - 1), is undefined.

Write *color* value to location specified by *coord.xy* in the 2D image identified by *coord.z* and mip-level *lod* in the 2D image array specified by *image*. Appropriate data format conversion to the specified image format is done before writing the color value. *coord.x*, *coord.y* and *coord.z* are considered to be unnormalized coordinates and must be in the range 0 .. image width of the mip-level specified by lod - 1, 0 .. image height -1 specified by lod - 1 and 0 .. image number of layers -1.

The behavior of **write_imagef**, **write_imagei** and **write_imageui** if (x, y, z) coordinate values are not in the range (0 ... image width of the mip-level specified by <math>lod - 1, 0 ... image height of the mip-level specified by <math>lod - 1, 0 ... image number of layers - 1), respectively or <math>lod value exceeds the (number of mip-levels in the image - 1), is undefined.

```
void write_imagef(
    write_only image3d_t image,
    int4 coord,
    int lod,
    float4 color)

void write_imagei(
    write_only image3d_t image,
    int4 coord,
    int lod,
    int4 color)

void write_imageui(
    write_only image3d_t image,
    int4 color)

void write_imageui(
    write_only image3d_t image,
    int4 coord,
    int lod,
    uint4 color)
```

Description

Write color value to location specified by coord.xyz and mip-level lod in the 3D image object specified by image. Appropriate data format conversion to the specified image format is done before writing the color value. coord.x, coord.y and coord.z are considered to be unnormalized coordinates and must be in the range 0 .. image width -1 specified by lod-1, 0 .. image height -1 specified by lod-1 and 0 .. image depth -1 specified by lod-1.

The behavior of **write_imagef**, **write_imagei** and **write_imageui** if (x, y, z) coordinate values are not in the range (0 ... image width of the mip-level specified by <math>lod - 1, 0 ... image height of the mip-level specified by <math>lod - 1, 0 ... image depth - 1), respectively or lod value exceeds the (number of mip-levels in the image - 1), is undefined.

The following new built-in functions are added to section 6.13.14.5.

Function

int get_image_num_mip_levels(image1d_t image) int get_image_num_mip_levels(image2d_t image) int get_image_num_mip_levels(image3d_t image) int get_image_num_mip_levels(image1d_array_t image) int get_image_num_mip_levels(image2d_array_t image) int get_image_num_mip_levels(image2d_depth_t image) int get_image_num_mip_levels(image2d_array_depth_t image)

Description

Return the number of mip-levels.

29.4. Additions to Creating OpenCL Memory Objects from OpenGL Objects

If both the <code>cl_khr_mipmap_image</code> and <code>cl_khr_gl_sharing</code> extensions are supported by the OpenCL device, the <code>cl_khr_gl_sharing</code> extension may also be used to create a mipmapped OpenCL image from a mipmapped OpenGL texture.

To create a mipmapped OpenCL image from a mipmapped OpenGL texture, pass a negative value as the *miplevel* argument to **clCreateFromGLTexture**. If *miplevel* is a negative value then an OpenCL mipmapped image object is created from a mipmapped OpenGL texture object, instead of an OpenCL image object for a specific miplevel of the OpenGL texture.

Note: For a detailed description of how the level of detail is computed, please refer to *section 3.9.7* of the OpenGL 3.0 specification.

Chapter 30. sRGB Image Writes

This section describes the cl_khr_srgb_image_writes extension.

This extension enables kernels to write to sRGB images using the **write_imagef** built-in function. The sRGB image formats that may be written to will be returned by **clGetSupportedImageFormats**.

When the image is an sRGB image, the **write_imagef** built-in function will perform the linear to sRGB conversion. Only the R, G, and B components are converted from linear to sRGB; the A component is written as-is.

30.1. General information

30.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

Chapter 31. Priority Hints

This section describes the **cl_khr_priority_hints** extension. This extension adds priority hints for OpenCL, but does not specify the scheduling behavior or minimum guarantees. It is expected that the user guides associated with each implementation which supports this extension will describe the scheduling behavior guarantees.

31.1. General information

31.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

31.2. Host-side API modifications

The function **clCreateCommandQueueWithProperties** (Section 5.1) is extended to support a priority value as part of the *properties* argument.

The priority property applies to OpenCL command queues that belong to the same OpenCL context.

The properties field accepts the CL_QUEUE_PRIORITY_KHR property, with a value of type cl_queue_priority_khr, which can be one of:

- CL QUEUE PRIORITY HIGH KHR
- CL_QUEUE_PRIORITY_MED_KHR
- CL_QUEUE_PRIORITY_LOW_KHR

If CL_QUEUE_PRIORITY_KHR is not specified then the default priority is CL_QUEUE_PRIORITY_MED_KHR.

To the error section for clCreateCommandQueueWithProperties, the following is added:

 CL_INVALID_QUEUE_PROPERTIES if the CL_QUEUE_PRIORITY_KHR property is specified and the queue is a CL_QUEUE_ON_DEVICE.

Chapter 32. Throttle Hints

This section describes the **cl_khr_throttle_hints** extension. This extension adds throttle hints for OpenCL, but does not specify the throttling behavior or minimum guarantees. It is expected that the user guide associated with each implementation which supports this extension will describe the throttling behavior guarantees.

Note that the throttle hint is orthogonal to functionality defined in **cl_khr_priority_hints** extension. For example, a task may have high priority (CL_QUEUE_PRIORITY_HIGH_KHR) but should at the same time be executed at an optimized throttle setting (CL_QUEUE_THROTTLE_LOW).

32.1. General information

32.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

32.2. Host-side API modifications

The function **clCreateCommandQueueWithProperties** (Section 5.1) is extended to support a new **CL_QUEUE_THROTTLE_KHR** value as part of the *properties* argument.

The properties field accepts the following values:

- CL_QUEUE_THROTTLE_HIGH_KHR (full throttle, i.e., OK to consume more energy)
- CL_QUEUE_THROTTLE_MED_KHR (normal throttle)
- CL_QUEUE_THROTTLE_LOW_KHR (optimized/lowest energy consumption)

If CL_QUEUE_THROTTLE_KHR is not specified then the default priority is CL_QUEUE_THROTTLE_MED_KHR.

To the error section for clCreateCommandQueueWithProperties, the following is added:

 CL_INVALID_QUEUE_PROPERTIES if the CL_QUEUE_THROTTLE_KHR property is specified and the queue is a CL_QUEUE_ON_DEVICE.

Chapter 33. Named Barriers for Subgroups

This section describes the **cl_khr_subgroup_named_barrier** extension. This extension adds barrier operations that cover subsets of an OpenCL work-group. Only the OpenCL API changes are described in this section. Please refer to the SPIR-V specification for information about using subgroups named barriers in the SPIR-V intermediate representation, and to the OpenCL C++ specification for descriptions of the subgroup named barrier built-in functions in the OpenCL C++ kernel language.

33.1. General information

33.1.1. Version history

Date	Version	Description
2020-04-21	1.0.0	First assigned version.

33.2. Changes to OpenCL specification

Add to table 4.3:

cl_device_info	Return Type	Description
CL_DEVICE_MAX_NAMED_BA RRIER_COUNT_KHR	cl_uint	Maximum number of named barriers in a work- group for any given kernel-instance running on the device. The minimum value is 8.

Chapter 34. Extended Async Copies (Provisional)

This section describes the **cl_khr_extended_async_copies** provisional extension. This extension augments built-in asynchronous copy functions to OpenCL C to support more patterns:

- 1. for async copy between 2D source and 2D destination.
- 2. for async copy between 3D source and 3D destination.

34.1. General information

34.1.1. Version history

Date	Version	Description
2020-04-21	0.9.0	First assigned version (provisional).

34.2. Additions to Chapter 6 of the OpenCL C Specification

The following new built-in functions are added to the *Async Copies from Global to Local Memory, Local to Global Memory, and Prefetch* functions described in *section 6.12.10* and *section 6.13.10* of the OpenCL 1.2 and OpenCL 2.0 C specifications.

Note that <code>async_work_group_strided_copy</code> is a special case of <code>async_work_group_copy_2D2D</code>, namely one which copies a single column to a single line or vice versa. For example: <code>async_work_group_strided_copy(dst, src, num_gentypes, src_stride)</code> is equal to <code>async_work_group_copy_2D2D(dst, src, 1, num_gentypes, src_stride-1, 1)</code>

These new built-in functions support the same gentype generic type names as the standard asynchronous copy functions unless otherwise stated.

```
event_t async_work_group_copy_2D2D(
    local gentype *dst,
    const __global gentype *src,
    size_t num_elements_per_line,
    size t num lines,
    size_t src_stride,
    size_t dst_stride,
    event_t event)
event_t async_work_group_copy_2D2D(
    __global gentype *dst,
    const __local gentype *src,
    size_t num_elements_per_line,
    size t num lines,
    size_t src_stride,
    size_t dst_stride,
    event_t event)
```

Description

Perform an asynchronous copy of *num_lines* lines from *src* to *dst*. Each line contains *num_elements_per_line* gentype elements. After each line of transfer, *src* address is incremented by (*src_stride + num_elements_per_line*) gentype elements, *dst* address is incremented by (*dst_stride + num_elements_per_line*) gentype elements for the next line of transfer.

For these functions, the stride describes the number of elements between the **end** of the current line and the **beginning** of the next line, i.e., without overlap.

Returns an event object that can be used by wait_group_events to wait for the async copy to finish. The *event* argument can also be used to associate the async_work_group_copy_2D2D with a previous async copy allowing an event to be shared by multiple async copies; otherwise *event* should be zero.

If *event* argument is non-zero, the event object supplied in *event* argument will be returned.

This function does not perform any implicit synchronization of source data such as using a **barrier** before performing the copy.

The behavior of async_work_group_copy_2D2D is undefined if the *num_elements_per_line* or *src_stride* or *dst_stride* values cause the *src* or *dst* addresses to exceed the upper bounds of the address space during the copy.

The async copy is performed by all work-items in a work-group and this built-in function must therefore be encountered by all work-items in a work-group executing the kernel with the same argument values; otherwise the results are undefined.

```
event_t async_work_group_copy_3D3D(
    local gentype *dst,
    const __global gentype *src,
    size_t num_elements_per_line,
    size_t num_lines,
    size_t src_line_stride,
    size_t dst_line_stride,
    size_t num_planes,
    size_t src_plane_stride,
    size_t dst_plane_stride,
    event_t event)
event_t async_work_group_copy_3D3D(
    __global gentype *dst,
    const __local gentype *src,
    size_t num_elements_per_line,
    size_t num_lines,
    size_t src_line_stride,
    size_t dst_line_stride,
    size_t num_planes,
    size_t src_plane_stride,
    size_t dst_plane_stride,
    event_t event)
```

Description

Perform an async copy of *num_planes* times *num_lines* lines from *src* to *dst* arranged in *num_planes* planes. Each plane contains *num_lines* lines. Each line contains *num_elements_per_line* gentype elements. After each line of transfer, *src* address is incremented by (*src_line_stride + num_elements_per_line*) gentype elements, *dst* address is incremented by (*dst_line_stride + num_elements_per_line*) gentype elements for the next line of transfer. For the last line of a plane, an additional *src_plane_stride* gentype elements is added to *src* address, and an additional *dst_plane_stride* gentype elements is address.

Returns an event object that can be used by wait_group_events to wait for the async copy to finish. The event argument can also be used to associate the async_work_group_copy_3D3D with a previous async copy allowing an event to be shared by multiple async copies; otherwise event should be zero.

If *event* argument is non-zero, the event object supplied in *event* argument will be returned.

This function does not perform any implicit synchronization of source data such as using a **barrier** before performing the copy.

The behavior of async_work_group_copy_3D3D is undefined if any of num_elements_per_line, src_line_stride, dst_line_stride, src_plane_stride or dst_plane_stride values cause the src or dst addresses to exceed the upper bounds of the address space during the copy.

The async copy is performed by all work-items in a work-group and this built-in function must therefore be encountered by all work-items in a work-group executing the kernel with the same argument values; otherwise the results are undefined.



This is a preview of an OpenCL provisional extension specification that has been Ratified under the Khronos Intellectual Property Framework. It is being made publicly available prior to being uploaded to the Khronos registry to enable review and feedback from the community. If you have feedback please create an issue on https://github.com/KhronosGroup/OpenCL-Docs/

Chapter 35. Async Work Group Copy Fence (Provisional)

This section describes the **cl_khr_async_work_group_copy_fence** provisional extension. The extension adds a new built-in function to OpenCL C to establish a memory synchronization ordering of asynchronous copies.

35.1. General information

35.1.1. Version history

Date	Version	Description
2020-04-21	0.9.0	First assigned version (provisional).

35.2. Additions to Chapter 6 of the OpenCL C Specification

The following new built-in function is added to the *Async Copies from Global to Local Memory, Local to Global Memory, and Prefetch* functions described in *section 6.12.10* and *section 6.13.10* of the OpenCL 1.2 and OpenCL 2.0 C specifications:

void async_work_group_copy_fence(
 cl_mem_fence_flags flags)

Description

copies preceding the

async_work_group_copy_fence must complete
their access to the designated memory or
memories, including both reads-from and
writes-to it, before async copies following the
fence are allowed to start accessing these
memories. In other words, every async copy
preceding the async_work_group_copy_fence
must happen-before every async copy following
the fence, with respect to the designated

Orders async copies produced by the work-items

of a work-group executing a kernel. Async

The *flags* argument specifies the memory address space and can be set to a combination of the following literal values:

CLK_LOCAL_MEM_FENCE CLK_GLOBAL_MEM_FENCE

memory or memories.

The async fence is performed by all work-items in a work-group and this built-in function must therefore be encountered by all work-items in a work-group executing the kernel with the same argument values; otherwise the results are undefined. This rule applies to ND-ranges implemented with uniform and non-uniform work-groups.



This is a preview of an OpenCL provisional extension specification that has been Ratified under the Khronos Intellectual Property Framework. It is being made publicly available prior to being uploaded to the Khronos registry to enable review and feedback from the community. If you have feedback please create an issue on https://github.com/KhronosGroup/OpenCL-Docs/

Chapter 36. Unique Device Identifiers

This section describes the cl_khr_device_uuid extension.

This extension adds the ability to query a universally unique identifier (UUID) for an OpenCL driver and OpenCL device. The UUIDs returned by the query may be used to identify drivers and devices across processes or APIs.

36.1. General information

36.1.1. Version history

Date	Version	Description
2020-08-27	1.0.0	First assigned version.

36.2. Additions to Chapter 4 of the OpenCL 3.0 API Specification

Add to Table 5 - OpenCL Device Queries:

Table 5. OpenCL Device Queries

cl_device_info	Return Type	Description
CL_DEVICE_UUID_KHR	<pre>cl_uchar[CL_UUID_S IZE_KHR]</pre>	Returns a universally unique identifier (UUID) for the device.
		Device UUIDs must be immutable for a given device across processes, driver APIs, driver versions, and system reboots.
CL_DRIVER_UUID_KHR	<pre>cl_uchar[CL_UUID_S IZE_KHR]</pre>	Returns a universally unique identifier (UUID) for the software driver for the device.
CL_DEVICE_LUID_VALID_KHR	cl_bool	Returns CL_TRUE if the device has a valid LUID and CL_FALSE otherwise.

cl_device_info	Return Type	Description
CL_DEVICE_LUID_KHR	<pre>cl_uchar[CL_LUID_S IZE_KHR]</pre>	Returns a locally unique identifier (LUID) for the device.
		It is not an error to query CL_DEVICE_LUID_KHR when CL_DEVICE_LUID_VALID_KHR returns CL_FALSE, but in this case the returned LUID value is undefined.
		When CL_DEVICE_LUID_VALID_KHR returns CL_TRUE, and the OpenCL device is running on the Windows operating system, the returned LUID value can be cast to an LUID object and must be equal to the locally unique identifier of an IDXGIAdapter1 object that corresponds to the OpenCL device.
CL_DEVICE_NODE_MASK_KHR	cl_uint	Returns a node mask for the device. It is not an error to query CL_DEVICE_NODE_MASK_KHR when CL_DEVICE_LUID_VALID_KHR returns CL_FALSE, but in this case the returned node mask is undefined.
		When CL_DEVICE_LUID_VALID_KHR returns CL_TRUE, the returned node mask must contain exactly one bit. If the OpenCL device is running on an operating system that supports the Direct3D 12 API and the OpenCL device corresponds to an individual device in a linked device adapter, the returned node mask identifies the Direct3D 12 node corresponding to the OpenCL device. Otherwise, the returned node mask must be 1.



While CL_DEVICE_UUID_KHR is specified to remain consistent across driver versions and system reboots, it is not intended to be usable as a serializable persistent identifier for a device. It may change when a device is physically added to, removed from, or moved to a different connector in a system while that system is powered down. Further, there is no reasonable way to verify with conformance testing that a given device retains the same UUID in a given system across all driver versions supported in that system. While implementations should make every effort to report consistent device UUIDs across driver versions, applications should avoid relying on the persistence of this value for uses other than identifying compatible devices for external object sharing purposes.

Chapter 37. Extended versioning

This extension introduces new platform and device queries that return detailed version information to applications. It makes it possible to return the exact revision of the specification or intermediate languages supported by an implementation. It also enables implementations to communicate a version number for each of the extensions they support and remove the requirement for applications to process strings to test for the presence of an extension or intermediate language or built-in kernel.

37.1. General information

37.1.1. Name Strings

cl_khr_extended_versioning

37.1.2. Contributors

Kévin Petit, Arm Ltd. Ben Ashbaugh, Intel Alastair Murray, Codeplay Software Ltd. Einar Hov, Arm Ltd.

37.1.3. Version history

Date	Version	Description
2020-02-12	1.0.0	Initial version.

37.1.4. Dependencies

This extension is written against the OpenCL Specification Version 2.2, Revision 11.

This extension requires OpenCL 1.0.

37.2. New API Types

37.2.1. Version

This extension introduces a new scheme to encode detailed (major, minor, patch/revision) version information into a single 32-bit unsigned integer:

- The major version is using bits 31-22
- The minor version is using bits 21-12
- The patch version is using bits 11-0

This scheme enables two versions to be ordered using the standard C/C++ operators. Macros are provided to extract individual fields or compose a full version from the individual fields.

```
typedef cl_uint cl_version_khr;
#define CL VERSION MAJOR BITS KHR (10)
#define CL_VERSION_MINOR_BITS_KHR (10)
#define CL_VERSION_PATCH_BITS_KHR (12)
#define CL_VERSION_MAJOR_MASK_KHR ((1 << CL_VERSION_MAJOR_BITS_KHR) - 1)
#define CL_VERSION_MINOR_MASK_KHR ((1 << CL_VERSION_MINOR_BITS_KHR) - 1)
#define CL VERSION PATCH MASK KHR ((1 << CL VERSION PATCH BITS KHR) - 1)
#define CL_VERSION_MAJOR_KHR(version) \
        ((version) >> (CL_VERSION_MINOR_BITS_KHR + CL_VERSION_PATCH_BITS_KHR))
#define CL_VERSION_MINOR_KHR(version) \
        (((version) >> CL_VERSION_PATCH_BITS_KHR) & CL VERSION MINOR MASK KHR)
#define CL_VERSION_PATCH_KHR(version) ((version) & CL_VERSION_PATCH_MASK_KHR)
#define CL MAKE VERSION KHR(major, minor, patch) \
    ((((major) & CL_VERSION_MAJOR_MASK_KHR) << (CL_VERSION_MINOR_BITS_KHR +</pre>
CL_VERSION_PATCH_BITS_KHR)) | \
     (((minor) & CL VERSION MINOR MASK KHR) << CL VERSION PATCH BITS KHR) | \
     ((patch) & CL_VERSION_PATCH_MASK_KHR))
```

37.2.2. Name and version

This extension adds a structure that can be used to describe a combination of a name alongside a version number:

```
#define CL_NAME_VERSION_MAX_NAME_SIZE_KHR 64

typedef struct _cl_name_version_khr {
    cl_version_khr version;
    char name[CL_NAME_VERSION_MAX_NAME_SIZE_KHR];
} cl_name_version_khr;
```

The name field is an array of CL_NAME_VERSION_MAX_NAME_SIZE_KHR bytes used as storage for a NUL-terminated string whose maximum length is therefore CL_NAME_VERSION_MAX_NAME_SIZE_KHR - 1.

37.3. New API Enums

Accepted value for the *param_name* parameter to **clGetPlatformInfo**:

```
CL_PLATFORM_NUMERIC_VERSION_KHR
CL_PLATFORM_EXTENSIONS_WITH_VERSION_KHR
```

Accepted value for the *param_name* parameter to **clGetDeviceInfo**:

CL_DEVICE_NUMERIC_VERSION_KHR
CL_DEVICE_OPENCL_C_NUMERIC_VERSION_KHR
CL_DEVICE_EXTENSIONS_WITH_VERSION_KHR
CL_DEVICE_ILS_WITH_VERSION_KHR
CL_DEVICE_BUILT_IN_KERNELS_WITH_VERSION_KHR

37.4. Modifications to the OpenCL API Specification

(Modify Section 4.1, Querying Platform Info)

(Add the following to Table 3, List of supported param_names by clGetPlatformInfo)

cl_platform_info	Return Type	Description
CL_PLATFORM_NUMERIC_VERSION_KHR	cl_version_khr	Returns detailed (major, minor, patch) numeric version information. The major and minor version numbers returned must match those returned via CL_PLATFORM_VERSION.
CL_PLATFORM_EXTENSIONS_WITH_VERSION_KHR	cl_name_version_khr[]	Returns an array of description (name and version) structures. The same extension name must not be reported more than once. The list of extensions reported must match the list reported via CL_PLATFORM_EXTENSIONS.

(Modify Section 4.2, Querying Devices)

(Add the following to Table 5, List of supported param_names by clGetDeviceInfo)

cl_device_info	Return Type	Description
CL_DEVICE_NUMERIC_VERSION_KHR	cl_version_khr	Returns detailed (major, minor, patch) numeric version information. The major and minor version numbers returned must match those returned via CL_DEVICE_VERSION.
CL_DEVICE_OPENCL_C_NUMERIC_VERSION _KHR	cl_version_khr	Returns detailed (major, minor, patch) numeric version information. The major and minor version numbers returned must match those returned via CL_DEVICE_OPENCL_C_VERSION.

cl_device_info	Return Type	Description
CL_DEVICE_EXTENSIONS_WITH_VERSION_ KHR	cl_name_version_khr[]	Returns an array of description (name and version) structures. The same extension name must not be reported more than once. The list of extensions reported must match the list reported via CL_DEVICE_EXTENSIONS.
CL_DEVICE_ILS_WITH_VERSION_KHR	cl_name_version_khr[]	Returns an array of descriptions (name and version) for all supported Intermediate Languages. Intermediate Languages with the same name may be reported more than once but each name and major/minor version combination may only be reported once. The list of intermediate languages reported must match the list reported via CL_DEVICE_IL_VERSION.
CL_DEVICE_BUILT_IN_KERNELS_WITH_VE RSION_KHR	cl_name_version_khr[]	Returns an array of descriptions for the built-in kernels supported by the device. Each built-in kernel may only be reported once. The list of reported kernels must match the list returned via CL_DEVICE_BUILT_IN_KERNELS.

37.5. Conformance tests

- 1. Each of the new queries described in this extension must be attempted and succeed.
- 2. It must be verified that the information returned by all queries that extend existing queries is consistent with the information returned by existing queries.
- 3. Some of the queries introduced by this extension impose uniqueness constraints on the list of returned values. It must be verified that these constraints are satisfied.

37.6. Issues

1. What compatibility policy should we define? e.g. a *revision* has to be backwards-compatible with previous ones

RESOLVED: No general rules as that wouldn't be testable. Here's a recommended policy:

- Patch version bump: only clarifications and small/obvious bugfixes.
- Minor version bump: backwards-compatible changes only.

- Major version bump: backwards compatibility may break.
- 2. Do we want versioning for built-in kernels as returned by CL_DEVICE_BUILT_IN_KERNELS?

RESOLVED: No immediate use-case for versioning but being able to get a list of individual kernels without parsing a string is desirable. Adding CL_DEVICE_BUILT_IN_KERNELS_WITH_VERSION_KHR.

3. What is the behaviour of the queries that return an array of structures when there are no elements to return?

RESOLVED: The query succeeds and the size returned is zero.

4. What value should be returned when version information is not available?

RESOLVED: If a patch version is not available, it should be reported as 0. If no version information is available, 0.0.0 should be reported. These values have been chosen as they are guaranteed to be lower than or equal to any other version.

5. Should we add a query to report SPIR-V extended instruction sets?

RESOLVED: It is unlikely that we will introduce many SPIR-V extended instruction sets without an accompanying API extension. Decided not to do this.

6. Should the queries for which the old-style query doesn't exist in a given OpenCL version be present (e.g. CL_DEVICE_BUILT_IN_KERNELS_WITH_VERSION_KHR prior to OpenCL 2.1 or without support for cl_khr_il_program or CL_DEVICE_OPENCL_C_NUMERIC_VERSION_KHR on OpenCL 1.0)?

RESOLVED: All the queries are always present. CL_DEVICE_BUILT_IN_KERNELS_WITH_VERSION_KHR returns an empty set when Intermediate Languages are not supported. CL_DEVICE_OPENCL_C_NUMERIC_VERSION_KHR always returns 1.0 on an OpenCL_1.0 platform.

7. Is reporting multiple Intermediate Languages with the same name and major/minor versions but differing patch versions allowed?

RESOLVED: No. This isn't aligned with the intended use for patch versions and makes it harder for implementations to guarantee consistency with the existing IL queries.

Chapter 38. Extended Subgroup Functions

38.1. Overview

This section describes a family of extensions that provide extended subgroup functionality. The extensions in this family are:

- cl_khr_subgroup_extended_types
- cl_khr_subgroup_non_uniform_vote
- cl_khr_subgroup_ballot
- cl_khr_subgroup_non_uniform_arithmetic
- cl_khr_subgroup_shuffle
- cl_khr_subgroup_shuffle_relative
- cl_khr_subgroup_clustered_reduce

The functionality added by these extensions includes:

- Additional data type support for subgroup broadcast, scan, and reduction functions;
- The ability to elect a single work item from a subgroup to perform a task;
- The ability to hold votes among work items in a subgroup;
- The ability to collect and operate on ballots from work items in the subgroup;
- The ability to use some subgroup functions, such as any, all, broadcasts, scans, and reductions within non-uniform flow control;
- Additional scan and reduction operators;
- Additional ways to exchange data among work items in a subgroup;
- Clustered reductions, that operate on a subset of work items in the subgroup.

This section describes changes to the OpenCL C Language for these extensions. There are no new API functions or enums added by these extensions.

38.2. General information

38.2.1. Version history

For all of the extensions described in this section:

Date	Version	Description
2020-12-15	1.0.0	First assigned version.

38.3. Summary of New OpenCL C Functions

```
// These functions are available to devices supporting
// cl khr subgroup extended types:
// Note: Existing functions supporting additional data types.
gentype sub_group_broadcast( gentype value, uint index )
gentype sub group reduce add( gentype value )
gentype sub_group_reduce_min( gentype value )
gentype sub_group_reduce_max( gentype value )
gentype sub_group_scan_inclusive_add( gentype value )
gentype sub_group_scan_inclusive_min( gentype value )
gentype sub_group_scan_inclusive_max( gentype value )
gentype sub_group_scan_exclusive_add( gentype value )
gentype sub_group_scan_exclusive_min( gentype value )
gentype sub_group_scan_exclusive_max( gentype value )
// These functions are available to devices supporting
// cl_khr_subgroup_non_uniform_vote:
int sub group elect()
int sub_group_non_uniform_all( int predicate )
int sub group non uniform any( int predicate )
int sub_group_non_uniform_all_equal( gentype value )
// These functions are available to devices supporting
// cl_khr_subgroup_ballot:
gentype sub_group_non_uniform_broadcast( gentype value, uint index )
gentype sub_group_broadcast_first( gentype value )
uint4 sub group ballot( int predicate )
int sub_group_inverse_ballot( uint4 value )
int sub_group_ballot_bit_extract( uint4 value, uint index )
uint sub_group_ballot_bit_count( uint4 value )
uint sub_group_ballot_inclusive_scan( uint4 value )
uint sub_group_ballot_exclusive_scan( uint4 value )
uint sub_group_ballot_find_lsb( uint4 value )
uint sub_group_ballot_find_msb( uint4 value )
uint4 get_sub_group_eq_mask()
uint4 get_sub_group_ge_mask()
uint4 get_sub_group_gt_mask()
uint4 get_sub_group_le_mask()
uint4 get_sub_group_lt_mask()
```

```
// These functions are available to devices supporting
// cl_khr_subgroup_non_uniform_arithmetic:
gentype sub_group_non_uniform_reduce_add( gentype value )
gentype sub group non uniform reduce mul( gentype value )
gentype sub_group_non_uniform_reduce_min( gentype value )
gentype sub_group_non_uniform_reduce_max( gentype value )
gentype sub group non uniform reduce and( gentype value )
gentype sub_group_non_uniform_reduce_or( gentype value )
gentype sub_group_non_uniform_reduce_xor( gentype value )
int
        sub group non uniform reduce logical and( int predicate )
int
        sub_group_non_uniform_reduce_logical_or( int predicate )
int
        sub_group_non_uniform_reduce_logical_xor( int predicate )
gentype sub_group_non_uniform_scan_inclusive_add( gentype value )
gentype sub_group_non_uniform_scan_inclusive_mul( gentype value )
gentype sub group non uniform scan inclusive min( gentype value )
gentype sub_group_non_uniform_scan_inclusive_max( gentype value )
gentype sub_group_non_uniform_scan_inclusive_and( gentype value )
gentype sub group non uniform scan inclusive or( gentype value )
gentype sub_group_non_uniform_scan_inclusive_xor( gentype value )
int
        sub_group_non_uniform_scan_inclusive_logical_and( int predicate )
        sub group non uniform scan inclusive logical or( int predicate )
int
int
        sub_group_non_uniform_scan_inclusive_logical_xor( int predicate )
gentype sub group non uniform scan exclusive add( gentype value )
gentype sub_group_non_uniform_scan_exclusive_mul( gentype value )
gentype sub group non uniform scan exclusive min( gentype value )
gentype sub group non uniform scan exclusive max( gentype value )
gentype sub_group_non_uniform_scan_exclusive_and( gentype value )
gentype sub group non uniform scan exclusive or( gentype value )
gentype sub group non uniform scan exclusive xor( gentype value )
        sub_group_non_uniform_scan_exclusive_logical_and( int predicate )
int
int
        sub group non uniform scan exclusive logical or( int predicate )
        sub group non uniform scan exclusive logical xor( int predicate )
int
// These functions are available to devices supporting
// cl_khr_subgroup_shuffle:
gentype sub group shuffle( gentype value, uint index )
gentype sub_group_shuffle_xor( gentype value, uint mask )
// These functions are available to devices supporting
// cl khr subgroup shuffle relative:
gentype sub_group_shuffle_up( gentype value, uint delta )
gentype sub_group_shuffle_down( gentype value, uint delta )
// These functions are available to devices supporting
// cl_khr_subgroup_clustered_reduce:
```

```
gentype sub_group_clustered_reduce_add( gentype value, uint clustersize )
gentype sub_group_clustered_reduce_mul( gentype value, uint clustersize )
gentype sub_group_clustered_reduce_min( gentype value, uint clustersize )
gentype sub_group_clustered_reduce_max( gentype value, uint clustersize )
gentype sub_group_clustered_reduce_and( gentype value, uint clustersize )
gentype sub_group_clustered_reduce_or( gentype value, uint clustersize )
gentype sub_group_clustered_reduce_xor( gentype value, uint clustersize )
int sub_group_clustered_reduce_logical_and( int predicate, uint clustersize )
int sub_group_clustered_reduce_logical_or( int predicate, uint clustersize )
int sub_group_clustered_reduce_logical_xor( int predicate, uint clustersize )
```

38.4. Extended Types

This section describes functionality added by cl_khr_subgroup_extended_types. This extension adds additional supported data types to the existing subgroup broadcast, scan, and reduction functions.

38.4.1. Modify the Existing Section Describing Subgroup Functions

Modify the first paragraph in this section that describes gentype type support for the subgroup broadcast, scan, and reduction functions to add scalar char, uchar, short, and ushort support, and to additionally add built-in vector type support for broadcast specifically. The functions in the table and their descriptions remain unchanged by this extension:

The table below describes OpenCL C programming language built-in functions that operate on a subgroup level. These built-in functions must be encountered by all work items in the subgroup executing the kernel. We use the generic type name gentype to indicate the built-in scalar data types char, uchar, short, ushort, int, uint, long, ulong, float, double (if double precision is supported), or half (if half precision is supported).

For the sub_group_broadcast function, the generic type name gentype may additionally be one of the supported built-in vector data types charn, ucharn, shortn, ushortn, intn, uintn, longn, ulongn, floatn, doublen (if double precision is supported), or halfn (if half precision is supported).

38.5. Votes and Elections

This section describes functionality added by cl_khr_subgroup_non_uniform_vote. This extension adds the ability to elect a single work item from a subgroup to perform a task and to hold votes among work items in a subgroup.

38.5.1. Add a new Section 6.15.X - Subgroup Vote and Elect Built-in Functions

The table below describes the OpenCL C programming language built-in functions to elect a single work item in a subgroup to perform a task and to collectively vote to determine a boolean condition for the subgroup. These functions need not be encountered by all work items in a subgroup executing the kernel. For the functions below, the generic type name gentype may be the one of the supported built-in scalar data types char, uchar, short, ushort, int, uint, long, ulong, float, double (if double precision is supported), or half (if half precision is supported).

Function	Description
<pre>int sub_group_elect()</pre>	Elects a single work item in the subgroup to perform a task. This function will return true (nonzero) for the active work item in the subgroup with the smallest subgroup local ID, and false (zero) for all other active work items in the subgroup.
<pre>int sub_group_non_uniform_all(int predicate)</pre>	Examines <i>predicate</i> for all active work items in the subgroup and returns a non-zero value if <i>predicate</i> is non-zero for all active work items in the subgroup and zero otherwise. Note: This behavior is the same as sub_group_all from cl_khr_subgroup and OpenCL 2.1, except this function need not be encountered by all work items in the subgroup executing the kernel.
<pre>int sub_group_non_uniform_any(int predicate)</pre>	Examines <i>predicate</i> for all active work items in the subgroup and returns a non-zero value if <i>predicate</i> is non-zero for any active work item in the subgroup and zero otherwise. Note: This behavior is the same as sub_group_any from cl_khr_subgroups and OpenCL 2.1, except this function need not be encountered by all work items in the subgroup executing the kernel.
<pre>int sub_group_non_uniform_all_equal(gentype value)</pre>	Examines <i>value</i> for all active work items in the subgroup and returns a non-zero value if <i>value</i> is equivalent for all active invocations in the subgroup and zero otherwise. Integer types use a bitwise test for equality. Floating-point types use an ordered floating-point test for equality.

38.6. Ballots

This section describes functionality added by cl_khr_subgroup_ballot. This extension adds the ability to collect and operate on ballots from work items in the subgroup.

38.6.1. Add a new Section 6.15.X - Subgroup Ballot Built-in Functions

The table below describes the OpenCL C programming language built-in functions to allow work items in a subgroup to collect and operate on ballots from work items in the subgroup. These

functions need not be encountered by all work items in a subgroup executing the kernel.

For the sub_group_non_uniform_broadcast and sub_group_broadcast_first functions, the generic type name gentype may be one of the supported built-in scalar data types char, uchar, short, ushort, int, uint, long, ulong, float, double (if double precision is supported), or half (if half precision is supported).

For the sub_group_non_uniform_broadcast function, the generic type name gentype may additionally be one of the supported built-in vector data types charn, ucharn, shortn, ushortn, intn, uintn, longn, ulongn, floatn, doublen (if double precision is supported), or halfn (if half precision is supported).

Function	Description
<pre>gentype sub_group_non_uniform_broadcast(gentype value, uint index)</pre>	Returns <i>value</i> for the work item with subgroup local ID equal to <i>index</i> . Behavior is undefined when the value of <i>index</i> is not equivalent for all active work items in the subgroup. The return value is undefined if the work item with subgroup local ID equal to <i>index</i> is inactive or if <i>index</i> is greater than or equal to the size of the subgroup.
<pre>gentype sub_group_broadcast_first(gentype value)</pre>	Returns <i>value</i> for the work item with the smallest subgroup local ID among active work items in the subgroup.
<pre>uint4 sub_group_ballot(int predicate)</pre>	Returns a bitfield combining the <i>predicate</i> values from all work items in the subgroup. Bit zero of the first vector component represents the subgroup local ID zero, with higher-order bits and subsequent vector components representing, in order, increasing subgroup local IDs. The representative bit in the bitfield is set if the work item is active and the <i>predicate</i> is non-zero, and is unset otherwise.

Function	Description
<pre>int sub_group_inverse_ballot(uint4 value)</pre>	Returns the predicate value for this work item in the subgroup from the bitfield <i>value</i> representing predicate values from all work items in the subgroup. The predicate return value will be non-zero if the bit in the bitfield <i>value</i> for this work item is set, and zero otherwise. Behavior is undefined when <i>value</i> is not equivalent for all active work items in the subgroup. This is a specialized function that may perform better than the equivalent
	<pre>sub_group_ballot_bit_extract on some implementations.</pre>
<pre>int sub_group_ballot_bit_extract(uint4 value, uint index)</pre>	Returns the predicate value for the work item with subgroup local ID equal to <i>index</i> from the bitfield <i>value</i> representing predicate values from all work items in the subgroup. The predicate return value will be non-zero if the bit in the bitfield <i>value</i> for the work item with subgroup
	local ID equal to <i>index</i> is set, and zero otherwise.
	The predicate return value is undefined if the work item with subgroup local ID equal to <i>index</i> is greater than or equal to the size of the subgroup.
uint sub_group_ballot_bit_count(uint4 value)	Returns the number of bits that are set in the bitfield <i>value</i> , only considering the bits in <i>value</i> that represent predicate values from all work items in the subgroup.
uint sub_group_ballot_inclusive_scan(uint4 value)	Returns the number of bits that are set in the bitfield <i>value</i> , only considering the bits in <i>value</i> representing work items with a subgroup local ID less than or equal to this work item's subgroup local ID.
uint sub_group_ballot_exclusive_scan(uint4 value)	Returns the number of bits that are set in the bitfield <i>value</i> , only considering the bits in <i>value</i> representing work items with a subgroup local ID less than this work item's subgroup local ID.

Function	Description
uint sub_group_ballot_find_lsb(uint4 value)	Returns the smallest subgroup local ID with a bit set in the bitfield <i>value</i> , only considering the bits in <i>value</i> that represent predicate values from all work items in the subgroup. If no bits representing predicate values from all work items in the subgroup are set in the bitfield <i>value</i> then the return value is undefined.
uint sub_group_ballot_find_msb(uint4 value)	Returns the largest subgroup local ID with a bit set in the bitfield <i>value</i> , only considering the bits in <i>value</i> that represent predicate values from all work items in the subgroup. If no bits representing predicate values from all work items in the subgroup are set in the bitfield <i>value</i> then the return value is undefined.
<pre>uint4 get_sub_group_eq_mask()</pre>	Generates a bitmask where the bit is set in the bitmask if the bit index equals the subgroup local ID and unset otherwise. Bit zero of the first vector component represents the subgroup local ID zero, with higher-order bits and subsequent vector components representing, in order, increasing subgroup local IDs.
<pre>uint4 get_sub_group_ge_mask()</pre>	Generates a bitmask where the bit is set in the bitmask if the bit index is greater than or equal to the subgroup local ID and less than the maximum subgroup size, and unset otherwise. Bit zero of the first vector component represents the subgroup local ID zero, with higher-order bits and subsequent vector components representing, in order, increasing subgroup local IDs.
<pre>uint4 get_sub_group_gt_mask()</pre>	Generates a bitmask where the bit is set in the bitmask if the bit index is greater than the subgroup local ID and less than the maximum subgroup size, and unset otherwise. Bit zero of the first vector component represents the subgroup local ID zero, with higher-order bits and subsequent vector components representing, in order, increasing subgroup local IDs.

Function	Description
<pre>uint4 get_sub_group_le_mask()</pre>	Generates a bitmask where the bit is set in the bitmask if the bit index is less than or equal to the subgroup local ID and unset otherwise. Bit zero of the first vector component represents the subgroup local ID zero, with higher-order bits and subsequent vector components representing, in order, increasing subgroup local IDs.
<pre>uint4 get_sub_group_lt_mask()</pre>	Generates a bitmask where the bit is set in the bitmask if the bit index is less than the subgroup local ID and unset otherwise. Bit zero of the first vector component represents the subgroup local ID zero, with higher-order bits and subsequent vector components representing, in order, increasing subgroup local IDs.

38.7. Non-Uniform Arithmetic

This section describes functionality added by cl_khr_subgroup_non_uniform_arithmetic. This extension adds the ability to use some subgroup functions within non-uniform flow control, including additional scan and reduction operators.

38.7.1. Add a new Section 6.15.X - Non Uniform Subgroup Scan and Reduction Built-in Functions

38.7.1.1. Arithmetic Operations

The table below describes the OpenCL C programming language built-in functions that perform simple arithmetic operations across work items in a subgroup. These functions need not be encountered by all work items in a subgroup executing the kernel. For the functions below, the generic type name gentype may be one of the supported built-in scalar data types char, uchar, short, ushort, int, uint, long, ulong, float, double (if double precision is supported), or half (if half precision is supported).

```
gentype sub_group_non_uniform_reduce_add(
    gentype value )
gentype sub_group_non_uniform_reduce_min(
    gentype value )
gentype sub_group_non_uniform_reduce_max(
    gentype value )
gentype sub_group_non_uniform_reduce_mul(
    gentype value )
```

```
gentype sub_group_non_uniform_scan_inclusive_add(
    gentype value )
gentype sub_group_non_uniform_scan_inclusive_min(
    gentype value )
gentype sub_group_non_uniform_scan_inclusive_max(
    gentype value )
gentype sub_group_non_uniform_scan_inclusive_mul(
    gentype value )
```

Description

Returns the summation, multiplication, minimum, or maximum of *value* for all active work items in the subgroup.

Note: This behavior is the same as the **add**, **min**, and **max** reduction built-in functions from cl_khr_subgroups and OpenCL 2.1, except these functions support additional types and need not be encountered by all work items in the subgroup executing the kernel.

Returns the result of an inclusive scan operation, which is the summation, multiplication, minimum, or maximum of *value* for all active work items in the subgroup with a subgroup local ID less than or equal to this work item's subgroup local ID.

Note: This behavior is the same as the add, min, and max inclusive scan built-in functions from cl_khr_subgroups and OpenCL 2.1, except these functions support additional types and need not be encountered by all work items in the subgroup executing the kernel.

```
gentype sub_group_non_uniform_scan_exclusive_add(
    gentype value )
gentype sub_group_non_uniform_scan_exclusive_min(
    gentype value )
gentype sub_group_non_uniform_scan_exclusive_max(
    gentype value )
gentype sub_group_non_uniform_scan_exclusive_mul(
    gentype value )
```

Description

Returns the result of an exclusive scan operation, which is the summation, multiplication, minimum, or maximum of *value* for all active work items in the subgroup with a subgroup local ID less than this work item's subgroup local ID.

If there is no active work item in the subgroup with a subgroup local ID less than this work item's subgroup local ID then an identity value I is returned. For add, the identity value is 0. For min, the identity value is the largest representable value for integer types, or +INF for floating point types. For max, the identity value is the minimum representable value for integer types, or -INF for floating point types. For mul, the identity value is 1.

Note: This behavior is the same as the add, min, and max exclusive scan built-in functions from cl_khr_subgroups and OpenCL 2.1, except these functions support additional types and need not be encountered by all work items in the subgroup executing the kernel.

Note: The order of floating-point operations is not guaranteed for the subgroup scan and reduction built-in functions that operate on floating point types, and the order of operations may additionally be non-deterministic for a given subgroup.

38.7.1.2. Bitwise Operations

The table below describes the OpenCL C programming language built-in functions that perform simple bitwise integer operations across work items in a subgroup. These functions need not be encountered by all work items in a subgroup executing the kernel. For the functions below, the generic type name gentype may be one of the supported built-in scalar data types char, uchar, short, ushort, int, uint, long, and ulong.

gentype sub_group_non_uniform_reduce_and(gentype value) gentype sub_group_non_uniform_reduce_or(gentype value) gentype sub_group_non_uniform_reduce_xor(

gentype value)

Description

Returns the bitwise **and**, **or**, or **xor** of *value* for all active work items in the subgroup.

```
gentype sub_group_non_uniform_scan_inclusive_and(
    gentype value )
gentype sub_group_non_uniform_scan_inclusive_or(
    gentype value )
gentype sub_group_non_uniform_scan_inclusive_xor(
    gentype value )
```

Returns the result of an inclusive scan operation, which is the bitwise **and**, **or**, or **xor** of *value* for all active work items in the subgroup with a subgroup local ID less than or equal to this work item's subgroup local ID.

```
gentype sub_group_non_uniform_scan_exclusive_and(
    gentype value )
gentype sub_group_non_uniform_scan_exclusive_or(
    gentype value )
gentype sub_group_non_uniform_scan_exclusive_xor(
    gentype value )
```

Returns the result of an exclusive scan operation, which is the bitwise **and**, **or**, or **xor** of *value* for all active work items in the subgroup with a subgroup local ID less than this work item's subgroup local ID.

If there is no active work item in the subgroup with a subgroup local ID less than this work item's subgroup local ID then an identity value I is returned. For **and**, the identity value is ~0 (all bits set). For **or** and **xor**, the identity value is 0.

38.7.1.3. Logical Operations

The table below describes the OpenCL C programming language built-in functions that perform simple logical operations across work items in a subgroup. These functions need not be encountered by all work items in a subgroup executing the kernel. For these functions, a non-zero *predicate* argument or return value is logically true and a zero *predicate* argument or return value is logically false.

Function Description

```
int sub_group_non_uniform_reduce_logical_and(
    int predicate )
int sub_group_non_uniform_reduce_logical_or(
    int predicate )
int sub_group_non_uniform_reduce_logical_xor(
    int predicate )
```

Returns the logical **and**, **or**, or **xor** of *predicate* for all active work items in the subgroup.

```
int sub_group_non_uniform_scan_inclusive_logical_and(
    int predicate )
int sub_group_non_uniform_scan_inclusive_logical_or(
    int predicate )
int sub_group_non_uniform_scan_inclusive_logical_xor(
    int predicate )
```

Returns the result of an inclusive scan operation, which is the logical **and**, **or**, or **xor** of *predicate* for all active work items in the subgroup with a subgroup local ID less than or equal to this work item's subgroup local ID.

```
int sub_group_non_uniform_scan_exclusive_logical_and(
    int predicate )
int sub_group_non_uniform_scan_exclusive_logical_or(
    int predicate )
int sub_group_non_uniform_scan_exclusive_logical_xor(
    int predicate )
```

Returns the result of an exclusive scan operation, which is the logical **and**, **or**, or **xor** of *predicate* for all active work items in the subgroup with a subgroup local ID less than this work item's subgroup local ID.

If there is no active work item in the subgroup with a subgroup local ID less than this work item's subgroup local ID then an identity value I is returned. For and, the identity value is true (non-zero). For or and xor, the identity value is false (zero).

38.8. General Purpose Shuffles

This section describes functionality added by cl_khr_subgroup_shuffle. This extension adds additional ways to exchange data among work items in a subgroup.

38.8.1. Add a new Section 6.15.X - Subgroup Shuffle Built-in Functions

The table below describes the OpenCL C programming language built-in functions that allow work items in a subgroup to exchange data. These functions need not be encountered by all work items in a subgroup executing the kernel. For the functions below, the generic type name gentype may be one of the supported built-in scalar data types char, uchar, short, ushort, int, uint, long, ulong, float,

double (if double precision is supported), or half (if half precision is supported).

Function	Description
<pre>gentype sub_group_shuffle(gentype value, uint index)</pre>	Returns <i>value</i> for the work item with subgroup local ID equal to <i>index</i> . The shuffle <i>index</i> need not be the same for all work items in the subgroup. The return value is undefined if the work item with subgroup local ID equal to <i>index</i> is inactive or if <i>index</i> is greater than or equal to the size of the subgroup.
<pre>gentype sub_group_shuffle_xor(gentype value, uint mask)</pre>	Returns <i>value</i> for the work item with subgroup local ID equal to this work item's subgroup local ID xor'd with <i>mask</i> . The shuffle <i>mask</i> need not be the same for all work items in the subgroup. The return value is undefined if the work item with subgroup local ID equal to the saleulated.
	with subgroup local ID equal to the calculated index is inactive or if the calculated index is greater than or equal to the size of the subgroup. This is a specialized function that may perform better than the equivalent sub_group_shuffle on some implementations.

38.9. Relative Shuffles

This section describes functionality added by cl_khr_subgroup_shuffle_relative. This extension adds specialized ways to exchange data among work items in a subgroup that may perform better on some implementations.

38.9.1. Add a new Section 6.15.X - Subgroup Relative Shuffle Built-in Functions

The table below describes specialized OpenCL C programming language built-in functions that allow work items in a subgroup to exchange data. These functions need not be encountered by all work items in a subgroup executing the kernel. For the functions below, the generic type name gentype may be one of the supported built-in scalar data types char, uchar, short, ushort, int, uint, long, ulong, float, double (if double precision is supported), or half (if half precision is supported).

Function	Description
<pre>gentype sub_group_shuffle_up(gentype value, uint delta)</pre>	Returns <i>value</i> for the work item with subgroup local ID equal to this work item's subgroup local ID minus <i>delta</i> . The shuffle <i>delta</i> need not be the same for all work items in the subgroup.
	The return value is undefined if the work item with subgroup local ID equal to the calculated index is inactive, or <i>delta</i> is greater than this work item's subgroup local ID.
	This is a specialized function that may perform better than the equivalent <pre>sub_group_shuffle</pre> on some implementations.
<pre>gentype sub_group_shuffle_down(gentype value, uint delta)</pre>	Returns <i>value</i> for the work item with subgroup local ID equal to this work item's subgroup local ID plus <i>delta</i> . The shuffle <i>delta</i> need not be the same for all work items in the subgroup.
	The return value is undefined if the work item with subgroup local ID equal to the calculated index is inactive, or this work item's subgroup local ID plus <i>delta</i> is greater than or equal to the size of the subgroup.
	This is a specialized function that may perform better than the equivalent <pre>sub_group_shuffle</pre> on some implementations.

38.10. Clustered Reductions

This section describes functionality added by cl_khr_subgroup_clustered_reduce. This extension adds support for clustered reductions that operate on a subset of work items in the subgroup.

38.10.1. Add a new Section 6.15.X - Subgroup Clustered Reduction Built-in Functions

This section describes arithmetic operations that are performed subset of work items in a subgroup, referred to as a cluster. A cluster is described by a specified cluster size. Work items in a subgroup are assigned to clusters such that for cluster size n, the n work items in the subgroup with the smallest subgroup local IDs are assigned to the first cluster, then the n remaining work items with the smallest subgroup local IDs are assigned to the next cluster, and so on. The specified cluster size must be an integer constant expression that is a power-of-two. Behavior is undefined if the specified cluster size is greater than the maximum size of a subgroup within the dispatch.

38.10.1.1. Arithmetic Operations

The table below describes the OpenCL C programming language built-in functions that perform simple arithmetic operations on a cluster of work items in a subgroup. These functions need not be encountered by all work items in a subgroup executing the kernel. For the functions below, the generic type name gentype may be one of the supported built-in scalar data types char, uchar, short, ushort, int, uint, long, ulong, float, double (if double precision is supported), or half (if half precision is supported).

Function | gentype sub_group_clustered_reduce_add(| gentype value, uint clustersize) | gentype sub_group_clustered_reduce_mul(| gentype value, uint clustersize) | gentype sub_group_clustered_reduce_min(| gentype value, uint clustersize) | gentype sub_group_clustered_reduce_max(| gentype value, uint clustersize)

Note: The order of floating-point operations is not guaranteed for the subgroup clustered reduction built-in functions that operate on floating point types, and the order of operations may additionally be non-deterministic for a given subgroup.

38.10.1.2. Bitwise Operations

The table below describes the OpenCL C programming language built-in functions to perform simple bitwise integer operations across a cluster of work items in a subgroup. These functions need not be encountered by all work items in a subgroup executing the kernel. For the functions below, the generic type name gentype may be the one of the supported built-in scalar data types char, uchar, short, ushort, int, uint, long, or ulong.

Function	Description
<pre>gentype sub_group_clustered_reduce_and(gentype value, uint clustersize) gentype sub_group_clustered_reduce_or(gentype value, uint clustersize) gentype sub_group_clustered_reduce_xor(gentype value, uint clustersize)</pre>	Returns the bitwise and , or , or xor of <i>value</i> for all active work items in the subgroup within a cluster of the specified <i>clustersize</i> .

38.10.1.3. Logical Operations

The table below describes the OpenCL C programming language built-in functions to perform simple logical operations across a cluster of work items in a subgroup. These functions need not be encountered by all work items in a subgroup executing the kernel. For these functions, a non-zero

predicate argument or return value is logically true and a zero predicate argument or return value is logically false.

Function	Description
<pre>int sub_group_clustered_reduce_logical_and(int predicate, uint clustersize) int sub_group_clustered_reduce_logical_or(int predicate, uint clustersize) int sub_group_clustered_reduce_logical_xor(int predicate, uint clustersize)</pre>	Returns the logical and , or , or xor of <i>predicate</i> for all active work items in the subgroup within a cluster of the specified <i>clustersize</i> .

38.11. Function Mapping and Capabilities

This section describes a possible mapping between OpenCL built-in functions and SPIR-V instructions and required SPIR-V capabilities.

This section is informational and non-normative.

OpenCL C Function	SPIR-V BuiltIn or Instruction	Enabling SPIR-V Capability
For OpenCL 2.1 or cl_khr_subgroup	oups:	
get_sub_group_size	SubgroupSize	Kernel
get_max_sub_group_size	SubgroupMaxSize	Kernel
get_num_sub_groups	NumSubgroups	Kernel
get_enqueued_num_sub_groups	NumEnqueuedSubgroups	Kernel
get_sub_group_id	SubgroupId	Kernel
get_sub_group_local_id	SubgroupLocalInvocationId	Kernel
sub_group_barrier	OpControlBarrier	None Needed
sub_group_all	OpGroupAll	Groups
sub_group_any	OpGroupAny	Groups
sub_group_broadcast	OpGroupBroadcast	Groups
sub_group_reduce_add	OpGroupIAdd, OpGroupFAdd	Groups
sub_group_reduce_min	OpGroupSMin, OpGroupUMin, OpGroupFMin	Groups
sub_group_reduce_max	OpGroupSMax, OpGroupUMax, OpGroupFMax	Groups
sub_group_scan_exclusive_add	OpGroupIAdd, OpGroupFAdd	Groups
sub_group_scan_exclusive_min	OpGroupSMin, OpGroupUMin, OpGroupFMin	Groups

OpenCL C Function	SPIR-V BuiltIn or Instruction	Enabling SPIR-V Capability
sub_group_scan_exclusive_max	OpGroupSMax, OpGroupUMax, OpGroupFMax	Groups
sub_group_scan_inclusive_add	OpGroupIAdd, OpGroupFAdd	Groups
sub_group_scan_inclusive_min	OpGroupSMin, OpGroupUMin, OpGroupFMin	Groups
sub_group_scan_inclusive_max	OpGroupSMax, OpGroupUMax, OpGroupFMax	Groups
sub_group_reserve_read_pipe	OpGroupReserveReadPipePac kets	Pipes
sub_group_reserve_write_pipe	OpGroupReserveReadWritePa ckets	Pipes
sub_group_commit_read_pipe	OpGroupCommitReadPipe	Pipes
sub_group_commit_write_pipe	OpGroupCommitWritePipe	Pipes
<pre>get_kernel_sub_group_count_ for_ndrange</pre>	OpGetKernelNDrangeSubGro upCount	DeviceEnqueue
<pre>get_kernel_max_sub_group_size_ for_ndrange</pre>	OpGetKernelNDrangeMaxSub GroupSize	DeviceEnqueue
For cl_khr_subgroup_extended_ty Note: This extension adds new ty	pes: pes to uniform subgroup operation	ons.
sub_group_broadcast	OpGroupBroadcast	Groups
sub_group_reduce_add	OpGroupIAdd, OpGroupFAdd	Groups
sub_group_reduce_min	OpGroupSMin, OpGroupUMin, OpGroupFMin	Groups
sub_group_reduce_max	OpGroupSMax, OpGroupUMax, OpGroupFMax	Groups
sub_group_scan_exclusive_add	OpGroupIAdd, OpGroupFAdd	Groups
sub_group_scan_exclusive_min	OpGroupSMin, OpGroupUMin, OpGroupFMin	Groups
sub_group_scan_exclusive_max	OpGroupSMax, OpGroupUMax, OpGroupFMax	Groups
sub_group_scan_inclusive_add	OpGroupIAdd, OpGroupFAdd	Groups
sub_group_scan_inclusive_min	OpGroupSMin, OpGroupUMin,	Groups

OpGroupFMin

OpenCL C Function	SPIR-V BuiltIn or Instruction	Enabling SPIR-V Capability
sub_group_scan_inclusive_max	OpGroupSMax, OpGroupUMax, OpGroupFMax	Groups
For cl_khr_subgroup_non_uniform	_vote:	
sub_group_elect	OpGroupNonUniformElect	GroupNonUniform
sub_group_non_uniform_all	OpGroupNonUniformAll	GroupNonUniformVote
sub_group_non_uniform_any	OpGroupNonUniformAny	GroupNonUniformVote
<pre>sub_group_non_uniform_all_ equal</pre>	OpGroupNonUniformAllEqua l	GroupNonUniformVote
For cl_khr_subgroup_ballot:		
<pre>sub_group_non_uniform_ broadcast</pre>	OpGroupNonUniformBroadca st	GroupNonUniformBallot
sub_group_broadcast_first	OpGroupNonUniformBroadca stFirst	GroupNonUniformBallot
sub_group_ballot	OpGroupNonUniformBallot	GroupNonUniformBallot
sub_group_inverse_ballot	OpGroupNonUniformInverse Ballot	GroupNonUniformBallot
sub_group_ballot_bit_extract	OpGroupNonUniformBallotBi tExtract	GroupNonUniformBallot
sub_group_ballot_bit_count	OpGroupNonUniformBallotBi tCount	GroupNonUniformBallot
<pre>sub_group_ballot_inclusive_ scan</pre>	OpGroupNonUniformBallotBi tCount	GroupNonUniformBallot
<pre>sub_group_ballot_exclusive_ scan</pre>	OpGroupNonUniformBallotBi tCount	GroupNonUniformBallot
sub_group_ballot_find_lsb	OpGroupNonUniformBallotFi ndLSB	GroupNonUniformBallot
sub_group_ballot_find_msb	OpGroupNonUniformBallotFi ndMSB	GroupNonUniformBallot
get_sub_group_eq_mask	SubgroupEqMask	GroupNonUniformBallot
get_sub_group_ge_mask	SubgroupGeMask	GroupNonUniformBallot
get_sub_group_gt_mask	SubgroupGtMask	GroupNonUniformBallot
get_sub_group_le_mask	SubgroupLeMask	GroupNonUniformBallot
get_sub_group_lt_mask	SubgroupLtMask	GroupNonUniformBallot
For cl_khr_subgroup_non_uniform	_arithmetic:	
<pre>sub_group_non_uniform_reduce_ add</pre>	OpGroupNonUniformIAdd, OpGroupNonUniformFAdd	GroupNonUniformArithmetic

OpenCL C Function	SPIR-V BuiltIn or Instruction	Enabling SPIR-V Capability
<pre>sub_group_non_uniform_reduce_ mul</pre>	OpGroupNonUniformIMul, OpGroupNonUniformFMul	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_reduce_ min</pre>	OpGroupNonUniformSMin, OpGroupNonUniformUMin, OpGroupNonUniformFMin	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_reduce_ max</pre>	OpGroupNonUniformSMax, OpGroupNonUniformUMax, OpGroupNonUniformFMax	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_reduce_ and</pre>	OpGroupNonUniformBitwise And	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_reduce_ or</pre>	OpGroupNonUniformBitwise Or	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_reduce_ xor</pre>	OpGroupNonUniformBitwise Xor	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_reduce_ logical_and</pre>	OpGroupNonUniformLogical And	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_reduce_ logical_or</pre>	OpGroupNonUniformLogical Or	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_reduce_ logical_xor</pre>	OpGroupNonUniformLogical Xor	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_scan_ inclusive_add</pre>	OpGroupNonUniformIAdd, OpGroupNonUniformFAdd	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_scan_ inclusive_mul</pre>	OpGroupNonUniformIMul, OpGroupNonUniformFMul	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_scan_ inclusive_min</pre>	OpGroupNonUniformSMin, OpGroupNonUniformUMin, OpGroupNonUniformFMin	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_scan_ inclusive_max</pre>	OpGroupNonUniformSMax, OpGroupNonUniformUMax, OpGroupNonUniformFMax	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_scan_ inclusive_and</pre>	OpGroupNonUniformBitwise And	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_scan_ inclusive_or</pre>	OpGroupNonUniformBitwise Or	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_scan_ inclusive_xor</pre>	OpGroupNonUniformBitwise Xor	GroupNonUniformArithmetic
<pre>sub_group_non_uniform_scan_ inclusive_logical_and</pre>	OpGroupNonUniformLogical And	GroupNonUniformArithmetic

OpenCL C Function	SPIR-V BuiltIn or Instruction	Enabling SPIR-V Capability	
<pre>sub_group_non_uniform_scan_ inclusive_logical_or</pre>	OpGroupNonUniformLogical Or	GroupNonUniformArithmetic	
<pre>sub_group_non_uniform_scan_ inclusive_logical_xor</pre>	OpGroupNonUniformLogical Xor	GroupNonUniformArithmetic	
<pre>sub_group_non_uniform_scan_ exclusive_add</pre>	OpGroupNonUniformIAdd, OpGroupNonUniformFAdd	GroupNonUniformArithmetic	
<pre>sub_group_non_uniform_scan_ exclusive_mul</pre>	OpGroupNonUniformIMul, OpGroupNonUniformFMul	GroupNonUniformArithmetic	
<pre>sub_group_non_uniform_scan_ exclusive_min</pre>	OpGroupNonUniformSMin, OpGroupNonUniformUMin, OpGroupNonUniformFMin	GroupNonUniformArithmetic	
<pre>sub_group_non_uniform_scan_ exclusive_max</pre>	OpGroupNonUniformSMax, OpGroupNonUniformUMax, OpGroupNonUniformFMax	GroupNonUniformArithmetic	
<pre>sub_group_non_uniform_scan_ exclusive_and</pre>	OpGroupNonUniformBitwise And	GroupNonUniformArithmetic	
<pre>sub_group_non_uniform_scan_ exclusive_or</pre>	OpGroupNonUniformBitwise Or	GroupNonUniformArithmetic	
<pre>sub_group_non_uniform_scan_ exclusive_xor</pre>	OpGroupNonUniformBitwise Xor	GroupNonUniformArithmetic	
<pre>sub_group_non_uniform_scan_ exclusive_logical_and</pre>	OpGroupNonUniformLogical And	GroupNonUniformArithmetic	
<pre>sub_group_non_uniform_scan_ exclusive_logical_or</pre>	OpGroupNonUniformLogical Or	GroupNonUniformArithmetic	
<pre>sub_group_non_uniform_scan_ exclusive_logical_xor</pre>	OpGroupNonUniformLogical Xor	GroupNonUniformArithmetic	
For cl_khr_subgroup_shuffle:			
sub_group_shuffle	OpGroupNonUniformShuffle	GroupNonUniformShuffle	
sub_group_shuffle_xor	OpGroupNonUniformShuffle Xor	GroupNonUniformShuffle	
For cl_khr_subgroup_shuffle_rel	ative:		
sub_group_shuffle_up	OpGroupNonUniformShuffle Up	GroupNonUniformShuffleRel ative	
sub_group_shuffle_down	OpGroupNonUniformShuffle Down	GroupNonUniformShuffleRel ative	
For cl_khr_subgroup_clustered_r	For cl_khr_subgroup_clustered_reduce:		
sub_group_reduce_clustered_add	OpGroupNonUniformIAdd, OpGroupNonUniformFAdd	GroupNonUniformClustered	

OpenCL C Function	SPIR-V BuiltIn or Instruction	Enabling SPIR-V Capability
sub_group_reduce_clustered_mul	OpGroupNonUniformIMul, OpGroupNonUniformFMul	GroupNonUniformClustered
sub_group_reduce_clustered_min	OpGroupNonUniformSMin, OpGroupNonUniformUMin, OpGroupNonUniformFMin	GroupNonUniformClustered
<pre>sub_group_reduce_clustered_max</pre>	OpGroupNonUniformSMax, OpGroupNonUniformUMax, OpGroupNonUniformFMax	GroupNonUniformClustered
sub_group_reduce_clustered_and	OpGroupNonUniformBitwise And	GroupNonUniformClustered
sub_group_reduce_clustered_or	OpGroupNonUniformBitwise Or	GroupNonUniformClustered
sub_group_reduce_clustered_xor	OpGroupNonUniformBitwise Xor	GroupNonUniformClustered
<pre>sub_group_reduce_clustered_ logical_and</pre>	OpGroupNonUniformLogical And	GroupNonUniformClustered
<pre>sub_group_reduce_clustered_ logical_or</pre>	OpGroupNonUniformLogical Or	GroupNonUniformClustered
<pre>sub_group_reduce_clustered_ logical_xor</pre>	OpGroupNonUniformLogical Xor	GroupNonUniformClustered

Chapter 39. PCI Bus Information Query

This extension adds a new query to obtain PCI bus information about an OpenCL device.

Not all OpenCL devices have PCI bus information, either due to the device not being connected to the system through a PCI interface or due to platform specific restrictions and policies. Thus this extension is only expected to be supported by OpenCL devices which can provide the information.

As a consequence, applications should always check for the presence of the extension string for each individual OpenCL device for which they intend to issue the new query for and should not have any assumptions about the availability of the extension on any given platform.

39.1. General information

39.1.1. Name Strings

```
cl_khr_pci_bus_info
```

39.1.2. Version History

Date	Version	Description
2021-04-19	1.0.0	Initial version.

39.1.3. Dependencies

This extension is written against the OpenCL API Specification Version V3.0.6.

This extension requires OpenCL 1.0.

39.2. New API Types

Structure returned by the device info query for CL_DEVICE_PCI_BUS_INFO_KHR:

```
typedef struct _cl_device_pci_bus_info_khr {
    cl_uint    pci_domain;
    cl_uint    pci_bus;
    cl_uint    pci_device;
    cl_uint    pci_function;
} cl_device_pci_bus_info_khr;
```

39.3. New API Enums

Accepted value for the *param_name* parameter to **clGetDeviceInfo**:

```
#define CL_DEVICE_PCI_BUS_INFO_KHR 0x410F
```

39.4. Modifications to the OpenCL API Specification

39.4.1. Section 4.2 - Querying Devices:

Add to Table 5 - OpenCL Device Queries:

Table 5. OpenCL Device Queries

cl_device_info	Return Type	Description
CL_DEVICE_PCI_BUS_INFO_KHR	<pre>cl_device_pci_bus_ info_khr</pre>	Returns PCI bus information for the device.
		The PCI bus information is returned as a single structure that includes the PCI bus domain, the PCI bus identifier, the PCI device identifier, and the PCI device function identifier.

Chapter 40. Extended Bit Operations

This extension adds OpenCL C functions for performing extended bit operations. Specifically, the following functions are added:

- bitfield insert: insert bits from one source operand into another source operand.
- bitfield extract: extract bits from a source operand, with sign- or zero-extension.
- bit reverse: reverse the bits of a source operand.

40.1. General Information

40.1.1. Name Strings

cl_khr_extended_bit_ops

40.1.2. Version History

Date	Version	Description
2021-04-22	1.0.0	Initial version.

40.1.3. Dependencies

This extension is written against the OpenCL 3.0 C Language Specification and the OpenCL SPIR-V Environment Specification Version V3.0.6.

This extension requires OpenCL 1.0.

40.2. New OpenCL C Functions

```
gentype bitfield_insert( gentype base, gentype insert, uint offset, uint count )
igentype bitfield_extract_signed( gentype base, uint offset, uint count )
ugentype bitfield_extract_unsigned( gentype base, uint offset, uint count )
gentype bit_reverse( gentype base )
```

40.3. Modifications to the OpenCL C Specification

40.3.1. Modify Section 6.15.3. Integer Functions:

Add a new Section 6.15.3.X. Extended Bit Operations:

The functions described in the following table can be used with built-in scalar or vector integer types to perform extended bit operations. The functions that operate on vector types operate component-wise. The description is per-component.

In the table below, the generic type name gentype refers to the built-in integer types char, charn,

uchar, ucharn, short, shortn, ushortn, int, intn, uint, uintn, long, longn, ulong, and ulongn. The generic type name igentype refers to the built-in signed integer types char, charn, short, shortn, int, intn, long, and longn. The generic type name ugentype refers to the built-in unsigned integer types uchar, ucharn, ushortn, uintn, uintn, ulong, and ulongn. n is 2, 3, 4, 8, or 16.

Table 47. Built-in Scalar and Vector Extended Bit Operations

Function

gentype bitfield_insert(
 gentype base, gentype insert,
 uint offset, uint count)

Description

Returns a copy of *base*, with a modified bitfield that comes from *insert*.

Any bits of the result value numbered outside [offset, offset + count - 1] (inclusive) will come from the corresponding bits in base.

Any bits of the result value numbered inside [offset, offset + count - 1] (inclusive) will come from the bits numbered [0, count - 1] (inclusive) of insert.

count is the number of bits to be modified. If *count* equals 0, the return value will be equal to *base*.

If *count* or *offset* or *offset* + *count* is greater than number of bits in <code>gentype</code> (for scalar types) or components of <code>gentype</code> (for vector types), the result is undefined.

igentype bitfield_extract_signed(
 gentype base,
 uint offset, uint count)

Returns an extracted bitfield from *base* with sign extension. The type of the return value is always a signed type.

The bits of *base* numbered in [*offset*, *offset* + *count* - 1] (inclusive) are returned as the bits numbered in [0, *count* - 1] (inclusive) of the result. The remaining bits in the result will be sign extended by replicating the bit numbered *offset* + *count* - 1 of *base*.

count is the number of bits to be extracted. If *count* equals 0, the result is 0.

If the *count* or *offset* or *offset* + *count* is greater than number of bits in gentype (for scalar types) or components of gentype (for vector types), the result is undefined.

Function Description Returns an extracted bitfield from base with ugentype bitfield_extract_unsigned(zero extension. The type of the return value is gentype base, always an unsigned type. uint offset, uint count) The bits of base numbered in [offset, offset + count - 1] (inclusive) are returned as the bits numbered in [0, count - 1] (inclusive) of the result. The remaining bits in the result will be zero. count is the number of bits to be extracted. If count equals 0, the result is 0. If the *count* or *offset* or *offset* + *count* is greater than number of bits in gentype (for scalar types) or components of gentype (for vector types), the result is undefined. Returns the value of *base* with reversed bits. gentype bit_reverse(That is, the bit numbered *n* of the result value gentype base) will be taken from the bit numbered width - n -1 of base (for scalar types) or a component of base (for vector types), where width is number of bits of gentype (for scalar types) or components of gentype (for vector types).

Chapter 41. Suggested Local Work Size Query

This extension adds the ability to query a suggested local work group size for a kernel running on a device for a specified global work size and global work offset. The suggested local work group size will match the work group size that would be chosen if the kernel were enqueued with the specified global work size and global work offset and a NULL local work size.

By using the suggested local work group size query an application has greater insight into the local work group size chosen by the OpenCL implementation, and the OpenCL implementation need not re-compute the local work group size if the same kernel is enqueued multiple times with the same parameters.

41.1. General Information

41.1.1. Name Strings

cl_khr_suggested_local_work_size

41.1.2. Version History

Date	Version	Description
2021-04-22	1.0.0	Initial version.

41.1.3. Dependencies

This extension is written against the OpenCL API Specification Version V3.0.6.

This extension requires OpenCL 1.0.

41.2. New API Functions

```
cl_int clGetKernelSuggestedLocalWorkSizeKHR(
    cl_command_queue command_queue,
    cl_kernel kernel,
    cl_uint work_dim,
    const size_t *global_work_offset,
    const size_t *global_work_size,
    size_t *suggested_local_work_size);
```

41.3. Modifications to the OpenCL API Specification

41.3.1. Section 5.9 - Kernel Objects:

41.3.1.1. New Section 5.9.4.X - Suggested Local Work Size Query

To query a suggested local work size for a kernel object, call the function

```
cl_int clGetKernelSuggestedLocalWorkSizeKHR(
    cl_command_queue command_queue,
    cl_kernel kernel,
    cl_uint work_dim,
    const size_t *global_work_offset,
    const size_t *global_work_size,
    size_t *suggested_local_work_size);
```

The returned suggested local work size is expected to match the local work size that would be chosen if the specified kernel object, with the same kernel arguments, were enqueued into the specified command queue with the specified global work size, specified global work offset, and with a NULL local work size.

- *command_queue* specifies the command queue and device for the query.
- *kernel* specifies the kernel object and kernel arguments for the query. The OpenCL context associated with *kernel* and *command_queue* must the same.
- work_dim specifies the number of work dimensions in the input global work offset and global work size, and the output suggested local work size.
- *global_work_offset* can be used to specify an array of at least *work_dim* global ID offset values for the query. This is optional and may be NULL to indicate there is no global ID offset.
- *global_work_size* is an array of at least *work_dim* values describing the global work size for the query.
- *suggested_local_work_size* is an output array of at least *work_dim* values that will contain the result of the query.

clGetKernelSuggestedLocalWorkSizeKHR returns **CL_SUCCESS** if the query executed successfully. Otherwise, it returns one of the following errors:

- CL_INVALID_COMMAND_QUEUE if command_queue is not a valid host command queue.
- CL_INVALID_KERNEL if *kernel* is not a valid kernel object.
- CL_INVALID_CONTEXT if the context associated with *kernel* is not the same as the context associated with *command_queue*.
- CL_INVALID_PROGRAM_EXECUTABLE if there is no successfully built program executable available for *kernel* for the device associated with *command_queue*.
- CL_INVALID_KERNEL_ARGS if all argument values for kernel have not been set.
- CL_MISALIGNED_SUB_BUFFER_OFFSET if a sub-buffer object is set as an argument to *kernel* and the offset specified when the sub-buffer object was created is not aligned to CL_DEVICE_MEM_BASE_ADDR_ALIGN for the device associated with *command_queue*.

- CL_INVALID_IMAGE_SIZE if an image object is set as an argument to *kernel* and the image dimensions are not supported by device associated with *command_queue*.
- CL_IMAGE_FORMAT_NOT_SUPPORTED if an image object is set as an argument to *kernel* and the image format is not supported by the device associated with *command_queue*.
- CL_INVALID_OPERATION if an SVM pointer is set as an argument to *kernel* and the device associated with *command_queue* does not support SVM or the required SVM capabilities for the SVM pointer.
- CL_INVALID_WORK_DIMENSION if *work_dim* is not a valid value (i.e. a value between 1 and CL_DEVICE_MAX_WORK_ITEM_DIMENSIONS).
- CL_INVALID_GLOBAL_WORK_SIZE if *global_work_size* is NULL or if any of the values specified in *global_work_size* are 0.
- CL_INVALID_GLOBAL_WORK_SIZE if any of the values specified in *global_work_size* exceed the maximum value representable by size_t on the device associated with *command_queue*.
- CL_INVALID_GLOBAL_OFFSET if the value specified in *global_work_size* plus the corresponding value in *global_work_offset* for dimension exceeds the maximum value representable by size_t on the device associated with *command_queue*.
- CL_OUT_OF_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- CL_OUT_OF_HOST_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the host.



These error conditions are consistent with error conditions for clEnqueueNDRangeKernel.

Chapter 42. Integer dot product

This extension adds support for SPIR-V instructions and OpenCL C built-in functions to compute the dot product of vectors of integers.

42.1. General Information

42.1.1. Name Strings

cl_khr_integer_dot_product

42.1.2. Version History

Date	Version	Description
2021-06-17	1.0.0	Initial version.

42.1.3. Dependencies

This extension is written against the OpenCL Specification Version 3.0.6, and OpenCL C Specification Version 3.0.6 and OpenCL Environment Specification Version 3.0.6.

This extension requires OpenCL 1.0.

42.1.4. Contributors

Kévin Petit, Arm Ltd.
Jeremy Kemp, Imagination Technologies
Ben Ashbaugh, Intel
Ruihao Zhang, Qualcomm
Stuart Brady, Arm Ltd
Balaji Calidas, Qualcomm
Ayal Zaks, Intel

42.2. New API Enums

Accepted value for the *param_name* parameter to **clGetDeviceInfo**:

```
CL_DEVICE_INTEGER_DOT_PRODUCT_INPUT_4x8BIT_PACKED_KHR (1 << 0)
CL_DEVICE_INTEGER_DOT_PRODUCT_INPUT_4x8BIT_KHR (1 << 1)

CL_DEVICE_INTEGER_DOT_PRODUCT_CAPABILITIES_KHR 0x1073
```

42.3. New OpenCL C Functions

This extension defines a number of new functions that operate on vectors of integers. The exact

function overloads available depend on the features supported by the target device.

```
uint dot(uchar4 a, uchar4 b);
int dot(char4 a, char4 b);
int dot(uchar4 a, char4 b);
int dot(char4 a, uchar4 b);
uint dot_acc_sat(uchar4 a, uchar4 b, uint acc);
int dot_acc_sat(char4 a, char4 b, int acc);
int dot_acc_sat(uchar4 a, char4 b, int acc);
int dot_acc_sat(char4 a, uchar4 b, int acc);
uint dot_4x8packed_uu_uint(uint a, uint b);
int dot_4x8packed_ss_int(uint a, uint b);
int dot_4x8packed_us_int(uint a, uint b);
int dot_4x8packed_su_int(uint a, uint b);
uint dot_acc_sat_4x8packed_uu_uint(uint a, uint b, uint acc);
int dot_acc_sat_4x8packed_ss_int(uint a, uint b, int acc);
int dot_acc_sat_4x8packed_us_int(uint a, uint b, int acc);
int dot_acc_sat_4x8packed_su_int(uint a, uint b, int acc);
```

42.4. Modifications to the OpenCL API Specification

(Modify Section 4.2, Querying Devices)

(Add the following to Table 4.3, Device Queries)

cl_device_info	Return Type	Description
CL_DEVICE_INTEGER_DOT_ PRODUCT_CAPABILITIES_ KHR	<pre>cl_device_integer_dot_ product_capabilities_ khr</pre>	Returns the integer dot product capabilities supported by the device.
		CL_DEVICE_INTEGER_DOT_PRODUCT_INPUT_4x8BIT_
		PACKED_KHR is always set indicating that all
		implementations that support
		cl_khr_integer_dot_product must support dot
		product built-in functions and, when SPIR-V is
		supported, SPIR-V instructions that take four-
		component vectors of 8-bit integers packed into
		32-bit integers as input.
		CL_DEVICE_INTEGER_DOT_PRODUCT_INPUT_4x8BIT_KHR
		is set when dot product built-in functions and,
		when SPIR-V is supported, SPIR-V instructions
		that take four-component of 8-bit elements as
		input are supported.

OpenCL 3 devices must report the following feature macros via CL_DEVICE_OPENCL_C_FEATURES when the corresponding bit is set in the bitfield returned for CL_DEVICE_INTEGER_DOT_PRODUCT_CAPABILITIES_KHR:

Feature bit	Feature macro
CL_DEVICE_INTEGER_DOT_PRODUCT_INPUT_4x8BIT_ PACKED_KHR	opencl_c_integer_dot_product_input_4x8bit_packed
CL_DEVICE_INTEGER_DOT_PRODUCT_INPUT_4x8BIT_KHR	opencl_c_integer_dot_product_input_4x8bit

42.5. Modifications to the OpenCL C Specification

(Modify section 6.13.3, Integer Functions)

The following built-in functions and preprocessor definitions are added:

```
#define cl khr integer dot product 1
if (CL_DEVICE_INTEGER_DOT_PRODUCT_INPUT_4x8BIT_PACKED_KHR) {
    #define __opencl_c_integer_dot_product_input_4x8bit_packed 1
    uint dot_4x8packed_uu_uint(uint a, uint b);
    int dot 4x8packed ss int(uint a, uint b);
    int dot_4x8packed_us_int(uint a, uint b);
    int dot_4x8packed_su_int(uint a, uint b);
    uint dot_acc_sat_4x8packed_uu_uint(uint a, uint b, uint acc);
    int dot_acc_sat_4x8packed_ss_int(uint a, uint b, int acc);
    int dot acc sat 4x8packed us int(uint a, uint b, int acc);
    int dot_acc_sat_4x8packed_su_int(uint a, uint b, int acc);
}
if (CL_DEVICE_INTEGER_DOT_PRODUCT_INPUT_4x8BIT_KHR) {
    #define __opencl_c_integer_dot_product_input_4x8bit 1
    uint dot(uchar4 a, uchar4 b);
    int dot(char4 a, char4 b);
    int dot(uchar4 a, char4 b);
    int dot(char4 a, uchar4 b);
    uint dot acc sat(uchar4 a, uchar4 b, uint acc);
    int dot_acc_sat(char4 a, char4 b, int acc);
    int dot_acc_sat(uchar4 a, char4 b, int acc);
    int dot_acc_sat(char4 a, uchar4 b, int acc);
}
```

- dot returns the dot product of the two input vectors a and b. The components of a and b are sign- or zero-extended to the width of the destination type and the vectors with extended components are multiplied component-wise. All the components of the resulting vectors are added together to form the final result.
- dot_acc_sat returns the saturating addition of the dot product of the two input vectors a and b and the accumulator acc:

```
product = dot(a,b);
result = add_sat(product, acc);
```

dot_*_4x8packed_XY_R returns the dot product of the two vectors packed into a and b (lowest component in least significant byte). The components are unpacked, sign- or zero-extended to the width of the destination type before the multiplications and additions. X represents the signedness of the components of a, Y that of the components of b. R is the return type.

42.6. Modifications to the OpenCL SPIR-V Environment Specification

See OpenCL SPIR-V Environment Specification.

42.7. Interactions with Other Extensions

If cl_khr_il_program is supported then the SPIR-V environment specification modifications described above apply.

Chapter 43. Extensions to the OpenCL SPIR-V Environment

An OpenCL SPIR-V environment may be modified by OpenCL extensions. Please refer to the OpenCL SPIR-V Environment Specification for descriptions how OpenCL extensions modify an OpenCL SPIR-V environment. In addition to the extensions described in this document, the OpenCL SPIR-V Environment Specification also describes how the following OpenCL extensions modify an OpenCL SPIR-V environment:

- cl_khr_spirv_no_integer_wrap_decoration
- cl_khr_spirv_extended_debug_info
- cl_khr_spirv_linkonce_odr

Index

C

clCreateEventFromEGLSyncKHR, 168
clCreateEventFromGLsyncKHR, 104
clCreateFromD3D10BufferKHR, 115
clCreateFromD3D10Texture2DKHR, 115
clCreateFromD3D10Texture3DKHR, 116
clCreateFromD3D11BufferKHR, 131
clCreateFromD3D11Texture2DKHR, 131
clCreateFromD3D11Texture3DKHR, 132
clCreateFromDX9MediaSurfaceKHR, 145
clCreateFromEGLImageKHR, 172
clCreateFromGLBuffer, 91
clCreateFromGLRenderbuffer, 96
clCreateFromGLTexture, 92
clCreateProgramWithILKHR, 188
clEnqueueAcquireD3D10ObjectsKHR, 119
clEnqueueAcquireD3D11ObjectsKHR, 135
clEnqueueAcquireDX9MediaSurfacesKHR, 147
clEnqueueAcquireEGLObjectsKHR, 174
clEnqueueAcquireGLObjects, 99
clEnqueueReleaseD3D10ObjectsKHR, 120
clEnqueueReleaseD3D11ObjectsKHR, 136
clEnqueueReleaseDX9MediaSurfacesKHR, 149
clEnqueueReleaseEGLObjectsKHR, 175
clEnqueueReleaseGLObjects, 100
clGetDeviceIDsFromD3D10KHR, 113
clGetDeviceIDsFromD3D11KHR, 129
$cl Get Device IDs From DX9 Media Adapter KHR, {\color{blue}144}$
$cl Get Extension Function Address For Platform, {\color{red} 4}$
clGetGLContextInfoKHR, 86
clGetGLObjectInfo, 97
clGetGLTextureInfo, 98
clGetKernelSubGroupInfoKHR, 197
clIcdGetPlatformIDsKHR, 10
clTerminateContextKHR, 183

Appendix A: Extensions Promoted to Core Features

A.1. For OpenCL 1.1:

- The functionality previously described by cl_khr_byte_addressable_store is now part of the core feature set.
- The functionality previously described by cl_khr_global_int32_base_atomics,
 cl_khr_global_int32_extended_atomics,
 cl_khr_local_int32_base_atomics,
 and cl_khr_local_int32_extended_atomics is now part of the core feature set.

A.2. For OpenCL 1.2:

• The functionality previously described by cl_khr_fp64 is now an optional core feature.

A.3. For OpenCL 2.0:

- The functionality described by cl_khr_3d_image_writes is part of the core feature set.
- The functionality described by **cl_khr_create_command_queue** is part of the core feature set.
- The functionality described by cl_khr_depth_images is now part of the core feature set.
- The functionality described by cl_khr_image2d_from_buffer is now part of the core feature set.

A.4. For OpenCL 2.1:

- The functionality described by cl_khr_il_program is now part of the core feature set.
- The API functionality described by **cl_khr_subgroups** is now part of the core API feature set, but the built-in functions described by **cl_khr_subgroups** must still be accessed as an extension to the OpenCL 2.0 C Language specification.

A.5. For OpenCL 3.0:

• The built-in functions described by **cl_khr_subgroups** are now supported in OpenCL C 3.0 when the <u>_opencl_c_subgroups</u> feature is supported.

Appendix B: Deprecated Extensions

B.1. For OpenCL 1.1:

• The **cl_khr_select_fprounding_mode** extension has been deprecated. Its use is no longer recommended.

Appendix C: Quick Reference

Extension Name	Brief Description	Status
cl_khr_3d_image_writes	Write to 3D images	Core Feature in OpenCL 2.0
cl_khr_async_work_group_copy_fence	Asynchronous Copy Fences	Provisional Extension
cl_khr_byte_addressable_store	Read and write from 8-bit and 16-bit pointers	Core Feature in OpenCL 1.1
cl_khr_create_command_queue	API to Create Command Queues with Properties	Core Feature in OpenCL 2.0
cl_khr_d3d10_sharing	Share Direct3D 10 Buffers and Textures with OpenCL	Extension
cl_khr_d3d11_sharing	Share Direct3D 11 Buffers and Textures with OpenCL	Extension
cl_khr_depth_images	Single Channel Depth Images	Core Feature in OpenCL 2.0
cl_khr_device_enqueue_local_arg_types	Pass Non-Void Local Pointers to Child Kernels	Extension
cl_khr_device_uuid	Unique Device and Driver Identifier Queries	Extension
cl_khr_dx9_media_sharing	Share DirectX 9 Media Surfaces with OpenCL	Extension
cl_khr_egl_event	Share EGL Sync Objects with OpenCL	Extension
cl_khr_egl_image	Share EGL Images with OpenCL	Extension
cl_khr_extended_async_copies	2D and 3D Async Copies	Provisional Extension
cl_khr_extended_bit_ops	Bit Insert, Extract, and Reverse Operations	Extension
cl_khr_extended_versioning	Extend versioning of platform, devices, extensions, etc.	Extension
cl_khr_fp16	Operations on 16-bit Floating-Point Values	Extension
cl_khr_fp64	Operations on 64-bit Floating-Point Values	Optional Core Feature in OpenCL 1.2
cl_khr_gl_depth_images	Share OpenGL Depth Images with OpenCL	Extension

Extension Name	Brief Description	Status
cl_khr_gl_event	Share OpenGL Fence Sync Objects with OpenCL	Extension
cl_khr_gl_msaa_sharing	Share OpenGL MSAA Textures with OpenCL	Extension
cl_khr_gl_sharing	Sharing OpenGL Buffers and Textures with OpenCL	Extension
cl_khr_global_int32_base_atomics	Basic Atomic Operations on 32-bit Integers in Global Memory	Core Feature in OpenCL 1.1
cl_khr_global_int32_extended_atomics	Extended Atomic Operations on 32-bit Integers in Global Memory	Core Feature in OpenCL 1.1
cl_khr_icd	Installable Client Drivers	Extension
cl_khr_il_program	Support for Intermediate Language (IL) Programs (SPIR-V)	Core Feature in OpenCL 2.1
cl_khr_image2d_from_buffer	Create 2D Images from Buffers	Core Feature in OpenCL 2.0
cl_khr_initialize_memory	Initialize Local and Private Memory on Allocation	Extension
cl_khr_int64_base_atomics	Basic Atomic Operations on 64-bit Integers in Global and Local Memory	Extension
cl_khr_int64_extended_atomics	Extended Atomic Operations on 64-bit Integers in Global and Local Memory	Extension
cl_khr_local_int32_base_atomics	Basic Atomic Operations on 32-bit Integers in Local Memory	Core Feature in OpenCL 1.1
cl_khr_local_int32_extended_atomics	Extended Atomic Operations on 32-bit Integers in Local Memory	Core Feature in OpenCL 1.1
cl_khr_integer_dot_product	Integer dot product operations	Extension
cl_khr_mipmap_image	Create and Use Images with Mipmaps	Extension
cl_khr_mipmap_image_writes	Write to Images with Mipmaps	Extension
cl_khr_pci_bus_info	Query PCI Bus Information for an OpenCL Device	Extension
cl_khr_priority_hints	Create Command Queues with Different Priorities	Extension
cl_khr_select_fprounding_mode	Set the Current Kernel Rounding Mode	DEPRECATED

Extension Name	Brief Description	Status
cl_khr_spir	Standard Portable Intermediate Representation Programs	Extension, Superseded by IL Programs / SPIR-V
cl_khr_srgb_image_writes	Write to sRGB Images	Extension
cl_khr_subgroups	Sub-Groupings of Work Items	Core Feature in OpenCL 2.1 (with minor changes)
cl_khr_subgroup_ballot	Exchange Ballots Among Sub- Groupings of Work Items	Extension
cl_khr_subgroup_clustered_reduce	Clustered Reductions for Sub- Groupings of Work Items	Extension
cl_khr_subgroup_extended_types	Additional Type Support for Sub- Group Functions	Extension
cl_khr_subgroup_named_barrier	Barriers for Subsets of a Work Group	Extension
cl_khr_subgroup_non_uniform_arithmetic	Sub-Group Arithmetic Functions in Non-Uniform Control Flow	Extension
cl_khr_subgroup_non_uniform_vote	Hold Votes Among Sub-Groupings of Work Items	Extension
cl_khr_subgroup_shuffle	General-Purpose Shuffles Among Sub-Groupings of Work Items	Extension
cl_khr_subgroup_shuffle_relative	Relative Shuffles Among Sub- Groupings of Work Items	Extension
cl_khr_suggested_local_work_size	Query a Suggested Local Work Size	Extension
cl_khr_terminate_context	Terminate an OpenCL Context	Extension
cl_khr_throttle_hints	Create Command Queues with Different Throttle Policies	Extension